

Lecture 3:

Parallel Programming Abstractions (and their corresponding HW/SW)

**Parallel Computer Architecture and Programming
CMU 15-418/15-618, Fall 2025**

Today's theme is a critical idea in this course.

And today's theme is:

Abstraction vs. Implementation

**Conflating abstraction with implementation
is a common cause for confusion in this**

An example: Programming with ISPC

ISPC

- Intel SPMD Program Compiler (**ISPC**)
- **SPMD: single program multiple data**
- <http://ispc.github.com/>

Recall: example program from

Compute $\sin(x)$ using Taylor expansion: $\sin(x) = x - x^3/3! + x^5/5! - x^7/7! + \dots$

```
void sinx(int N, int terms, float* x, float* result)
{
    for (int i=0; i<N; i++)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        int denom = 6; // 3!
        int sign = -1;

        for (int j=1; j<=terms; j++)
        {
            value += sign * numer / denom;
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }

        result[i] = value;
    }
}
```

sin(x) in ISPC

Compute $\sin(x)$ using Taylor expansion: $\sin(x) = x - x^3/3! + x^5/5! - \dots$

C++ code:

```
#include "sinx_ispc.h"

int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

ISPC code:

```
export void sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assume N % programCount = 0
    for (uniform int i=0; i<N; i+=programCount)
    {
        int idx = i + programIndex;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[idx] * x[idx];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

SPMD programming abstraction:

Call to ISPC function spawns “gang” of ISPC
“program instances”

$\sin(x)$ in ISPC

Compute $\sin(x)$ using Taylor expansion: $\sin(x) = x - x^3/3! + x^5/5! -$

C++ code:

```
#include "sinx_ispc.h"

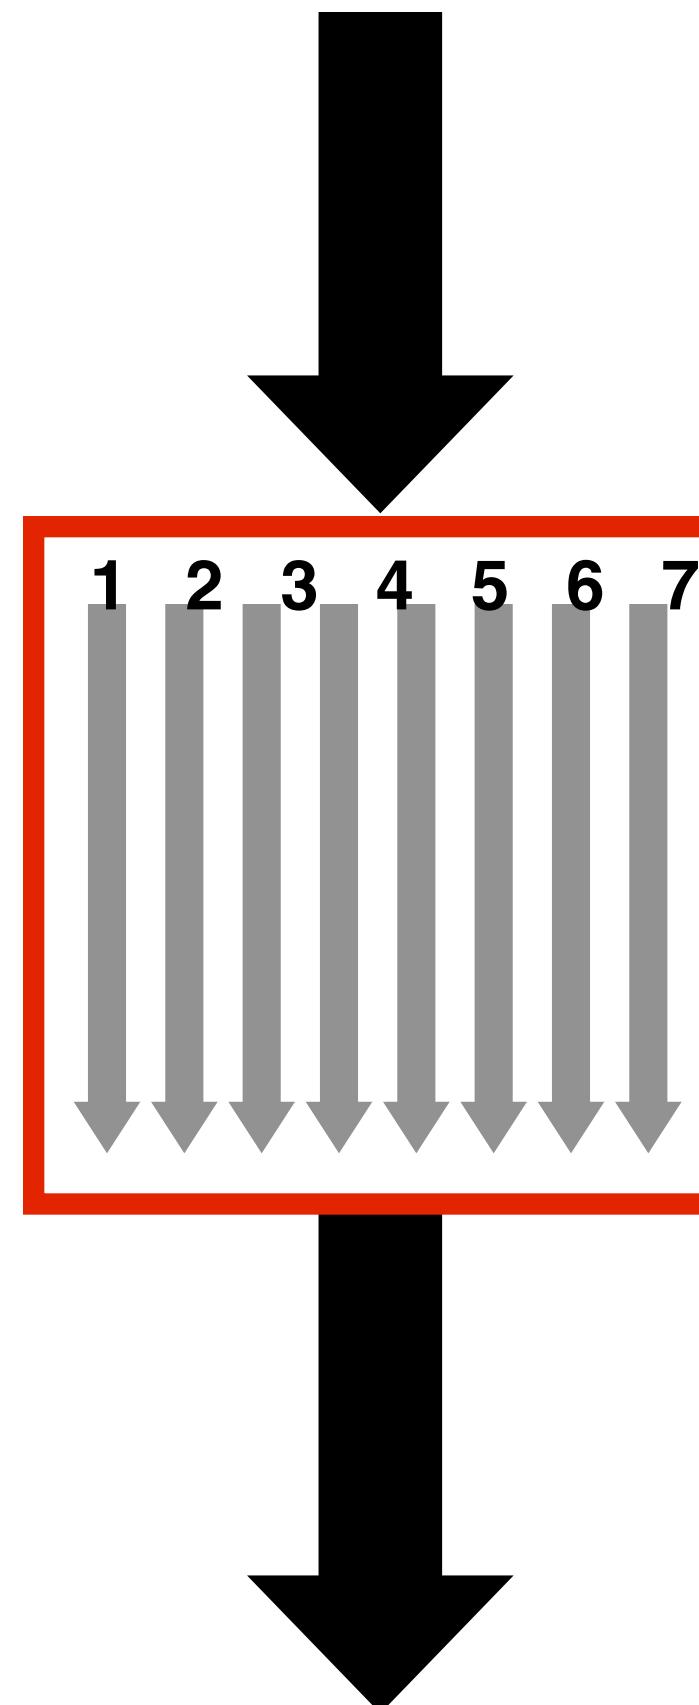
int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

SPMD programming abstraction:

Call to ISPC function spawns “**gang**” of ISPC
“**program instances**”



Sequential execution

Call to sinx()
Begin executing
programCount instances
of sinx() (ISPC code)

sinx() returns.
Completion of ISPC
program instances.

Sequential execution
(C code)

$\sin(x)$ in ISPC

“Interleaved” assignment of array elements to program instances

C++ code:

```
#include "sinx_ispc.h"

int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

ISPC Keywords:

programCount: number of simultaneously executing instances in the gang (uniform value)

programIndex: id of the current instance in the gang. (a non-uniform value: “varying”)

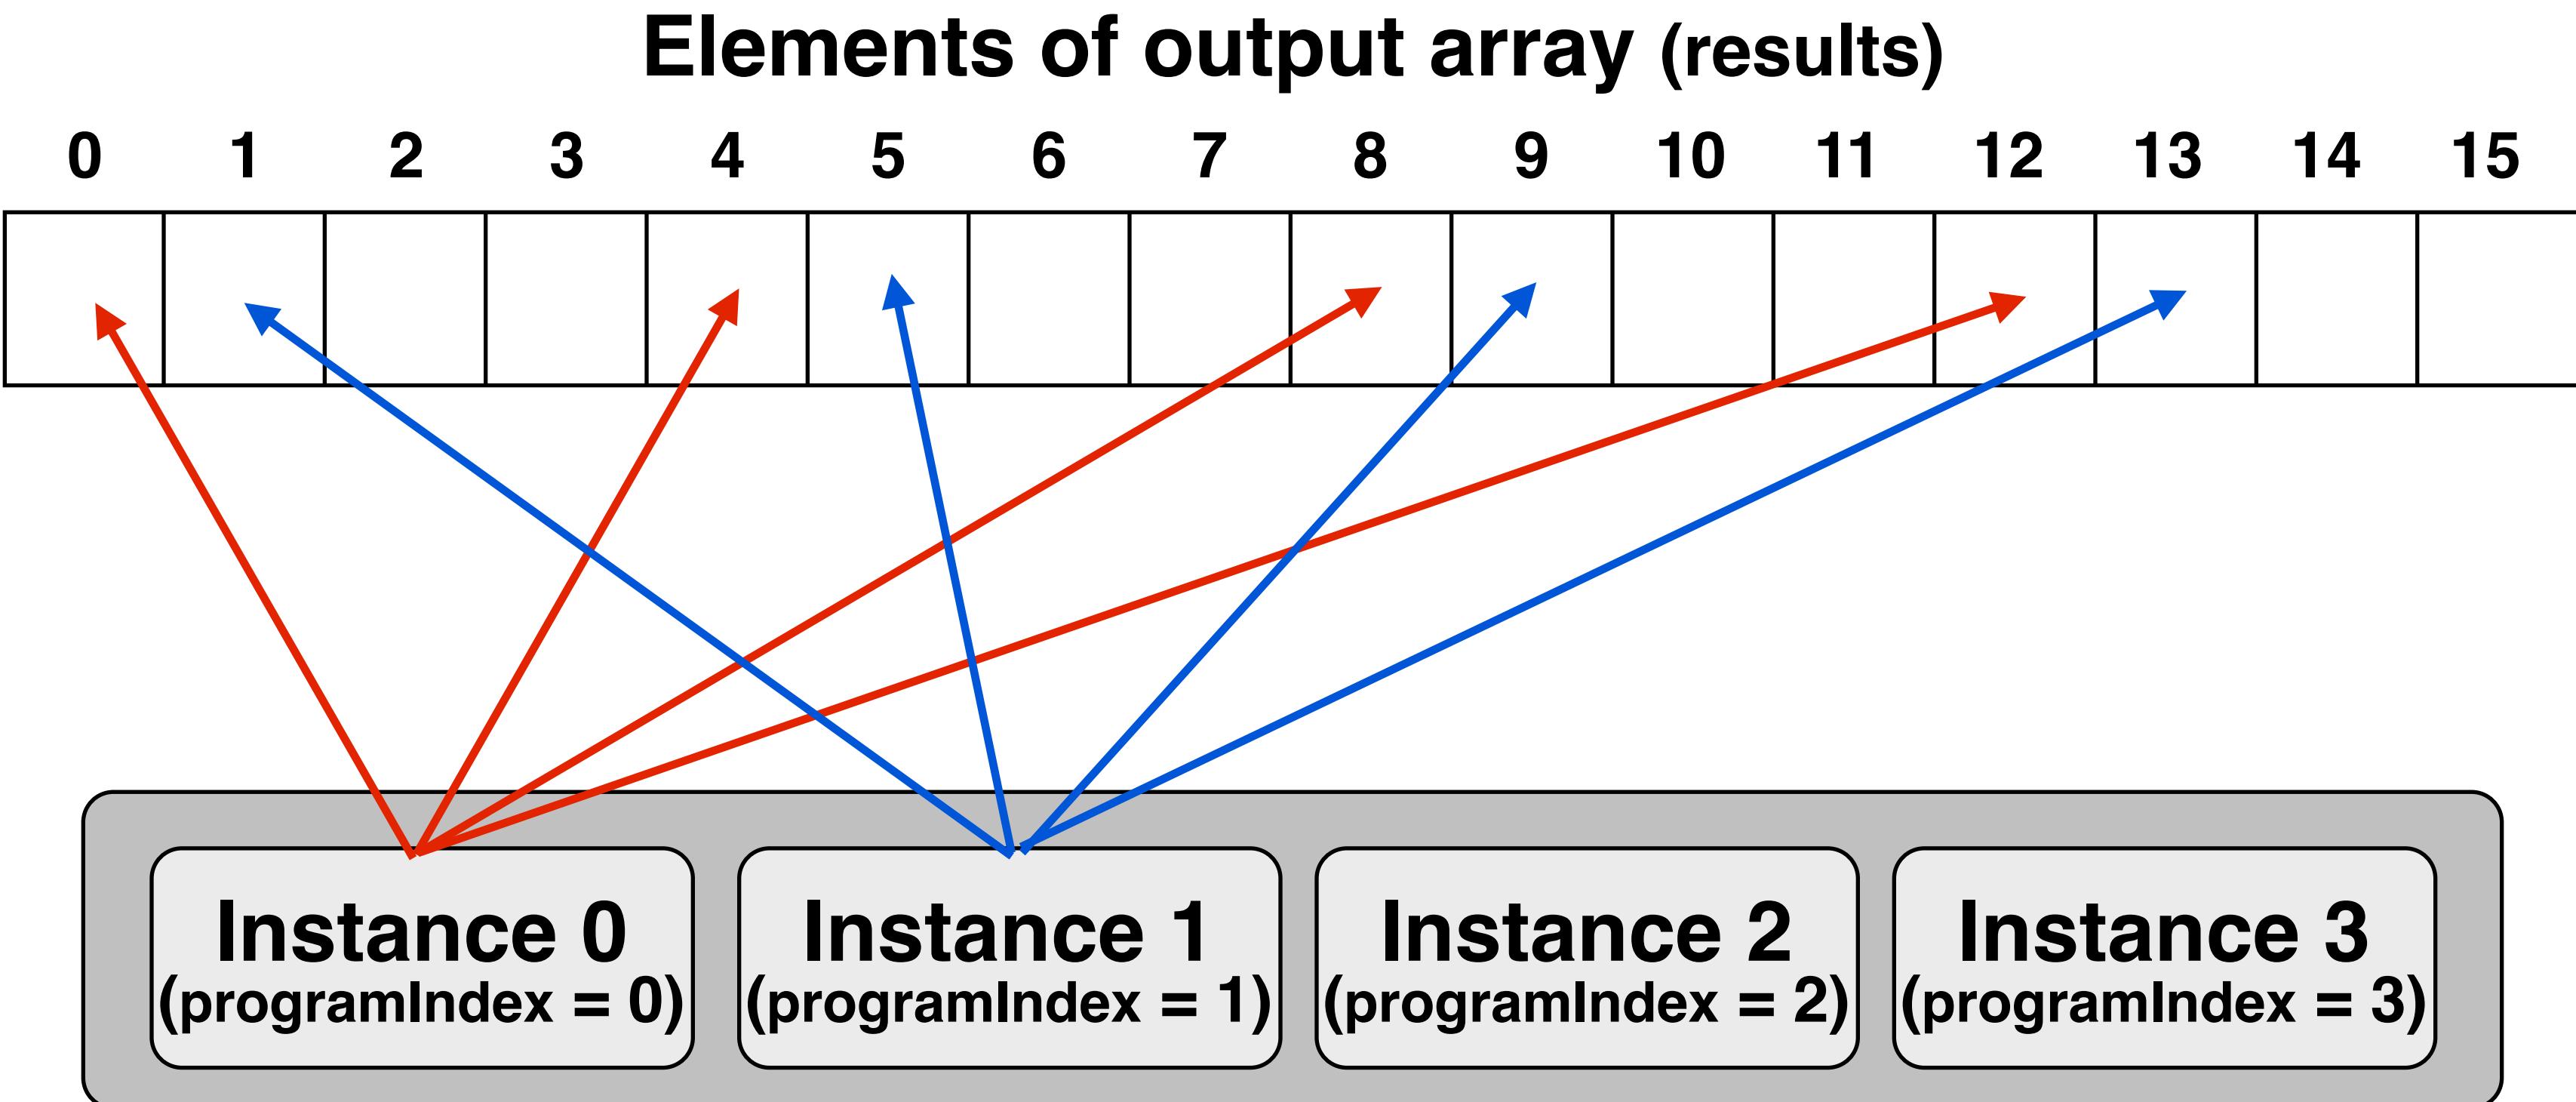
uniform: A type modifier. All instances

ISPC code:

```
export void sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assumes N % programCount = 0
    for (uniform int i=0; i<N; i+=programCount)
    {
        int idx = i + programIndex;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[idx] * x[idx];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

Interleaved assignment of program instances to loop



“Gang” of ISPC program instances

In this illustration: gang contains four instances: `programCount = 4`

ISPC implements the gang abstraction using SIMD

C++ code:

```
#include "sinx_ispc.h"

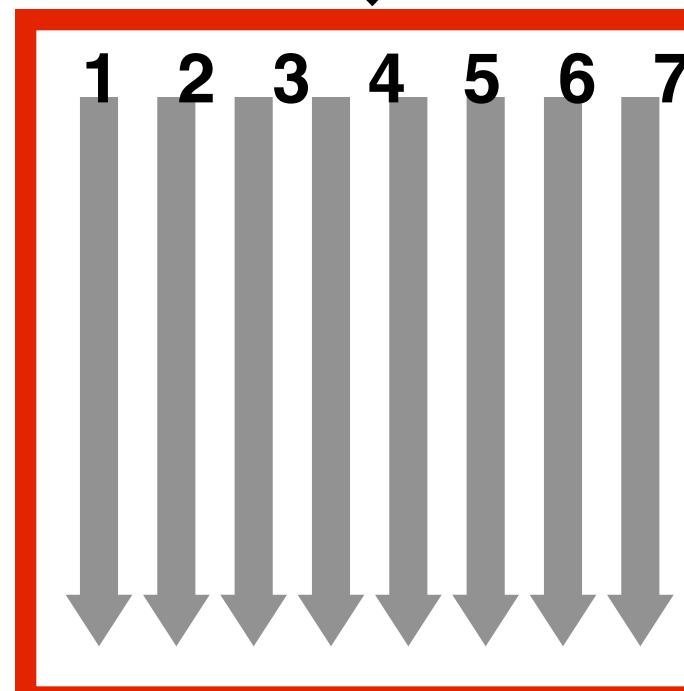
int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```



Sequential execution



Call to sinx()
Begin executing
programCount instances
of sinx() (ISPC code)



sinx() returns.
Completion of ISPC
program instances.

SPMD programming abstraction:

Call to ISPC function spawns “gang” of ISPC “program instances”

Sequential execution
(C code)

ISPC compiler generates SIMD implementation:

Number of instances in a gang is the **SIMD width** of the hardware (or a small multiple of SIMD width)

ISPC compiler generates binary (.o) with SIMD instructions

$\sin(x)$ in ISPC: version 2

“Blocked” assignment of elements to

C++ code:

```
#include "sinx_ispc.h"

int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

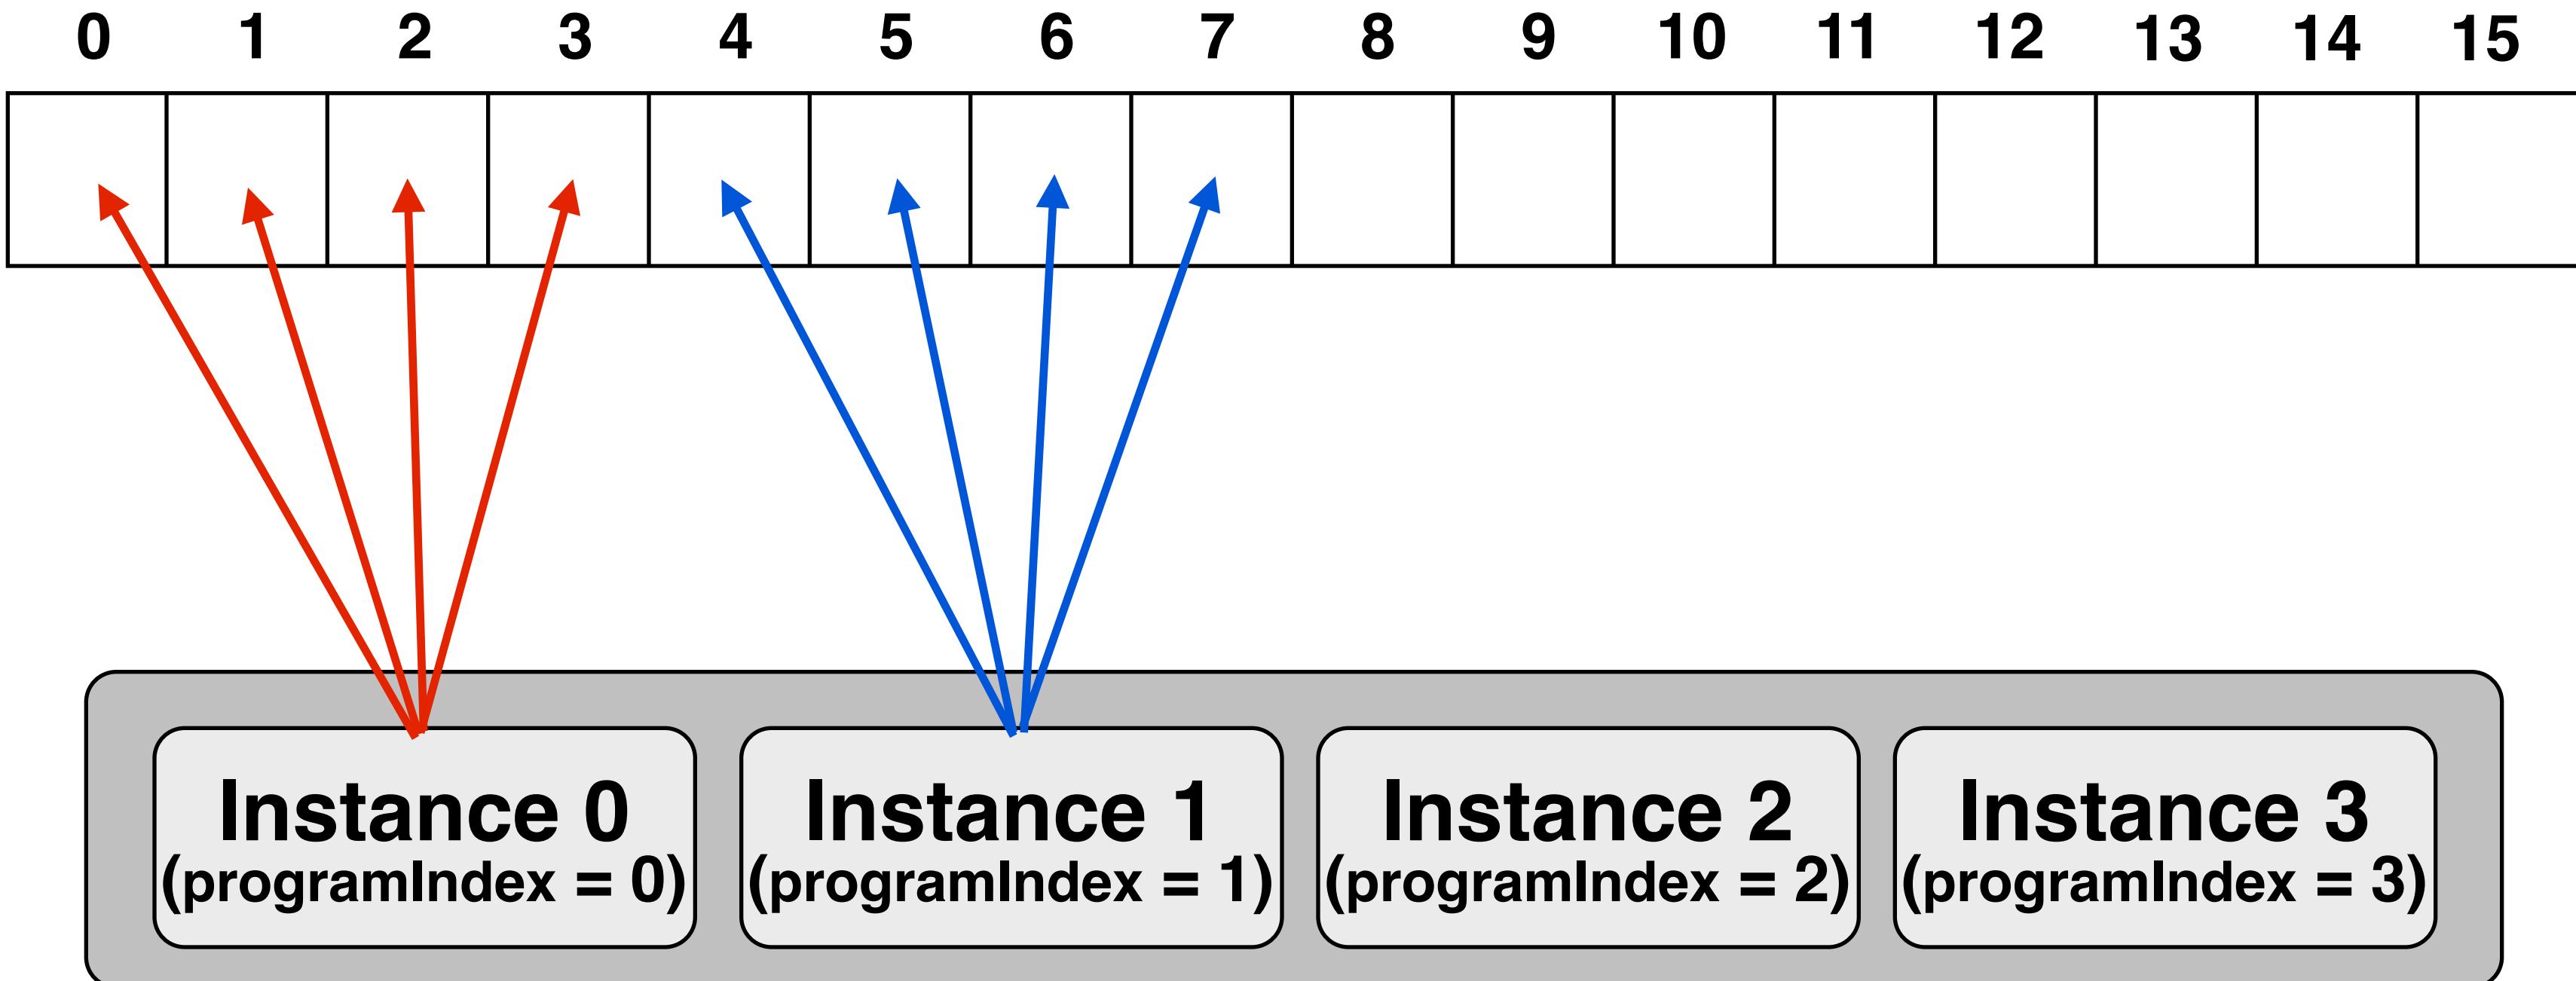
ISPC code:

```
export void sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)
{
    // assume N % programCount = 0
    uniform int count = N / programCount;
    int start = programIndex * count;
    for (uniform int i=0; i<count; i++)
    {
        int idx = start + i;
        float value = x[idx];
        float numer = x[idx] * x[idx] * x[idx];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[idx] * x[idx];
            denom *= (j+3) * (j+4);
            sign *= -1;
        }
        result[idx] = value;
    }
}
```

Blocked assignment of program instances to loop iterations

Elements of output array (results)



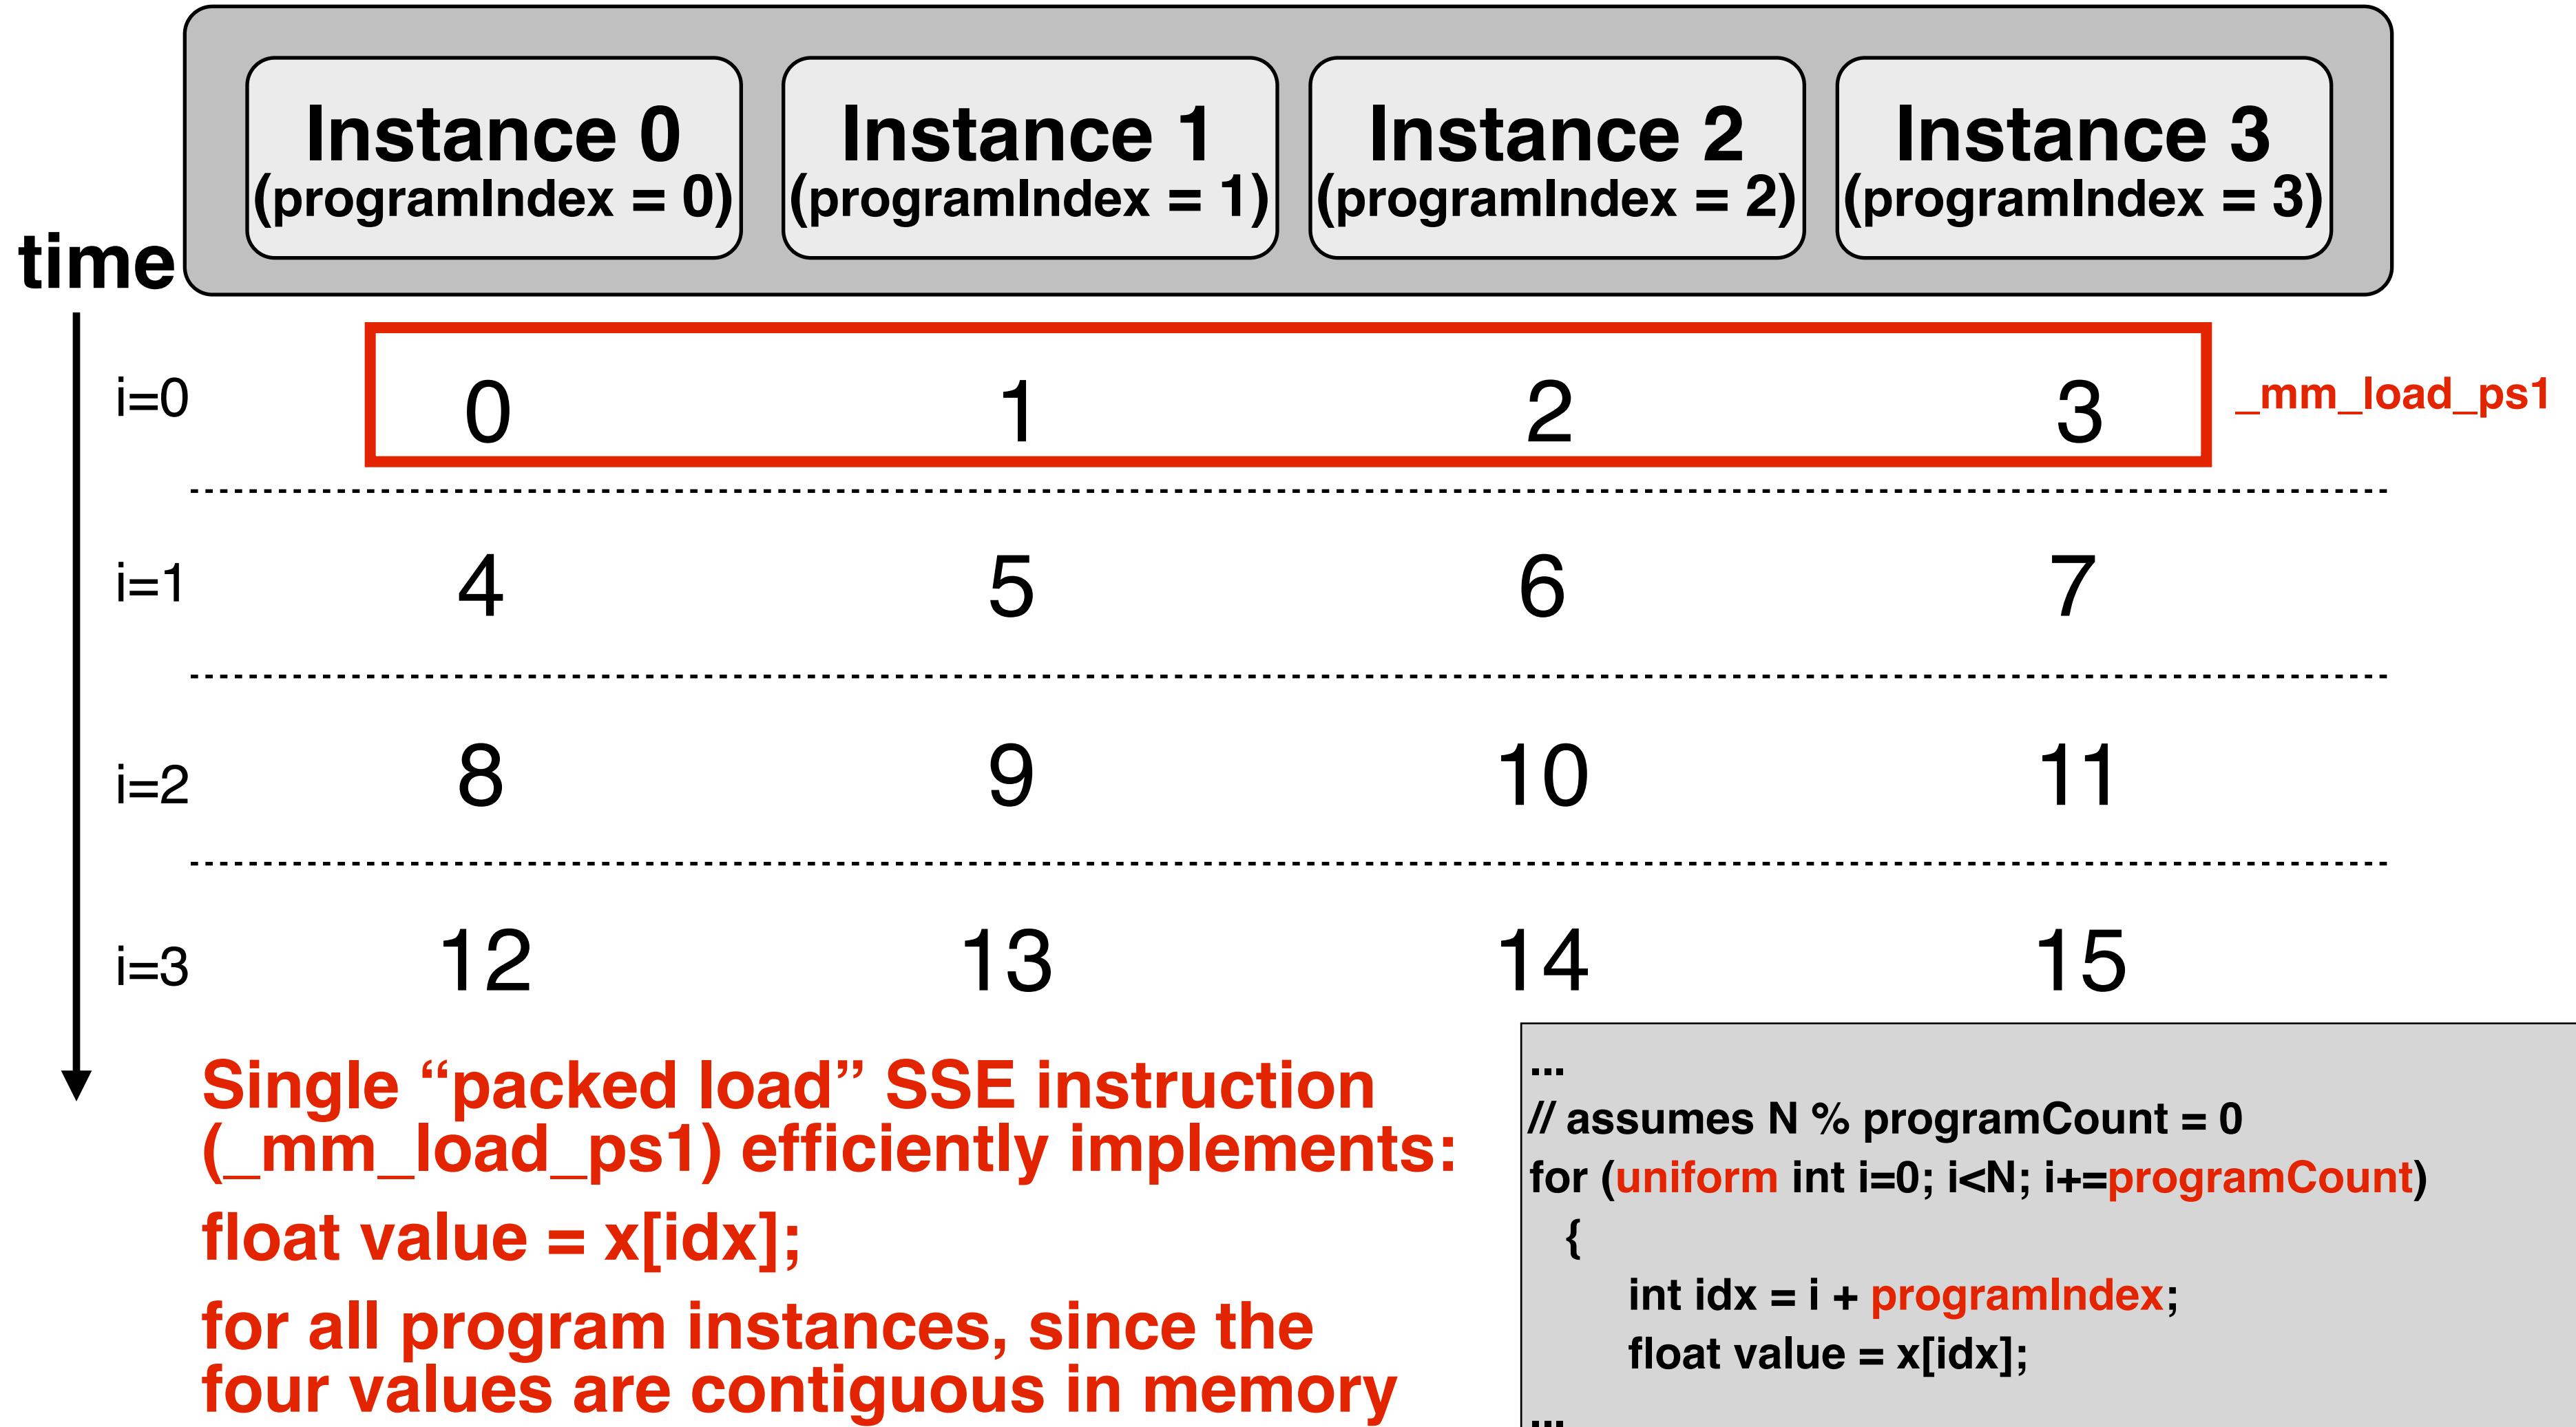
“Gang” of ISPC program instances

In this illustration: gang contains four instances: $\text{programCount} = 4$

Schedule: interleaved

“Gang” of ISPC program instances

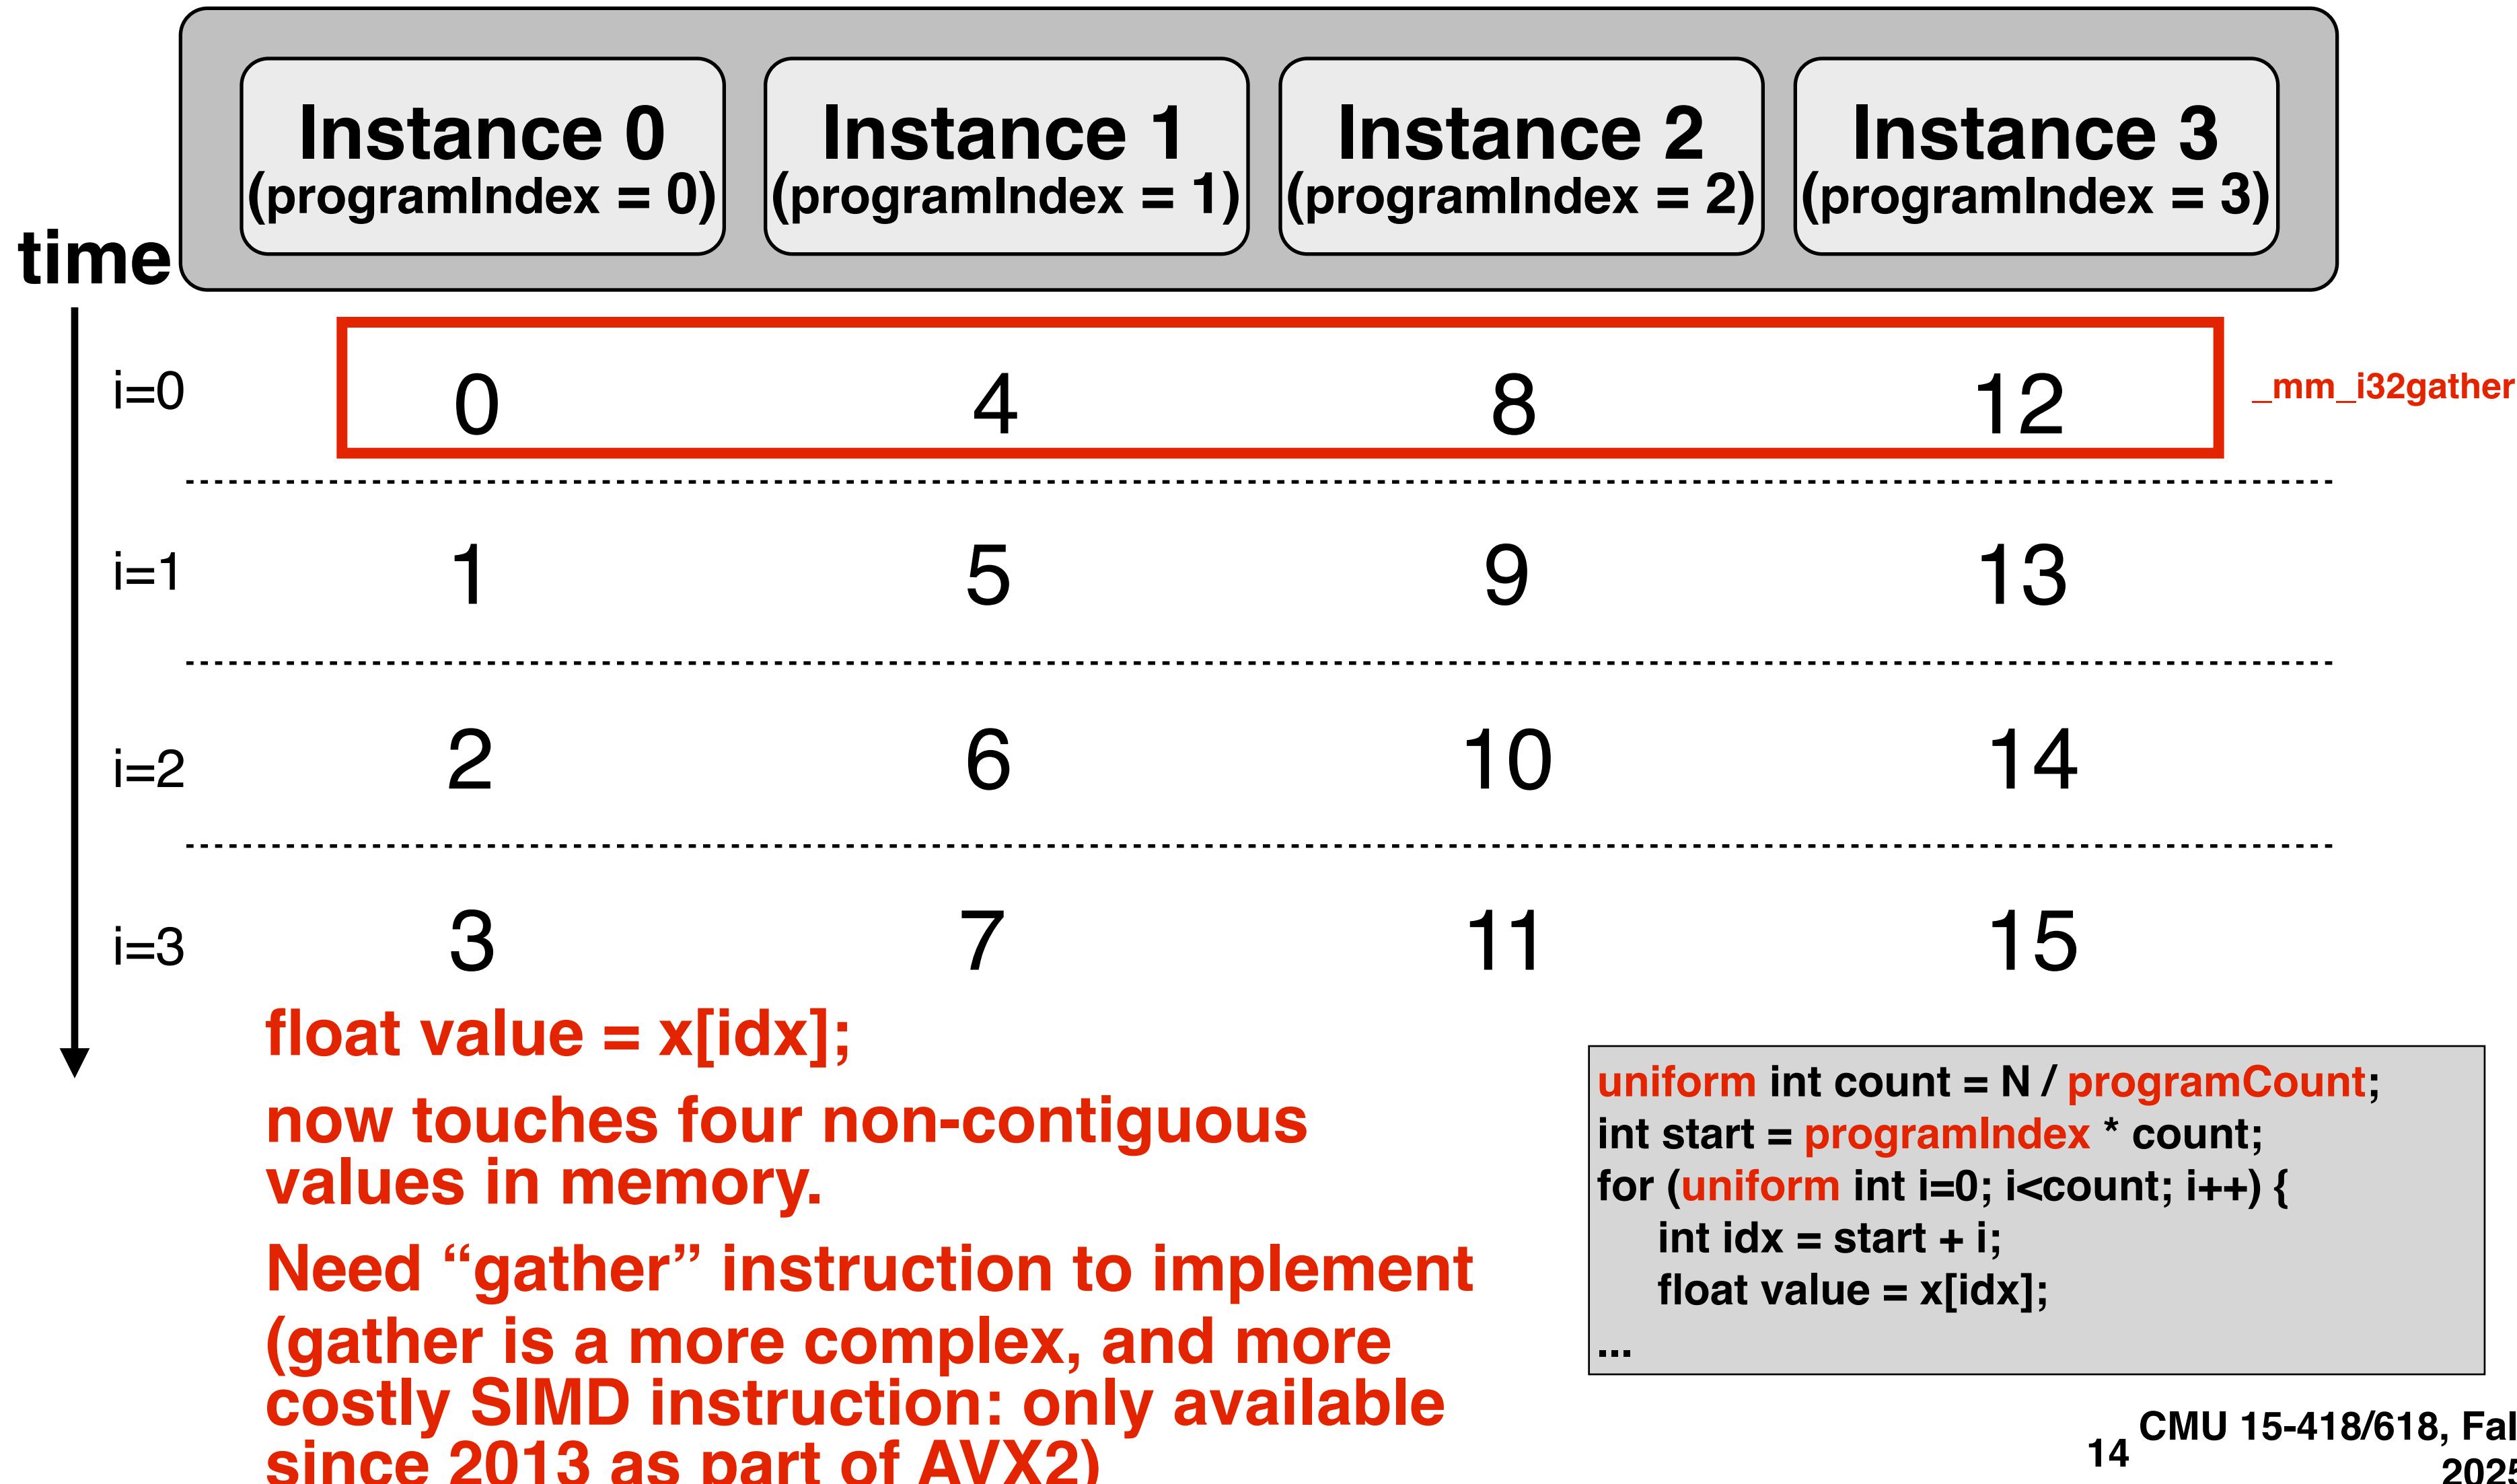
Gang contains four instances: `programCount = 4`



Schedule: blocked assignment

“Gang” of ISPC program instances

Gang contains four instances: `programCount = 4`



Raising level of abstraction

Compute $\sin(x)$ using Taylor expansion: $\sin(x) = x - x^3/3! + x^5/5! - \dots$

C++ code:

```
#include "sinx_ispc.h"

int N = 1024;
int terms = 5;
float* x = new float[N];
float* result = new float[N];

// initialize x here

// execute ISPC code
sinx(N, terms, x, result);
```

foreach: key ISPC language construct

- **foreach declares parallel loop iterations**
 - Programmer says: these are the iterations the instances in a gang must cooperatively perform
- **ISPC implementation assigns iterations to program instances in**

ISPC code:

```
export void sinx(
    uniform int N,
    uniform int terms,
    uniform float* x,
    uniform float* result)

{
    foreach (i = 0 ... N)
    {
        float value = x[i];
        float numer = x[i] * x[i] * x[i];
        uniform int denom = 6; // 3!
        uniform int sign = -1;

        for (uniform int j=1; j<=terms; j++)
        {
            value += sign * numer / denom
            numer *= x[i] * x[i];
            denom *= (2*j+2) * (2*j+3);
            sign *= -1;
        }
        result[i] = value;
    }
}
```

ISPC: abstraction vs.

- **Single program, multiple data (SPMD) programming model**
 - Programmer “thinks”: running a gang is spawning `programCount` logical instruction streams (each with a different value of `programIndex`)
 - This is the programming abstraction
 - Program is written in terms of this abstraction
- **Single instruction, multiple data (SIMD) implementation**
 - ISPC compiler emits vector instructions (SSE4 or AVX) that carry out the logic performed by a ISPC gang
 - ISPC compiler handles mapping of conditional control flow to vector instructions (by masking vector lanes, etc.)

ISPC discussion: sum

Compute the sum of all array elements in parallel

```
export uniform float sumall1(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        sum += x[i];  
    }  
  
    return sum;  
}
```

```
export uniform float sumall2(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum;  
    float partial = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        partial += x[i];  
    }  
  
    // from ISPC math library  
    sum = reduce_add(partial);  
  
    return sum;  
}
```

Correct ISPC solution

sum is of type **uniform float** (one copy of variable for all program instances)

x[i] is not a **uniform** expression (different value for each program instance)

ISPC discussion: sum

Compute the sum of all array elements in parallel

Each instance accumulates a **private partial sum** (no communication)

Partial sums are added together using the **reduce_add()** cross-instance communication primitive. The result is the same total sum for all program instances (reduce_add() returns a uniform float)

```
float sumall2(int N, float* x) {  
  
    float tmp[8]; // assume 16-byte alignment  
    __mm256 partial = __mm256_broadcast_ss(0.0f);  
  
    for (int i=0; i<N; i+=8)  
        partial = __mm256_add_ps(partial, __mm256_load_ps(&x[i]));  
  
    __mm256_store_ps(tmp, partial);  
  
    float sum = 0.f;  
    for (int i=0; i<8; i++)  
        sum += tmp[i];  
  
    return sum;  
}
```

```
float sumall2(  
    uniform int N,  
    uniform float* x)  
{  
    uniform float sum;  
    float partial = 0.0f;  
    foreach (i = 0 ... N)  
    {  
        partial += x[i];  
    }  
  
    // from ISPC math library  
    sum = reduce_add(partial);  
  
    return sum;  
}
```

* Self-test: If you understand why this implementation complies with the semantics of the ISPC gang abstraction, then you've got good

ISPC tasks

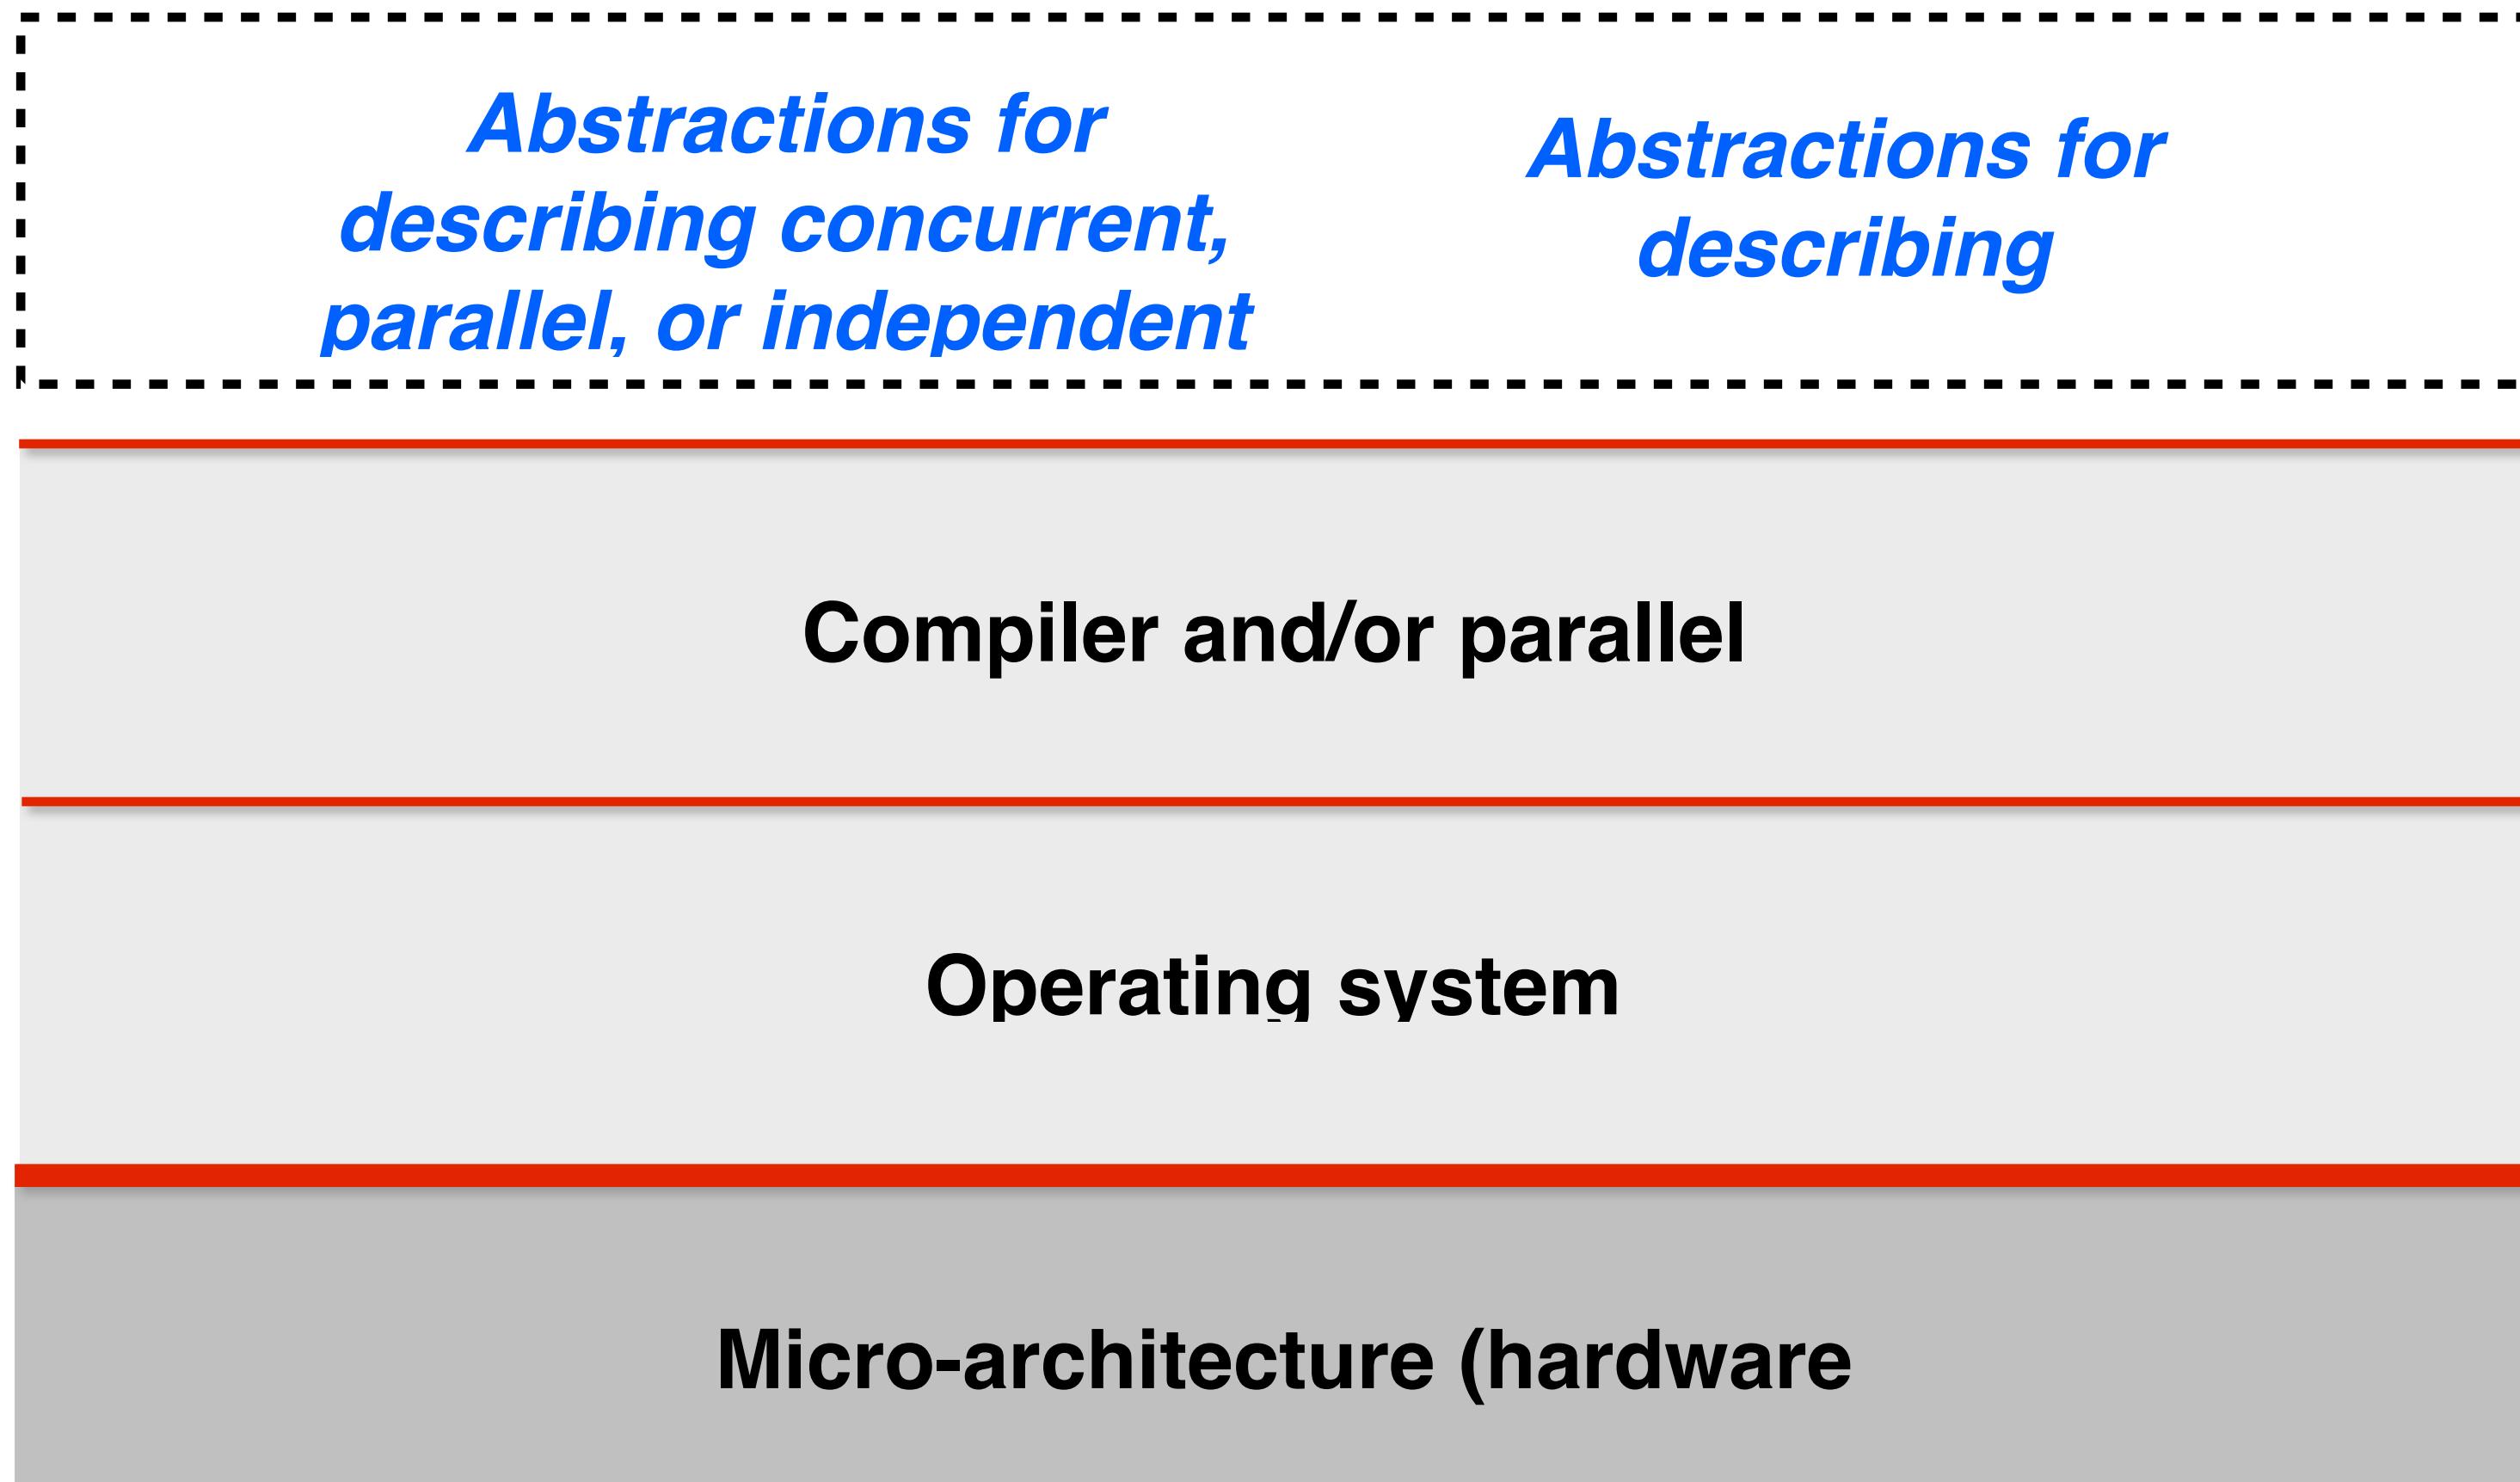
- The ISPC **gang abstraction** is implemented by SIMD instructions on one core.
- So... all the code I've shown you in the previous slides would have executed **on only one of the four cores** of the GHC machines.
- ISPC contains **another abstraction**: a “**task**” that is used to **achieve multi-core execution**. I'll let you read up about that.

The second half of today's

- **Four parallel programming models**
 - That differ in communication abstractions presented to the programmer
 - Programming models influence how programmers think when writing programs
- **Four machine architectures**
 - Abstraction presented by the hardware to low-level software
 - Typically reflect hardware implementation's capabilities
- **We'll focus on differences in communication and cooperation**

System layers: interface,

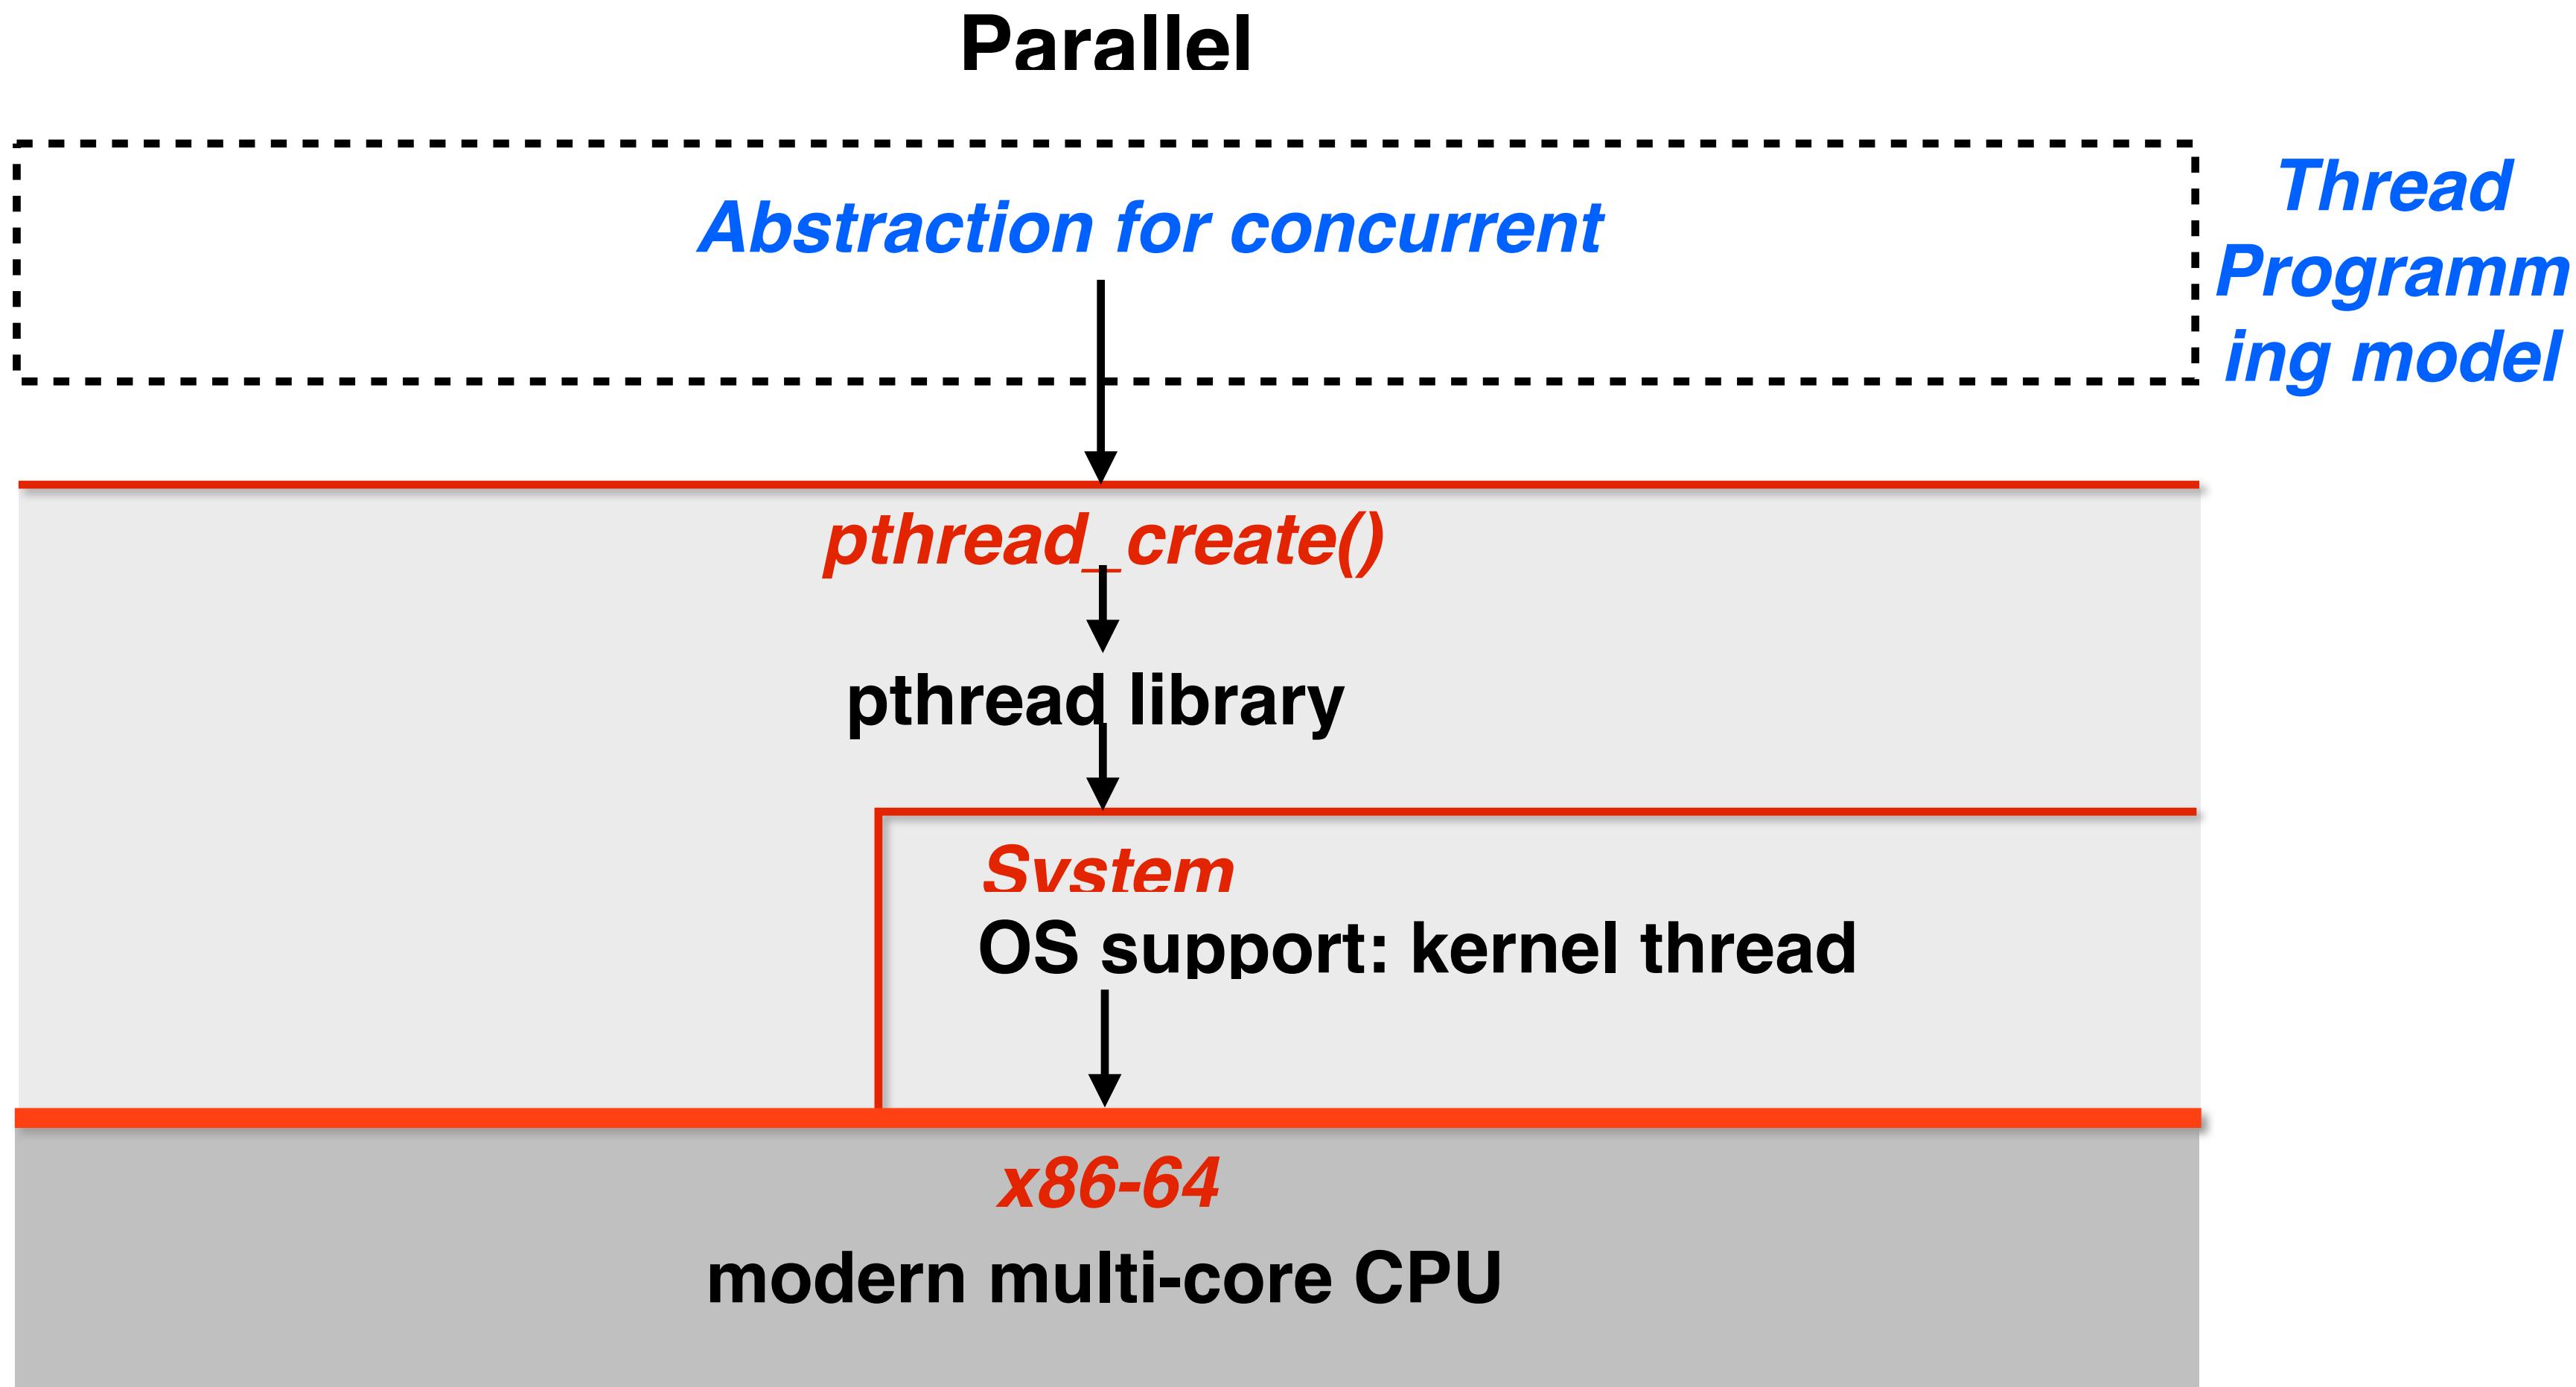
Parallel



*Blue italic text: abstraction/
concept*

Black text: system

Example: expressing parallelism

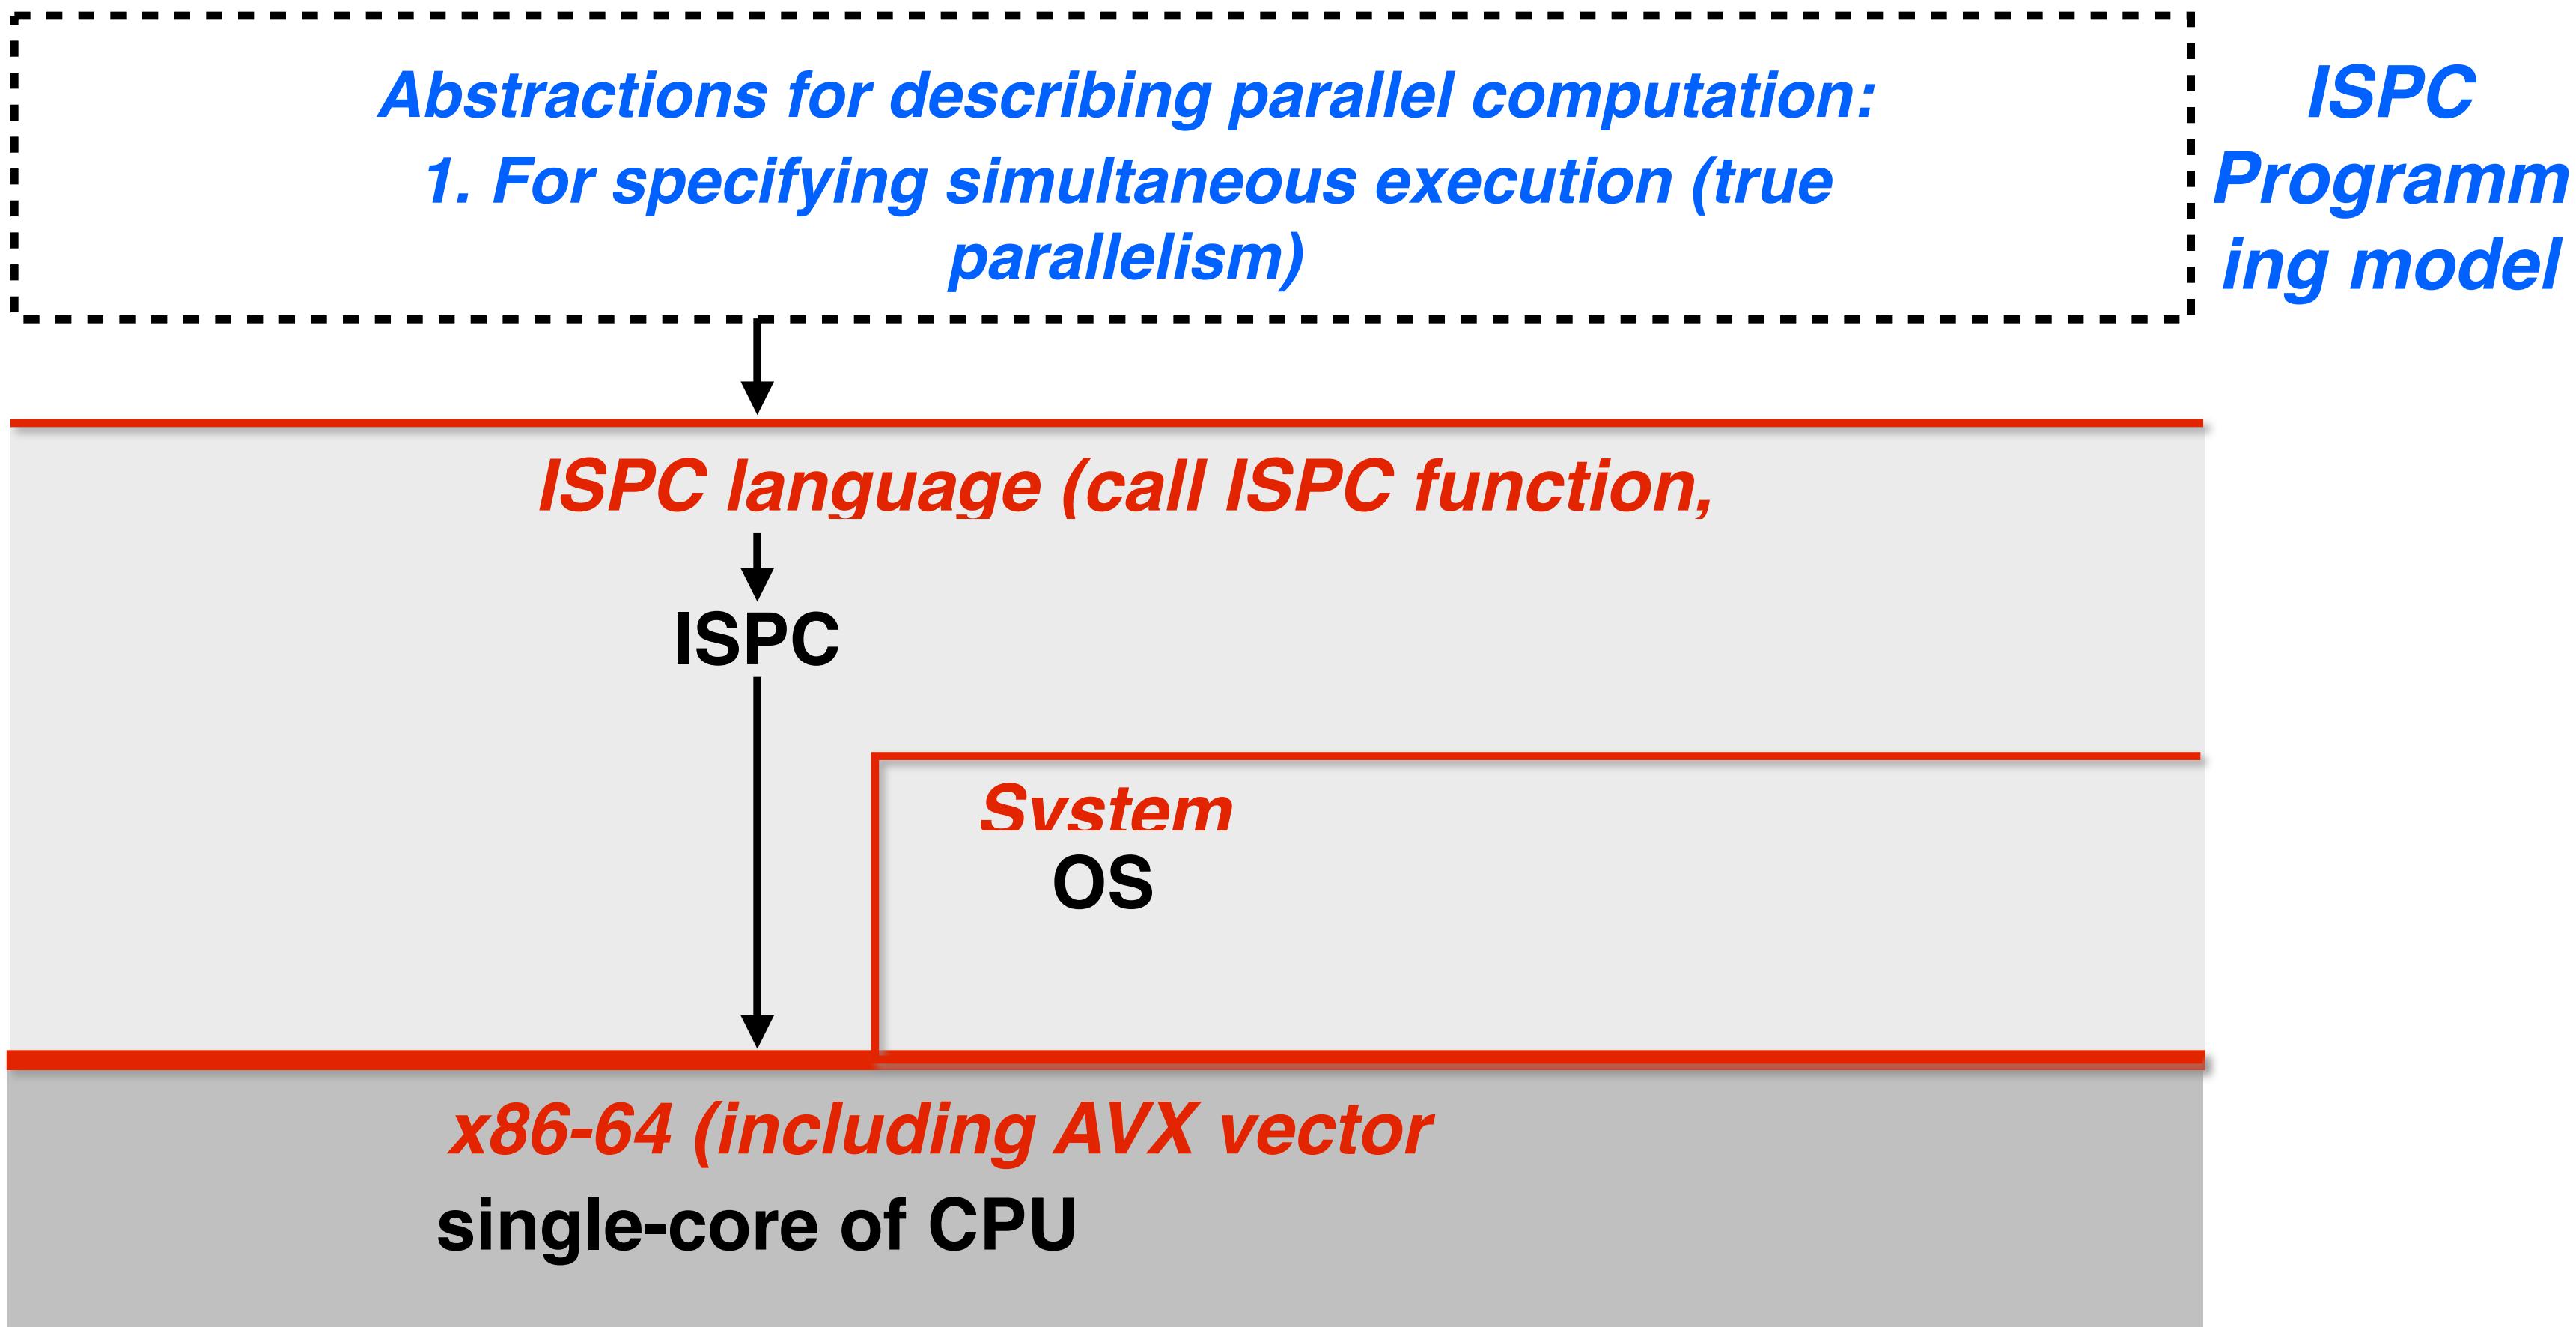


Blue italic text: abstraction/concept

Black text: system

Example: expressing parallelism

Parallel



Note: This diagram is specific to the ISPC gang abstraction. ISPC also has the “task” language primitive for multi-core execution.

Four models of communication (abstractions)

1. Shared address space
2. Message passing
3. Data parallel
4. Systolic arrays

Shared address space model of

Shared address space model

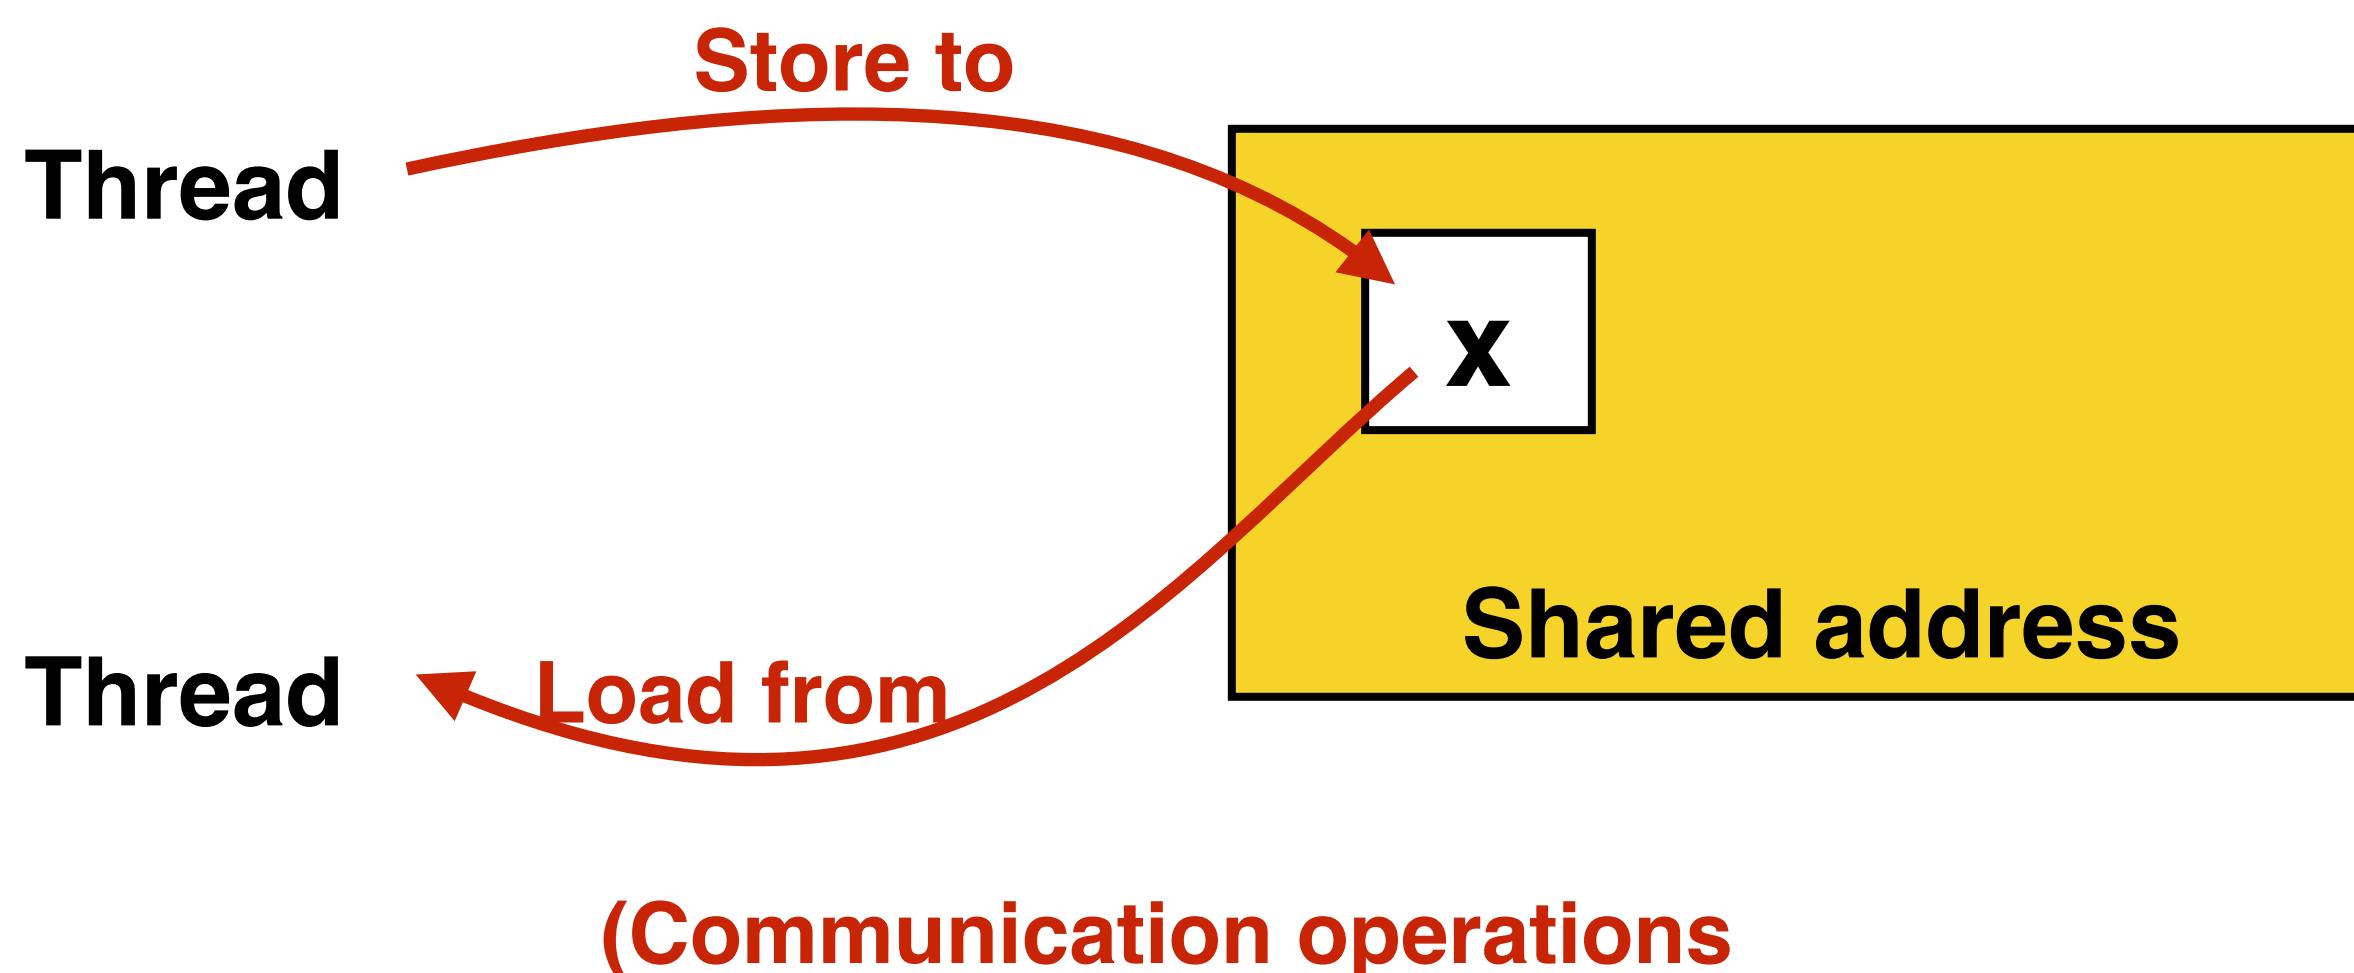
- Threads communicate by reading/writing to shared variables
- Shared variables are like a big bulletin board
 - Any thread can read or write to shared variables

Thread 1:

```
int x = 0;  
spawn_thread(foo, &x);  
x = 1;
```

Thread 2:

```
void foo(int* x) {  
    while (x == 0) {}  
    print x;  
}
```



Shared address space model

Synchronization primitives are **also shared variables**:
e.g., **locks**

Thread 1:

```
int x = 0;  
Lock my_lock;  
  
spawn_thread(foo, &x, &my_lock);
```

```
mylock.lock();  
x++;  
mylock.unlock();
```

Thread 2:

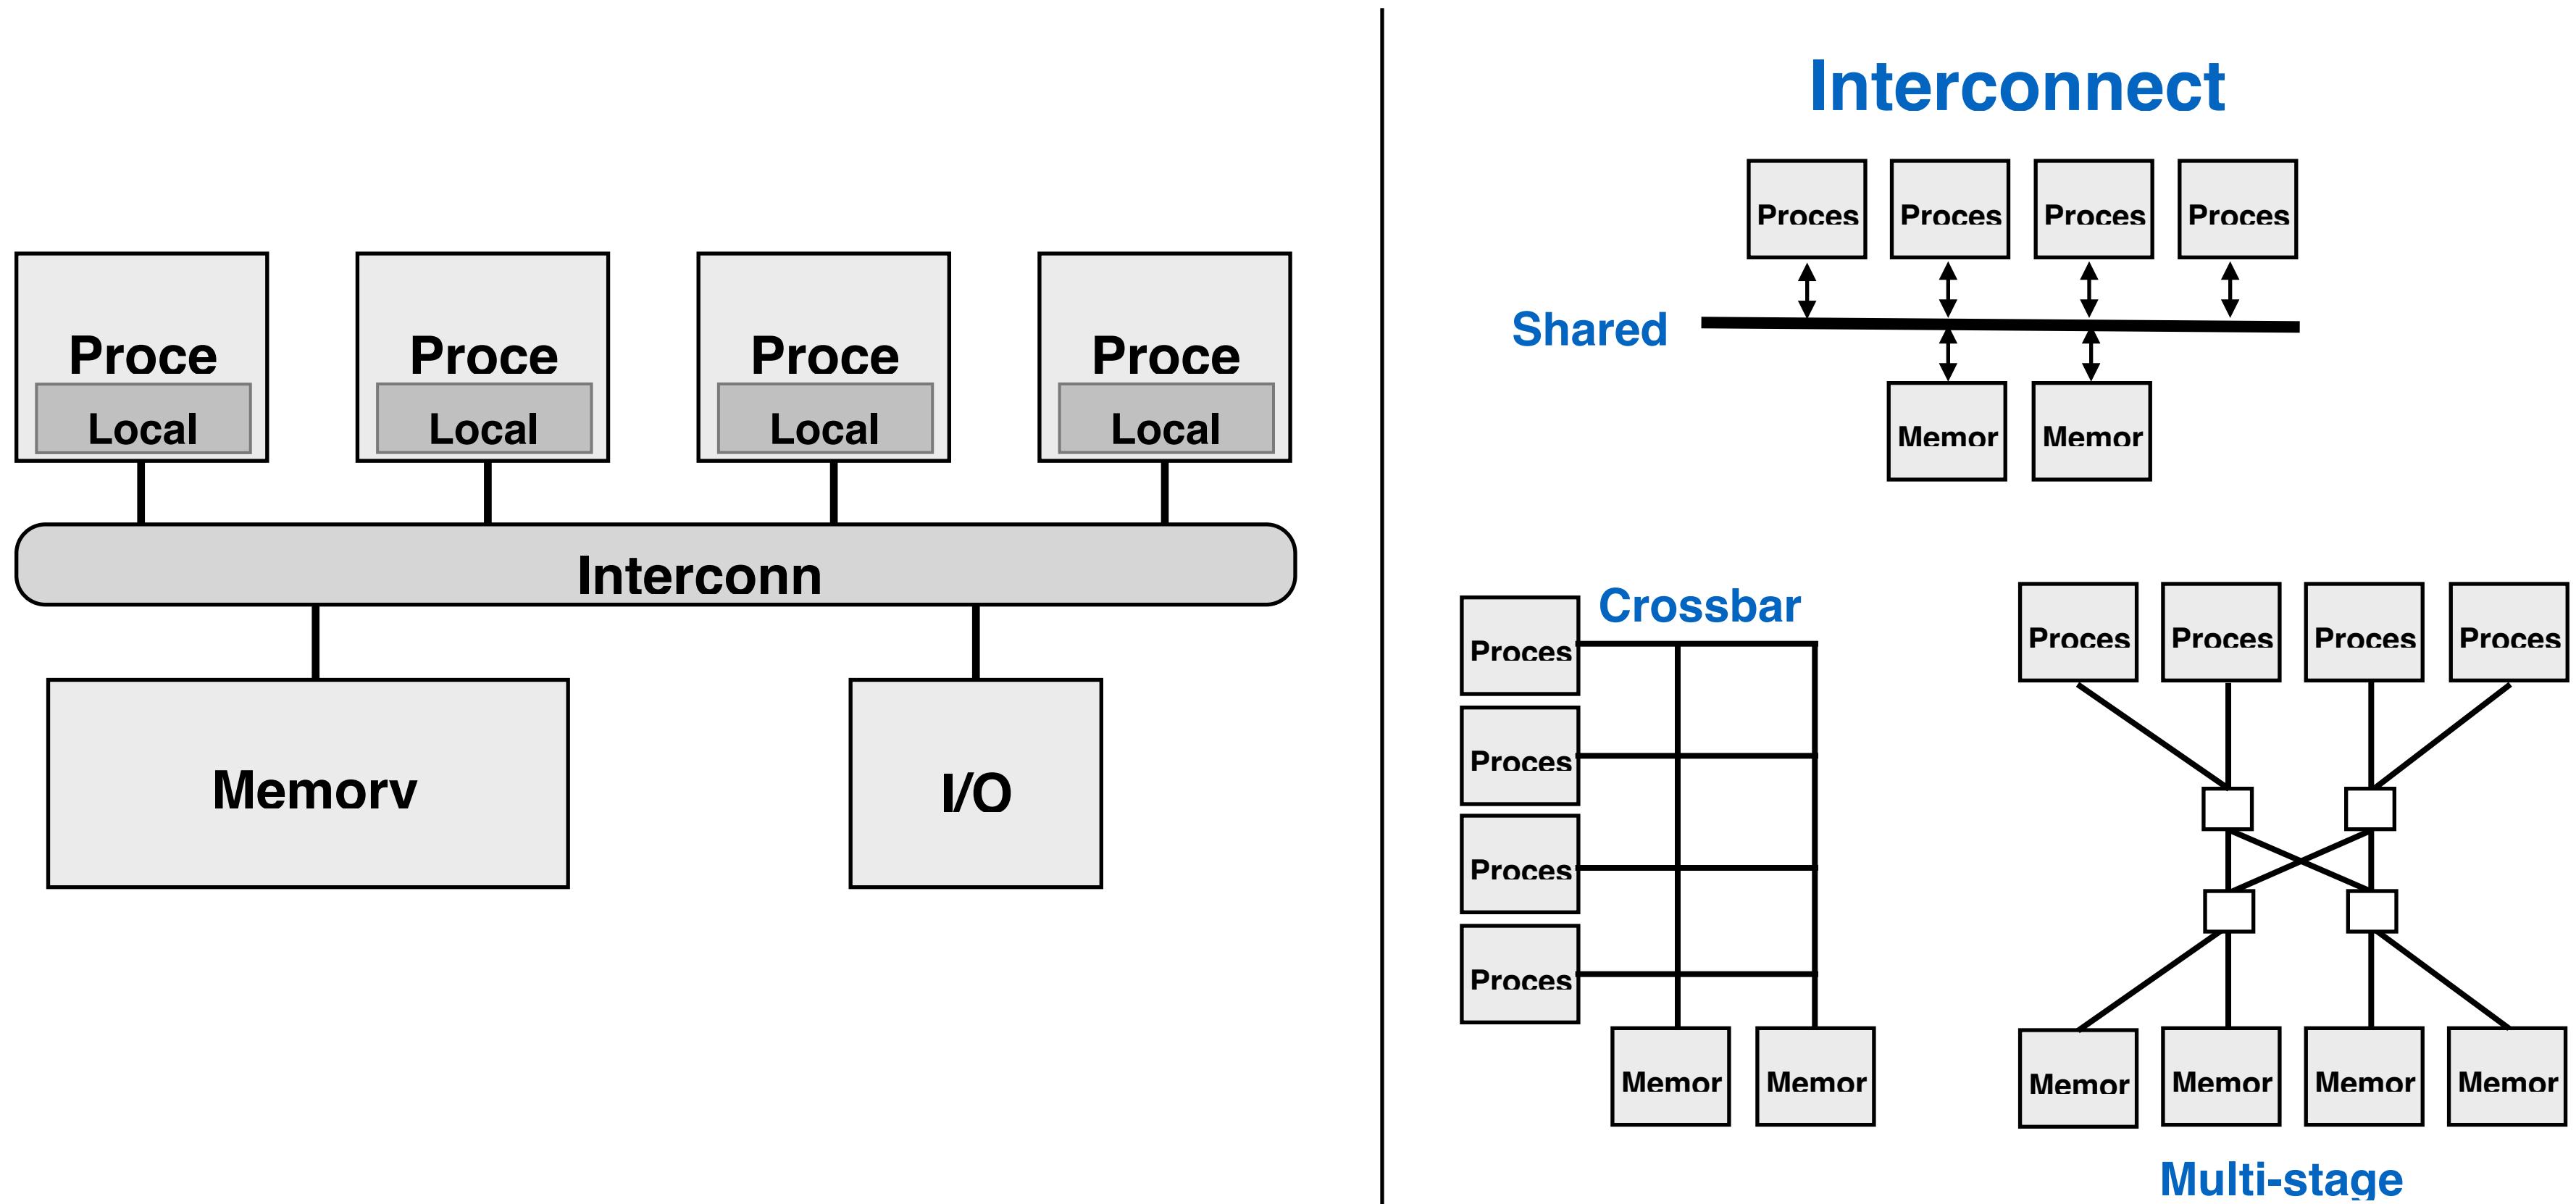
```
void foo(int* x, lock* my_lock)  
{  
    my_lock->lock();  
    x++;  
    my_lock->unlock();  
  
    print x;  
}
```

Shared address space model

- Threads **communicate** by:
 - **Reading/writing to shared variables**
 - Inter-thread communication is implicit in memory operations
 - Thread 1 stores to X
 - Later, thread 2 reads X (and observes update of value by thread 1)
 - **Manipulating synchronization primitives**
 - e.g., ensuring mutual exclusion via use of locks
- This is a natural extension of sequential programming
 - In fact, all our discussions in class have assumed a shared address space so far!

HW implementation of a shared

Key idea: any processor can directly reference any memory



Symmetric (shared-memory) multi-processor (SMP):

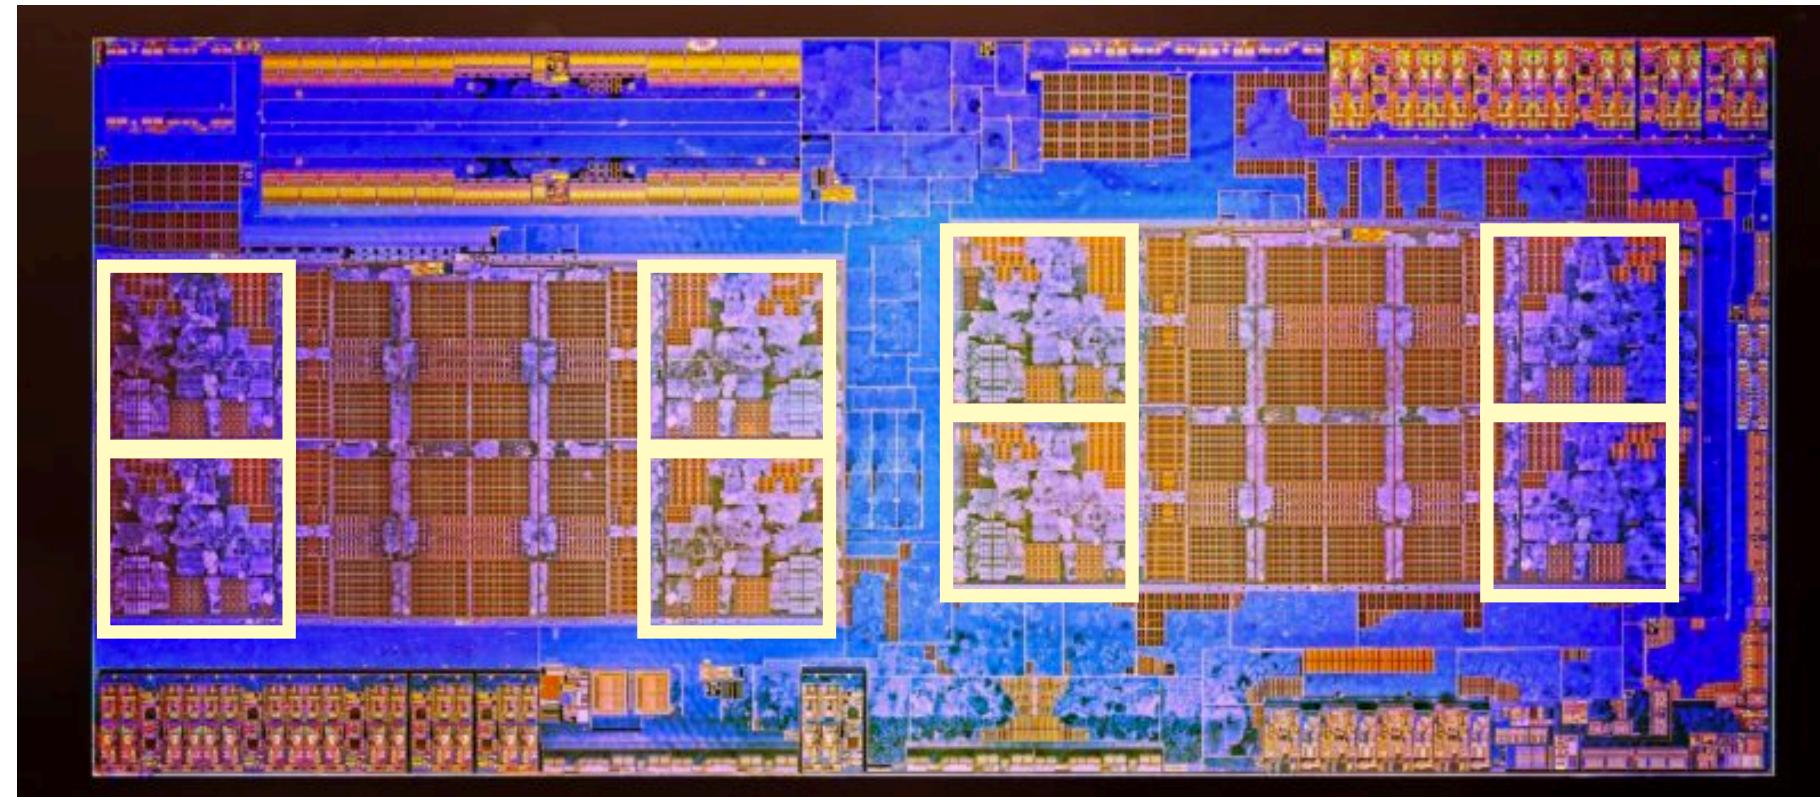
* caching introduces non-uniform access times, but we'll talk about that later

Shared address space HW

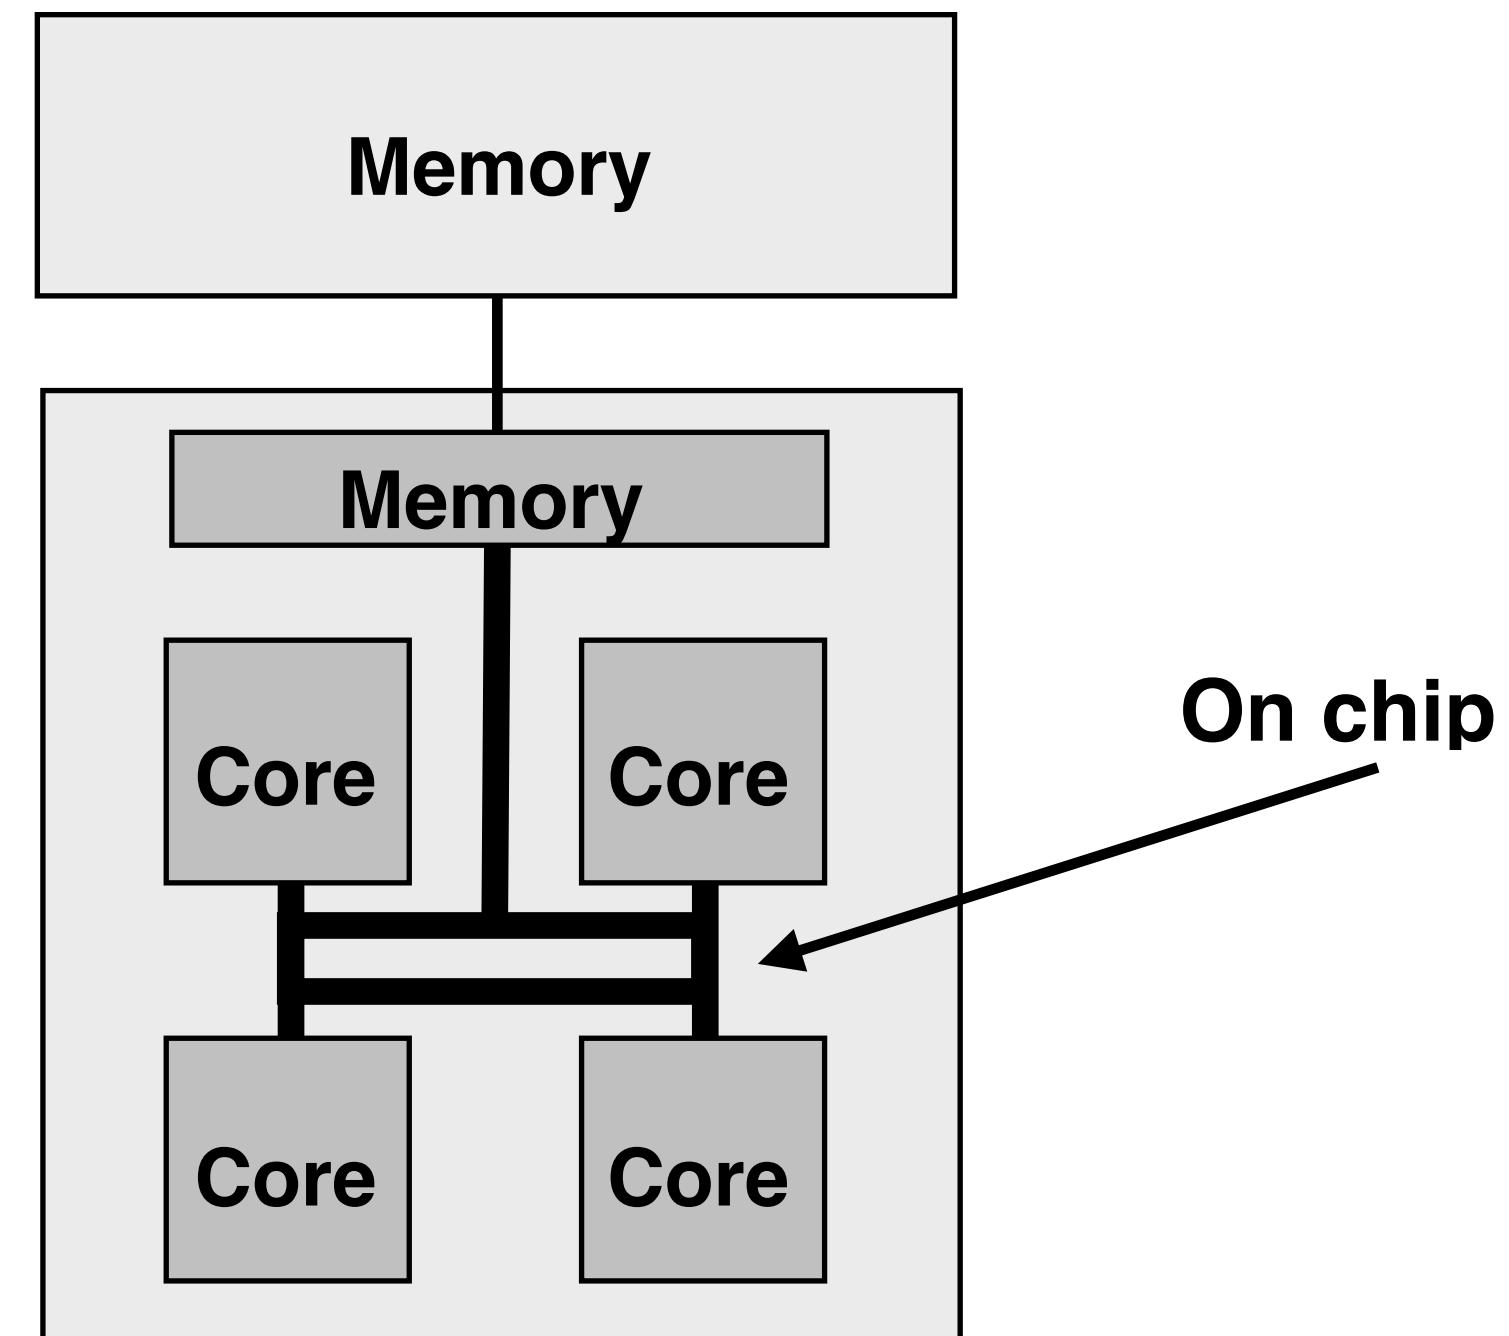
Commodity x86 examples



Intel Alder Lake-S, 16 cores (8+8) (interconnect is a ring)

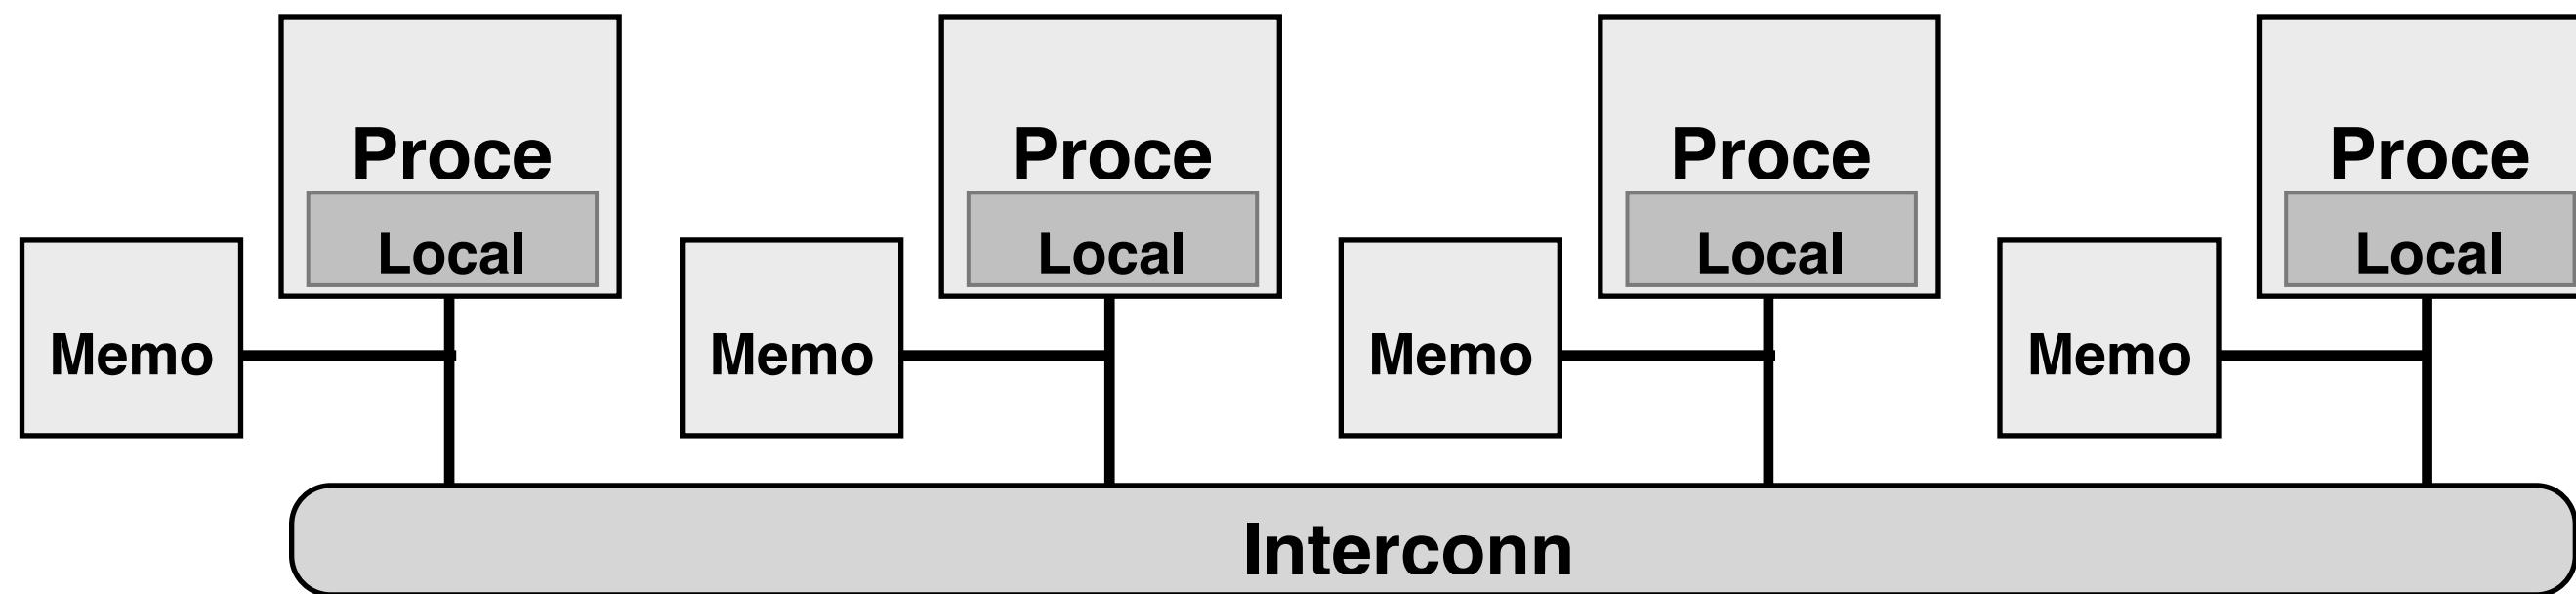


AMD Ryzen (8



Non-uniform memory access

All processors can access any memory location, but... the cost of memory access (latency and/or bandwidth) is

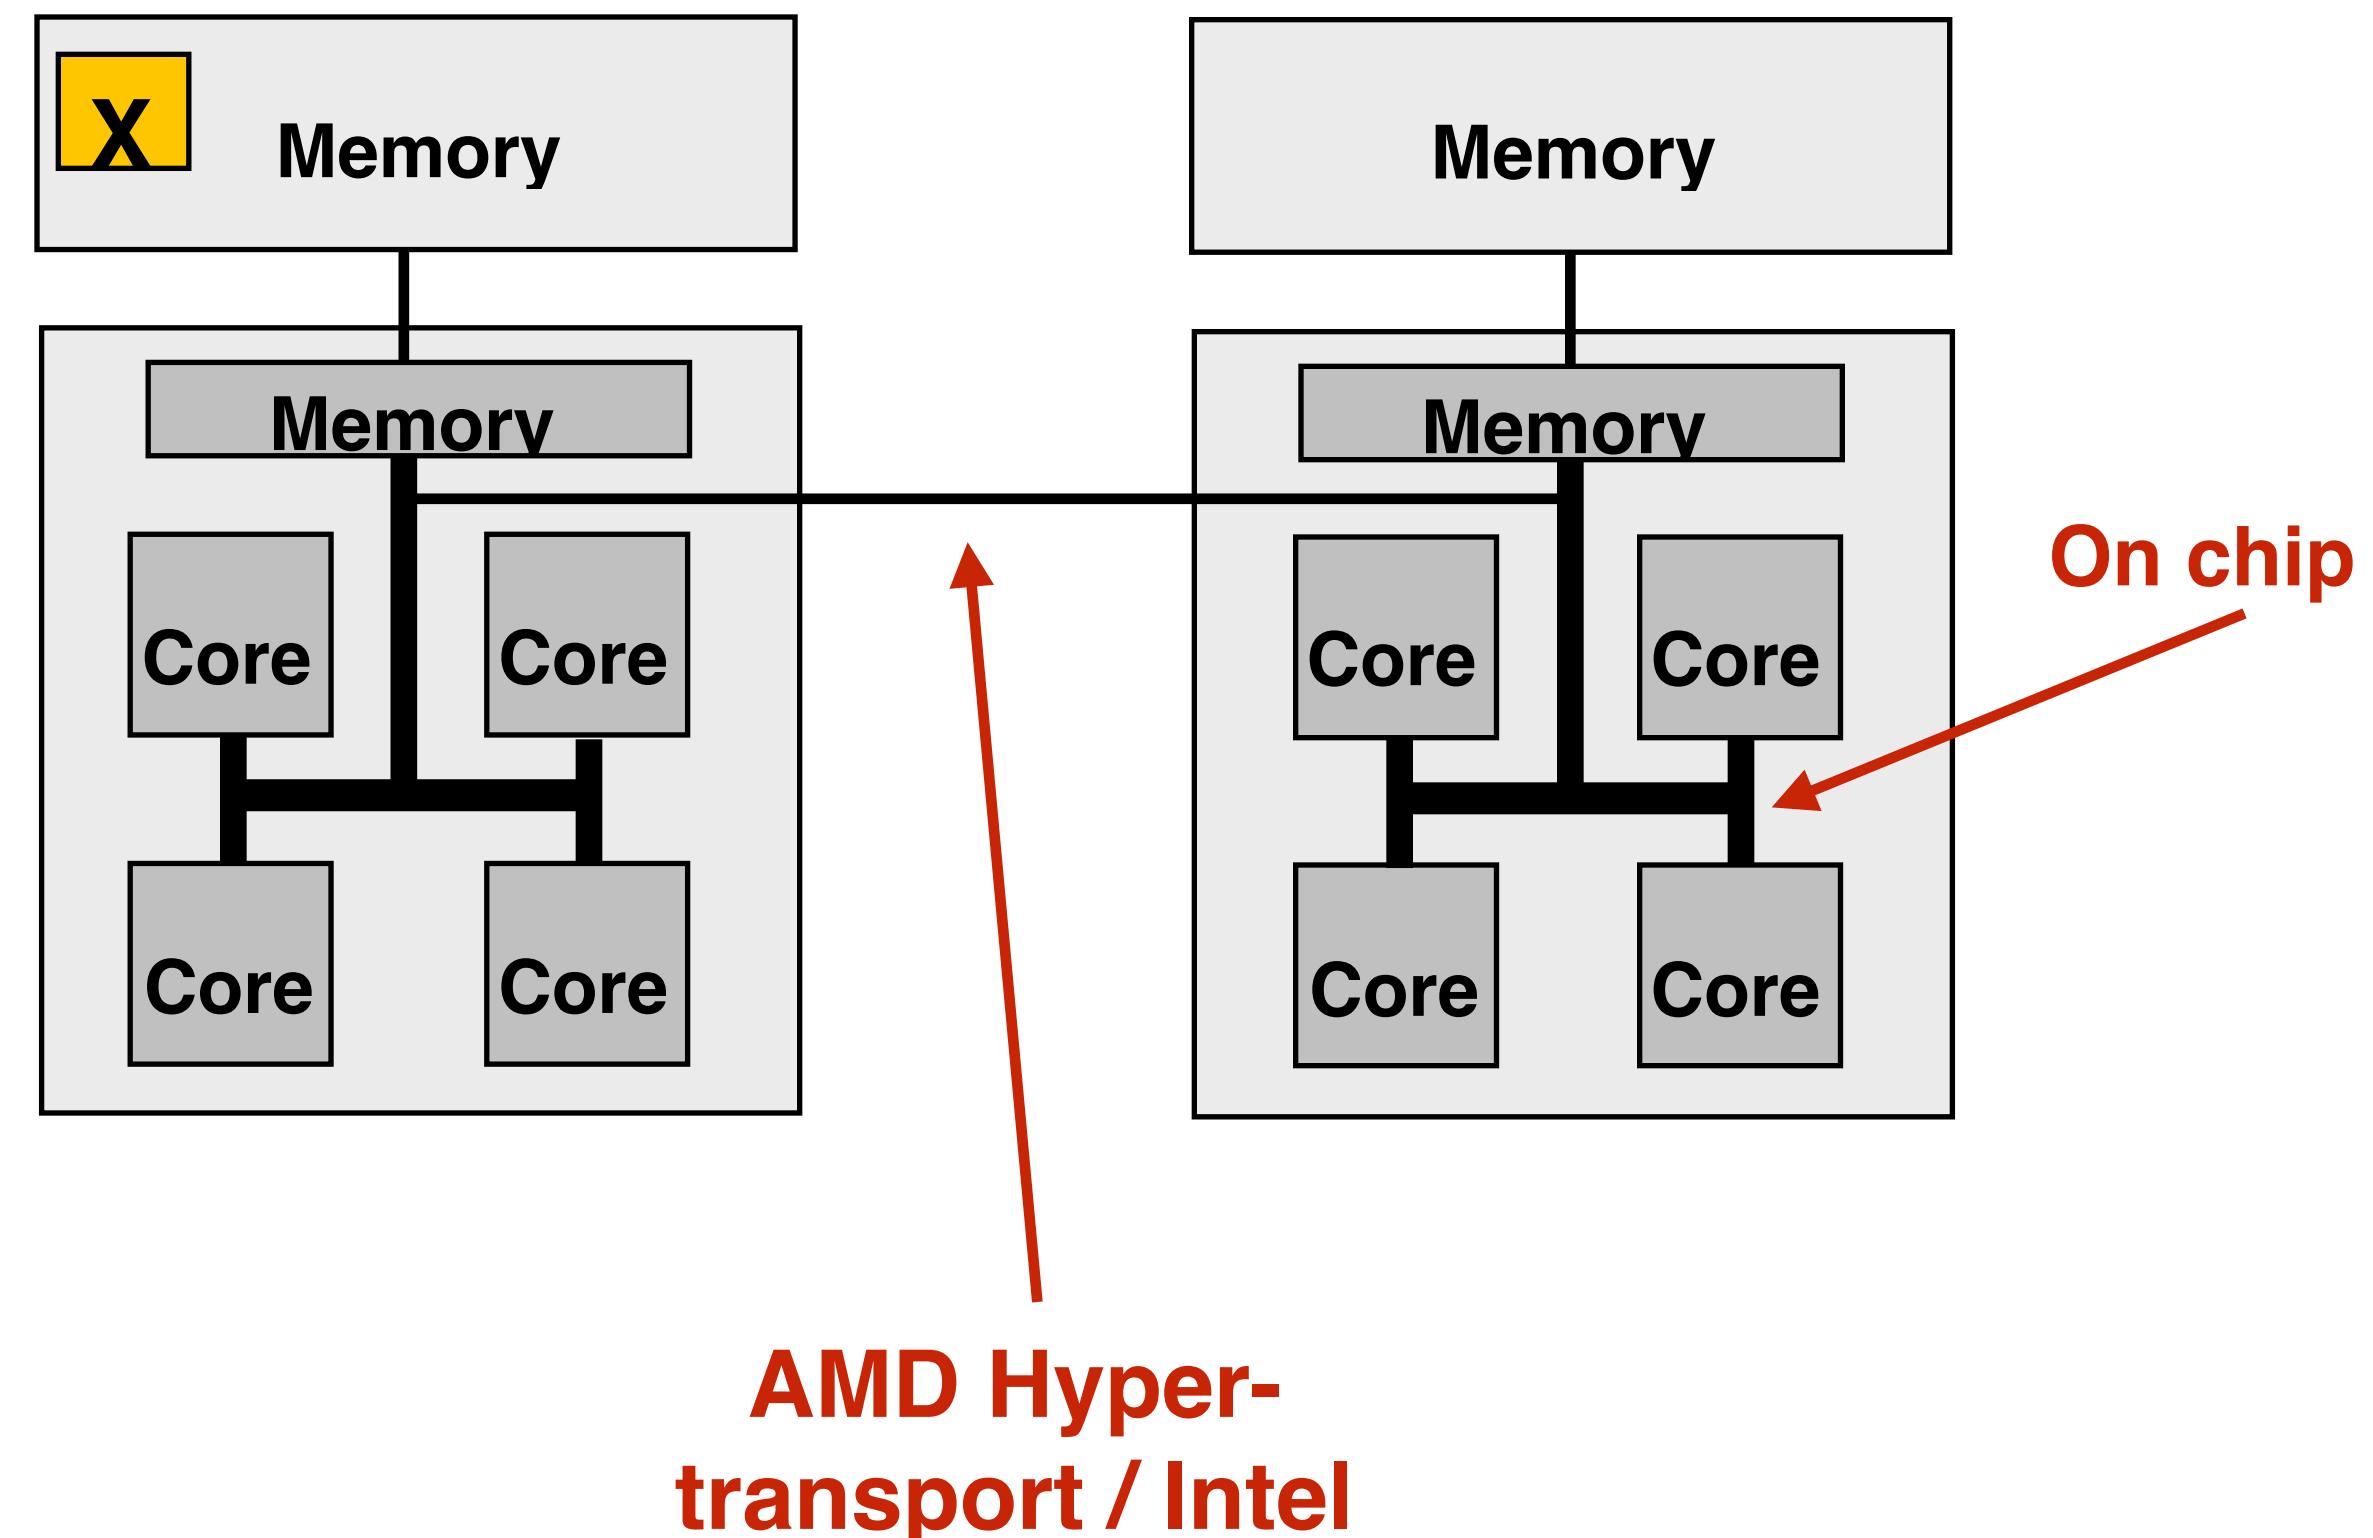
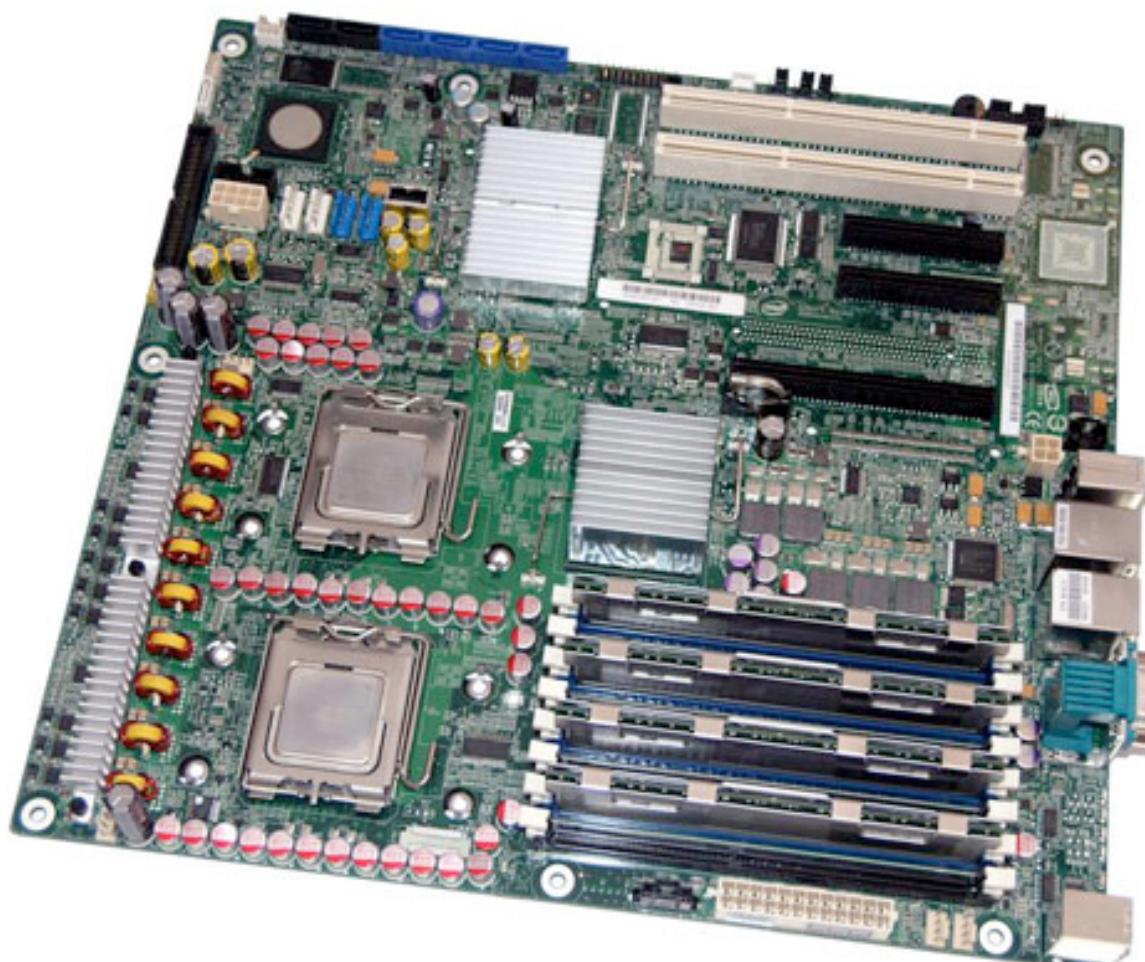


- Problem with preserving uniform access time in a system: **scalability**
 - **GOOD:** costs are uniform, **BAD:** they are **uniformly bad** (memory is uniformly far away)
- **NUMA designs are more scalable**
 - **Low latency and high bandwidth to local memory**
- Cost is increased programmer effort for performance tuning
 - **Finding, exploiting locality** is important to performance

Non-uniform memory access

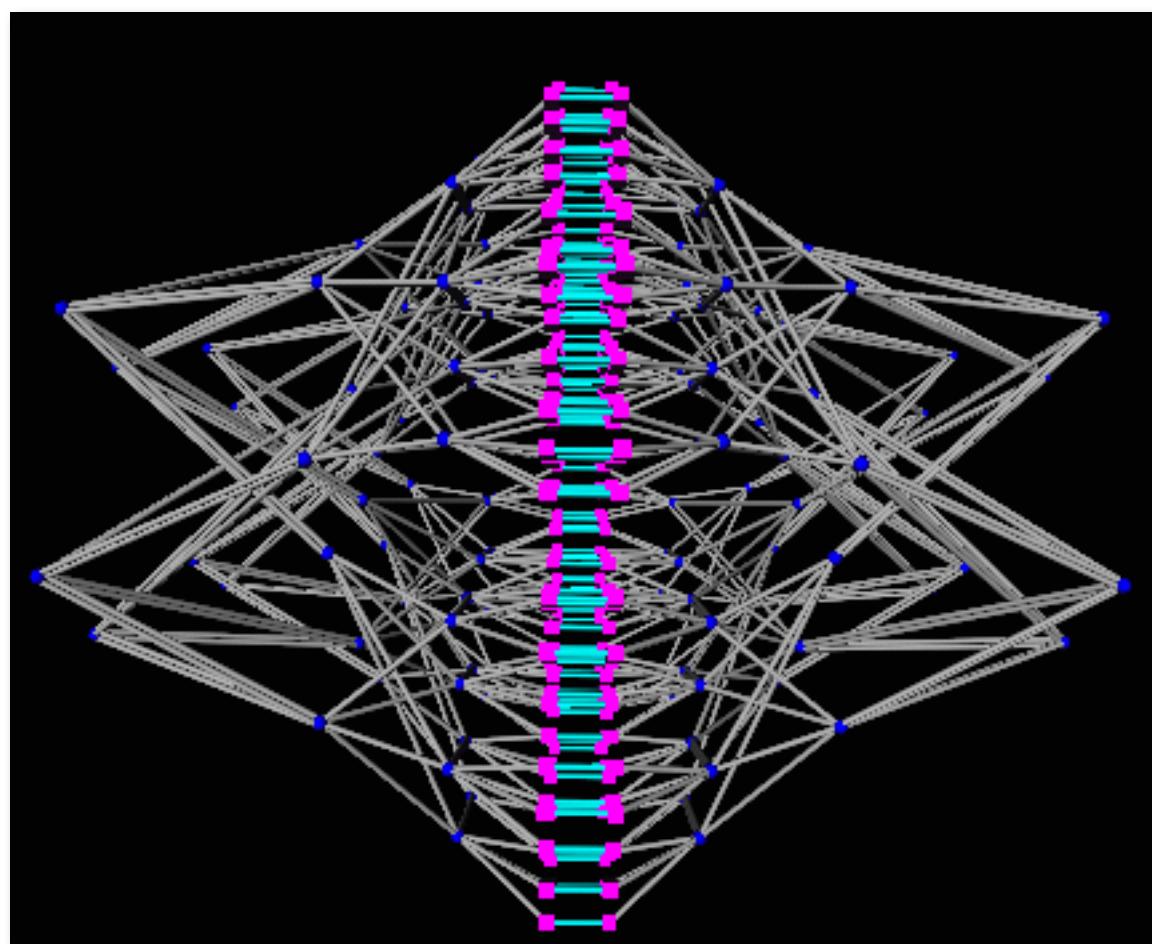
Example: latency to access address x is higher from

Example: modern dual-socket



SGI Altix UV 1000

- 256 blades, 2 CPUs per blade, 8 cores per CPU = **4096 cores**
- **Single shared address space**
- Interconnect: fat tree



Fat tree topology



Image credit: Pittsburgh Supercomputing Center

Summary: shared address

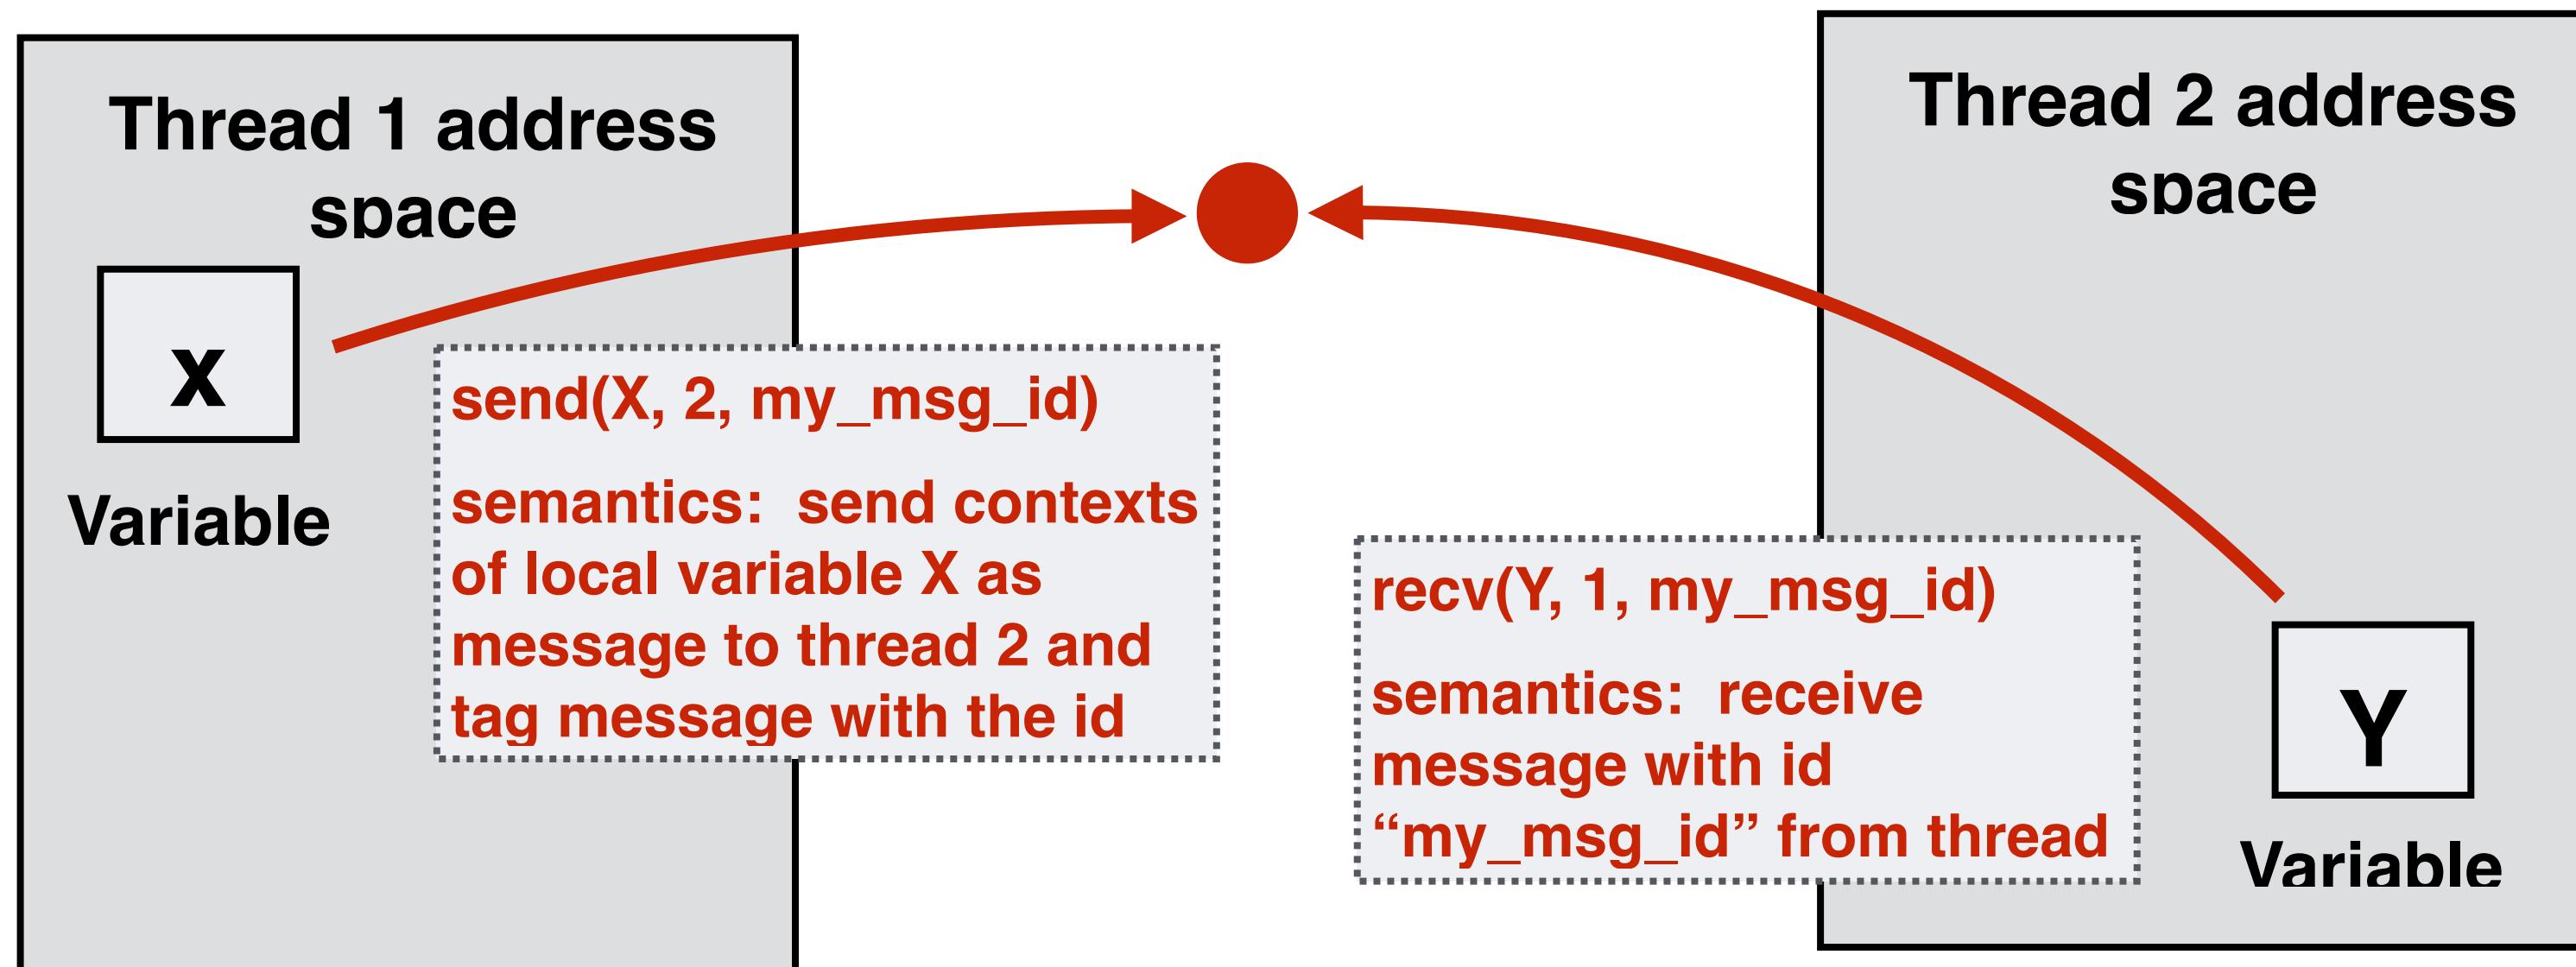
- **Communication abstraction**
 - Threads read/write **shared variables**
 - Threads manipulate **synchronization** primitives: locks, semaphors, etc.
 - Logical extension of uniprocessor programming *
- **Requires hardware support to implementation efficiently**
 - Any processor can load and store from any address (its shared address space!)

* But NUMA implementation requires reasoning about locality for performance

Message passing model of

Message passing model

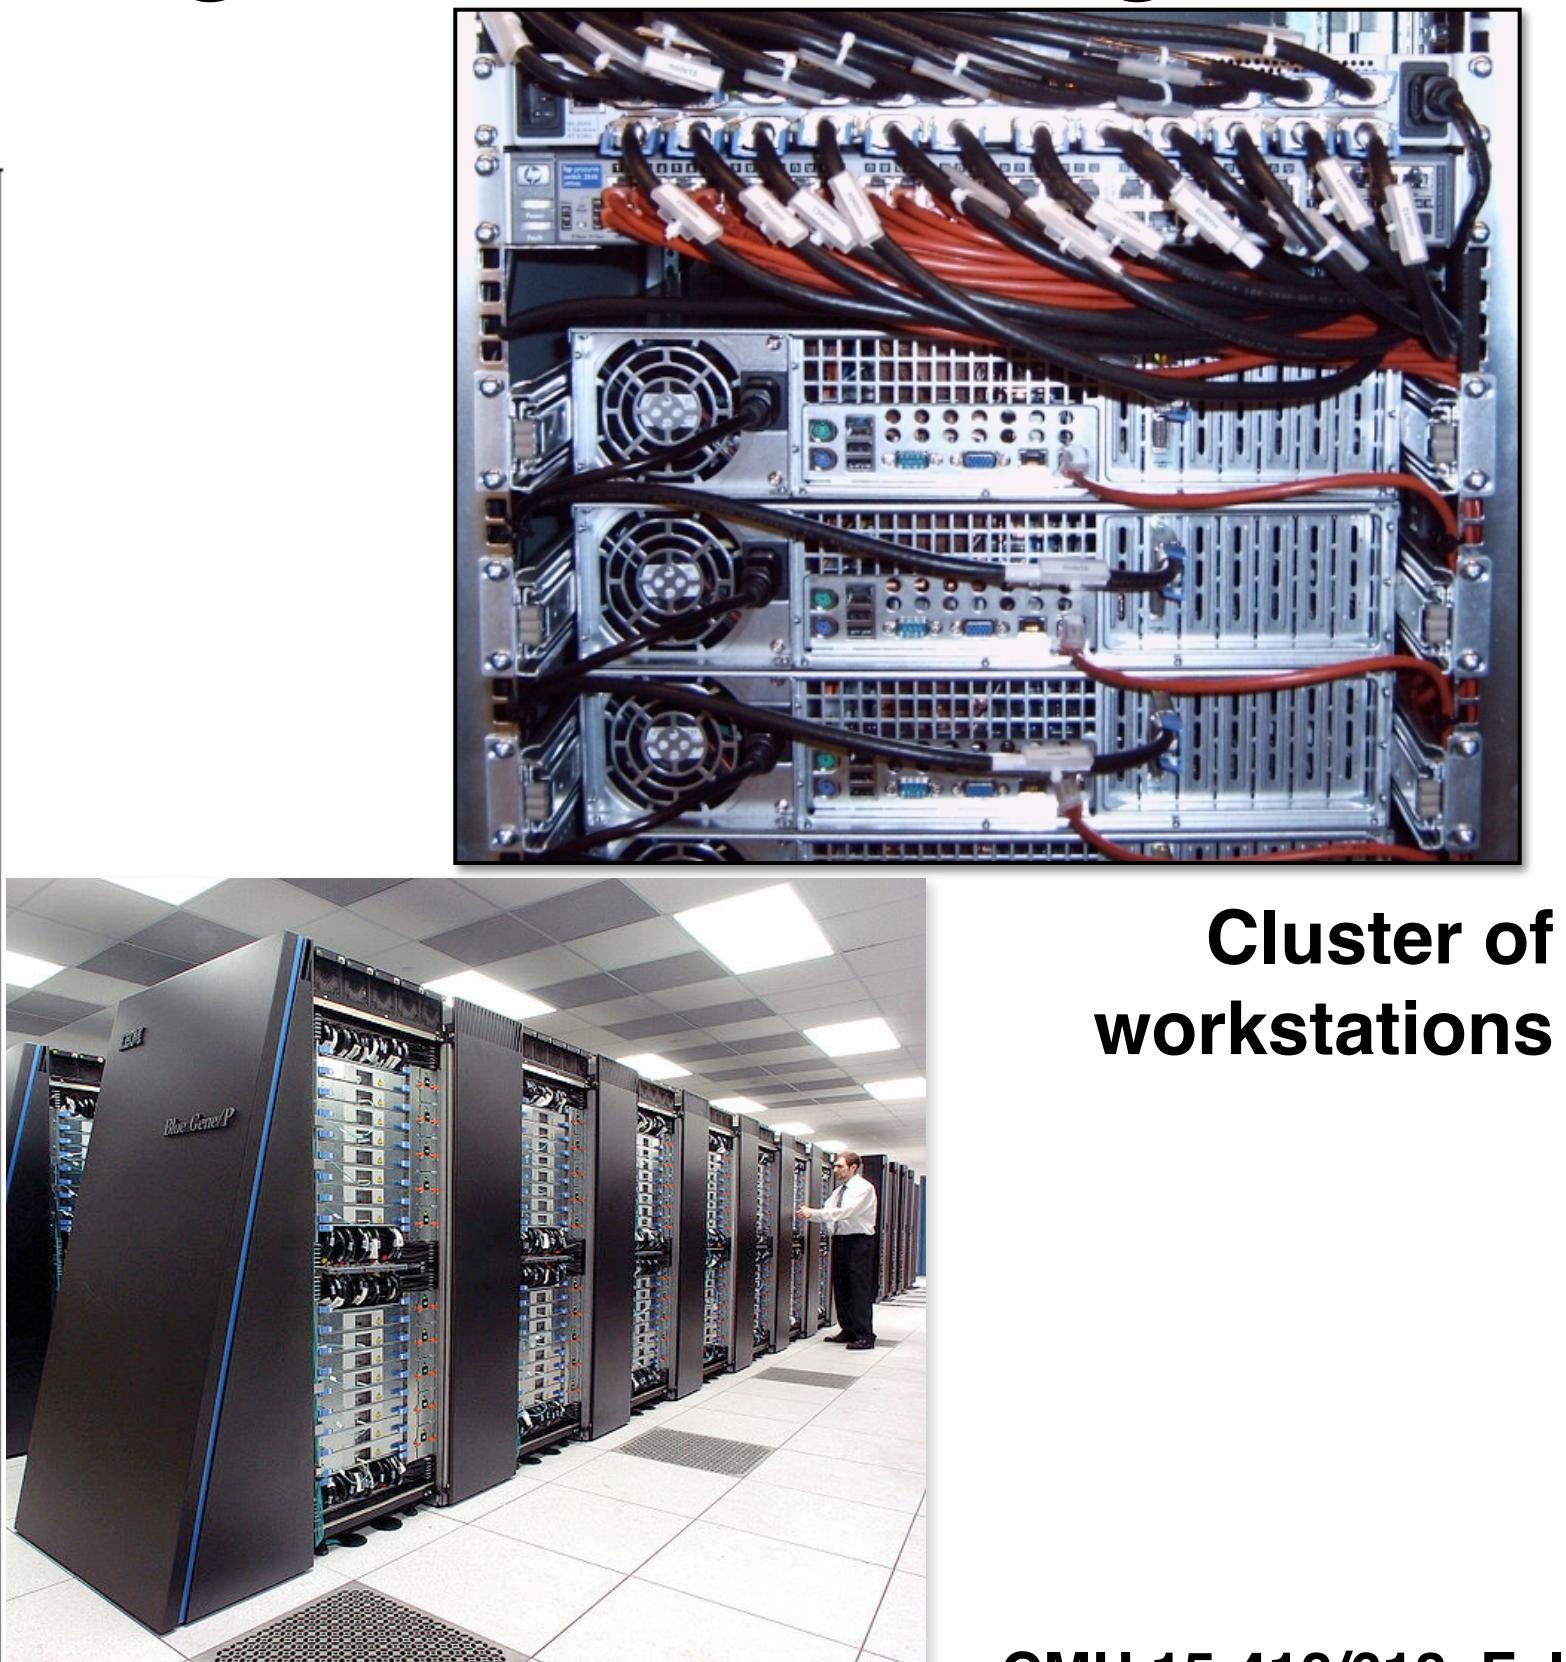
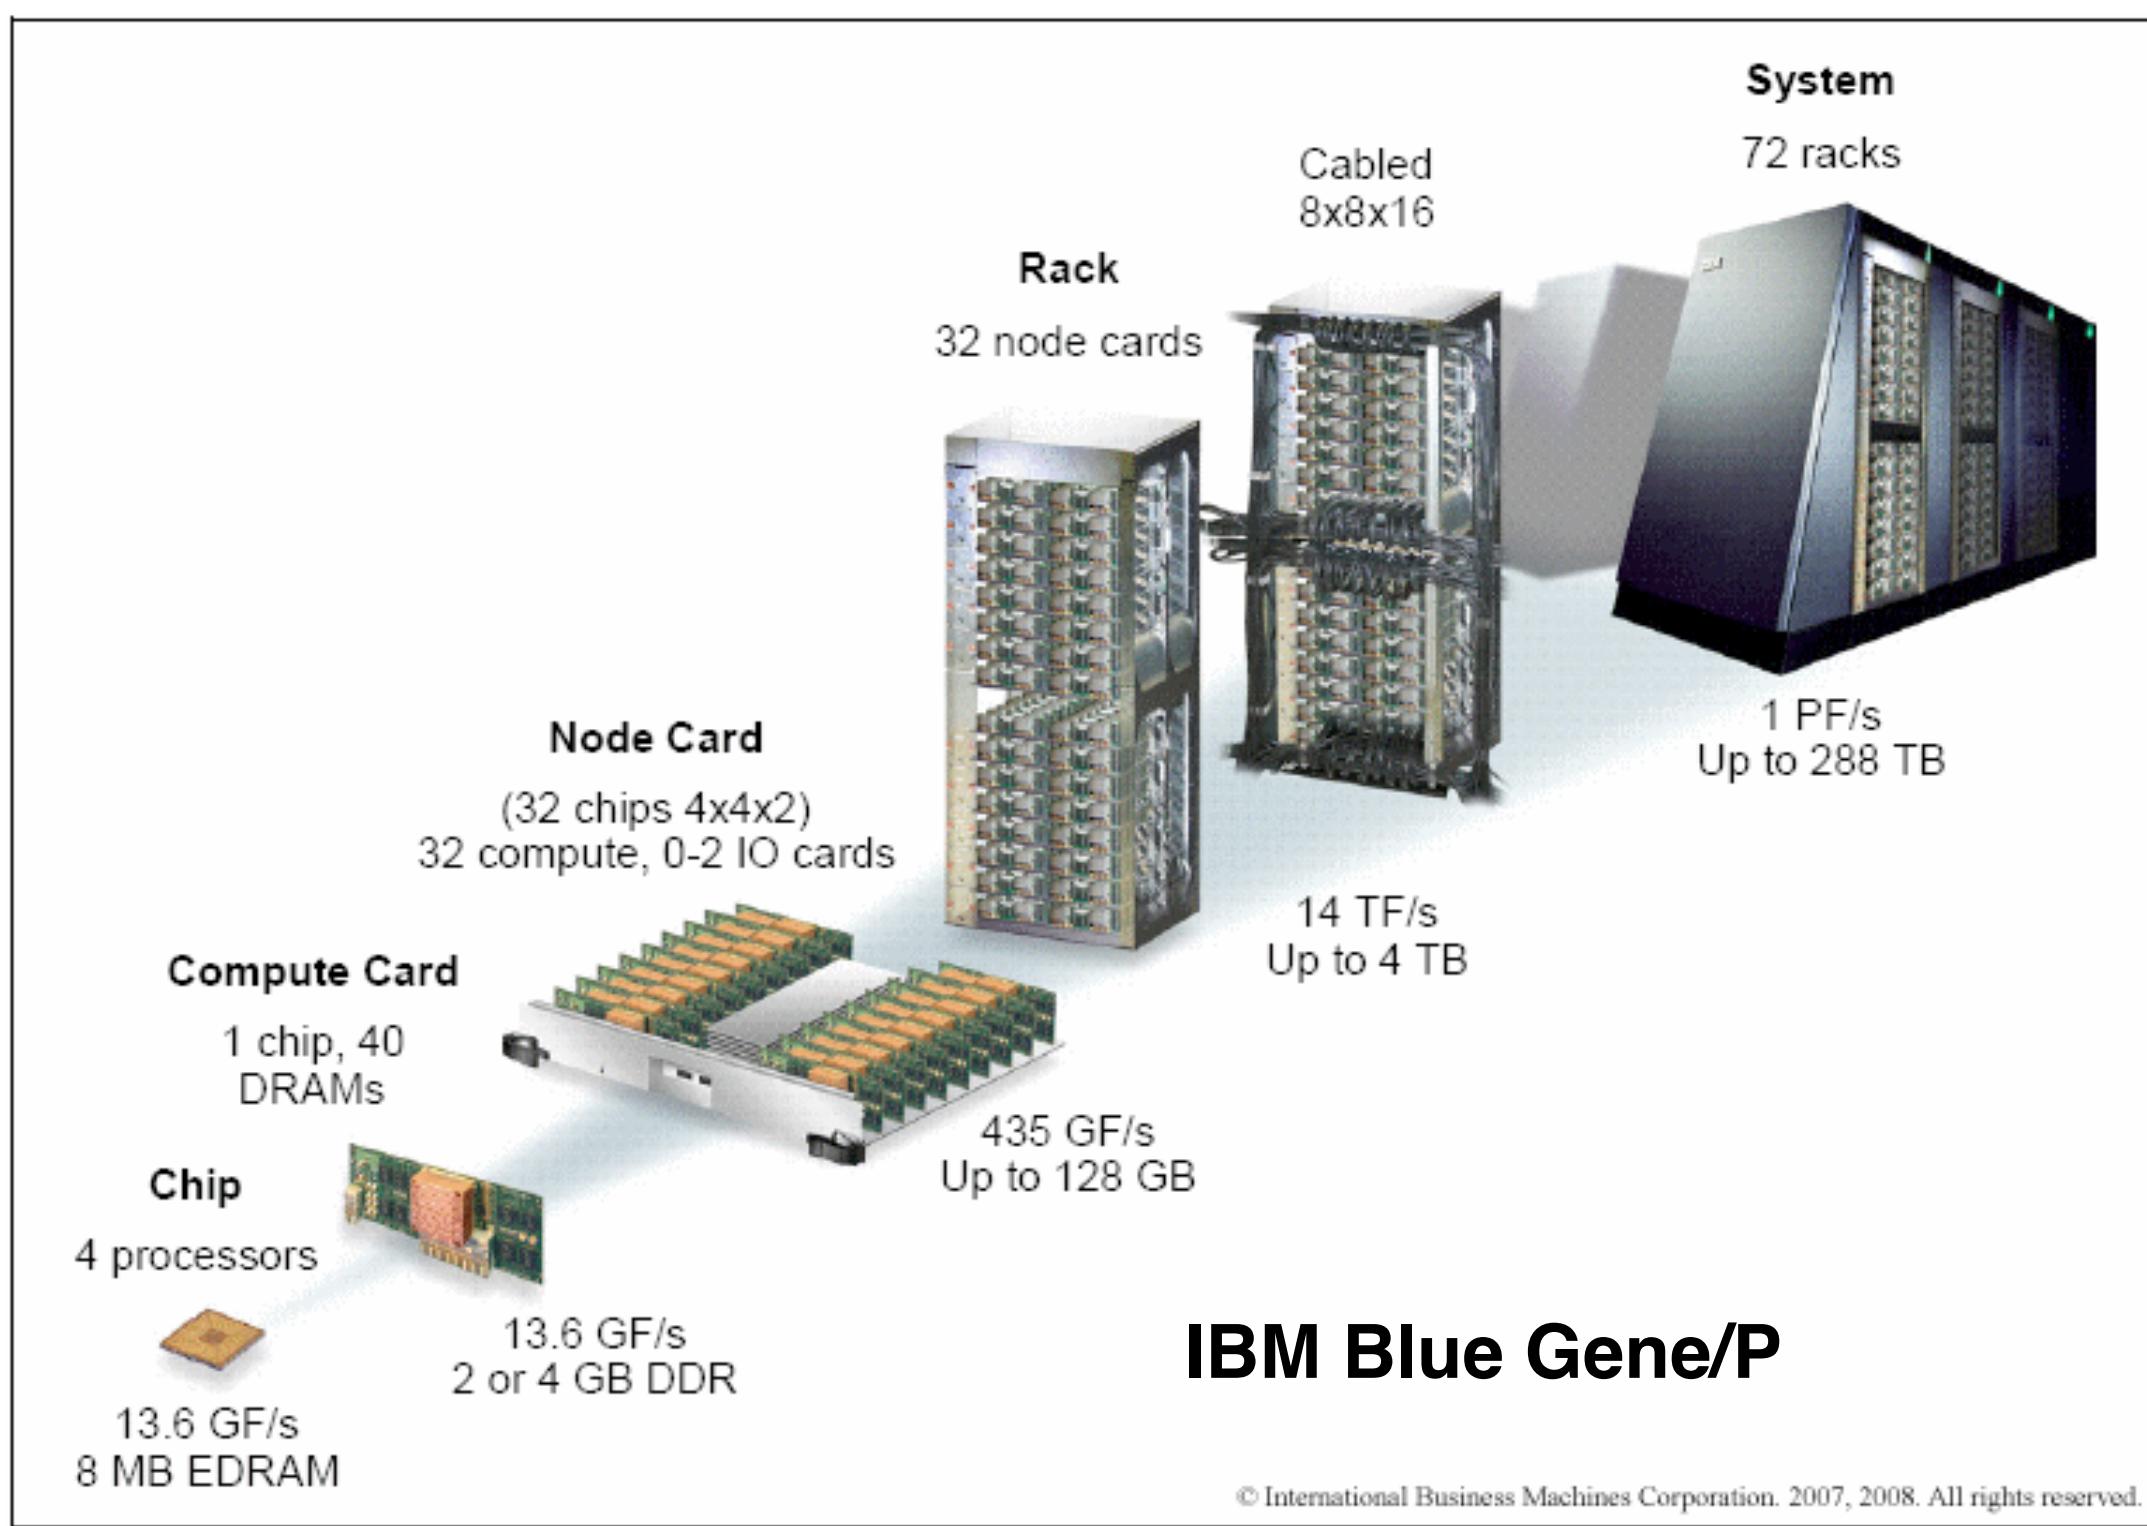
- Threads operate within their own **private address spaces**
- Threads **communicate** by **sending/receiving messages**
 - send: specifies recipient, buffer to be transmitted, and optional message identifier (“tag”)
 - receive: sender, specifies buffer to store data, and optional



(Communication operations shown

Messaeae bassina

- Popular software library: **MPI** (message passing interface)
- Hardware need not implement system-wide loads and stores to execute message passing programs (need only be able to communicate messages)
 - Can connect commodity systems together to form large



15-418/618 “latedavs” cluster

- 15 node cluster *
 - 1 head node, 14 worker nodes accessible via a batch job queue
- Each node contains:
 - Two, six-core Xeon e5-2620 v3 processors (844 GFLOPS peak)
 - 2.4 GHz, 15MB L3 cache, hyper-threading, AVX2 (“8-wide”) instruction support
 - 16 GB RAM (60 GB/sec bandwidth)
 - NVIDIA K40 GPU (4.2 TFLOPS peak)
 - One Xeon Phi 5110p co-processor board (2 TFLOPS peak)
 - 60 “simple” 1 GHz x86 cores
 - 4-threads per core, AVX512 (“16-wide”) instruction support
 - 8 GB RAM (320 GB/sec bandwidth)

The correspondence between programming models and

- Common to **implement message passing abstractions** on machines that **implement a shared address space in hardware**
 - “**Sending message**” = **copying memory from message library buffers**
 - “**Receiving message**” = **copy data from message library buffers**
- Can **implement shared address space abstraction** on machines that **do not support it in HW** (via less efficient SW solution)
 - Mark all pages with shared variables as invalid

The data-parallel model

Recall: programming models impose structure on programs

- **Shared address space: very little structure**
 - All threads can read and write to all shared variables
 - Pitfall: due to implementation: not all reads and writes have the same cost (and that cost is not apparent in program text)
- **Message passing: highly structured communication**
 - All communication occurs in the form of messages (can read program and see where the communication is)
- **Data-parallel: very rigid computation**

Data-parallel model

- Historically: same operation on each element of an array
 - Matched capabilities SIMD supercomputers of 80's
 - Connection Machine (CM-1, CM-2): thousands of processors, one instruction decode unit
 - Cray supercomputers: vector processors
 - $\text{add}(A, B, n) \leftarrow$ this was one instruction on vectors A, B of length n
- Matlab is another good example: $C = A + B$ (A, B, and C are vectors of same length)
- Today: often takes form of SPMD programming
 - **map(function, collection)**
 - Where function is applied to each element of collection independently

Gather/scatter: two key data-parallel communication primitives

Map absolute_value onto stream produced by

```
const int N = 1024;
stream<float> input(N);
stream<int> indices;
stream<float> tmp_input(N);
stream<float> output(N);
```

```
stream_gather(input, indices, tmp_input);
absolute_value(tmp_input, output);
```

Map absolute_value onto stream,
gather:
scatter results:

```
const int N = 1024;
stream<float> input(N);
stream<int> indices;
stream<float> tmp_output(N);
stream<float> output(N);
```

```
absolute_value(input, tmp_output);
stream_scatter(tmp_output, indices, output);
```

ISPC equivalent:

```
export void absolute_value(
    uniform float N,
    uniform float* input,
    uniform float* output,
    uniform int* indices)
{
    foreach (i = 0 ... n)
    {
        float tmp = input[indices[i]];
        if (tmp < 0)
            output[i] = -tmp;
        else
            output[i] = tmp;
    }
}
```

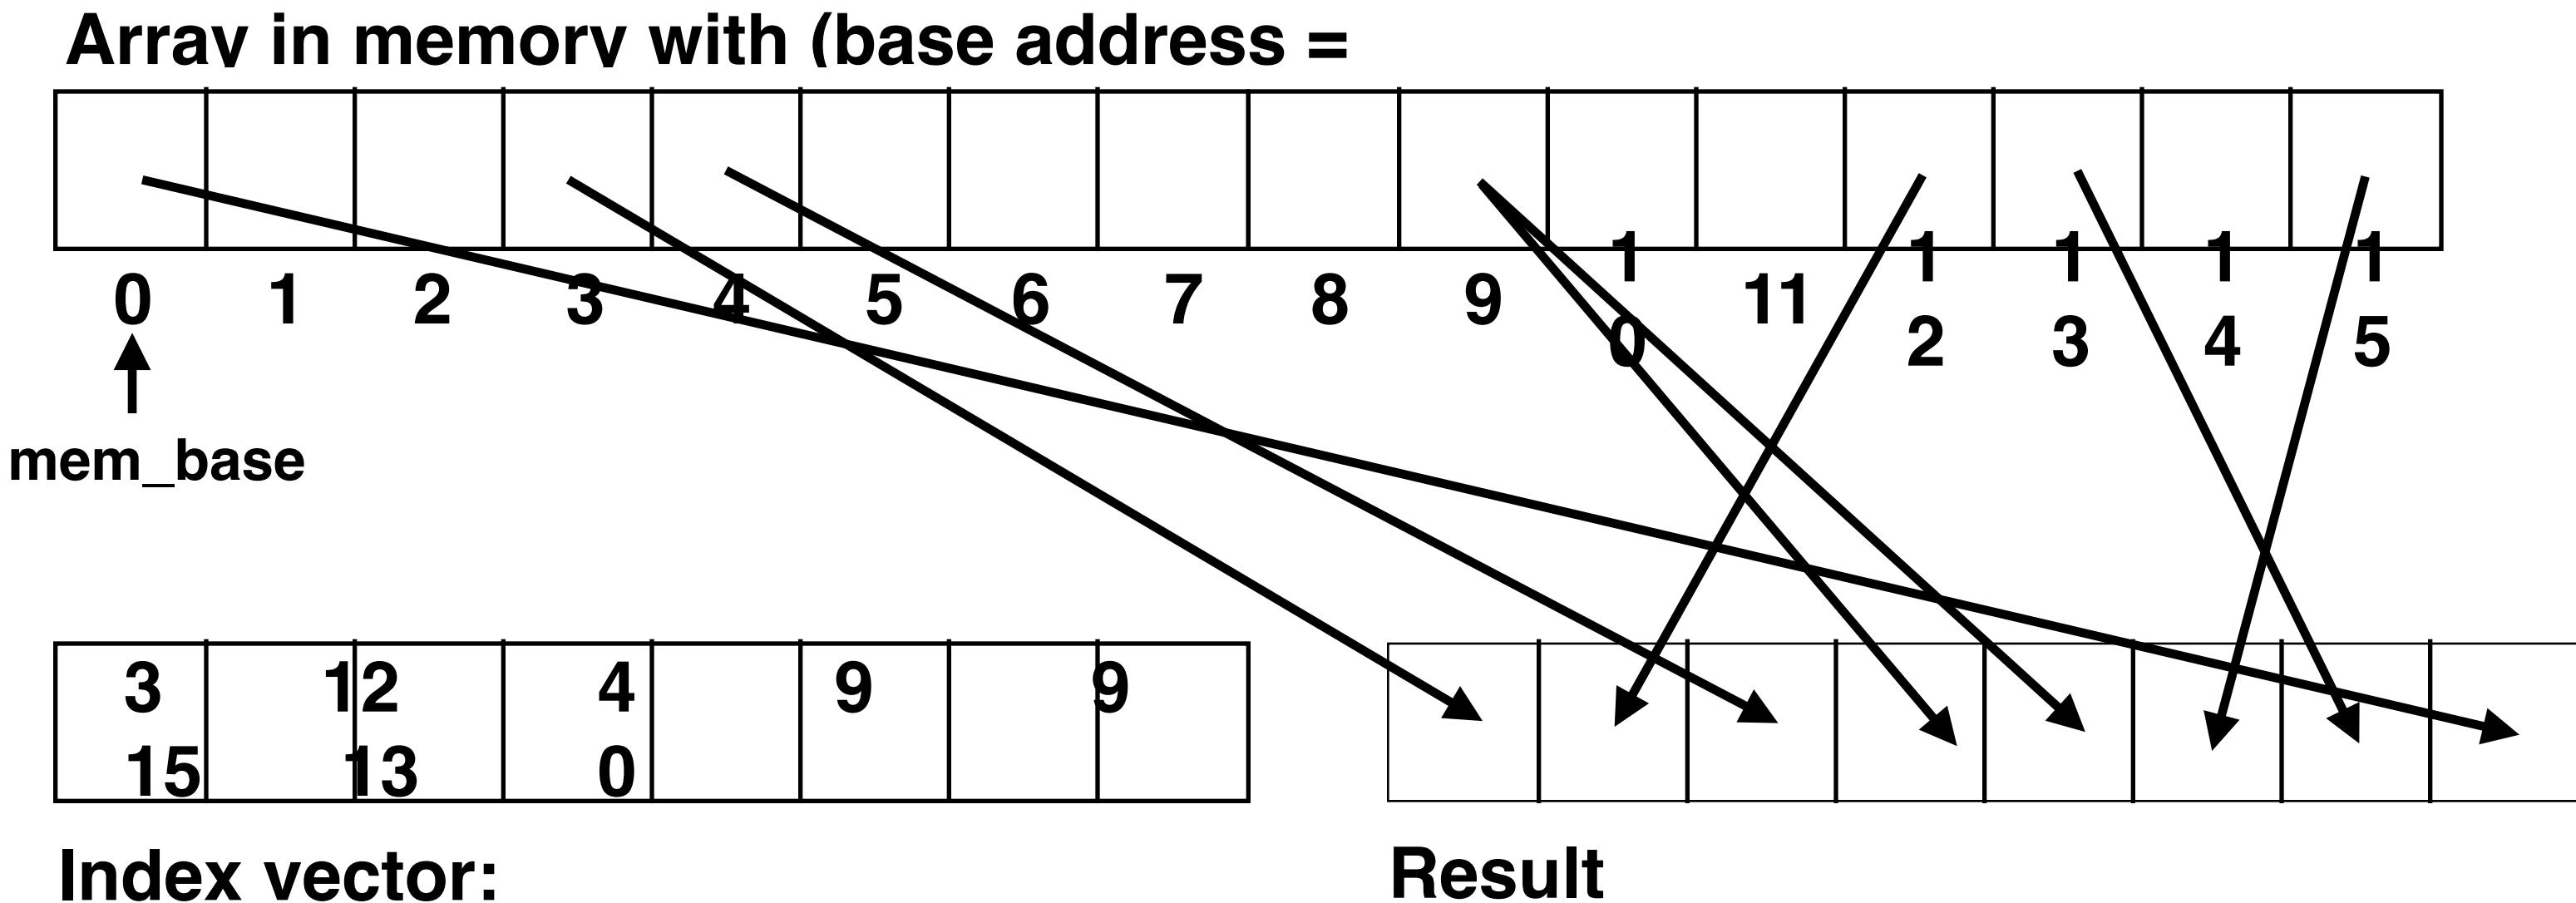
ISPC equivalent:

```
export void absolute_value(
    uniform float N,
    uniform float* input,
    uniform float* output,
    uniform int* indices)
{
    foreach (i = 0 ... n)
    {
        if (input[i] < 0)
            output[indices[i]] = -input[i];
        else
            output[indices[i]] = input[i];
    }
}
```

Gather instruction

gather(R1, R0, mem_base);

“Gather from buffer `mem_base` into `R1` according to



Gather supported with AVX2 in 2013

But AVX2 does not support SIMD scatter (must implement as scalar loop)

Scatter instruction exists in AVX512

Hardware supported gather/scatter does exist on GPUs.

Summary: data-parallel model

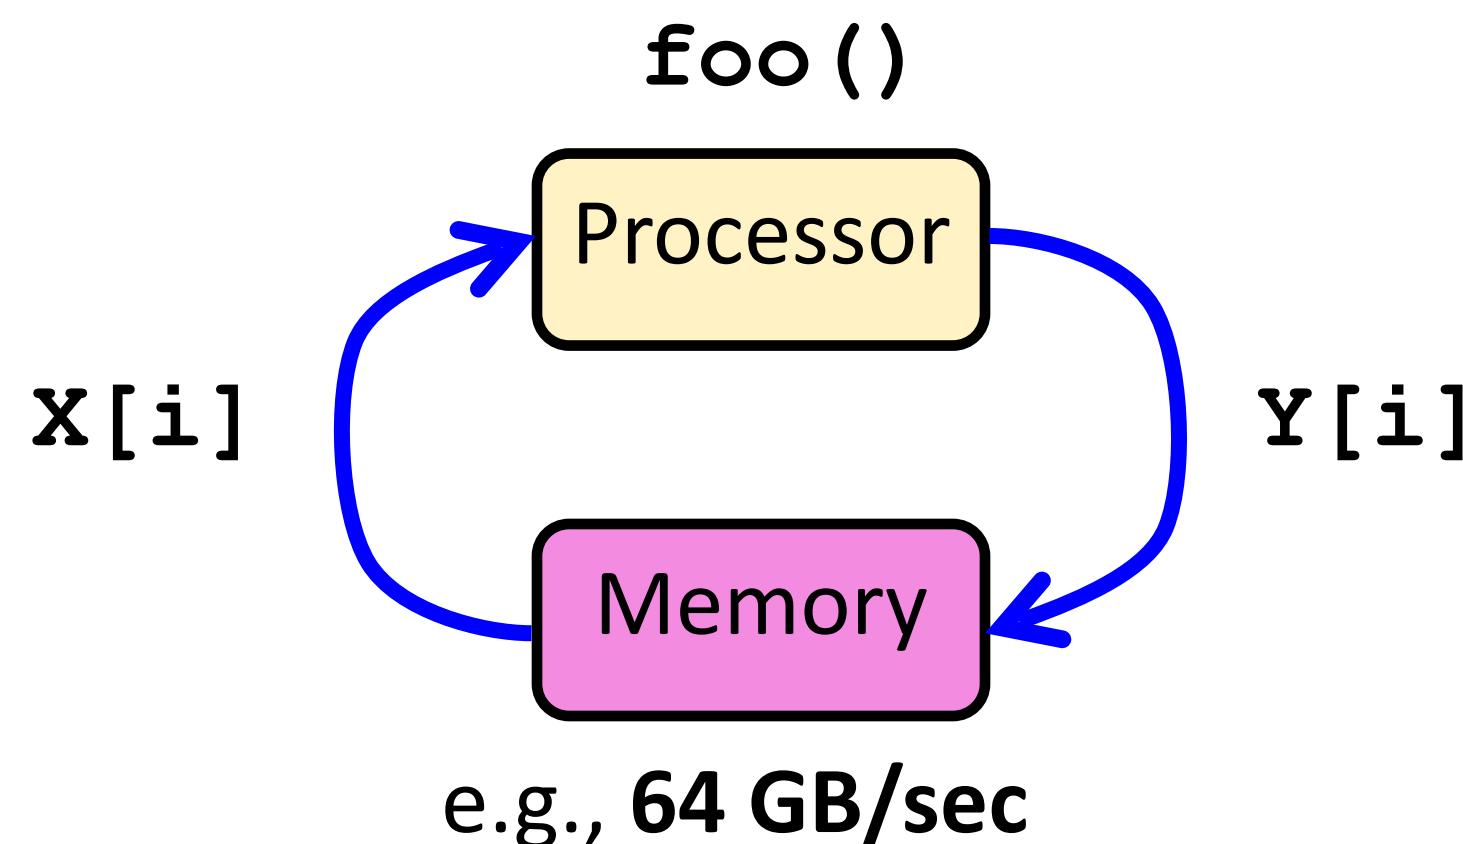
- Data-parallelism is about **imposing rigid program structure** to facilitate simple programming and advanced optimizations
- Basic structure: **map a function onto a large collection of data**
 - Functional: side-effect free execution
 - No communication among distinct function invocations (allow invocations to be scheduled in any order, including in parallel)
- In practice that's how many simple programs work
- But... many modern performance-oriented data-parallel languages do not strictly enforce this structure

The systolic arrays model

Motivation for Systolic Arrays

- Target domain:
 - data-intensive applications where **memory bandwidth** is the **bottleneck**
- Simple bandwidth limited example:

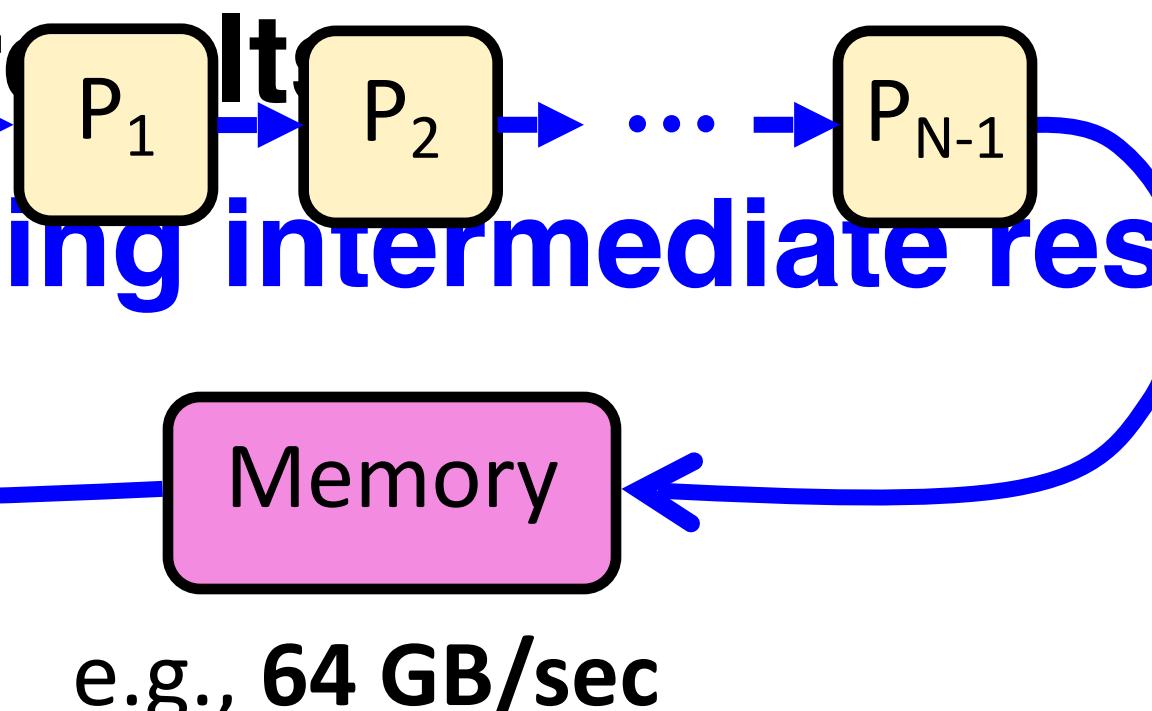
```
for i = 0 to N-1
    Y[i] = foo(X[i]);
```



- Due to lack of temporal locality, caches do not solve the problem

Systolic Arrays: Key Insights

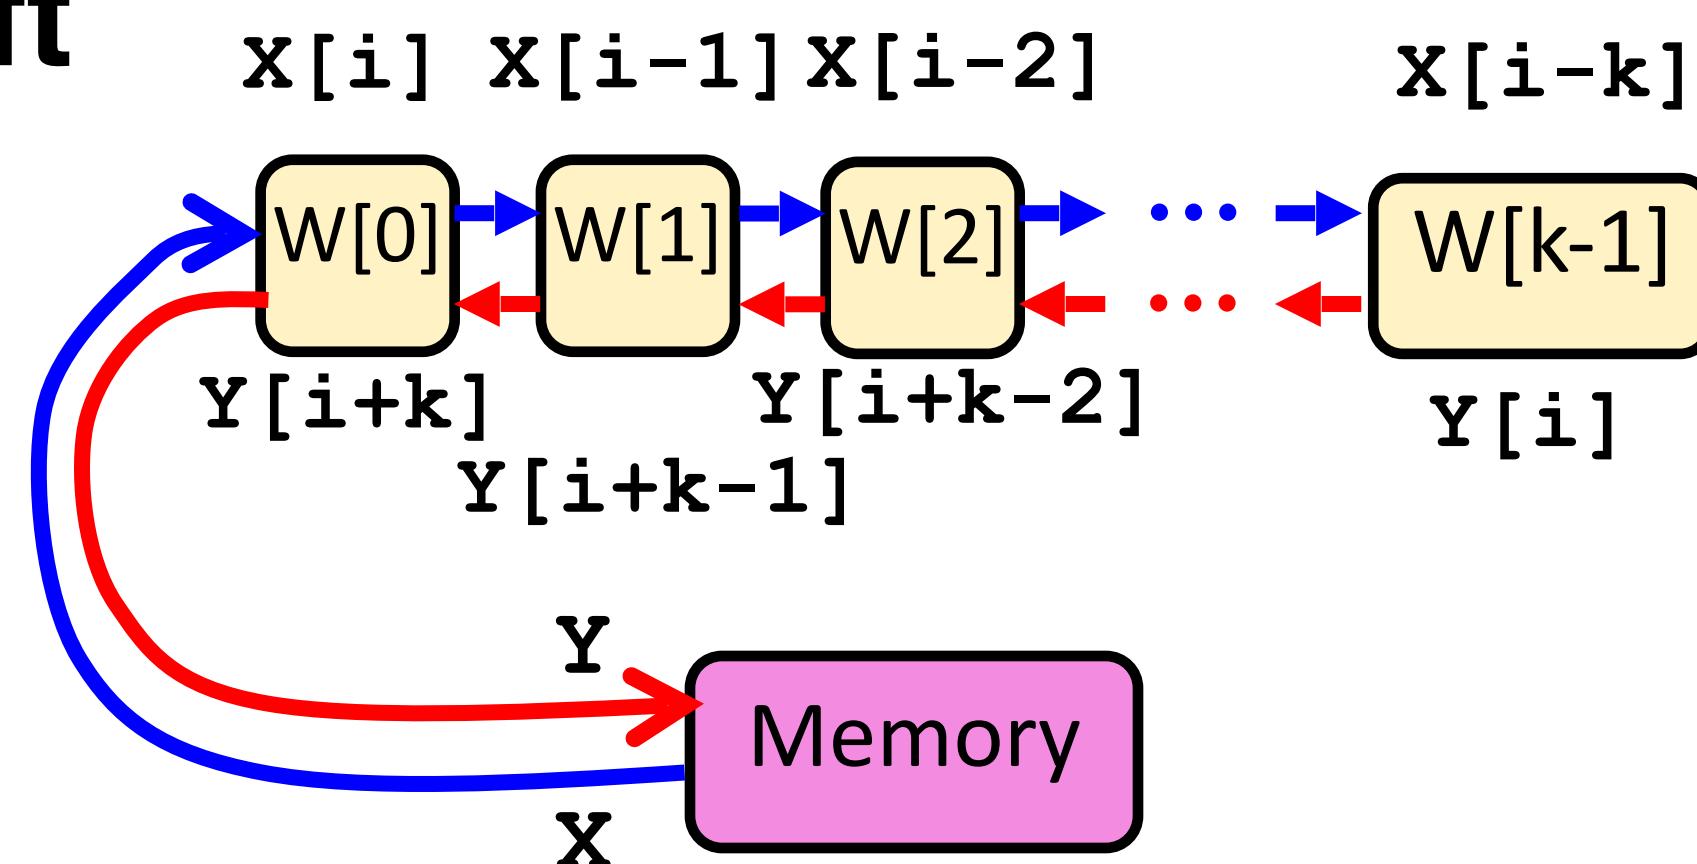
- Because memory bandwidth is the critical resource:
 - we really want to **avoid accessing memory unnecessarily**
- Hence whenever we read data from memory:
 - **perform all necessary computation** on it before writing back a result
 - by **directly pipelining intermediate results** between processors



Systolic Arrays Example: 1D

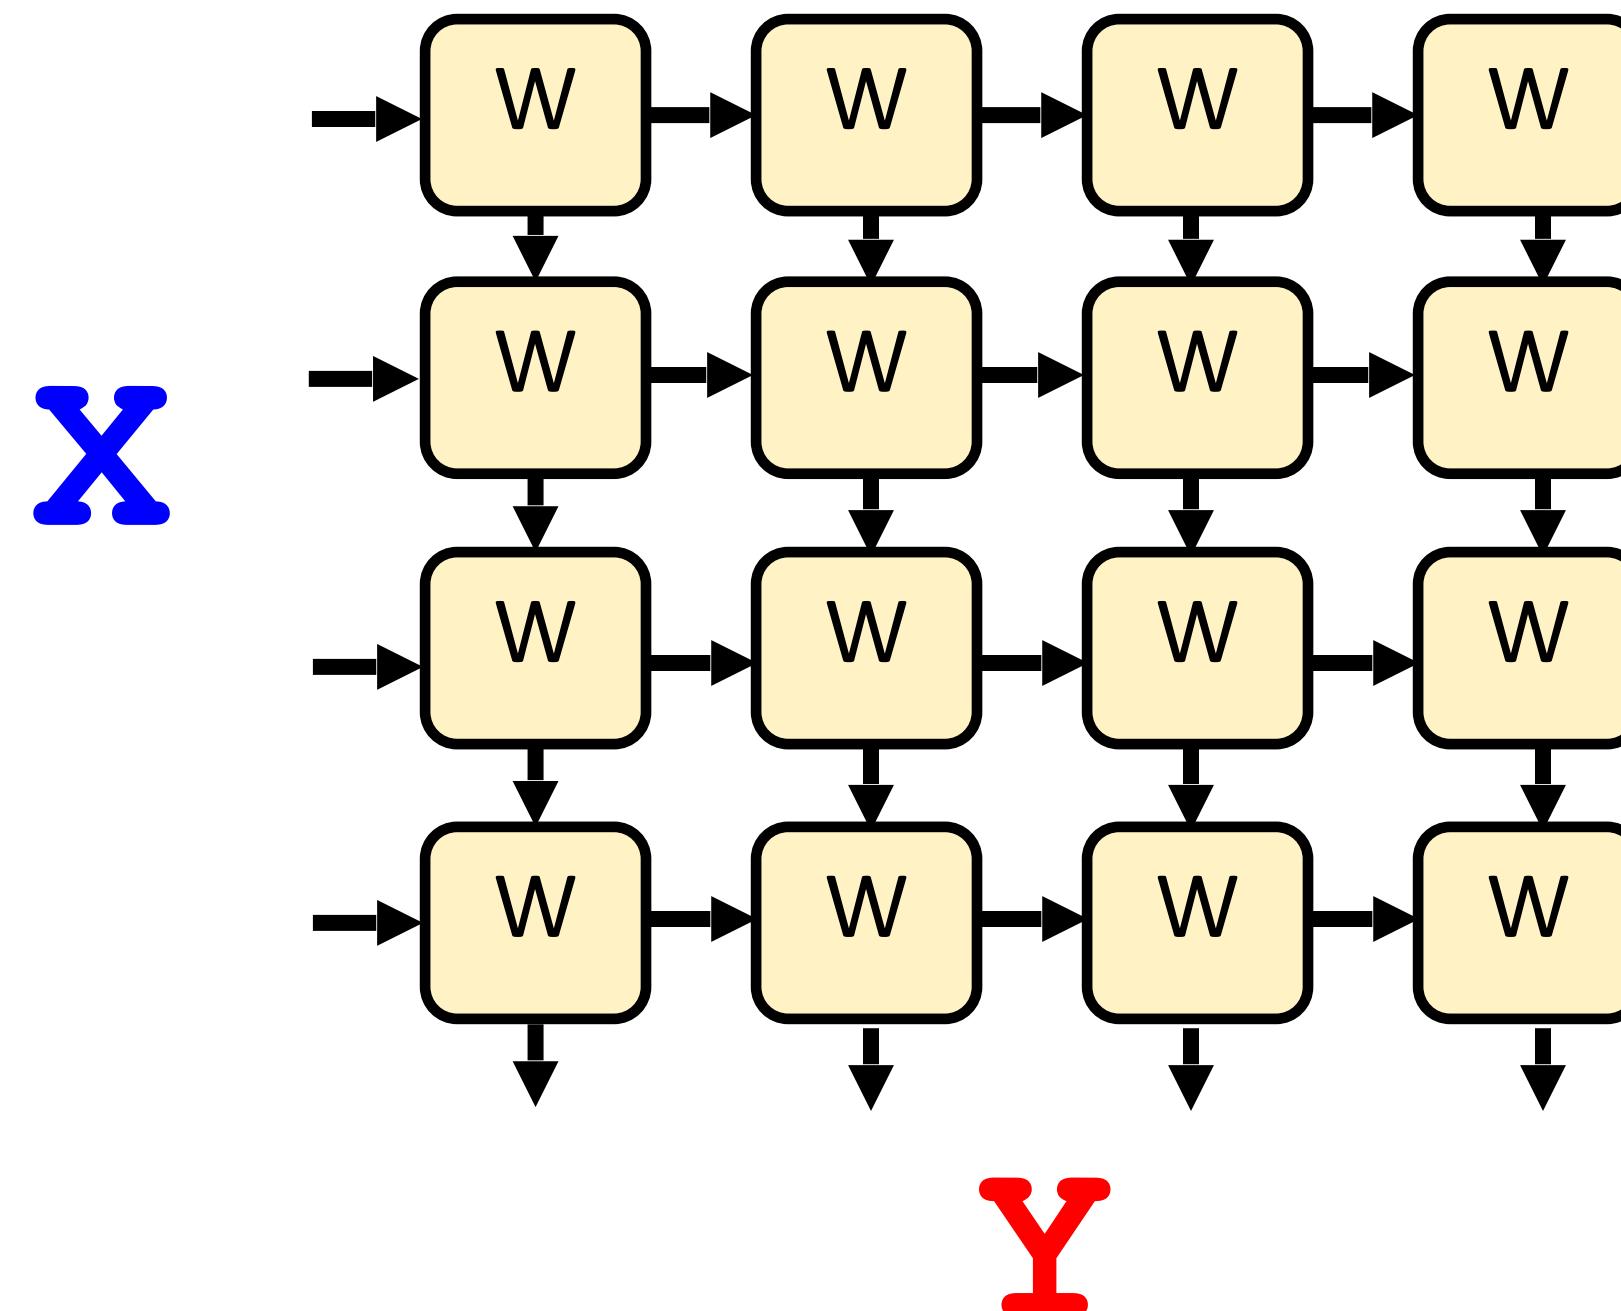
$$Y[i] = \sum_{k=0}^{K-1} W[k] \cdot X[i - k]$$

- pre-load the weights ($W[i]$) onto each processor
- input data ($X[i]$) flows from left to right
- partial results (eventually forming $Y[i]$) flow from right to left



Systolic Arrays Today: ML

$$Y_{i,j} = \sum_{k=1}^n X_{i,k} W_{k,j}$$



- Basic idea behind AWS Trainium. Google

Four parallel programming

- **Shared address space**
 - Communication is unstructured, implicit in loads and stores
 - Natural way of programming, but can shoot yourself in the foot easily
 - Program might be correct, but not perform well
- **Message passing**
 - Structure all communication as messages
 - Often harder to get first correct program than shared address space
 - Structure often helpful in getting to first correct, scalable program
- **Data parallel**
 - Structure computation as a big “map” over a collection

Modern practice: mixed

- Use **shared address space** programming **within a multi-core node** of a cluster, use **message passing between nodes**
 - Very, very common in practice
 - Use convenience of shared address space where it can be implemented efficiently (within a node), require explicit communication elsewhere
- Data-parallel-ish programming models support shared-memory style synchronization primitives in kernels
 - Permit limited forms of inter-iteration communication (e.g., CUDA, OpenCL)
- In a future lecture... CUDA/OpenCL use data-

Summary

- Programming models provide a way to think about the organization of parallel programs. They provide **abstractions** that admit many possible **implementations**.
- Restrictions imposed by these abstractions are designed to reflect realities of parallelization and communication costs
 - **Shared address space machines**: hardware supports any processor accessing any address
 - **Messaging passing machines**: may have hardware to accelerate message send/receive/buffering
 - It is desirable to keep “abstraction distance” low so programs have predictable performance, but want it high enough for code flexibility/portability