Lecture 22:

Domain-Specific Programming Systems

**Parallel Computer Architecture and Programming
CMU 15-418/15-618, Fall 2025**

Logistics

- **Reminder: Course feedback poll is live:**
  https://docs.google.com/forms/d/e/1FAIpQLSf8QuPJL20ZX6EFCsxsaWmJ5UVBtXgNYDC8SNoKf3GWfqv7zQ/viewform?usp=dialog **(also Piazza)**

  - **Small number of responses so far**

  - **Please take a few minutes to share your feedback!**

- **Special lecturer on Friday: Nathan Serafin, CMU PhD candidate**

Review: The inexorable rise of hardware specialization

- **Hardware perspective: It's always best to specialize for efficiency**
  - Eliminate "waste" of programmable design (can be >99% of energy)
  - Parallelism: Replicate efficient custom circuit many times
  - Data movement: Put exactly the memories you want, exactly where they need to be

- **Software perspective: It's always best to be programmable**
  - Let the programmer/compiler decide what to do
    - Enable innovation, eliminate unnecessary computation
  - What am I supposed to do with specialized hardware?!?

Course themes:

Designing computer systems that <u>scale</u>
(running faster given more resources)

Designing computer systems that are <u>efficient</u>
(running faster under constraints on resources)

Techniques discussed:
    Exploiting parallelism in applications
   **Exploiting locality in applications**
    Leveraging hardware specialization (earlier lecture)

Claim: most software uses modern hardware resources inefficiently

- **Consider a piece of sequential C code**
  - Call the performance of this code our "baseline performance"
- **Well-written sequential C code: ~ 5-10x faster**
- **Assembly language program:** *maybe* **another small constant factor faster**
- **Java, Python, PHP, etc. ??**

Code performance: relative to C (single core)

GCC -O3 (no manual vector optimizations)

Slowdown (Compared to C++)

Legend:
- = NBody (blue)
- = Mandlebrot (green)
- = Tree Alloc/Delloc (orange)
- = Power method (compute eigenvalue) (red)

Values above bars: 51, 40/57/53, 47, 44/114x

Categories: Java, Scala, C# (Mono), Haskell, Go, Javascript (V8), Lua, PHP, Python 3, Ruby (JRuby)

"no data" appears for Scala (Tree Alloc) position and Javascript (V8) Mandlebrot position.

Data from: The Computer Language Benchmarks Game:
http://shootout.alioth.debian.org

Even good C code is inefficient

Recall Assignment 1's Mandelbrot program

Consider execution on a high-end laptop: quad-core, Intel Core i7, AVX instructions...

Single core, with AVX vector instructions: 5.8x speedup over C implementation

Multi-core + hyper-threading + AVX instructions: 21.7x speedup

Conclusion: basic C implementation compiled with -O3 leaves <u>a lot</u> of performance on the table

Making efficient use of modern machines is challenging

(proof by assignments 2, 3, and 4)

In our assignments, you only programmed homogeneous parallel computers.
...And parallelism even in that context was not easy.

Assignment 2: GPU cores only

Assignments 3 & 4: shared memory / message passing

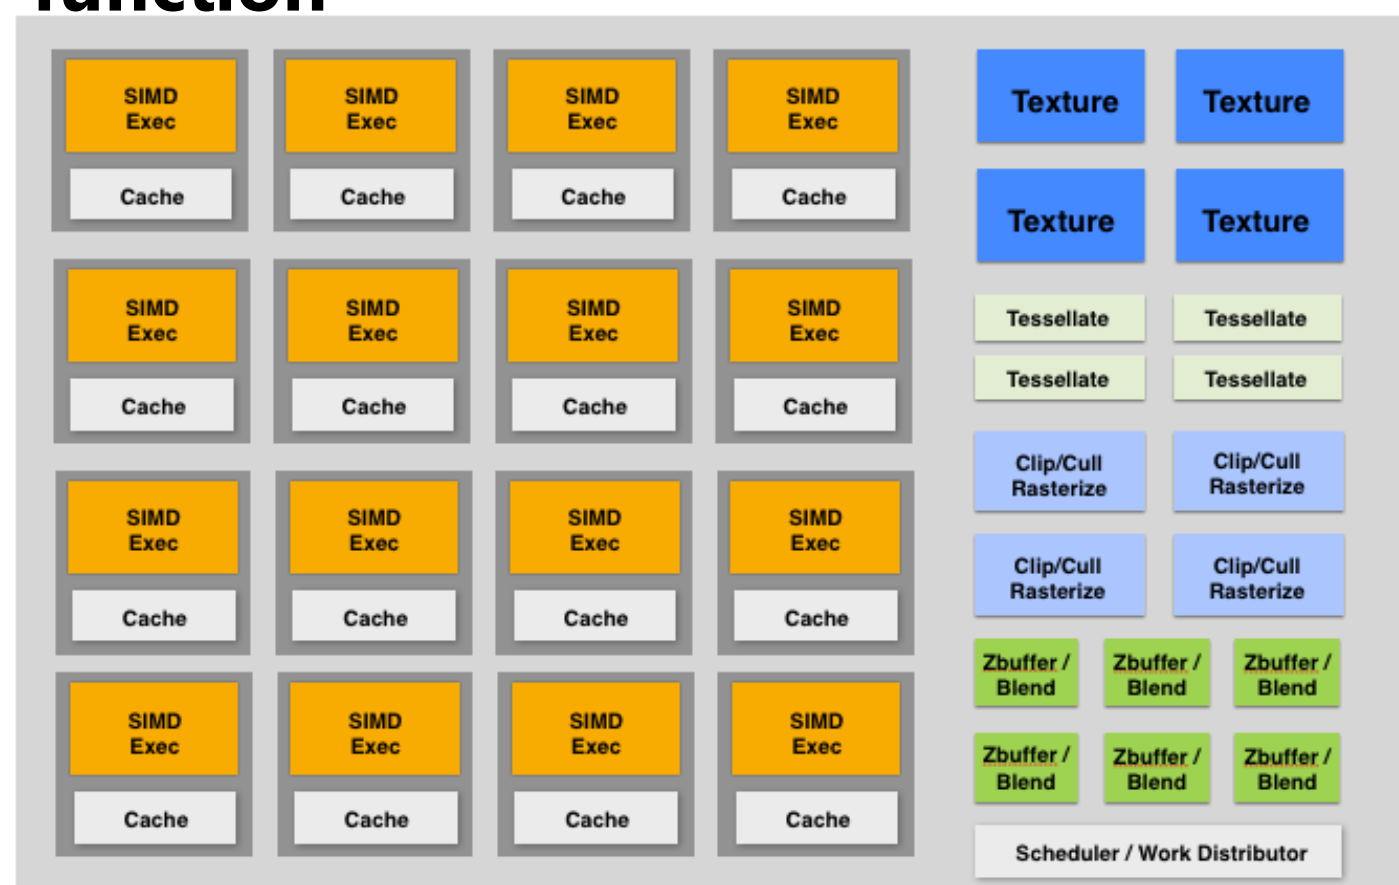# Recall: need for efficiency leading to heterogeneous parallel platforms


Integrated CPU + GPU

CPU+data-parallel accelerator

GPU:
throughput cores + fixed-function

Qualcomm Snapdragon SoC
800 PROCESSOR

Krait 400 CPU
features 28HPm process technology
superior
2GHz+ performance

Adreno 330 for
advanced graphics

Hexagon QDSP6
for ultra low power
applications and custom
programmability

Integrated LTE³, 802.11ac³, USB 3.0
and BT 4.0 offers broad array
of high speed connectivity

Ultra HD Capture
and Playback
DTS-HD and Dolby
Digital Plus audio
Expanded Gestures

55MP with dual ISP

Support for up
to 2560x2048 display
Miracast 1080p
HD support

IZat GNSS with
support for three
GPS constellations

Mobile system-on-a-chip:
CPU+GPU+media processing

Hardware trend: specialization of execution

- **Multiple forms of parallelism**
  - **SIMD/vector processing** $\longrightarrow$ **Fine-granularity parallelism: perform same logic on different data**
  - **Multi-threading** $\longrightarrow$ **Mitigate inefficiencies (stalls) caused by unpredictable data access**
  - **Multi-core**
  - **Multiple node** $\longrightarrow$ **Varying scales of coarse-granularity parallelism**
  - **Multiple server**

- **Heterogeneous execution capability**

  - **Programmable, latency-centric** (e.g., "CPU-like" cores)

  - **Programmable, throughput-optimized** (e.g., "GPU-like" cores)

  - **Fixed-function, application-specific** (e.g., image/video/audio processing)

Motivation for parallelism and specialization: maximize compute capability given constraints on chip area, chip energy consumption. Result: amazingly high compute capability in a wide range of devices!

Hardware diversity is a huge challenge

- **Machines with very different performance characteristics**
- **Even worse: different technologies and performance characteristics within <u>the same machine</u> at different scales**
  - **Within a core: SIMD, multi-threading: fine-granularity sync and communication**
  - **Across cores: coherent shared memory via fast on-chip network**
  - **Hybrid CPU+GPU multi-core: incoherent (potentially) shared memory**
  - **Across racks: distributed memory, multi-stage network**

Variety of programming models to abstract HW

- **Machines with very different performance characteristics**

- **Worse: different technologies and performance characteristics within <u>the same machine</u> at different scales**

  - **Within a core**: SIMD, multi-threading: fine grained sync and comm

    - <u>Abstractions</u>: **SPMD programming (ISPC, Cuda, OpenCL, Metal, Renderscript)**

  - **Across cores**: coherent shared memory via fast on-chip network

    - <u>Abstractions: </u>**OpenMP pragma, Cilk, TBB**

  - **Hybrid CPU+GPU** multi-core: incoherent (potentially) shared memory

    - **Abstractions: OpenCL, CUDA, HIP**

This is a huge challenge

- **Machines with very different performance characteristics**

- **Worse: different performance characteristics within <u>the same machine</u> at different scales**

- **To be efficient, software must be optimized for HW characteristics**

  - **Difficult even in the case of one level of one machine**

  - **Combinatorial complexity of optimizations when considering a complex machine, or different machines**

  - **Loss of software portability**

Open computer science question:

How do we enable programmers to productively write software that efficiently uses current and future heterogeneous, parallel machines?

The [magical] ideal parallel programming language

High Performance
(software is scalable and efficient)

??

Productivity
(ease of development)

Completeness
(applicable to most problems we want to write a program for)
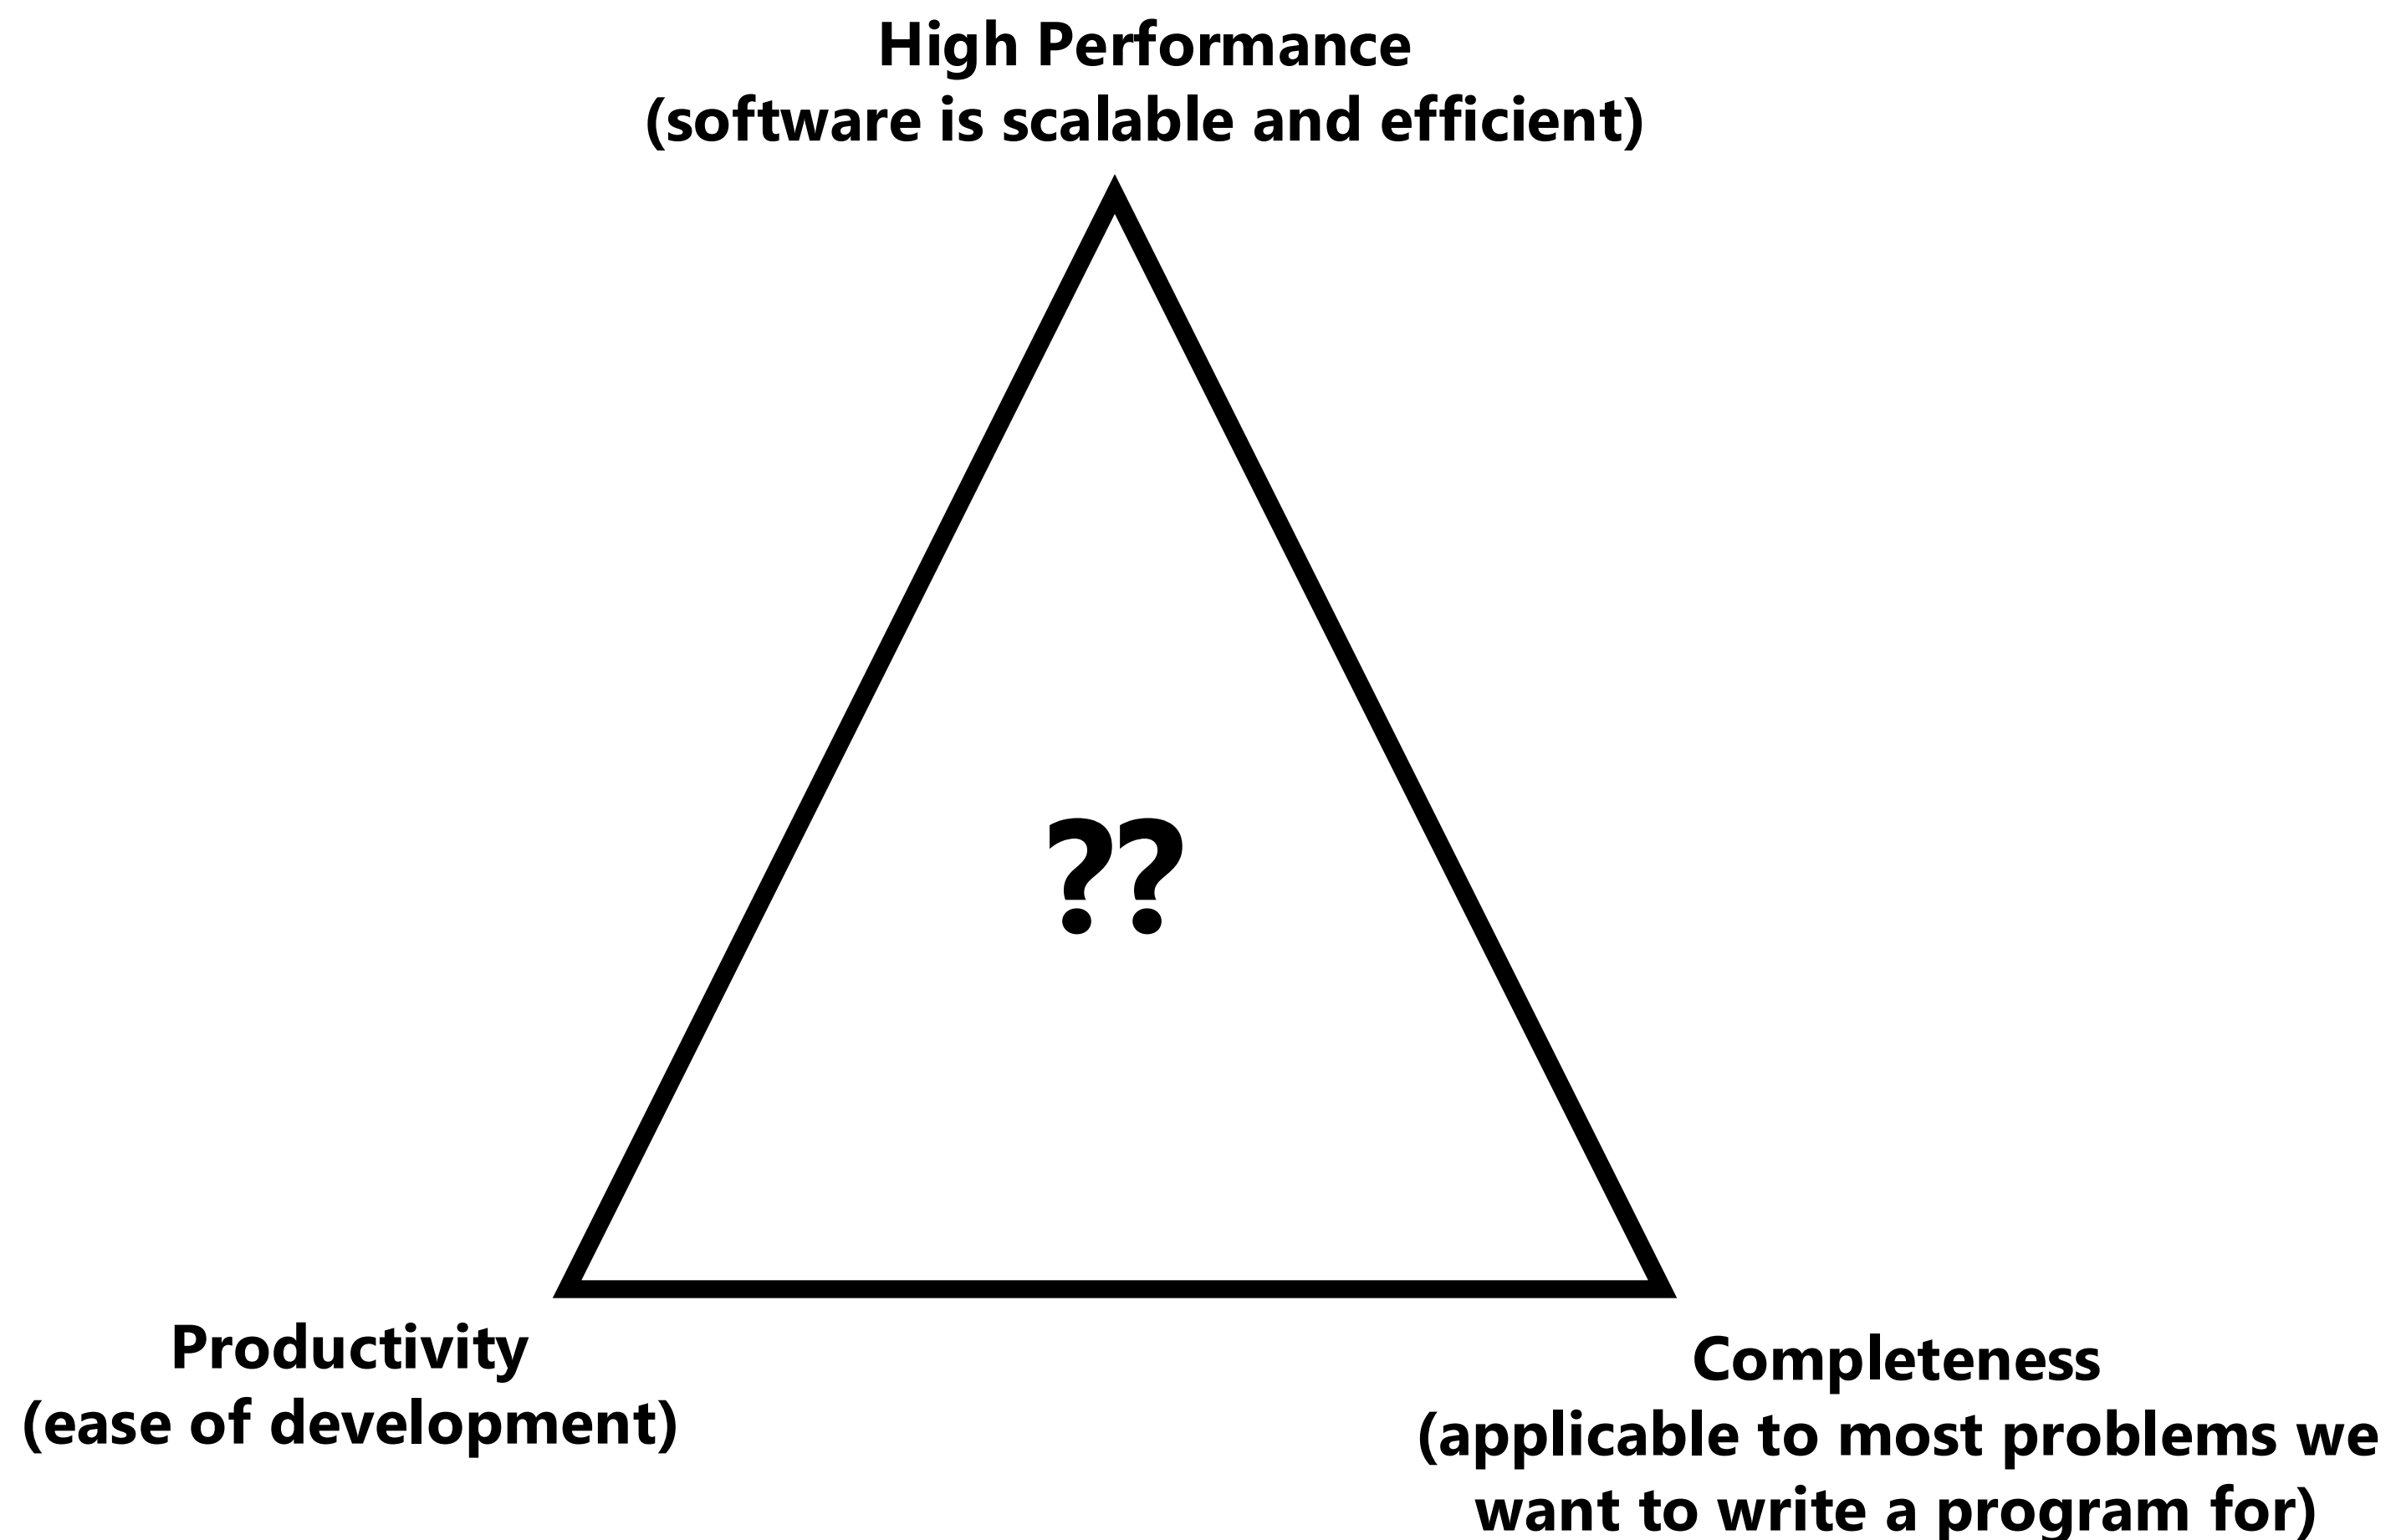
Successful programming languages
(Success = widely used)

High Performance
(software is scalable and efficient)

C/C++

??

Productivity
(ease of development)

Completeness
(applicable to most problems we want to write a program for)

python
JavaScript
Ruby
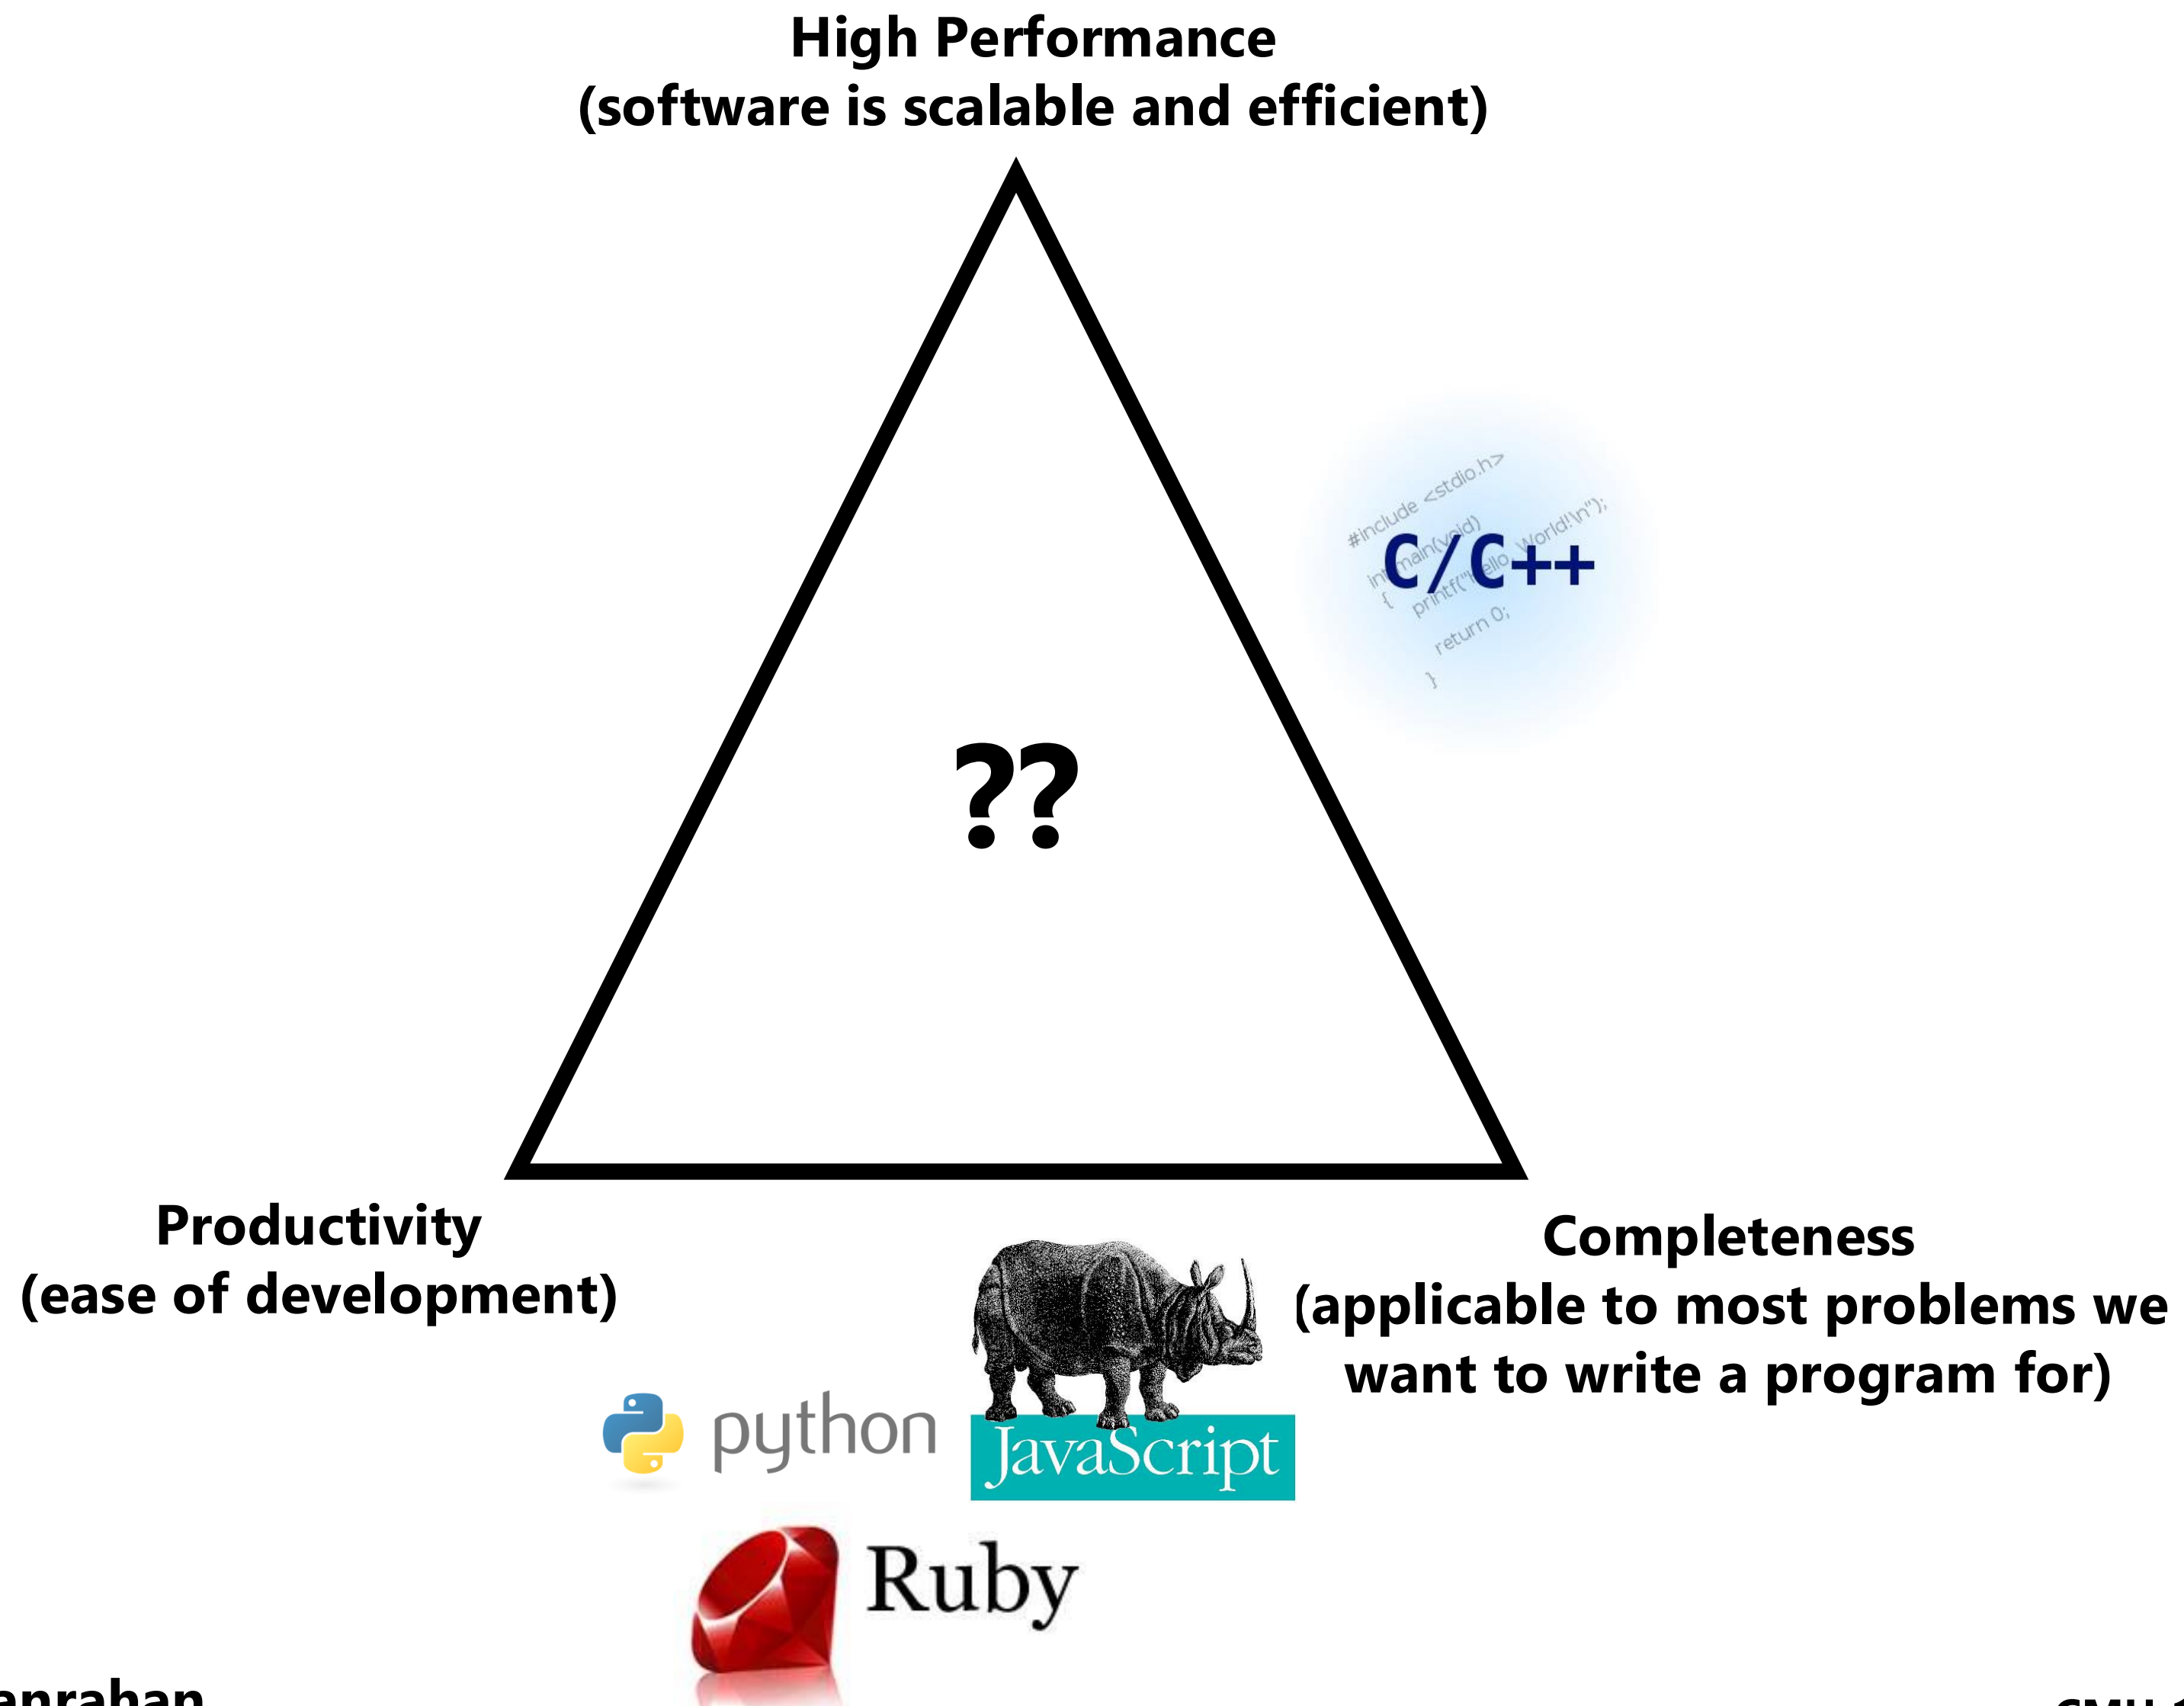
Growing interest in domain-specific programming systems
To realize high performance and productivity: willing to sacrifice completeness

High Performance
(software is scalable and efficient)

Domain-specific languages (DSL) and programming frameworks

C/C++

??

Productivity
(ease of development)

Completeness
(applicable to most problems we want to write a program for)

python   JavaScript

Ruby
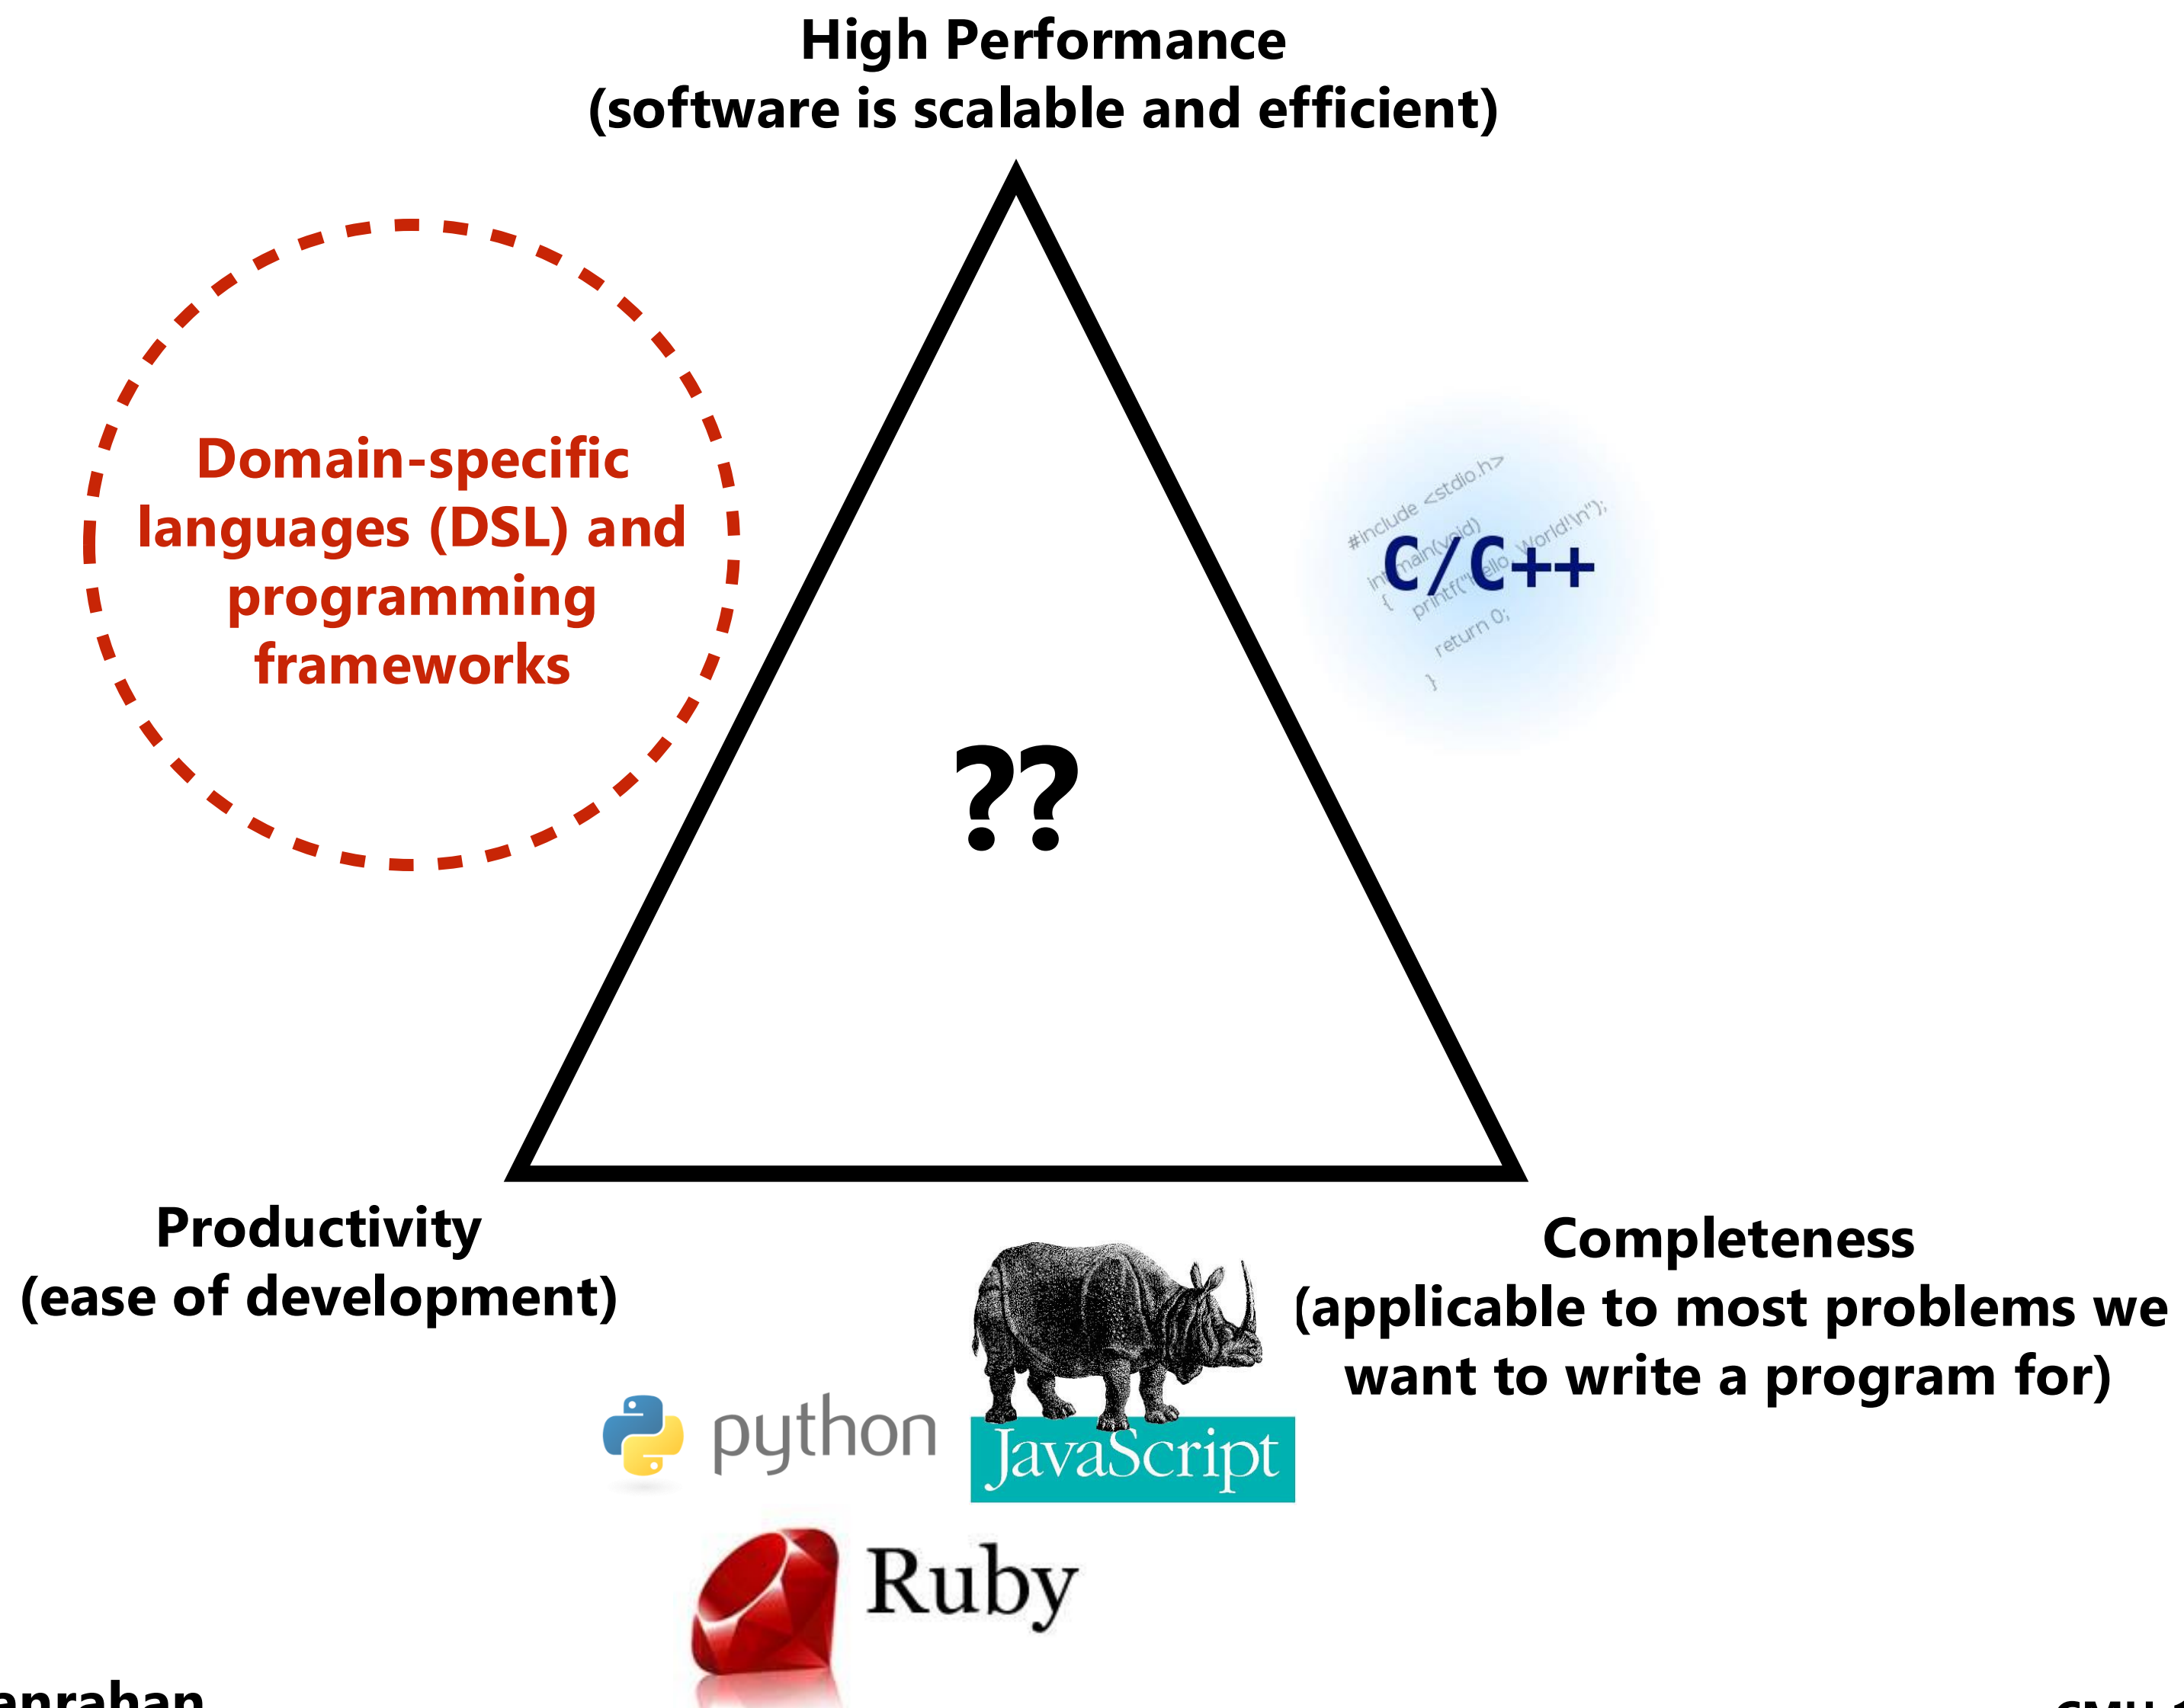
Domain-specific programming systems

- **Main idea:** raise level of abstraction for expressing programs

- **Introduce high-level programming primitives specific to an application domain**

  - **<u>Productive</u>: intuitive to use, portable across machines, primitives correspond to behaviors frequently used to solve problems in targeted domain**

  - **<u>Performant</u>: system uses domain knowledge to provide efficient, optimized implementation(s)**

    - **Given a machine: system knows what algorithms to use, parallelization strategies to employ for this domain**

    - **Optimization goes beyond efficient mapping of software to hardware! The hardware platform itself can be optimized to the abstractions as well**

- **<u>Cost</u>: loss of generality/completeness**

Two domain-specific programming examples

1. **Liszt**: for scientific computing on meshes

2. **Halide**: for image processing

What are other domain specific languages?

(SQL is another good example)

Example 1:
Lizst: a language for solving PDE's on meshes

**[DeVito et al. Supercomputing 11,
SciDac '11]**

**Slide credit for this section of lecture:
Pat Hanrahan and Zach Devito (Stanford)**

http://liszt.stanford.edu/

What a Liszt program does

A Liszt program is run on a mesh
A Liszt program defines, and compute the value of, fields defined on the mesh

Position is a field defined at each mesh vertex.
The field's value is represented by a 3-vector.

```
val Position = FieldWithConst[Vertex,Float3](0.f, 0.f, 0.f)
val Temperature = FieldWithConst[Vertex,Float](0.f)
val Flux = FieldWithConst[Vertex,Float](0.f)
val JacobiStep = FieldWithConst[Vertex,Float](0.f)
```

Color key:
Fields
Mesh entity

Notes:
Fields are a higher-kinded type
(special function that maps a type to a new type)

Liszt program: heat conduction on mesh

Program computes the value of fields defined on meshes

Set flux for all vertices to 0.f;

```
var i = 0;
while ( i < 1000 ) {
    Flux(vertices(mesh)) = 0.f;
    JacobiStep(vertices(mesh)) = 0.f;
    for (e <- edges(mesh)) {
        val v1 = head(e)
        val v2 = tail(e)
        val dP = Position(v1) - Position(v2)
        val dT = Temperature(v1) - Temperature(v2)
        val step = 1.0f/(length(dP))
        Flux(v1) += dT*step
        Flux(v2) -= dT*step
        JacobiStep(v1) += step
        JacobiStep(v2) += step
    }
    i += 1
}
```

Color key:
Fields
Mesh
Topology functions
Iteration over set

Independently, for each edge in the mesh

Given edge, loop body accesses/modifies field values at adjacent mesh vertices

Access value of field at mesh vertex v2

Liszt's topological operators

Used to access mesh elements relative to some input vertex, edge, face, etc.
Topological operators are the <u>only way</u> to access mesh data in a Liszt program
Notice how many operators return sets (e.g., "all edges of this face")

```
BoundarySet¹[ME <: MeshElement](name : String) : Set[ME]
vertices(e : Mesh) : Set[Vertex]
cells(e : Mesh) : Set[Cell]
edges(e : Mesh) : Set[Edge]
faces(e : Mesh) : Set[Face]
```

```
vertices(e : Vertex) : Set[Vertex]
cells(e : Vertex) : Set[Cell]
edges(e : Vertex) : Set[Edge]
faces(e : Vertex) : Set[Face]
```

```
cells(e : Cell) : Set[Cell]
vertices(e : Cell) : Set[Vertex]
faces(e : Cell) : Set[Face]
edges(e : Cell) : Set[Edge]
```

```
vertices(e : Edge) : Set[Vertex]
facesCCW²(e : Edge) : Set[Face]
cells(e : Edge) : Set[Cell]
head(e : Edge) : Vertex
tail(e : Edge) : Vertex
flip⁴(e : Edge) : Edge
towards⁵(e : Edge, t : Vertex) : Edge
```

```
cells(e : Face) : Set[Cell]
edgesCCW²(e : Face) : Set[Edge]
vertices(e : Face) : Set[Vertex]
inside³(e : Face) : Cell
outside³(e : Face) : Cell
flip⁴(e : Face) : Face
towards⁵(e : Face, t : Cell) : Face
```

Liszt programming

- **A Liszt program describes operations on fields of an abstract mesh representation**

- **Application specifies type of mesh (regular, irregular) and its topology**

- **Mesh representation is chosen by Liszt (not by the programmer)**
  - **Based on mesh type, program behavior, and target machine**

Well, that's interesting. I write a program, and the compiler decides what data structure it should use based on what operations my code performs.

Compiling to parallel computers

Recall challenges you have faced in your assignments

1. Identify **parallelism**
2. Identify **data locality**
3. Reason about required **synchronization**

Now consider how to automate this process in the Liszt compiler.

Key: determining program dependencies

1. Identify **parallelism:**
    Absence of dependencies implies code can be executed in parallel
2. Identify **data locality:**
    Partition data based on dependencies (localize dependent computations for faster synchronization)
3. Reason about required **synchronization:**
    Synchronization is needed to respect dependencies (must wait until the values a computation depends on are known)

In general programs, compilers are unable to infer dependencies at global scale: `a[f(i)] += b[i]`   (must execute `f(i)` to know if dependency exists across loop iterations `i`)

Liszt is constrained to allow dependency analysis

Lizst infers "stencils": "stencil" = mesh elements accessed in an iteration of loop
= dependencies for the iteration

Statically analyze code to find stencil of each top-level for loop

- Extract nested mesh element reads
- Extract field operations

```
for (e <- edges(mesh)) {
  val v1 = head(e)
  val v2 = tail(e)
  val dP = Position(v1) - Position(v2)
  val dT = Temperature(v1) - Temperature(v2)
  val step = 1.0f/(length(dP))
  Flux(v1) += dT*step
  Flux(v2) -= dT*step
  JacobiStep(v1) += step
  JacobiStep(v2) += step
}
…
```
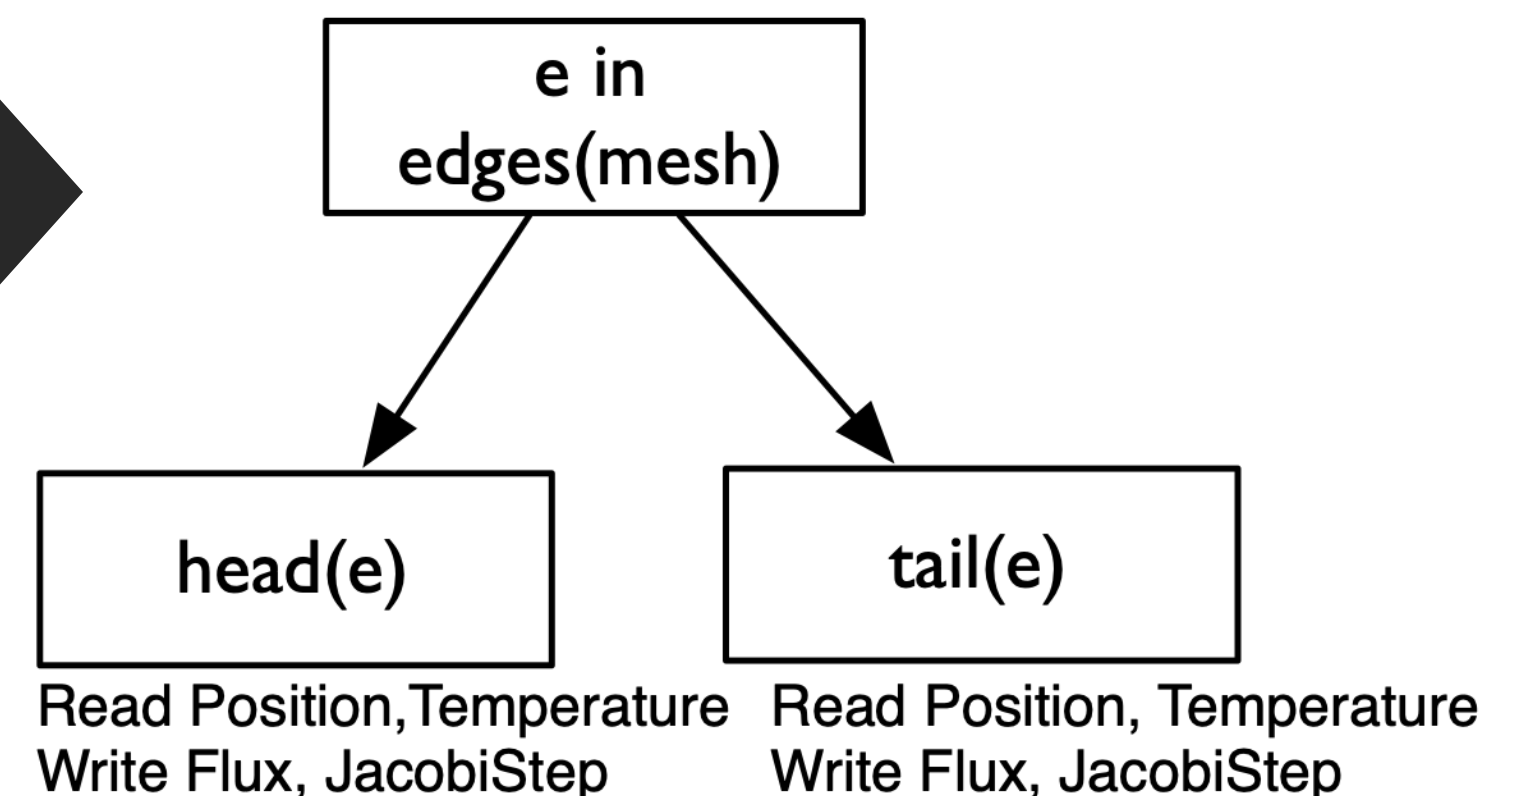
Edge 6's read stencil is D and F

e in edges(mesh)

head(e)
Read Position,Temperature
Write Flux, JacobiStep

tail(e)
Read Position, Temperature
Write Flux, JacobiStep

Restrict language for dependency analysis

Language restrictions:

- **Mesh elements are only accessed through built-in topological functions:**

```
cells(mesh), …
```

- **Single static assignment:**

```
val v1 = head(e)
```

- **Data in fields can only be accessed using mesh elements:**

```
Pressure(v)
```

- **No recursive functions**

Restrictions allow compiler to automatically infer stencil for a loop iteration.

Portable parallelism: use dependencies to implement different parallel execution strategies

I'll discuss two strategies...

Strategy 1: mesh partitioning

Strategy 2: mesh coloring

Ghost Cell

Owned Cell

Schedule set of nonconflicting threads per color

Batch 1	Batch 2	Batch 3	Batch 4

1	3	8	11	0	5	7	10	4	9	2
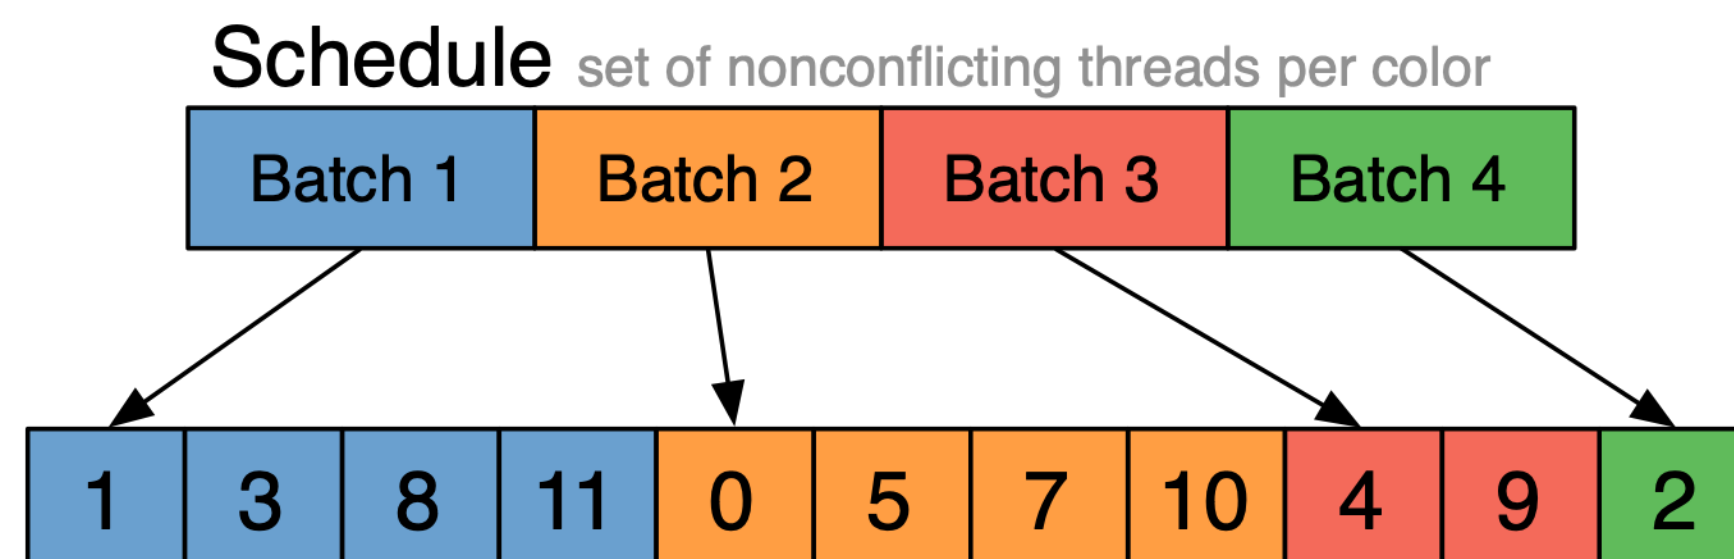
Imagine compiling a Lizst program to the (entire) Latedays cluster

(multiple nodes, distributed address space)

How might Liszt distribute a graph across these nodes?

Distributed memory implementation of Liszt
Mesh + Stencil → Graph → Partition

```
for(f <- faces(mesh)) {
  rhoOutside(f) =
    calc_flux(f, rho(outside(f))) +
    calc_flux(f, rho(inside(f)))
}
```

Initial Partition
(by ParMETIS)
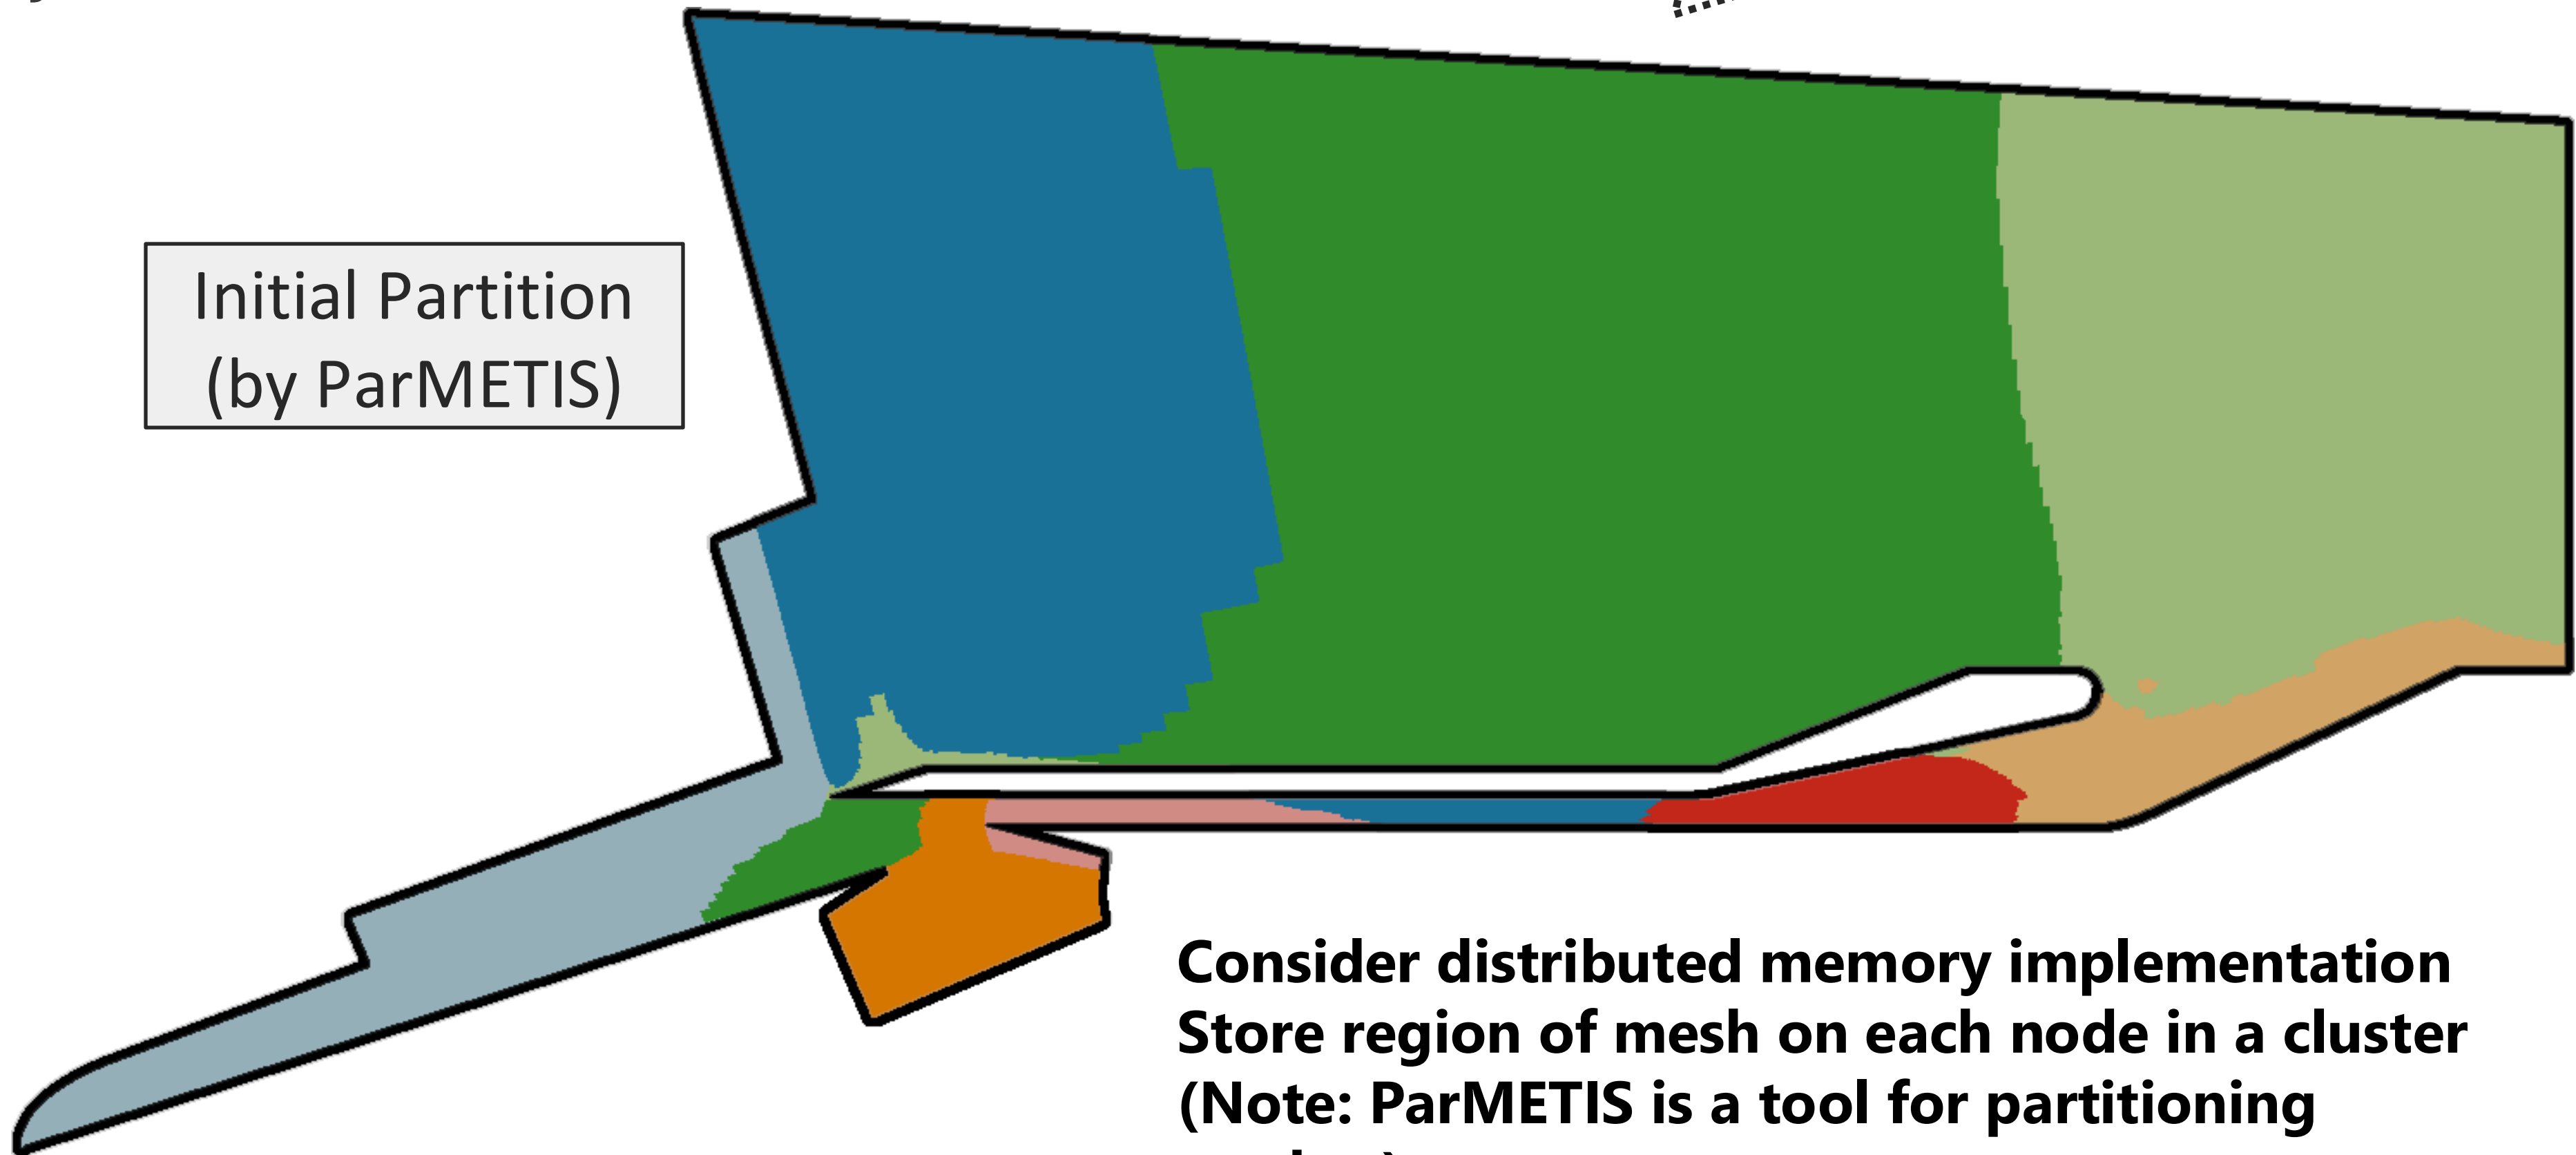
Consider distributed memory implementation
Store region of mesh on each node in a cluster
(Note: ParMETIS is a tool for partitioning
meshes)

Ghost Cells

Each processor also needs data for neighboring cells to perform computation ("ghost cells") Listz allocates ghost region storage and emits required communication to implement topological operators.

Imagine compiling a Lizst program to a GPU

(single address space, many tiny threads)

GPU implementation: parallel reductions

In previous example, one region of mesh assigned per processor (or node in MPI cluster)

On GPU, natural parallelization is one edge per CUDA thread

Threads (each edge assigned to 1 CUDA thread)

0	1	2	3	4	5	6	7	8	9	10	11

A	B	C	D	E	F	G	H

Flux field values (per vertex)

```
for (e <- edges(mesh)) {
    …
    Flux(v1) += dT*step
    Flux(v2) -= dT*step
    …
}
```

Different edges share a vertex: requires atomic update of per-vertex field data

GPU implementation: conflict graph

Threads (each edge assigned to 1 CUDA thread)

| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 |

| A | B | C | D | E | F | G | H |

Flux field values (per vertex)

Identify mesh edges with colliding writes (lines in graph indicate presence of collision)

Can simply run program once to get this information.
(results valid for subsequent executions provided mesh does not change)

GPU implementation: conflict graph

Threads (each edge assigned to 1 CUDA thread)

0	1	2	3	4	5	6	7	8	9	10	11

| A	B	C	D		E		F	G		H

Flux field values (per vertex)

"Color" nodes in graph such that no connected nodes have the same color

Can execute on GPU in parallel, without atomic operations, by running all nodes with the same color in a single CUDA launch.

Cluster performance of Lizst program

256 nodes, 8 cores per node (message-passing implemented using MPI)

Important: performance portability!
Same Lizst program also runs with high efficiency on GPU (results not shown here).
But uses a <u>different algorithm</u> when compiled to GPU! (graph coloring)

Liszt summary

- **Productivity:**
  - Abstract representation of mesh: vertices, edges, faces, fields (concepts that a scientist thinks about already!)
  - Intuitive topological operators

- **Portability**
  - Same code runs on large cluster of CPUs (MPI) and GPUs (and combinations thereof!)

- **High-performance**
  - Language is constrained to allow compiler to track dependencies
  - Used for locality-aware partitioning in distributed memory implementation
  - Used for graph coloring in GPU implementation
  - Compiler knows how to chooses different parallelization strategies for different platforms
  - Underlying mesh representation can be customized by system based on usage and platform (e.g, don't store edge pointers if code doesn't need it, choose struct of arrays vs. array of structs for per-vertex fields)

Example 2:
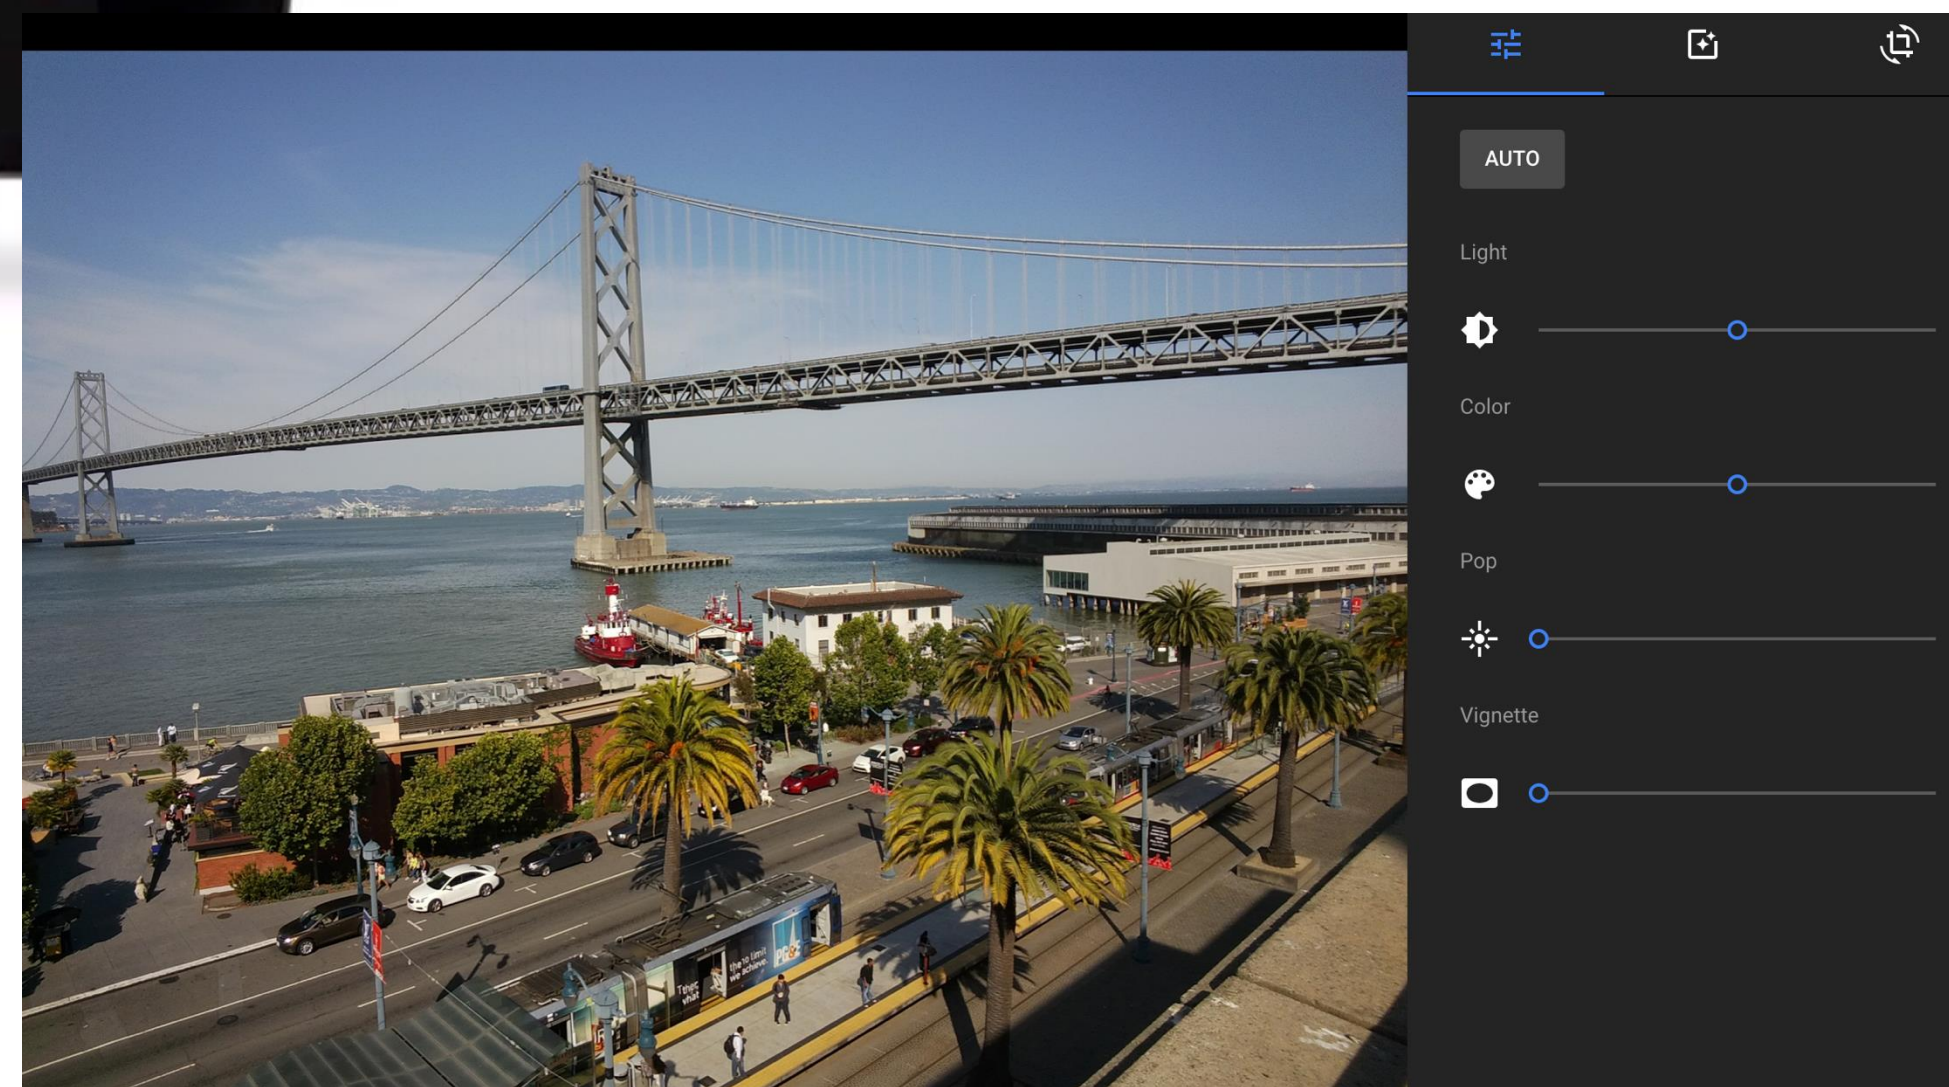# Halide: a domain-specific language for image processing

Jonathan Ragan-Kelley, Andrew Adams et al.
[SIGGRAPH 2012, PLDI 13]

Halide used in practice

- Halide used to implement Android HDR+ app

- Halide code used to process all images uploaded to Google Photos

A quick tutorial on high-performance image processing

What does this C code do?

```c
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9};


for (int j=0; j<HEIGHT; j++) {
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      for (int ii=0; ii<3; ii++)
        tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];
      output[j*WIDTH + i] = tmp;
  }
}
```

3x3 box blur

(Zoom view)

3x3 image blur

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9,
                   1.0/9, 1.0/9, 1.0/9};


for (int j=0; j<HEIGHT; j++) {
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      for (int ii=0; ii<3; ii++)
        tmp += input[(j+jj)*(WIDTH+2) + (i+ii)] * weights[jj*3 + ii];
    output[j*WIDTH + i] = tmp;
  }
}
```

Total work per image = 9 x WIDTH x HEIGHT

For NxN filter: N^2 x WIDTH x HEIGHT

Two-pass 3x3 blur

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];


float weights[] = {1.0/3, 1.0/3, 1.0/3};


for (int j=0; j<(HEIGHT+2); j++)
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int ii=0; ii<3; ii++)
      tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
    tmp_buf[j*WIDTH + i] = tmp;
  }


for (int j=0; j<HEIGHT; j++) {
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
    output[j*WIDTH + i] = tmp;
  }
}
```
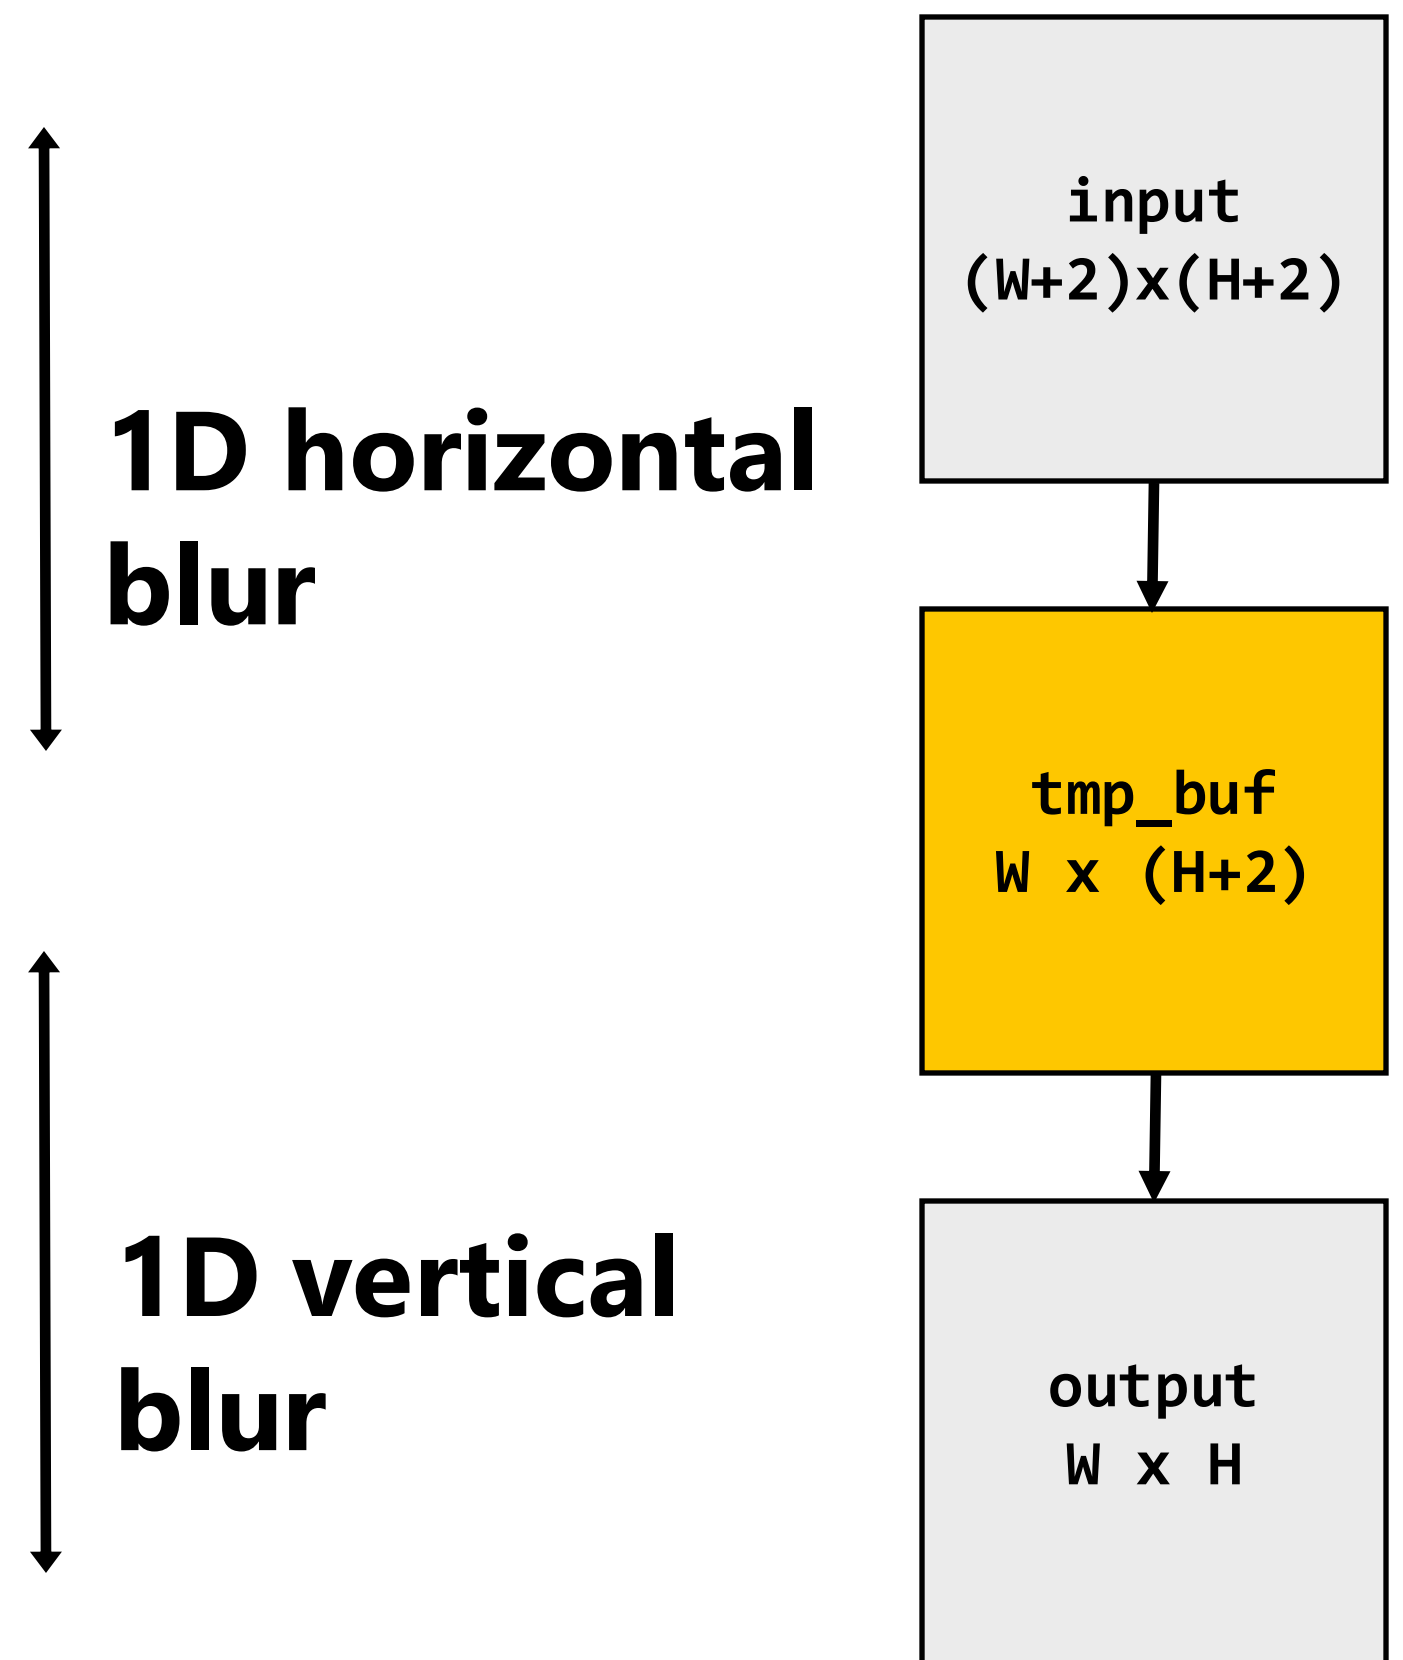
Total work per image = 6 x WIDTH x HEIGHT

For NxN filter: 2N x WIDTH x HEIGHT

WIDTH x HEIGHT extra storage
3X lower arithmetic intensity than 3D blur

1D horizontal blur

1D vertical blur

input
(W+2)x(H+2)

tmp_buf
W x (H+2)

output
W x H

Two-pass image blur: locality

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (HEIGHT+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<(HEIGHT+2); j++)
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int ii=0; ii<3; ii++)
      tmp += input[j*(WIDTH+2) + i+ii] * weights[ii];
    tmp_buf[j*WIDTH + i] = tmp;
  }

for (int j=0; j<HEIGHT; j++) {
  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      tmp += tmp_buf[(j+jj)*WIDTH + i] * weights[jj];
    output[j*WIDTH + i] = tmp;
  }
}
```

Intrinsic bandwidth requirements of algorithm:
Application must read each element of input image and must write each element of output image.

Data from `input` reused three times. (immediately reused in next two i-loop iterations after first load, never loaded again.)
- **Perfect cache behavior: never load required data more than once**
- **Perfect use of cache lines (don't load unnecessary data into cache)**

Two pass: loads/stores to `tmp_buf` are overhead (this memory traffic is an artifact of the two-pass implementation: it is not intrinsic to computation being performed)

Data from `tmp_buf` reused three times (but three rows of image data are accessed in between)
- **Never load required data more than once… if cache has capacity for <u>three rows of image</u>**
- **Perfect use of cache lines (don't load unnecessary data into cache)**

Two-pass image blur, "chunked" (version 1)

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * 3];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<HEIGHT; j++) {

  for (int j2=0; j2<3; j2++)
    for (int i=0; i<WIDTH; i++) {
      float tmp = 0.f;
      for (int ii=0; ii<3; ii++)
        tmp += input[(j+j2)*(WIDTH+2) + i+ii] * weights[ii];
      tmp_buf[j2*WIDTH + i] = tmp;

  for (int i=0; i<WIDTH; i++) {
    float tmp = 0.f;
    for (int jj=0; jj<3; jj++)
      tmp += tmp_buf[jj*WIDTH + i] * weights[jj];
    output[j*WIDTH + i] = tmp;
  }
}
```

Only 3 rows of intermediate buffer need to be allocated

input
(W+2)x(H+2)

tmp_buf (Wx3)

output
W x H

Produce 3 rows of tmp_buf (only what's needed for one row of output)

Combine them together to get one row of output

Total work per row of output:
 - **step 1: 3 x 3 x WIDTH work**
 - **step 2: 3 x WIDTH work**

Total work per image = 12 x WIDTH x HEIGHT   ????

Loads from tmp_buffer are cached (assuming tmp_buffer fits in cache)

Two-pass image blur, "chunked" (version 2)

```
int WIDTH = 1024;
int HEIGHT = 1024;
float input[(WIDTH+2) * (HEIGHT+2)];
float tmp_buf[WIDTH * (CHUNK_SIZE+2)];
float output[WIDTH * HEIGHT];

float weights[] = {1.0/3, 1.0/3, 1.0/3};

for (int j=0; j<HEIGHT; j+CHUNK_SIZE) {

  for (int j2=0; j2<CHUNK_SIZE+2; j2++)
    for (int i=0; i<WIDTH; i++) {
      float tmp = 0.f;
      for (int ii=0; ii<3; ii++)
        tmp += input[(j+j2)*(WIDTH+2) + i+ii] * weights[ii];
      tmp_buf[j2*WIDTH + i] = tmp;

  for (int j2=0; j2<CHUNK_SIZE; j2++)
    for (int i=0; i<WIDTH; i++) {
      float tmp = 0.f;
      for (int jj=0; jj<3; jj++)
        tmp += tmp_buf[(j2+jj)*WIDTH + i] * weights[jj];
      output[(j+j2)*WIDTH + i] = tmp;
    }
}
```
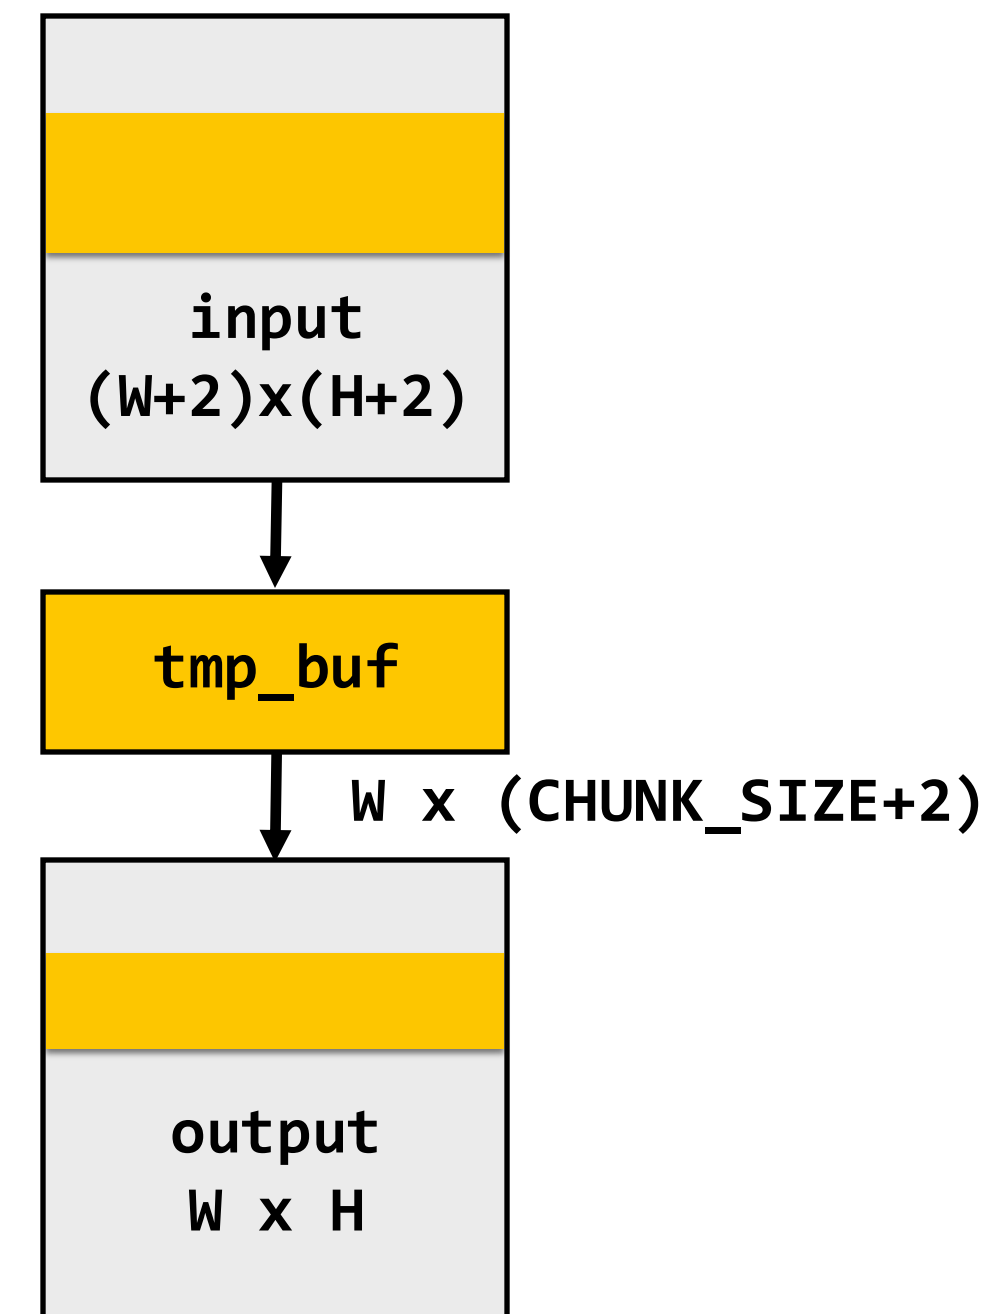
Sized to fit in cache (capture all producer-consumer locality)

Produce enough rows of tmp_buf to produce a CHUNK_SIZE number of rows of output

Produce CHUNK_SIZE rows of output

input
(W+2)x(H+2)

tmp_buf

W x (CHUNK_SIZE+2)

output
W x H

Total work per chunck of output:
(assume CHUNK_SIZE = 16)
- **Step 1: 18 x 3 x WIDTH work**
- **Step 2: 16 x 3 x WIDTH work**
Total work per image: (34/16) x 3 x WIDTH x HEIGHT
= 6.4 x WIDTH x HEIGHT

Trends to ideal 6 x WIDTH x HEIGHT as CHUNK_SIZE is increased!

Conflicting goals (once again…)

- **Want to be work efficient (perform fewer operations)**
- **Want to take advantage of locality when present**
  - **Otherwise work-efficient code will be bandwidth bound**
  - **Ideally: bandwidth cost of implementation is very close to intrinsic cost of algorithm: data is loaded from memory once and reused as much as needed prior to being discarded from processor's cache**
- **Want to execute in parallel (multi-core, SIMD within core)**

Optimized C++ code: 3x3 image blur

Good: 10x faster: on a quad-core CPU than original two-pass code

Bad: specific to SSE (not AVX2), CPU-code only, hard to tell what is going on at all!

```cpp
void fast_blur(const Image &in, Image &blurred) {
  __m128i one_third = _mm_set1_epi16(21846);
  #pragma omp parallel for
  for (int yTile = 0; yTile < in.height(); yTile += 32) {
    __m128i a, b, c, sum, avg;
    __m128i tmp[(256/8)*(32+2)];
    for (int xTile = 0; xTile < in.width(); xTile += 256) {
      __m128i *tmpPtr = tmp;
      for (int y = -1; y < 32+1; y++) {
        const uint16_t *inPtr = &(in(xTile, yTile+y));
        for (int x = 0; x < 256; x += 8) {
          a = _mm_loadu_si128((__m128i*)(inPtr-1));
          b = _mm_loadu_si128((__m128i*)(inPtr+1));
          c = _mm_load_si128((__m128i*)(inPtr));
          sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
          avg = _mm_mulhi_epi16(sum, one_third);
          _mm_store_si128(tmpPtr++, avg);
          inPtr += 8;
        }}
      tmpPtr = tmp;
      for (int y = 0; y < 32; y++) {
        __m128i *outPtr = (__m128i *)(&(blurred(xTile, yTile+y)));
        for (int x = 0; x < 256; x += 8) {
          a = _mm_load_si128(tmpPtr+(2*256)/8);
          b = _mm_load_si128(tmpPtr+256/8);
          c = _mm_load_si128(tmpPtr++);
          sum = _mm_add_epi16(_mm_add_epi16(a, b), c);
          avg = _mm_mulhi_epi16(sum, one_third);
          _mm_store_si128(outPtr++, avg);
}}}}}
```

Multi-core execution (partition image vertically)

Modified iteration order: 256x32 block-major iteration (to maximize cache hit rate)

use of SIMD vector intrinsics

two passes fused into one: tmp data read from cache

Halide blur (algorithm description)

```
// Halide 3x3 blur program definition
Func halide_blur(Func in) {

   Func blurx, out;
   Var  x, y;




   blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
   out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;
   return out;
}


// top-level calling code
Image<uint8_t> input = load_image("myimage.png");    // define input image
Func my_program = halide_blur(input);                 // define pipeline
Image<uint8_t> output = my_program.realize(input.width(), input.height(),
                                   input.channels()); // execute pipeline

output.save("myblurredimage.png");
```

Images are pure functions
Functions map integer coordinates (in up to a 4D domain) to values (e.g., colors of corresponding pixels)
(`in`, `blurx` and `out` are functions)

Algorithms are a series of functions (think: pipeline stages)

Value of `blurx` at coordinate (x,y) is given by expression accessing three values of `in`

NOTE: execution order and storage are unspecified by the abstraction. The implementation can evaluate, reevaluate, cache individual points as desired!

Think of a Halide program as a pipeline

```
// Halide 3x3 blur program definition
Func halide_blur(Func in) {

  Func blurx, out;
  Var  x, y;

  blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
  out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;
  return out;
}
```

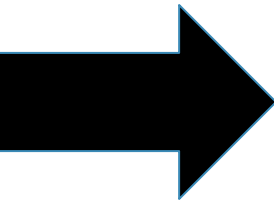# Halide schedule describes <u>how</u> to execute a pipeline

```
// Halide program definition
Func halide_blur(Func in) {

   Func blurx, out;
   Var  x, y, xi, yi

   // the "algorithm description"  (what to do)
   blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
   out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

   // "the schedule" (how to do it)
   out.tile(x, y, xi, yi, 256, 32).vectorize(xi,8).parallel(y);


   blurx.chunk(x).vectorize(x, 8);


   return out;
}
```

When evaluating `out`, use 2D tiling order (loops named by x, y, xi, yi). Use tile size 256 x 32.

Vectorize the xi loop (8-wide)

Use threads to parallelize the y loop

Produce only chunks of `blurx` at a time. Vectorize the x (innermost) loop

Halide schedule describes <u>how</u> to execute a pipeline

```
// Halide program definition
Func halide_blur(Func in) {

  Func blurx, out;
  Var  x, y, xi, yi

  // the "algorithm description"  (what to do)
  blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
  out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

  // "the schedule" (how to do it)
  out.tile(x, y, xi, yi, 256, 32).vectorize(xi,8).parallel(y);
  blurx.chunk(x).vectorize(x, 8);
  return out;
}
```

Given a schedule, Halide carries out mechanical process of implementing the specified schedule

```
void halide_blur(uint8_t* in, uint8_t* out) {
    #pragma omp parallel for
    for (int y=0; y<HEIGHT; y+=32) {      // tile loop
      for (int x=0; y<WIDTH; x+=256) {  // tile loop

        // buffer
        uint8_t* blurx[34 * 256];

        // produce intermediate buffer
        for (int yi=0; yi<34; yi++) {
          // SIMD vectorize this loop (not shown)
          for (int xi=0; xi<256; xi++) {
            blurx[yi*256+xi] =
                (in[(y+yi-1)*WIDTH+x+xi-1] +
                 in[(y+yi-1)*WIDTH+x+xi] +
                 in[(y+yi-1)*WIDTH+x+xi+1]) / 3.0;
          }
        }

        // consume intermediate buffer
        for (int yi=0; yi<32; yi++) {
          // SIMD vectorize this loop (not shown)
          for (int xi=0; xi<256; xi++) {
            out[(y+yi)*256+(x+xi)] =
                (blurx[yi*256+xi] +
                 blurx[(yi+1)*256+xi] +
                 blurx[(yi+2)*256+xi]) / 3.0;
          }
        }
      } // loop over tiles
    }  // loop over tiles
}
```

Halide: two domain-specific co-languages

- **Functional language** for **describing image processing operations**

- **Domain-specific language** for **describing schedules**

- **<u>Design principle</u>: separate "algorithm specification" from its schedule**

    - Programmer's responsibility: provide a high-performance schedule

    - Compiler's responsibility: carry out mechanical process of generating threads, SIMD instructions, managing buffers, etc.

    - <u>Result</u>: **enable programmer to rapidly explore space of schedules**

        - (e.g., "tile these loops", "vectorize this loop", "parallelize this loop across cores")

- **Domain scope:**

    - All computation on regular N-D coordinate spaces

    - Only feed-forward pipelines (includes special support for reductions and fixed recursion depth)

   - All dependencies inferable by compiler

Producer/consumer scheduling primitives

**Four basic scheduling primitives shown below**

in   tmp   blurred

**breadth first:** each function is entirely evaluated before the next one.

"Root"

in   tmp   blurred

total fusion: values are computed on the fly each time that they are needed.

"Inline"

in   tmp   blurred

sliding window: values are computed when needed then stored until not useful anymore.

"Sliding Window"

in   tmp   blurred

tiles: overlapping regions are processed in parallel, functions are evaluated one after another.

"Chunked"

Producer/consumer scheduling primitives

```
// Halide program definition
Func halide_blur(Func in) {

  Func blurx, out;
  Var  x, y, xi, yi

  // the "algorithm description"  (what to do)
  blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
  out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

  // "the schedule" (how to do it)
  blurx.compute_at(ROOT);
  return out;

}
```

**"Root":
compute all points of the
producer, then run consumer
(minimal locality)**

```
void halide_blur(uint8_t* in, uint8_t* out) {
    uint8_t blurx[WIDTH * HEIGHT];

    for (int y=0; y<HEIGHT; y++) {
      for (int x=0; y<WIDTH; x++) {
          blurx[] = ...


    for (int y=0; y<HEIGHT; y++) {
      for (int x=0; y<WIDTH; x++) {
          out[] = ...
}
```

```
// Halide program definition
Func halide_blur(Func in) {

  Func blurx, out;
  Var  x, y, xi, yi

  // the "algorithm description"  (what to do)
  blurx(x,y) = (in(x-1, y) + in(x,y) + in(x+1,y)) / 3.0f;
  out(x,y)   = (blurx(x,y-1) + blurx(x,y) + blurx(x,y+1)) / 3.0f;

  // "the schedule" (how to do it)
  blurx.inline();
  return out;

}
```

**"Inline":
revaluate producer at every
use site in consumer
(maximal locality)**

```
void halide_blur(uint8_t* in, uint8_t* out) {
    for (int y=0; y<HEIGHT; y++) {
      for (int x=0; y<WIDTH; x++) {
          out[] = (((in[(y-1)*WIDTH+x-1] +
                     in[(y-1)*WIDTH+x] +
                     in[(y-1)*WIDTH+x+1]) / 3) +
                   ((in[y*WIDTH+x-1] +
                     in[y*WIDTH+x] +
                     in[y*WIDTH+x+1]) / 3) +
                   ((in[(y+1)*WIDTH+x-1] +
                     in[(y+1)*WIDTH+x] +
                     in[(y+1)*WIDTH+x+1]) / 3));
}
```

Domain iteration primitives

1	2	3	4	5	6
7	8	9	10	11	12
13	14	15	16	17	18
19	20	21	22	23	24
25	26	27	28	29	30
31	32	33	34	35	36

serial y, serial x

1	7	13	19	25	31
2	8	14	20	26	32
3	9	15	21	27	33
4	10	16	22	28	34
5	11	17	23	29	35
6	12	18	24	30	36

serial x, serial y

Specify both order and how to parallelize (multi-thread, SIMD vector)

1	2
3	4
5	6
7	8
9	10
11	12

serial y
vectorized x

1	2
1	2
1	2
1	2
1	2
1	2

parallel y
vectorized x

1	2	5	6	9	10
3	4	7	8	11	12
13	14	17	18	21	22
15	16	19	20	23	24
25	26	29	30	33	34
27	28	31	32	35	36

split x into $2x_o+x_i$,
split y into $2y_o+y_i$,
serial y_o, x_o, y_i, x_i

2D blocked iteration order

Example Halide results

- **Camera RAW processing pipeline**

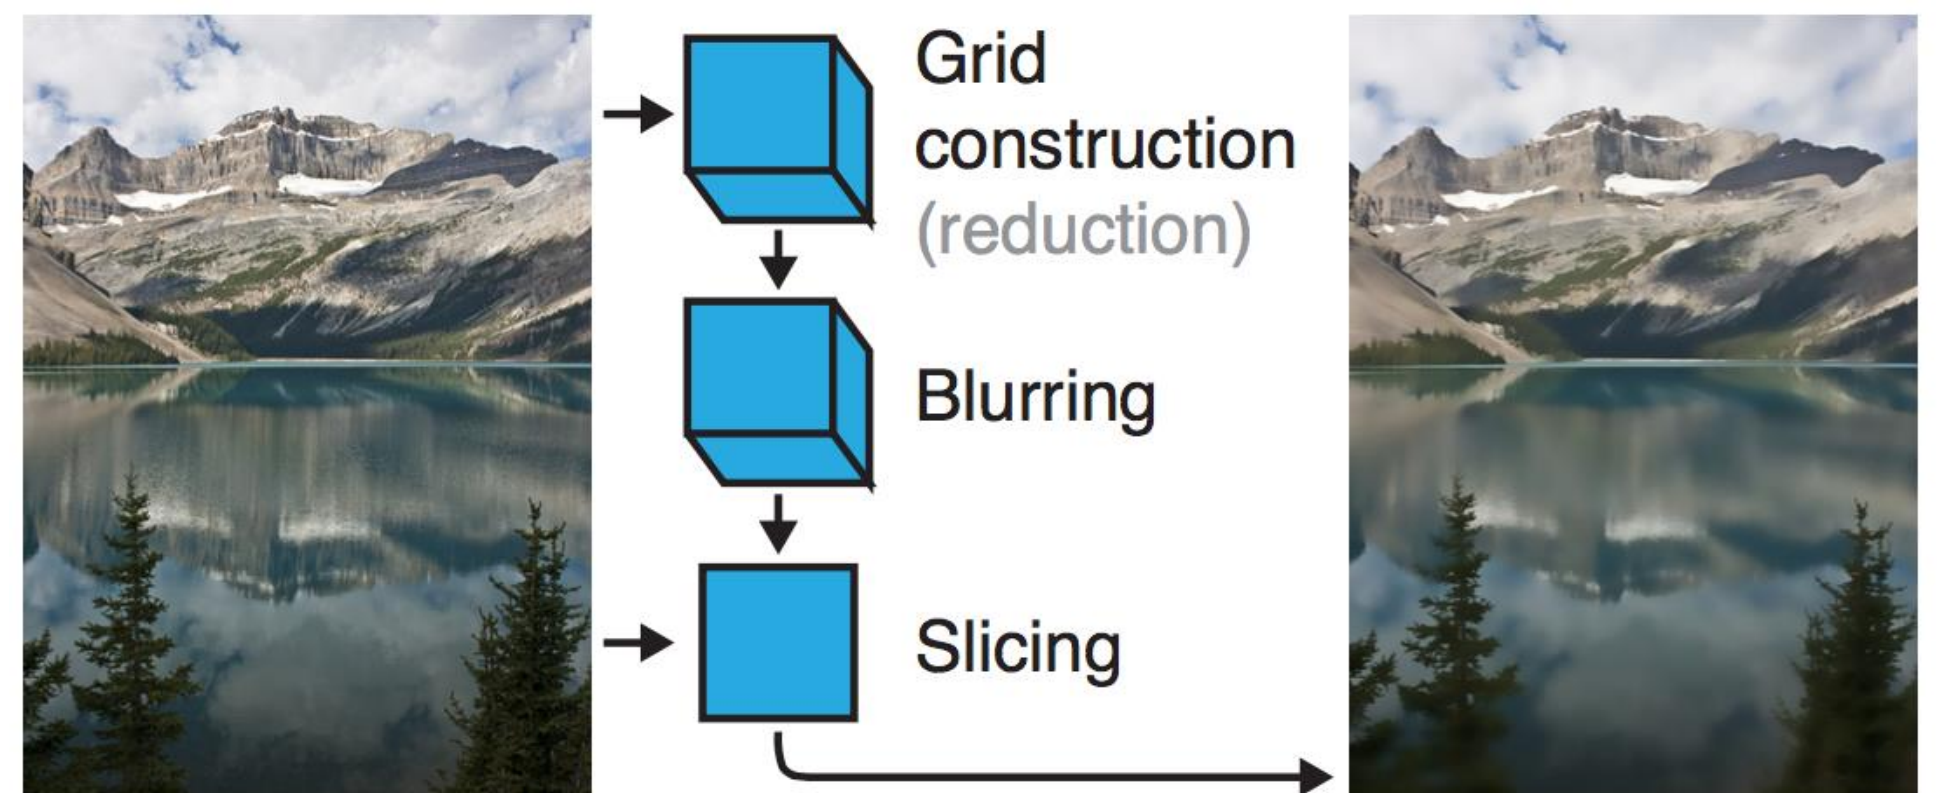 (Convert RAW sensor data to RGB image)

  - **Original: 463 lines of hand-tuned ARM NEON assembly**
  - **Halide: 2.75x less code, 5% faster**

- **Bilateral filter**

  (Common image filtering operation used in many applications)

  - **Original 122 lines of C++**
  - **Halide: 34 lines algorithm + 6 lines schedule**
    - **CPU implementation: 5.9x faster**
    - **GPU implementation: 2x faster than hand-written CUDA**

Stepping back: what is Halide?

- **Halide is a DSL for helping good developers optimize image processing code more rapidly**
  - Halide doesn't decide how to optimize a program for a novice programmer
  - Halide provides primitives for a programmer (that has strong knowledge of code optimization, such as a 418 student) to rapidly express what optimizations the system should apply
  - Halide carries out the nitty-gritty of mapping that strategy to a machine

Automatically generating Halide schedules

[Mullapudi, CMU 2016]

Extend Halide compiler to automatically generate schedule for programmer
- Compiler input: Halide program + size of expected input/output images
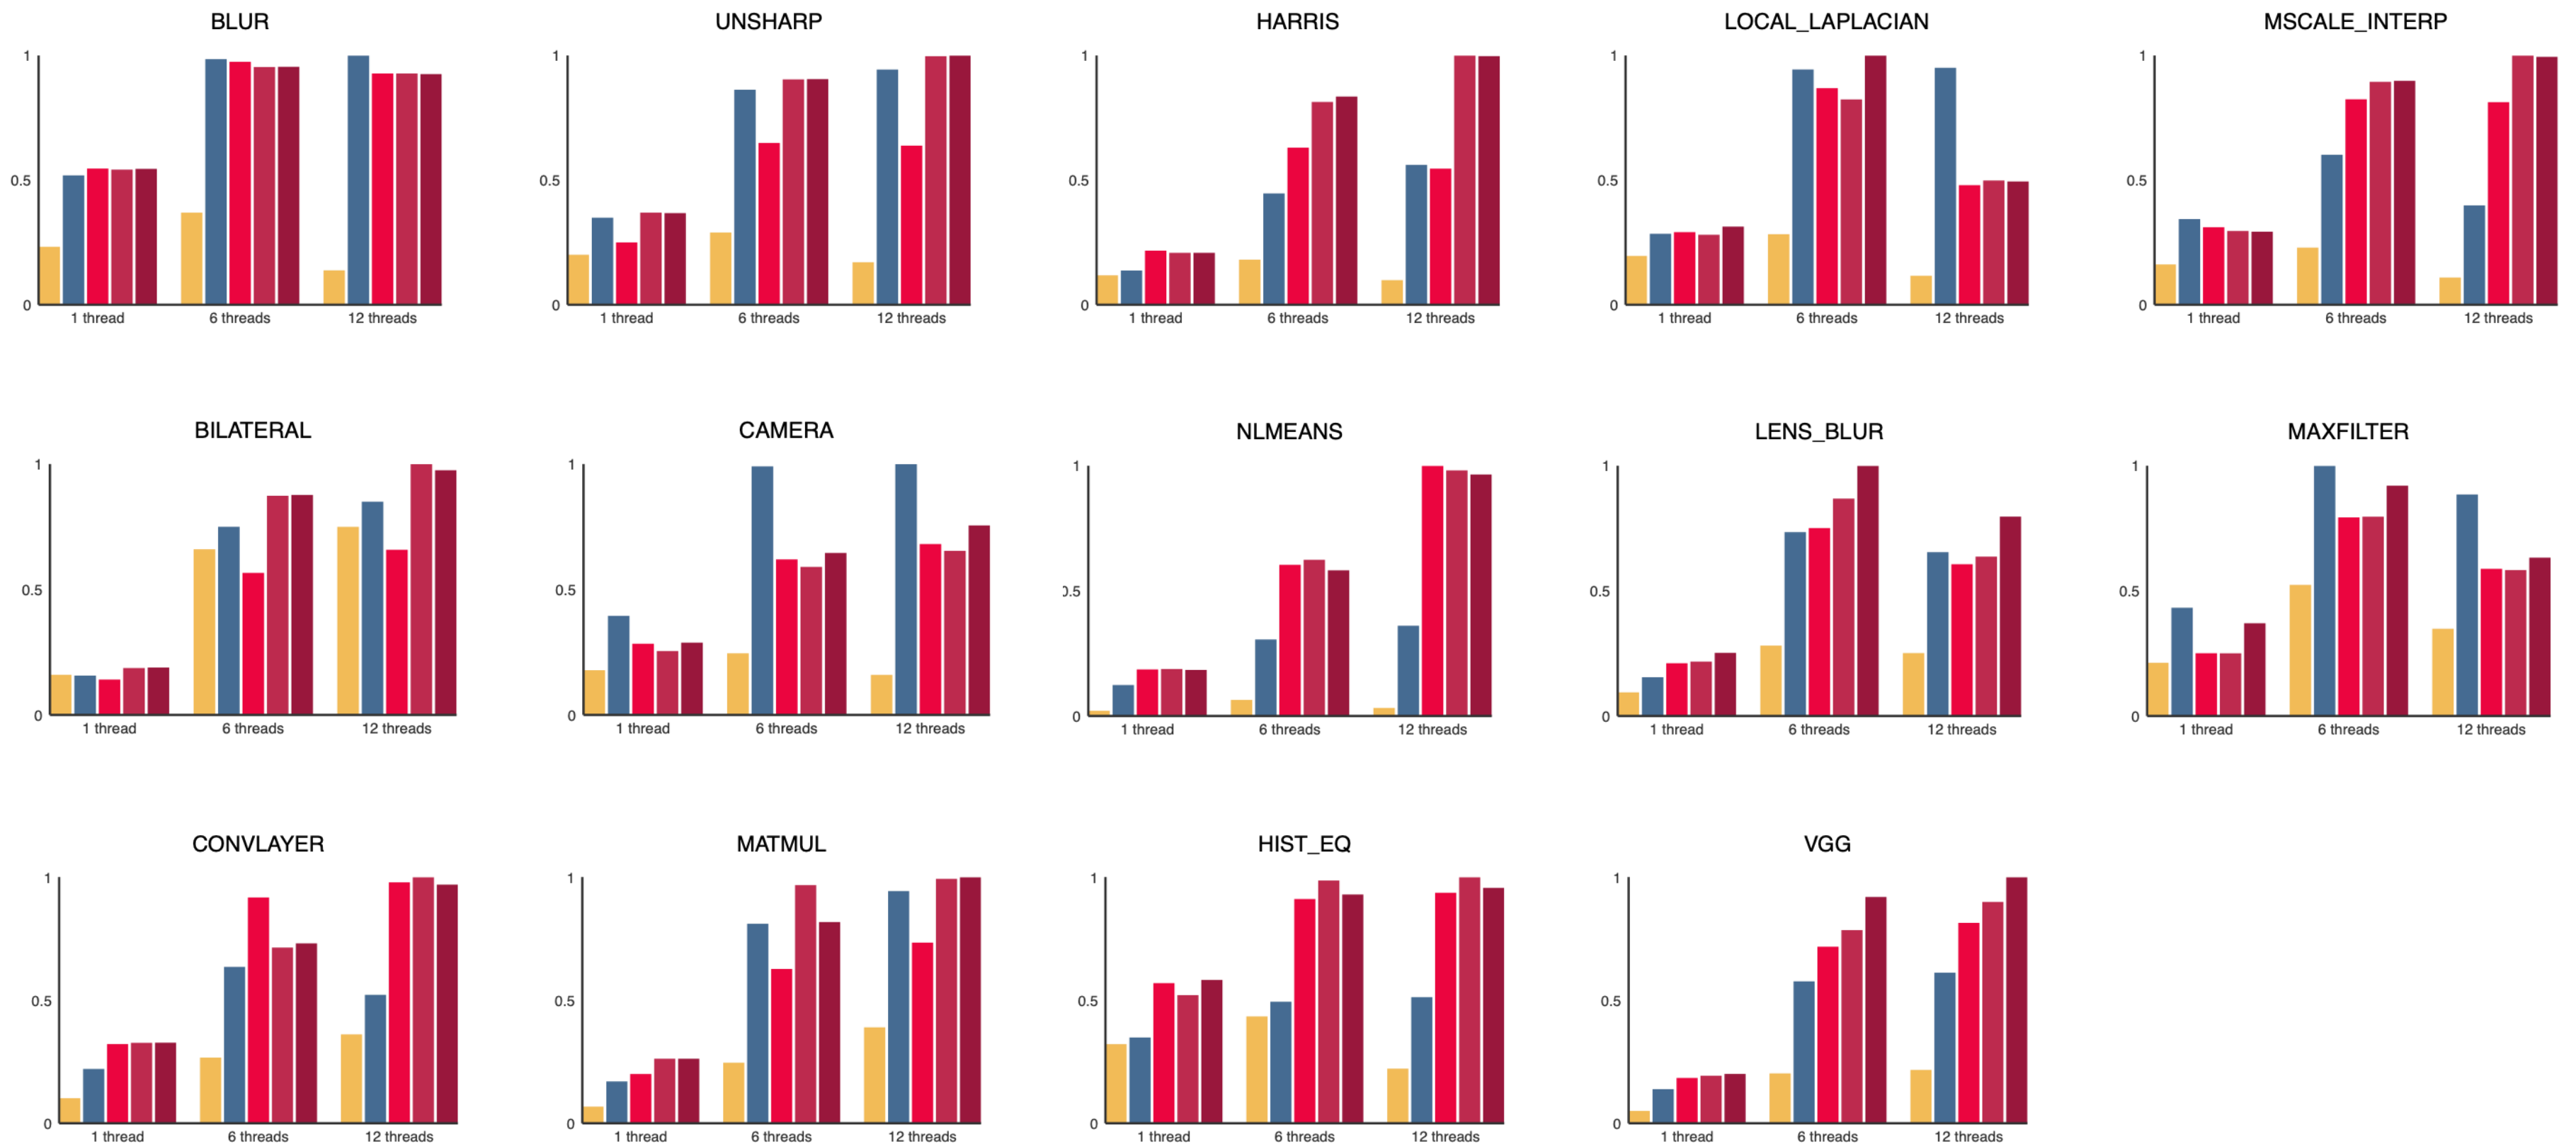
■ = Naive schedule

■ = Expert manual schedule
  (best human-created schedule)

■ = Automatically generated schedule (no autotuning, ~ seconds)

■ = Automatically generated, with auto-tuning (~ 10 minutes)

■ = Automatically generated, auto-tuning over 3 days   CMU 15-418/618, Fall 2025

"Racing" top Halide programmers

Halide auto-scheduler produced schedules that were better than those of expert Google Halide programmers in two of three cases (it got beat in one!)

Optimization of Manually Authored Schedules

LENSBLUR

MAXFILTER

NL MEANS

Throughput (1/ms)

Schedule development time (minutes)

■ = Auto-scheduler   ■ = Programmer 1   ■ = Programmer 2

Darkroom/Rigel

- **Directly synthesize FGPA implementation of image processing pipeline from a high-level description (a constrained "Halide-like" language)**


```
bx = im(x,y)
  (I(x-1,y) +
   I(x,y) +
   I(x+1,y))/3
end
by = im(x,y)
  (bx(x,y-1) +
   bx(x,y) +
   bx(x,y+1))/3
end
sharpened = im(x,y)
  I(x,y) + 0.1*
  (I(x,y) - by(x,y))
end
```
Stencil Language

- **Goal: ultra high efficiency image processing**

Many other recent domain-specific programming systems

Less domain specific than examples given today, but still designed specifically for:
data-parallel computations on big data for distributed systems ("Map-Reduce")

DSL for graph-based machine learning computations
Also see Green-Marl, Ligra (DSLs for describing operations on graphs)

Model-view-controller paradigm for web-applications

Ongoing efforts in many domains...

Simit: a language for physical simulation [MIT]

Domain-specific programming system development

- **Can develop DSL as a stand-alone language**
  - Graphics shading languages
  - MATLAB, SQL
- **"Embed" DSL in an existing generic language**
  - e.g., C++ library (GraphLab, OpenGL host-side API, Map-Reduce)
  - Lizst syntax above was all valid Scala code
- **Active research idea:**
  - Design generic languages that have facilities that assist rapid embedding of new domain-specific languages
  - "What is a good language for rapidly making new DSLs?"

Summary

- **Modern machines: parallel and heterogeneous**
  - Only way to increase compute capability in energy-constrained world
- **Most software uses small fraction of peak capability of machine**
  - Very challenging to tune programs to these machines
  - Tuning efforts are not portable across machines
- **Domain-specific programming environments trade-off generality to achieve productivity, performance, and portability**
  - Case studies today: Liszt, Halide
  - <u>Common trait</u>: languages provide abstractions that make dependencies known
  - Understanding dependencies is necessary but not sufficient: need domain restrictions and domain knowledge for system to synthesize efficient implementations