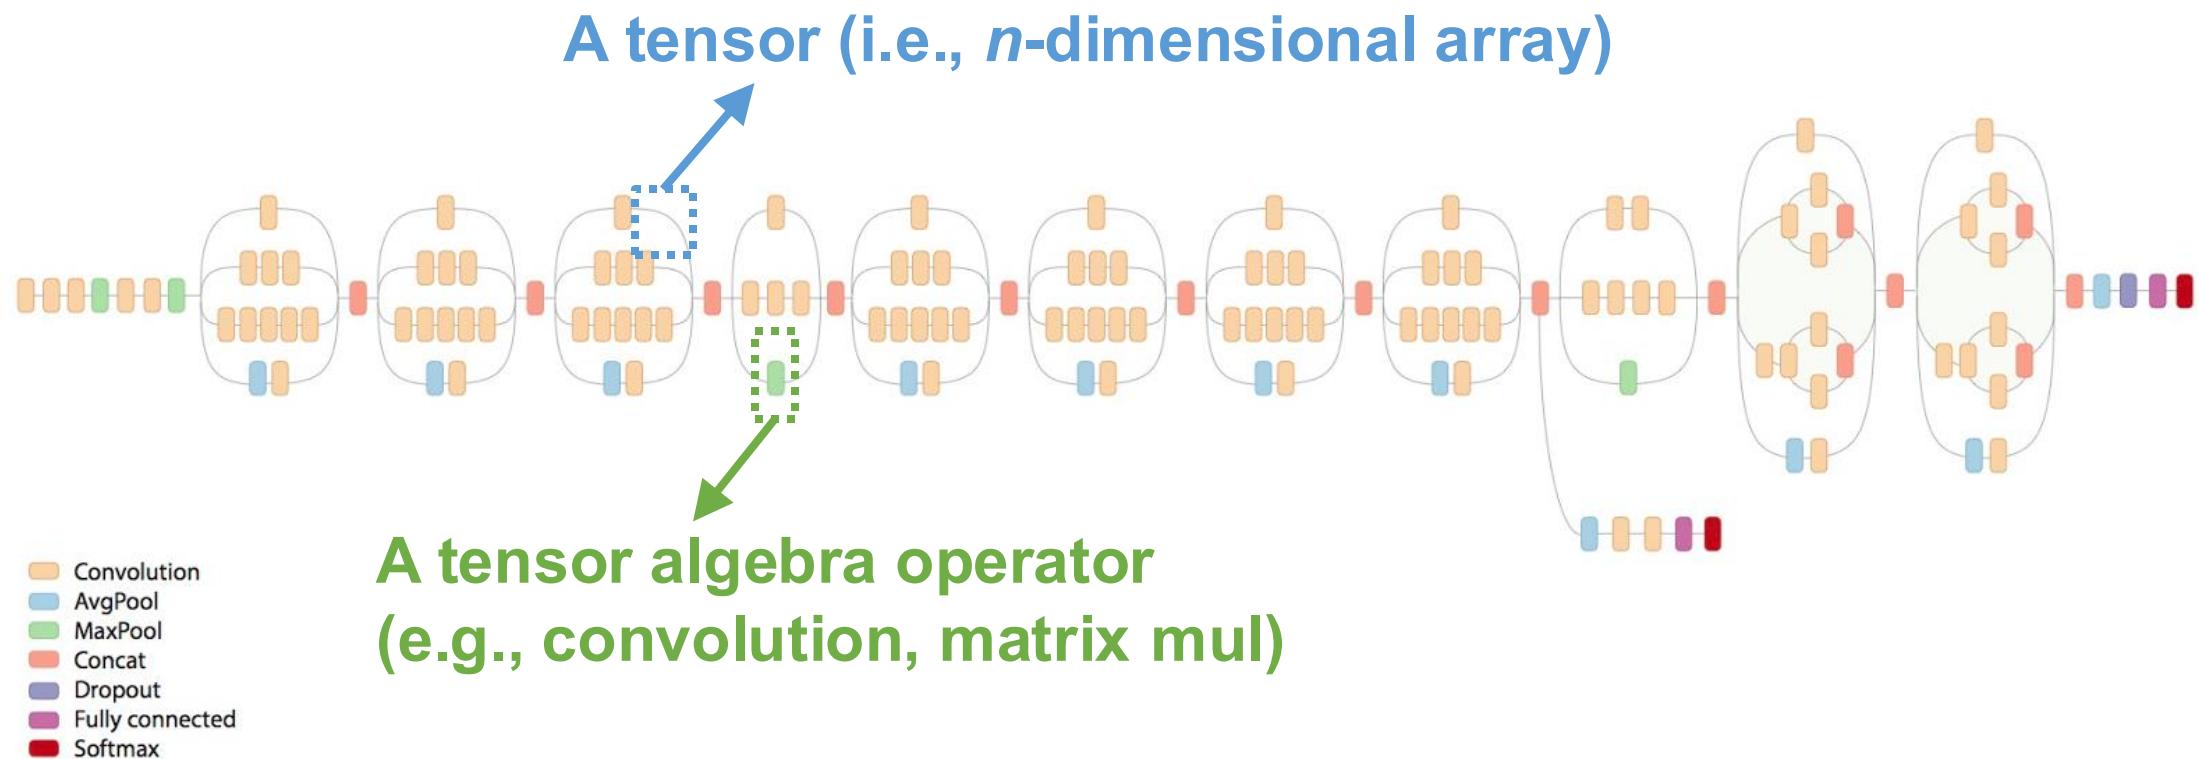


# **Lecture 24: Parallel Deep Learning (Data Parallelism)**

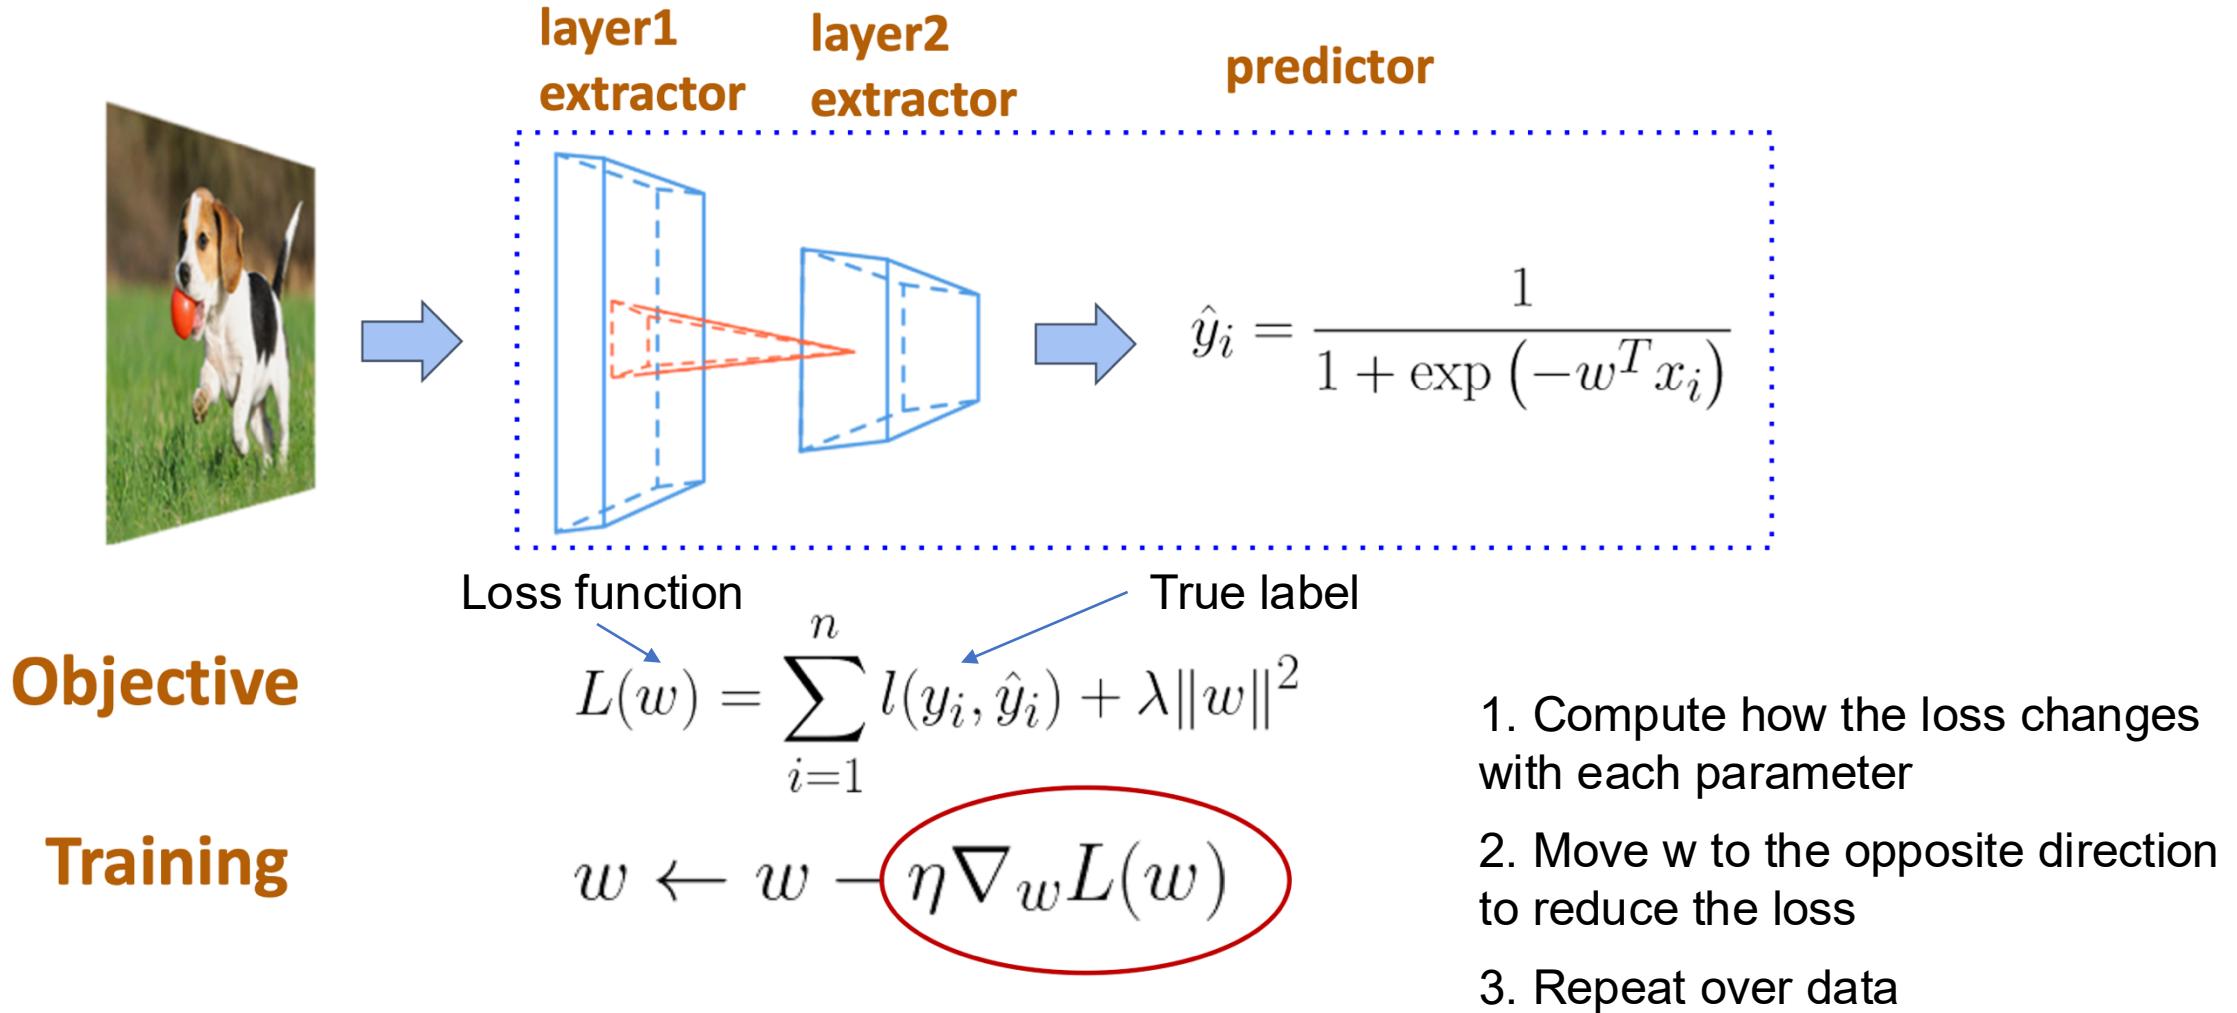
**Parallel Computer Architecture and Programming  
CMU 15-418/15-618, Fall 2025**

# Recap: Deep Neural Network

- Collection of simple trainable mathematical units that work together to solve complicated tasks



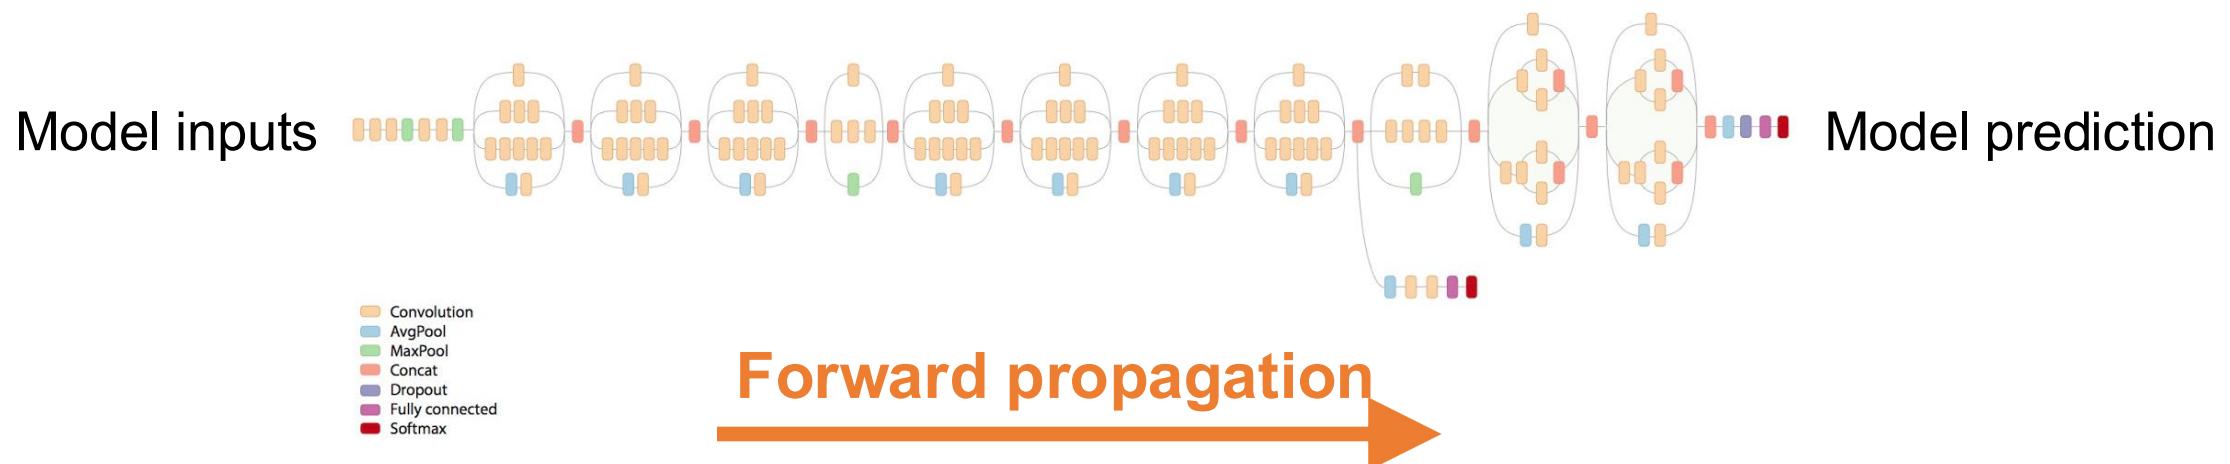
# Recap: DNN Training Overview



# DNN Training Process

Train ML models through many iterations of 3 stages

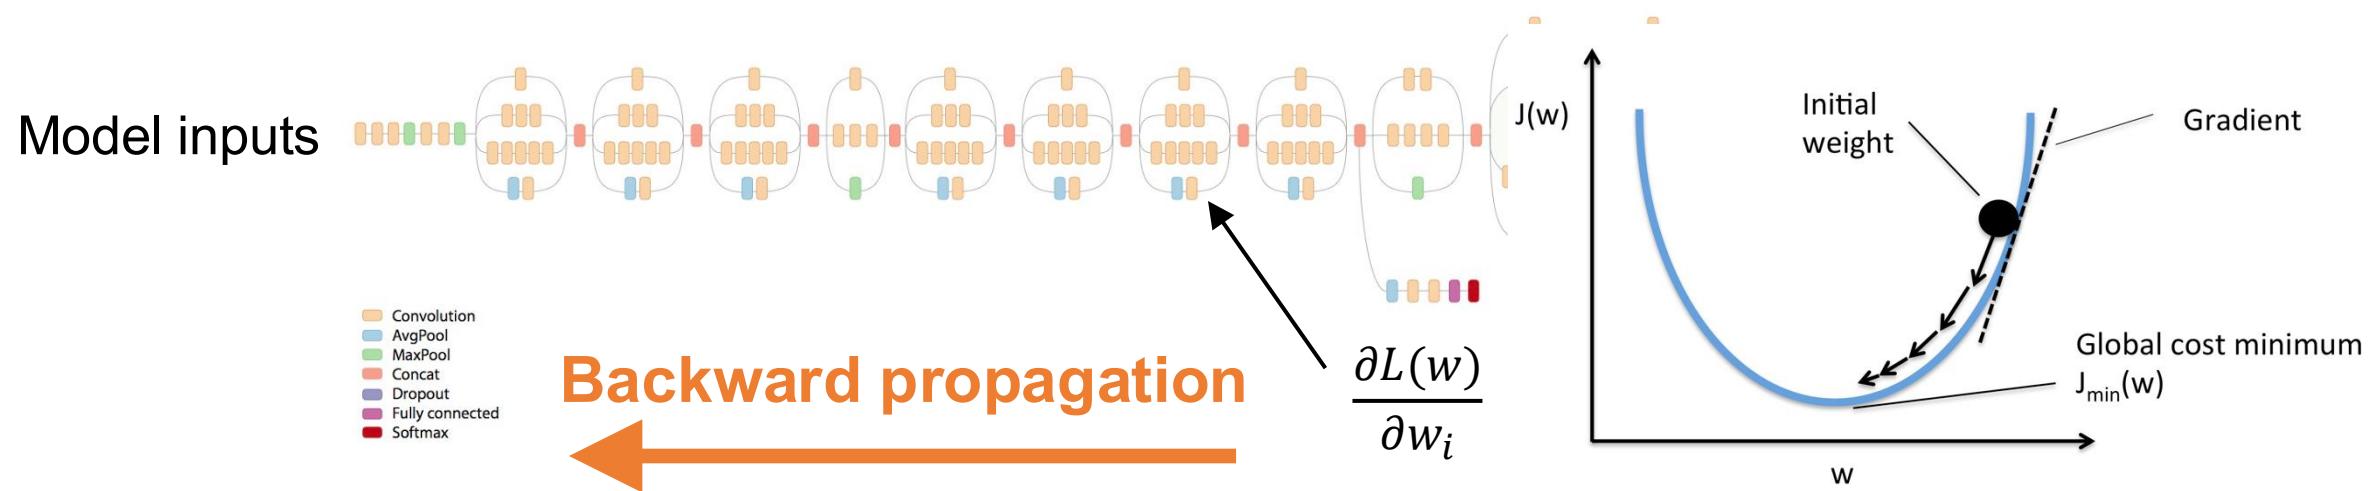
1. **Forward propagation**: apply model to a batch of input samples and run calculations through operators to produce a prediction
2. **Backward propagation**: run the model in reverse to produce error for each trainable weight
3. **Weight update**: use the loss value to update model weights



# DNN Training Process

Train ML models through many iterations of 3 stages

1. **Forward propagation**: apply model to a batch of input samples and run calculations through operators to produce a prediction
2. **Backward propagation**: run the model in reverse to produce a gradient for each trainable weight
3. **Weight update**: use the loss value to update model weights



# DNN Training Process

Train ML models through many iterations of 3 stages

1. **Forward propagation**: apply model to a batch of input samples and run calculations through operators to produce a prediction
2. **Backward propagation**: run the model in reverse to produce a gradient for each trainable weight
3. **Weight update**: use the gradients to update model weights

$$w_i := w_i - \gamma \frac{\partial L(w)}{\partial w_i} = w_i - \frac{\gamma}{n} \sum_{j=1}^n \boxed{\frac{\partial l_j(w)}{\partial w_i}}$$

Gradients of individual samples

# How can we parallelize DNN training?

The gradient of the loss function wrt the parameters

$$w_i := w_i - \gamma \nabla L(w_i) = w_i - \frac{\gamma}{n} \sum_{j=1}^n \nabla L_j(w_i)$$

Parameter vector → weights at iter i

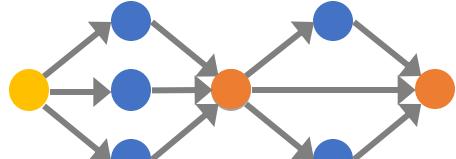
Learning rate → step size

Batch Size

Loss function for  $w_i$

```
graph TD; A[Parameter vector → weights at iter i] --> w_i[w_i]; B[Learning rate → step size] --> grad["γ ∇L(w_i)"]; C[Batch Size, Loss function for w_i] --> sum["n ∑ j=1 ∇L_j(w_i)"];
```

# Data Parallelism



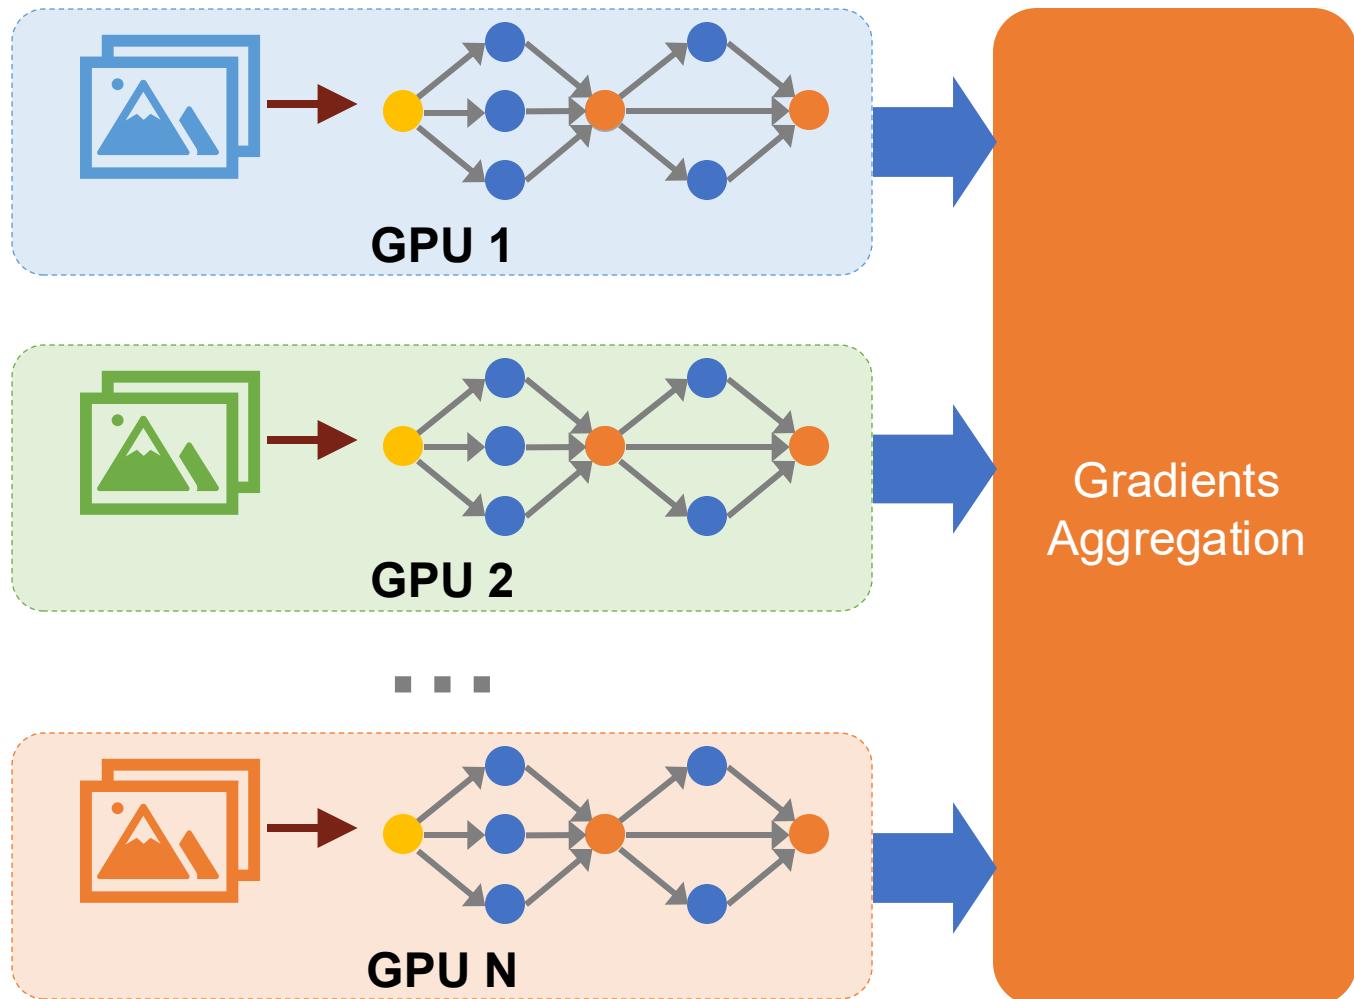
ML Model



Training Dataset

$$w_i := w_i - \gamma \nabla L(w_i) = w_i - \frac{\gamma}{n} \sum_{j=1}^n \nabla L_j(w_i)$$

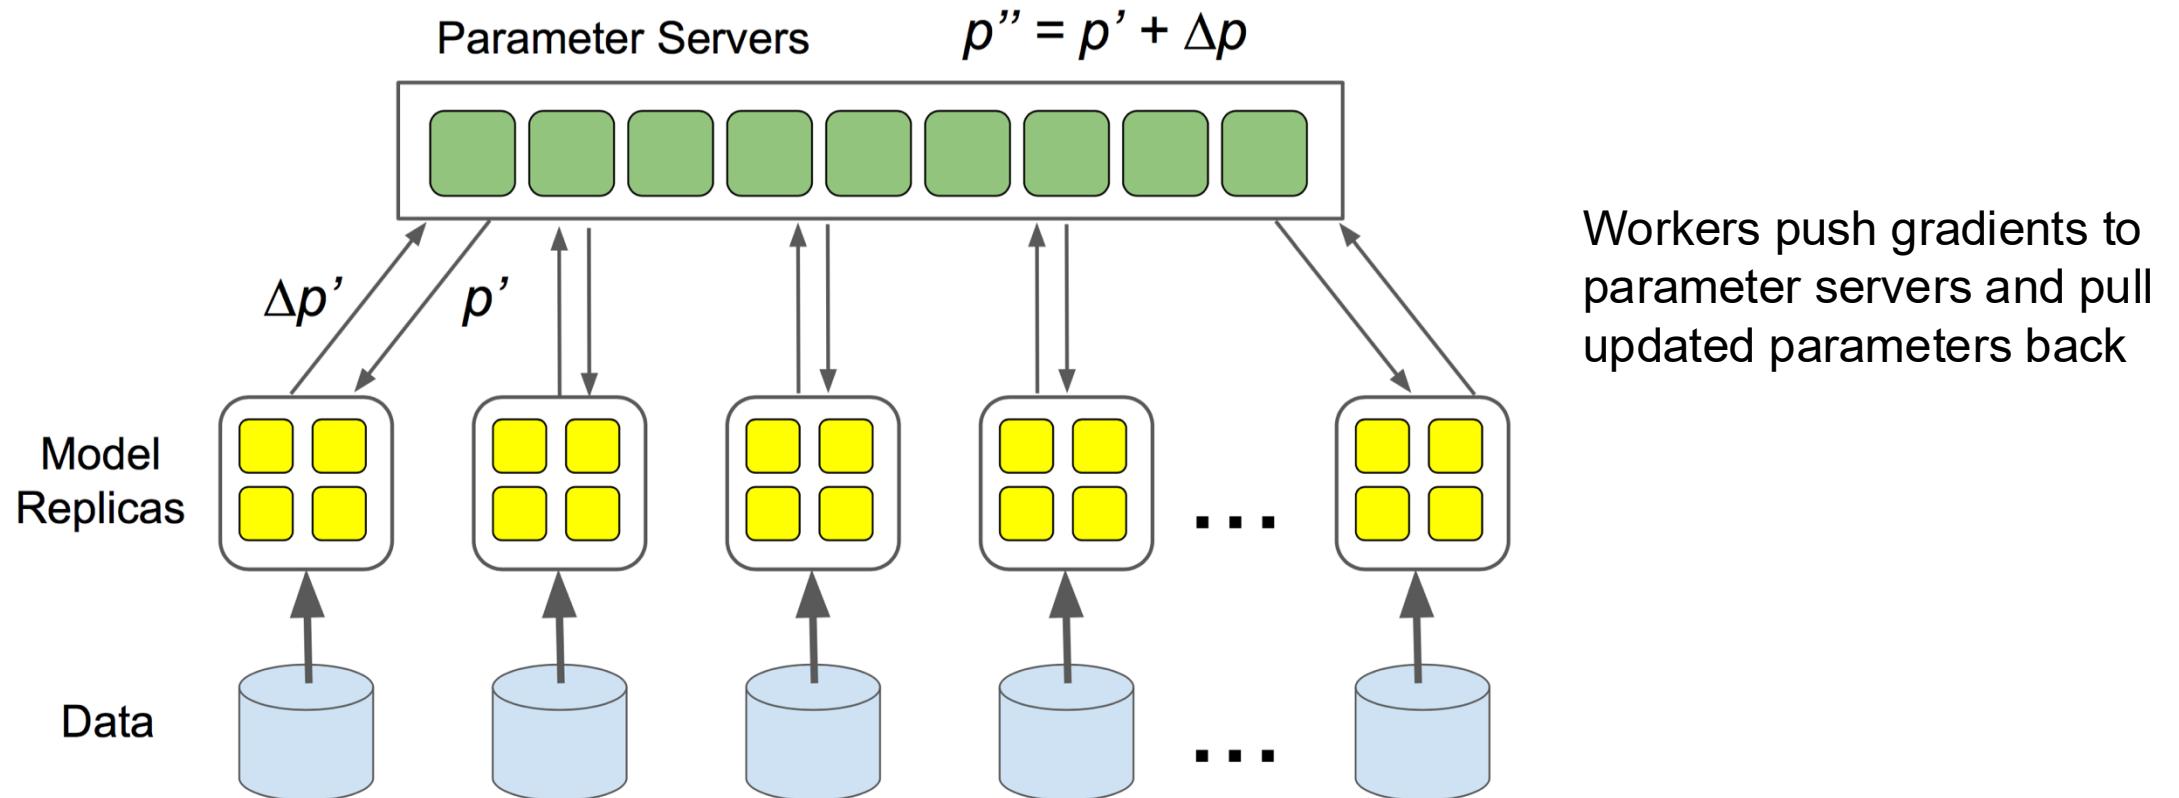
1. Partition training data into batches



2. Compute the gradients of each batch on a GPU

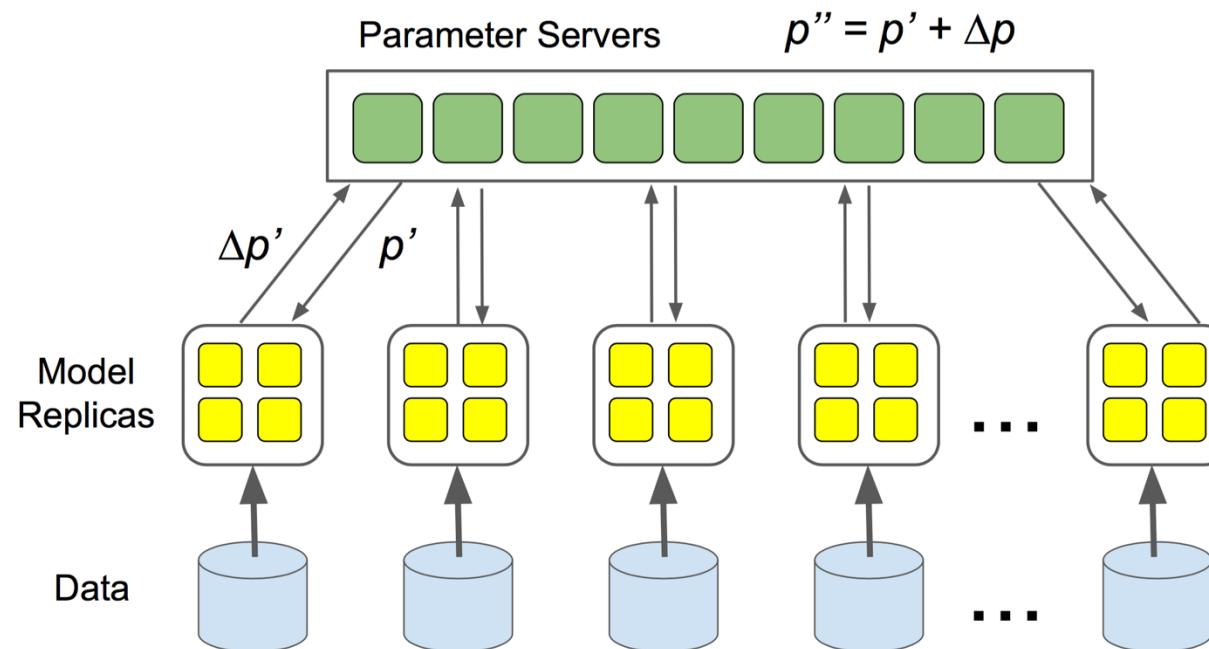
3. Aggregate gradients across GPUs

# Data Parallelism: Parameter Server



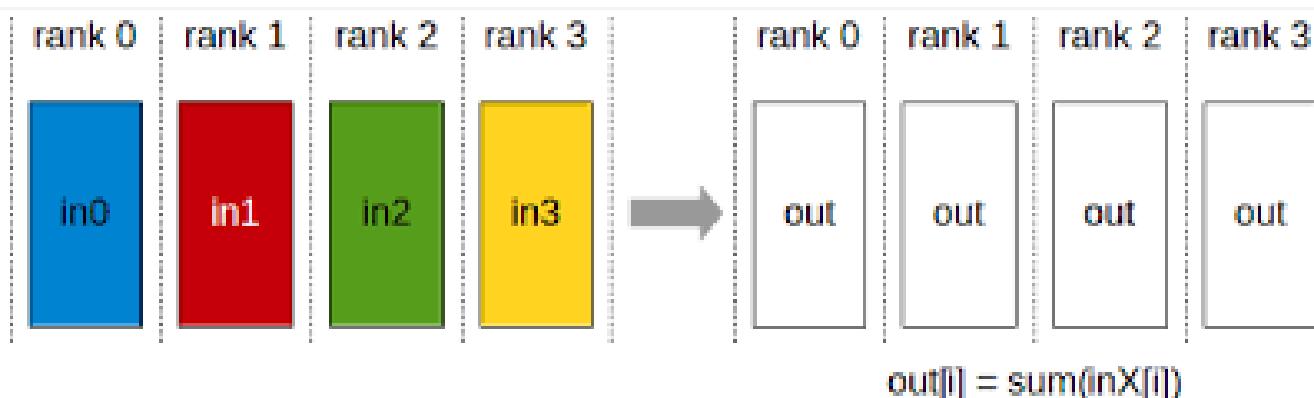
# Inefficiency of Parameter Server

- **Centralized communication:** all workers communicate with parameter servers for weights update; cannot scale to large numbers of workers
- How can we decentralize communication in DNN training?



# Inefficiency of Parameter Server

- **Centralized communication**: all workers communicate with parameter servers for weights update; cannot scale to large numbers of workers
- How can we decentralize communication in DNN training?
- **AllReduce**: perform element-wise reduction across multiple devices

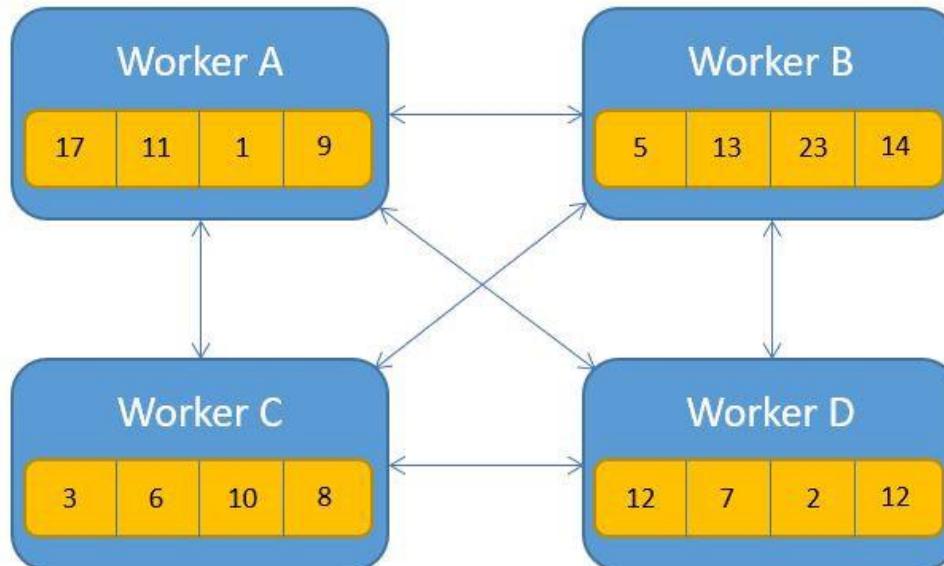


# Different Ways to Perform AllReduce

- Naïve AllReduce
- Ring AllReduce
- Tree AllReduce
- Butterfly AllReduce

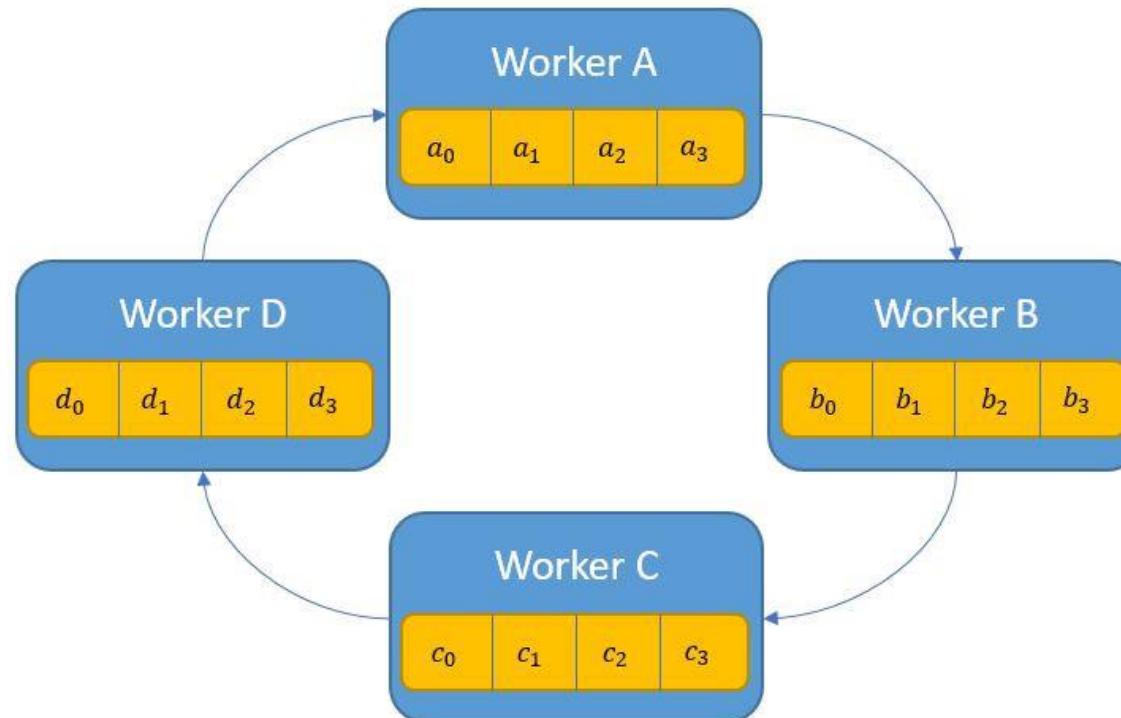
# Naïve AllReduce

- Each worker can send its local gradients to all other workers
- If we have N workers and each worker contains M parameters
- Overall communication:  $N * (N-1) * M$  parameters
- **Issue:** each worker communicates with all other workers; have the same scalability issue as parameter server



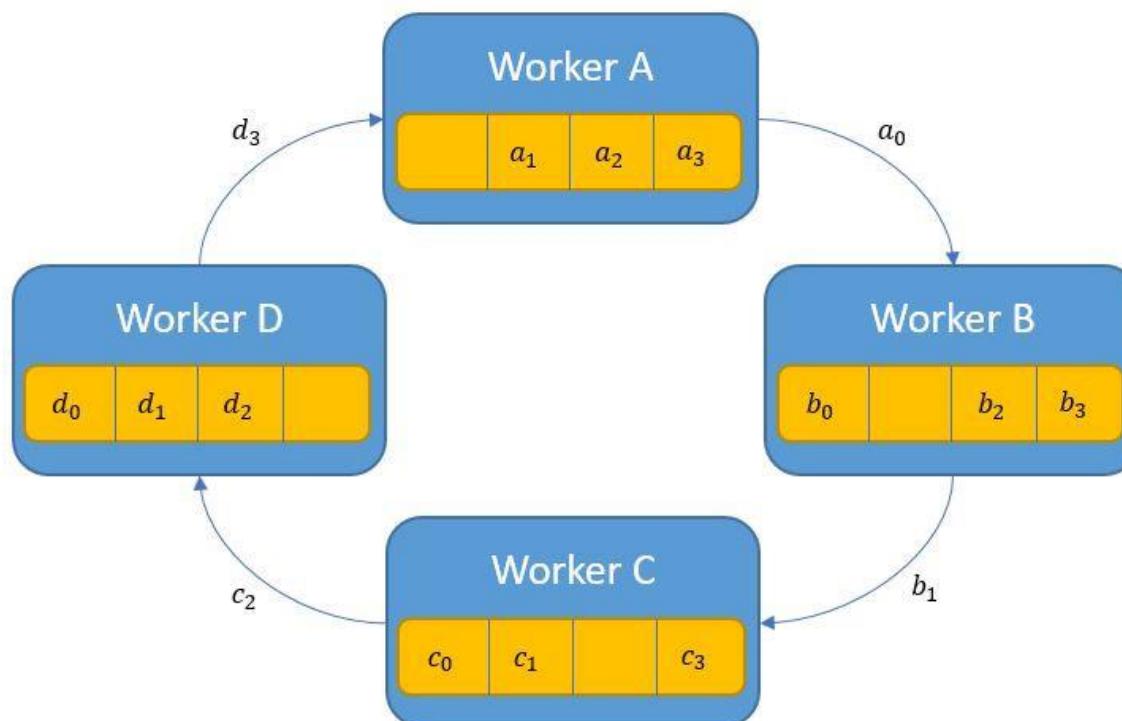
# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat N times



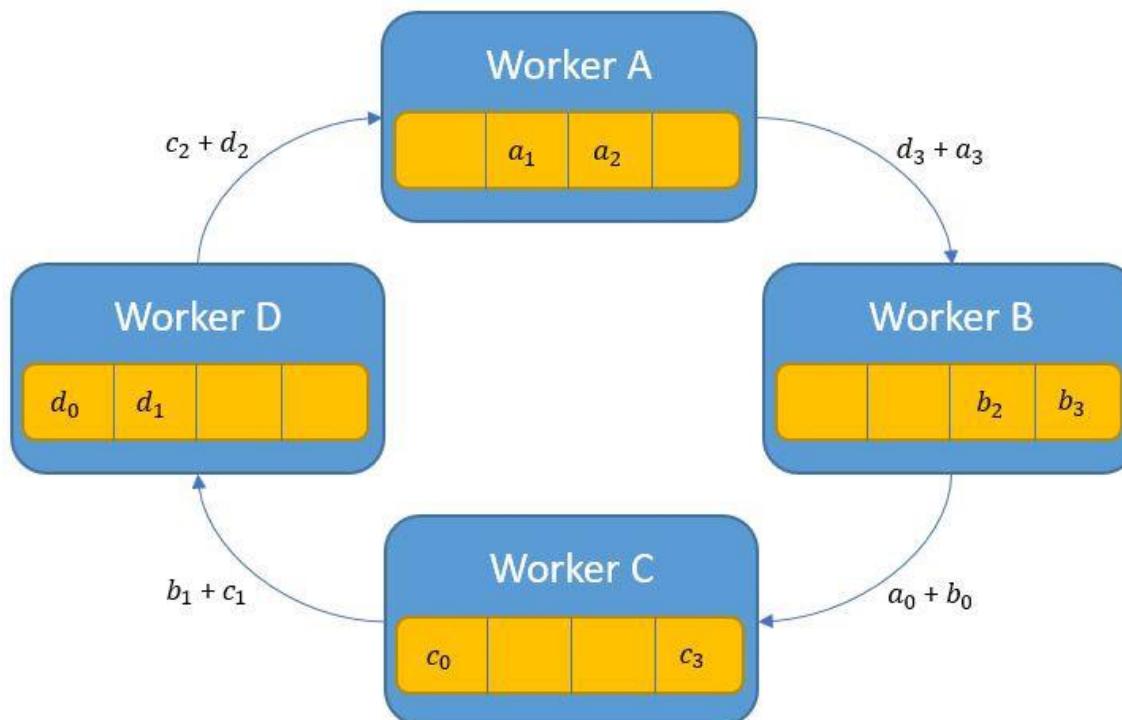
# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat N times



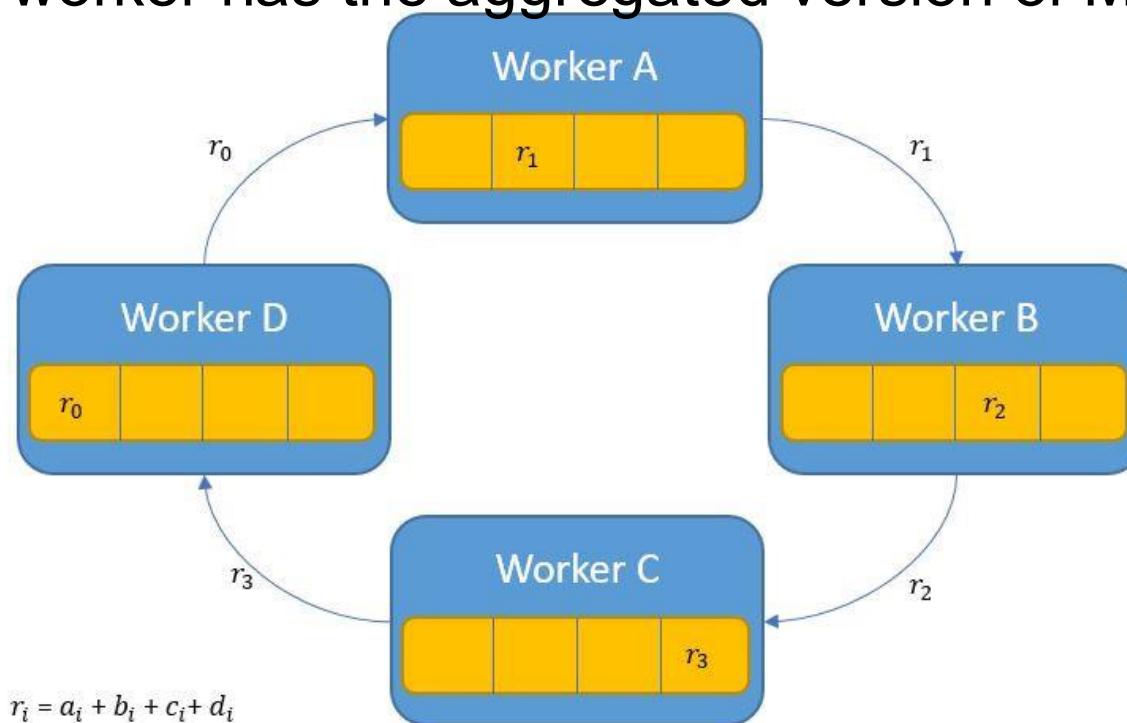
# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat N times



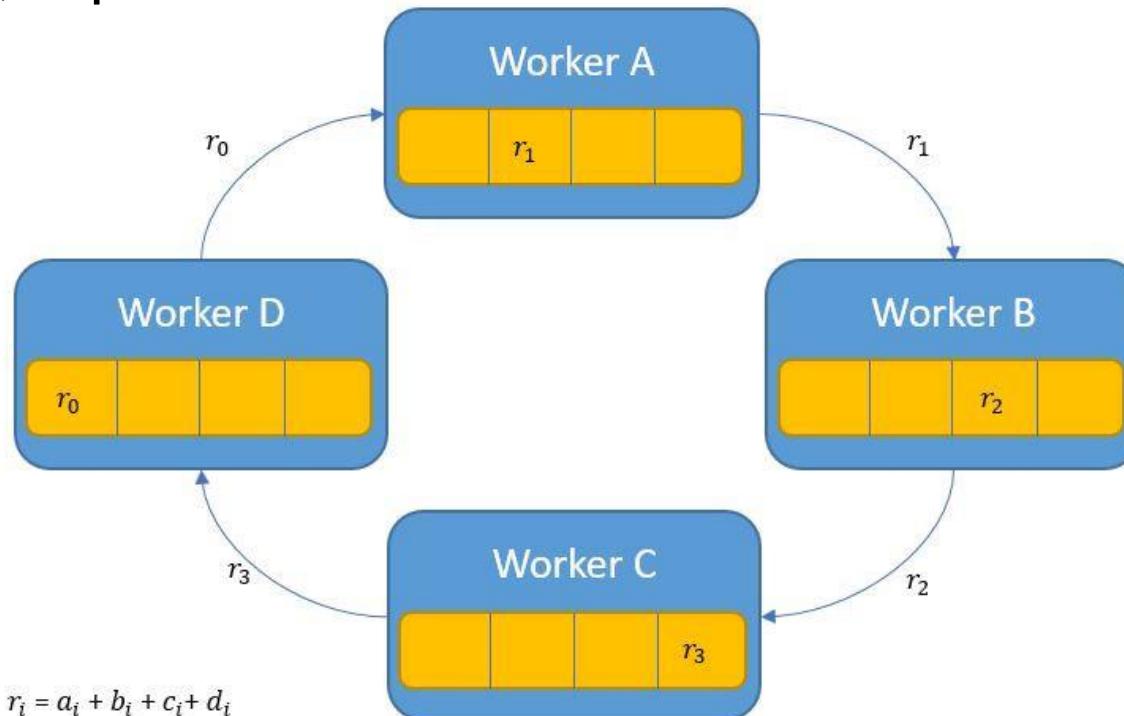
# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat  $N$  times
- After step 1, each worker has the aggregated version of  $M/N$  parameters



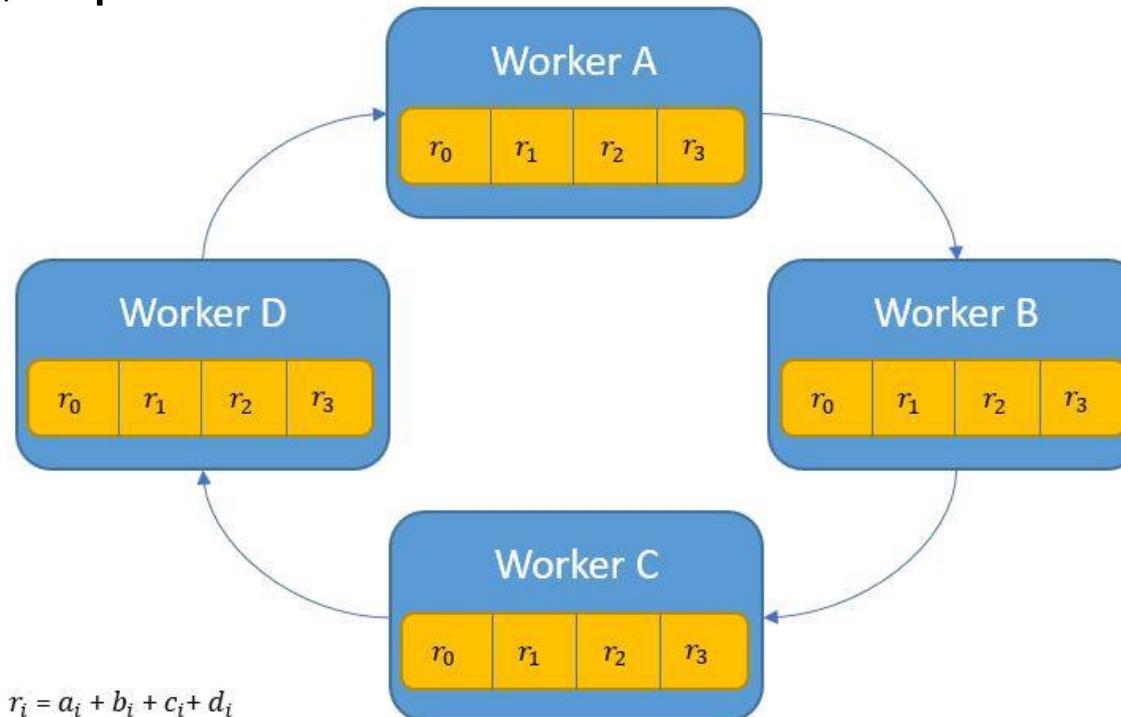
# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat N times
- Step 2 (Broadcast): each worker send one slice of aggregated parameters to the next worker; repeat N times



# Ring AllReduce

- Construct a ring of N workers, divide M parameters into N slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat N times
- Step 2 (Broadcast): each worker send one slice of aggregated parameters to the next worker; repeat N times

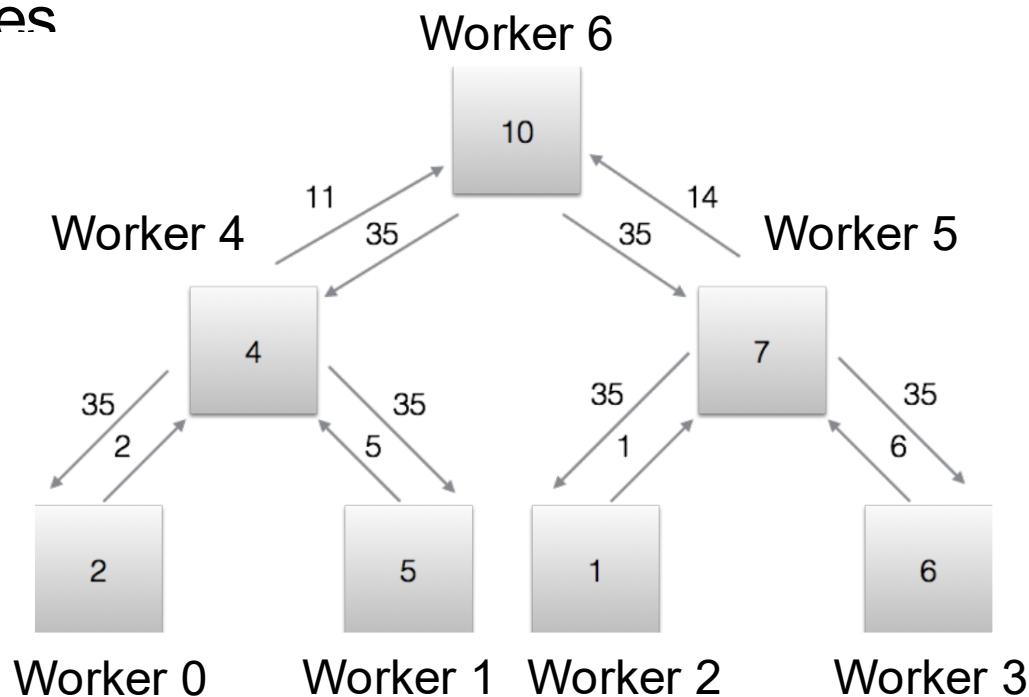


# Ring AllReduce

- Construct a ring of  $N$  workers, divide  $M$  parameters into  $N$  slices
- Step 1 (Aggregation): each worker send one slice ( $M/N$  parameters) to the next worker on the ring; repeat  $N$  times
- Step 2 (Broadcast): each worker send one slice of aggregated parameters to the next worker; repeat  $N$  times
- Overall communication:  $2 * M * N$  parameters
  - Aggregation:  $M * N$  parameters
  - Broadcast:  $M * N$  parameters

# Tree AllReduce

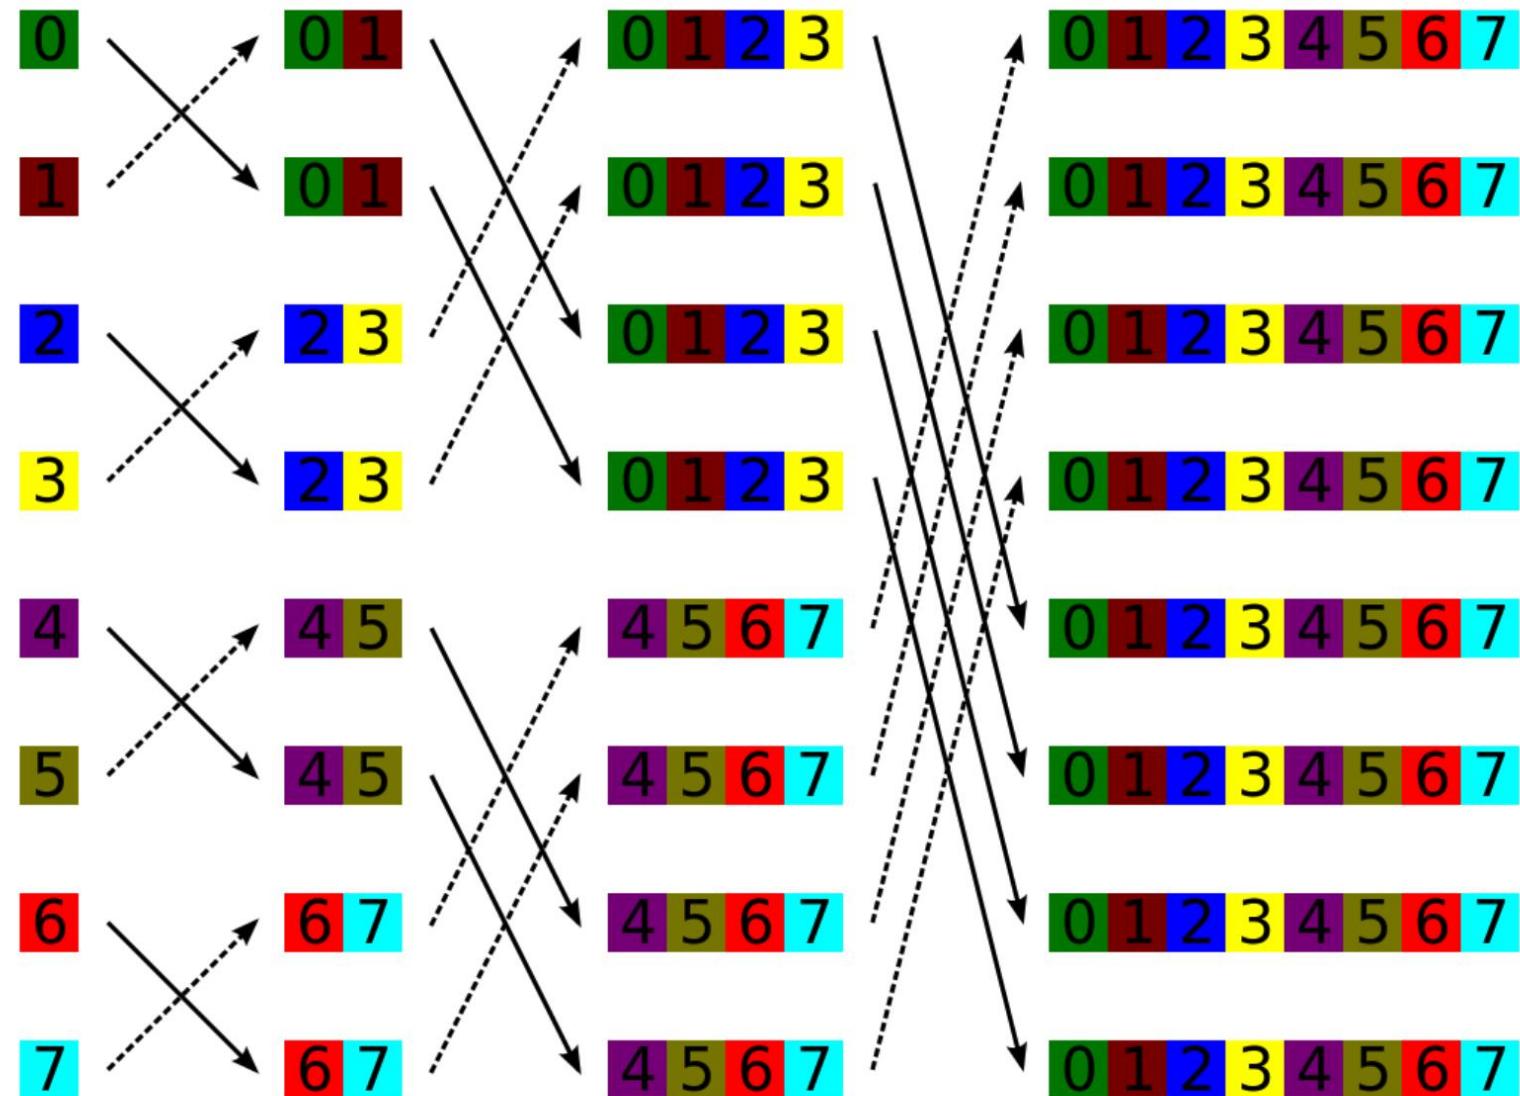
- Construct a tree of N workers;
- Step 1 (Aggregation): each worker sends M parameters to its parent; repeat  $\log(N)$  times
- Step 2 (Broadcast): each worker sends M parameters to its children; repeat  $\log(N)$  times



# Tree AllReduce

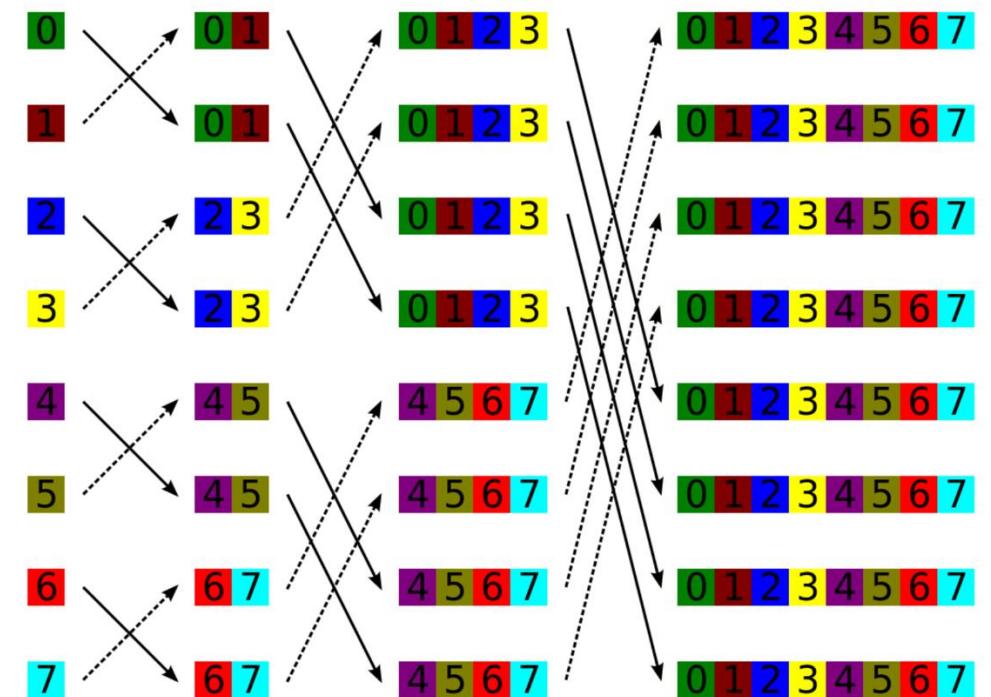
- Construct a tree of  $N$  workers;
- Step 1 (Aggregation): each worker sends  $M$  parameters to its parent; repeat  $\log(N)$  times
- Step 2 (Broadcast): each worker sends  $M$  parameters to its children; repeat  $\log(N)$  times
- Overall communication:  $2 * N * M$  parameters
  - Aggregation:  $M * N$  parameters
  - Broadcast:  $M * N$  parameters

# Butterfly Network



# Butterfly AllReduce

- Repeat  $\log(N)$  times:
  1. Each worker sends  $M$  parameters to its target node in the butterfly network
  2. Each worker aggregates gradients locally
- Overall communication:  $N * M * \log(N)$  parameters



# Comparing different AllReduce Methods

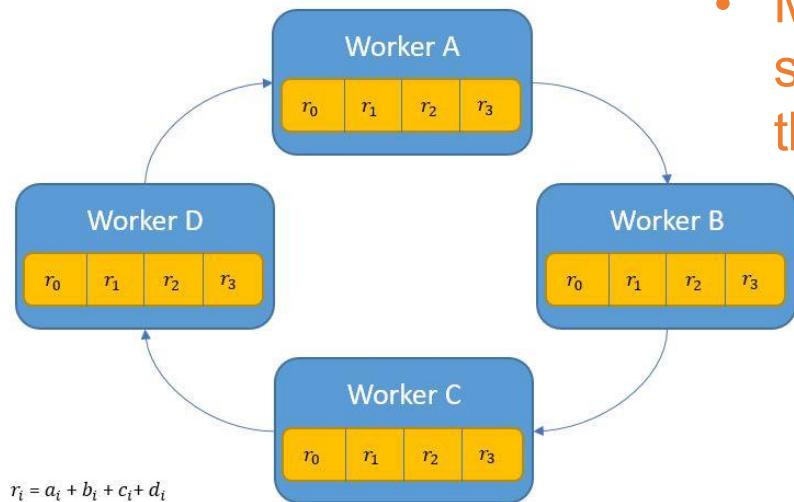
	Parameter Server	Naïve AllReduce	Ring AllReduce	Tree AllReduce	Butterfly AllReduce
Overall communication	$2 \times N \times M$	$N^2 \times M$	$2 \times N \times M$	$2 \times N \times M$	$N \times M \times \log N$

Question: Ring AllReduce is more efficient and scalable than Tree AllReduce and Parameter Server, why?

# Ring AllReduce v.s. Tree AllReduce v.s. Parameter Server

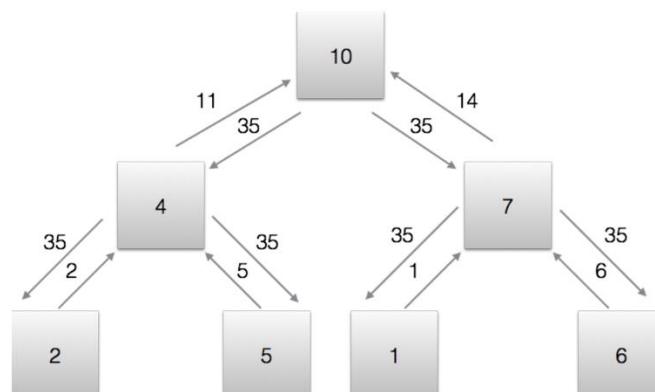
## Ring AllReduce:

- Best latency
- Balanced workload across workers
- More scalable since each worker sends  $2 \cdot M$  parameters (independent to the number of workers)



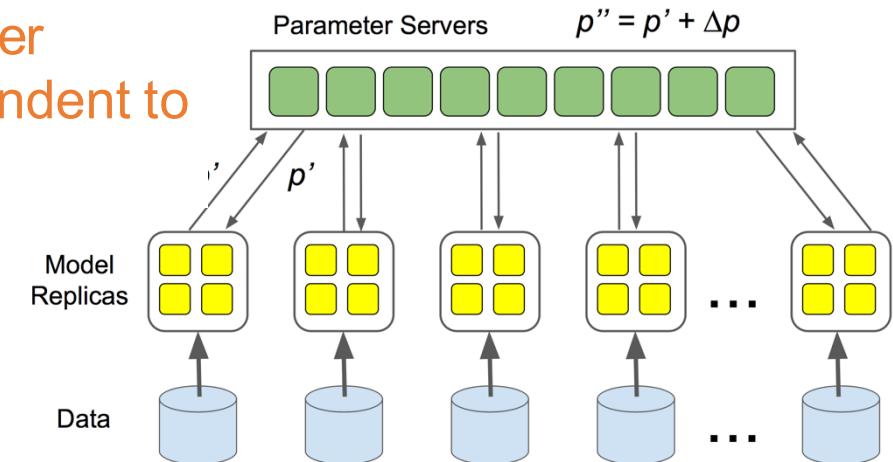
Each worker sends **M/N** parameters per iteration; repeat for **2\*N** iterations

**Latency**:  $M/N * (2*N) / \text{bandwidth}$



Each worker sends **M** parameters per iteration; repeat for **2\*log(N)** iterations

**Latency**:  $M * 2 * \log(N) / \text{bandwidth}$



All workers send **M** parameters to parameter servers and receive **M** parameters from servers

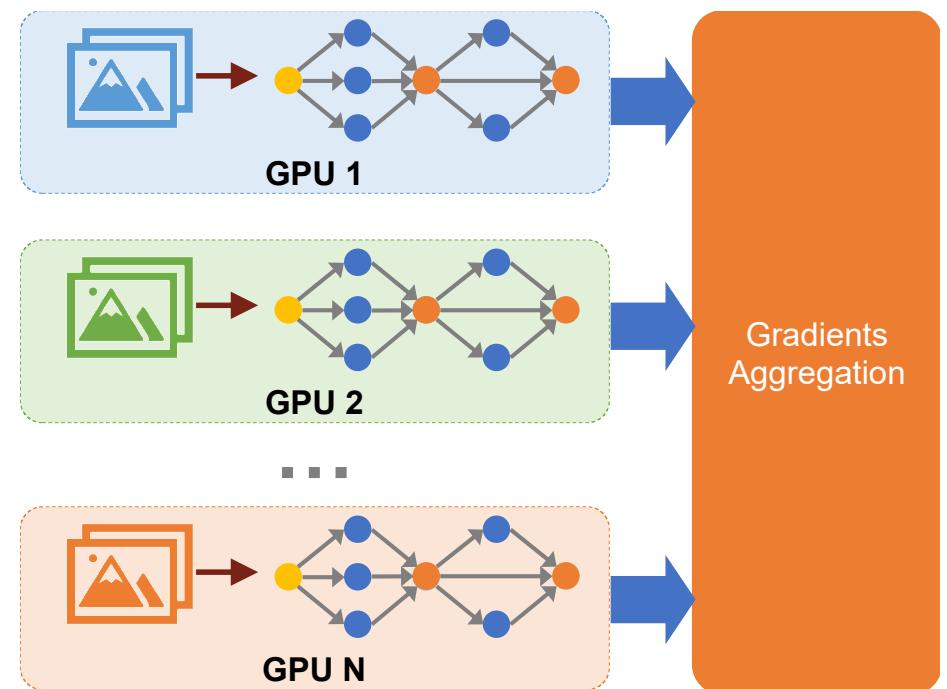
**Latency**:  $M * N / \text{bandwidth}$

# Recap: Data Parallelism (Quiz)

Each worker keeps a replica of the entire model and communicates with other workers to synchronize weights updates

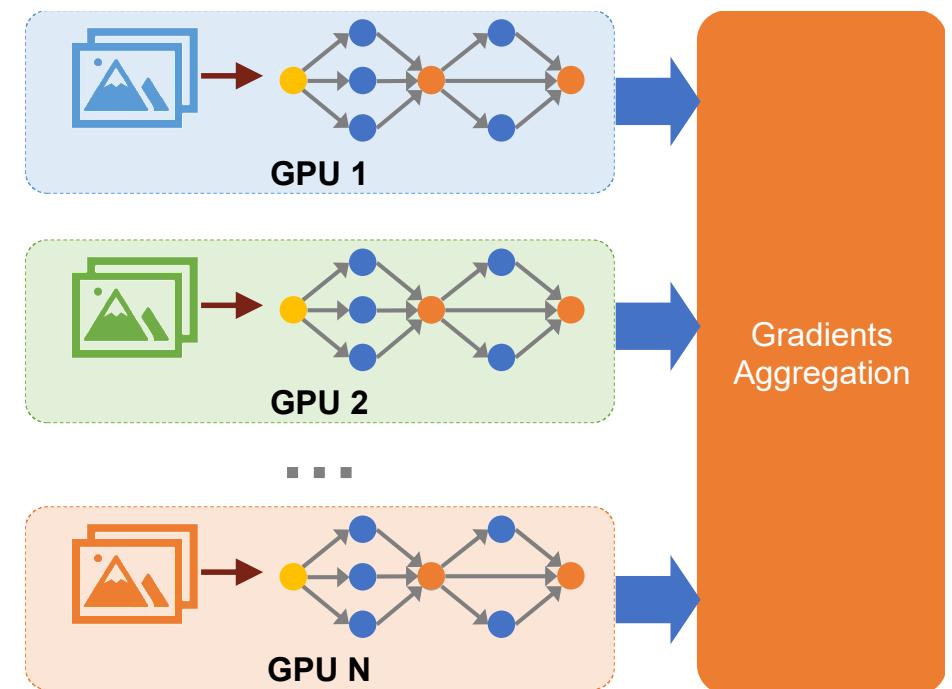
Gradients aggregation methods:

- Parameter Server
- Ring AllReduce
- Tree AllReduce
- Butterfly AllReduce
- Etc.



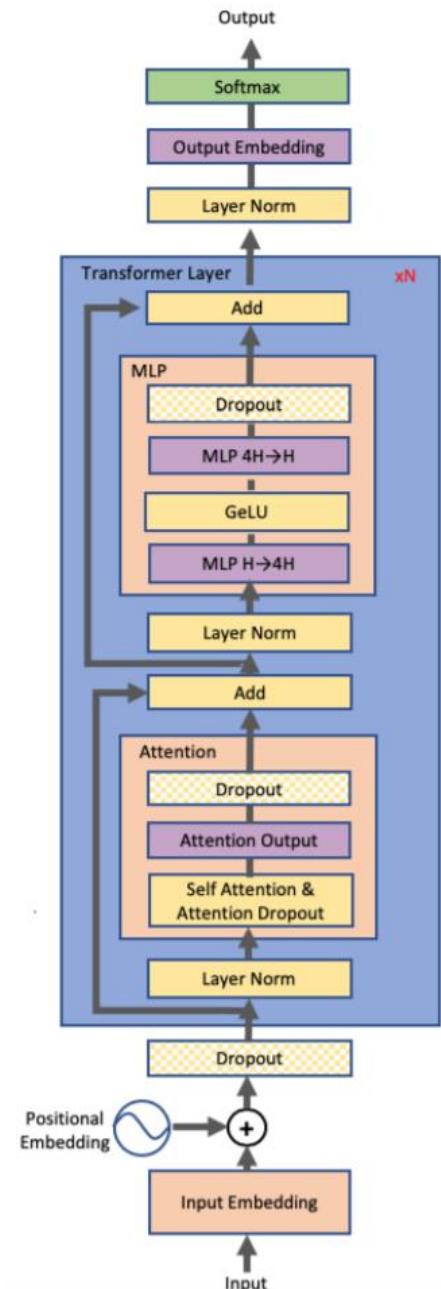
# An Issue with Data Parallelism

- Each GPU saves a replica of the entire model
- Cannot train large models that exceed GPU device memory



# Large Model Training Challenges

	Bert-Large	GPT-2	Turing 17.2 NLG	GPT-3
Parameters	0.32B	1.5B	17.2B	175B
Layers	24	48	78	96
Hidden Dimension	1024	1600	4256	12288
Relative Computation	1x	4.7x	54x	547x
Memory Footprint	5.12GB	24GB	275GB	2800GB

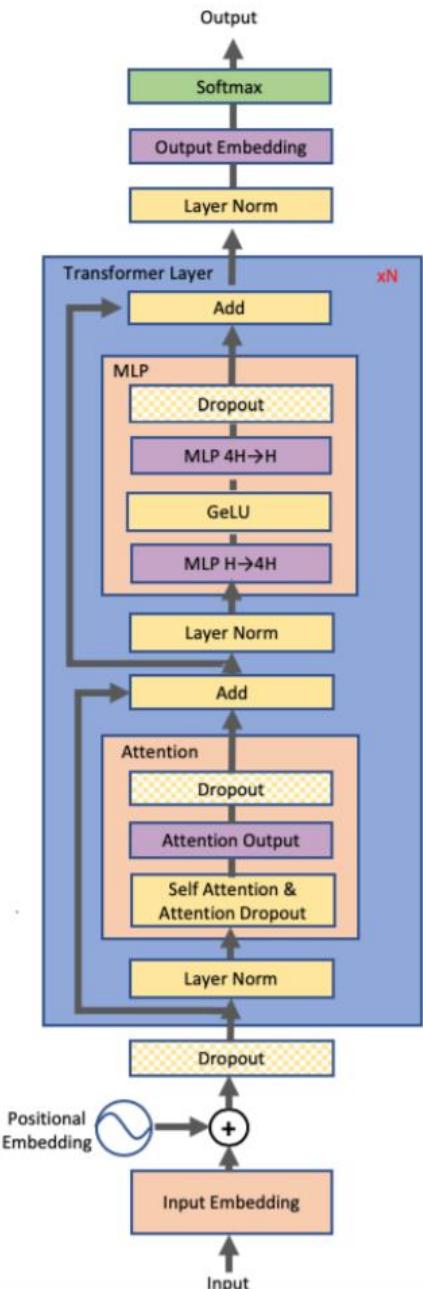


# Large Model Training Challenges

	Bert-Large	GPT-2	Turing 17.2 NLG	GPT-3
Parameters	0.32B	1.5B	17.2B	175B
Layers	24	48	78	96
Hidden Dimension	1024	1600	4256	12288
Relative Computation	1x	4.7x	54x	547x
Memory Footprint	5.12GB	24GB	275GB	2800GB

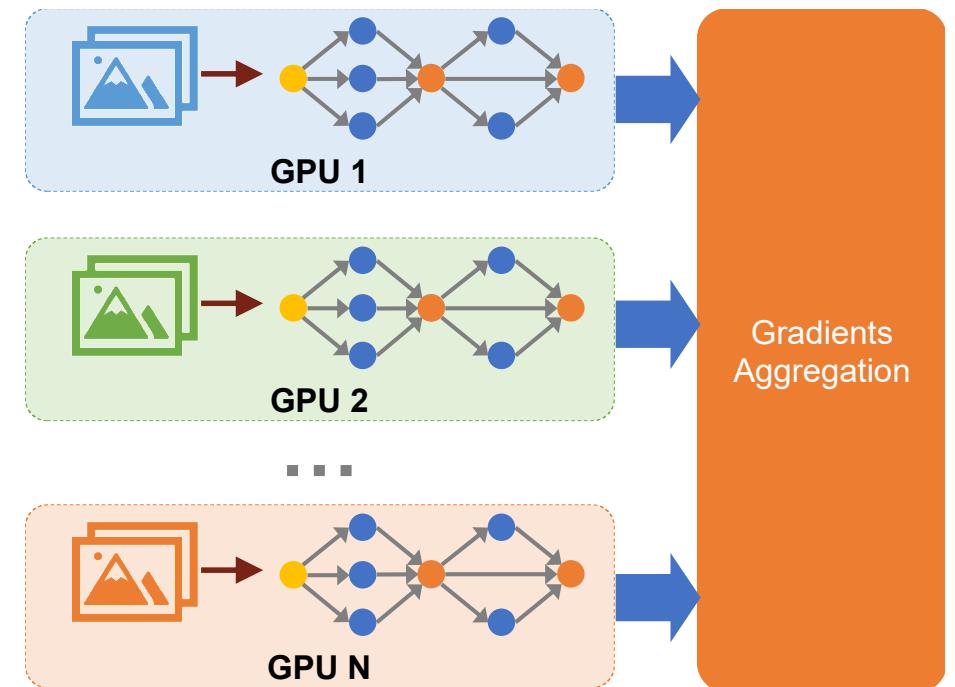
NVIDIA V100 GPU memory capacity: 16G/32G  
NVIDIA A100 GPU memory capacity: 40G/80G

Out of Memory



# ZeRO: Zero Redundancy Optimizer

- Eliminating data redundancy in data parallel training
- A widely used technique for data parallel training of large models



# Revisit: Stochastic Gradient Descent

```
For t = 1 to T
     $\Delta w = \eta \times \frac{1}{b} \sum_{i=1}^b \nabla \left( loss(f_w(x_i, y_i)) \right)$  // compute derivative and update
    w -= Δw // apply update
End
```

Backward pass      Forward pass

The diagram illustrates the flow of data and gradients in a stochastic gradient descent step. A blue arrow labeled "Backward pass" points to the gradient operator  $\nabla$  in the update equation. A green arrow labeled "Forward pass" points to the loss function  $loss(f_w(x_i, y_i))$  within the gradient calculation.

# Adaptive Learning Rates (Adam)

For  $t = 1$  to  $T$

$$g = \frac{1}{b} \sum_{i=1}^b \nabla \left( \text{loss}(f_w(x_i, y_i)) \right)$$

$$\Delta w = \text{adam}(g)$$

$w -= \Delta w$  // apply update

End

$$\nu_t = \beta_1 * \nu_{t-1} - (1 - \beta_1) * g_t$$
$$s_t = \beta_2 * s_{t-1} - (1 - \beta_2) * g_t^2$$

$$\Delta \omega_t = -\eta \frac{\nu_t}{\sqrt{s_t + \epsilon}} * g_t$$

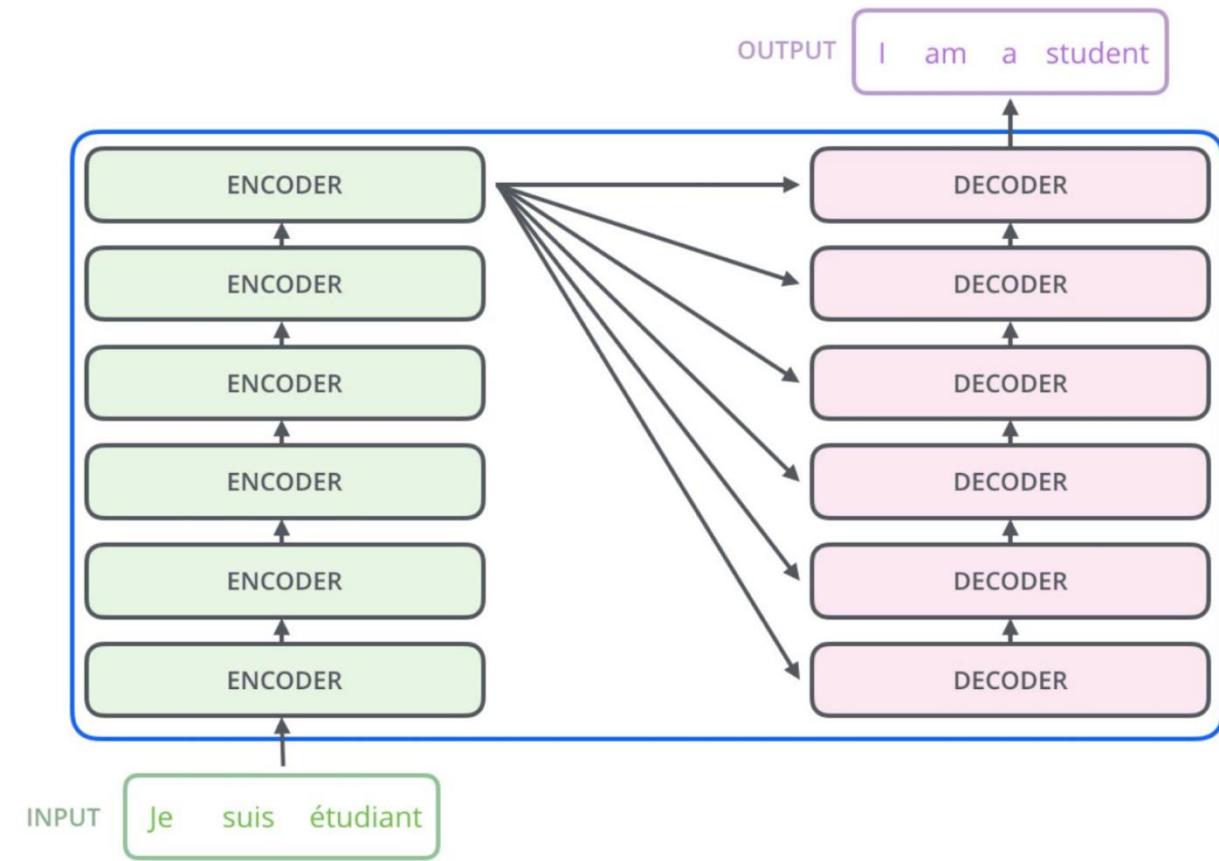
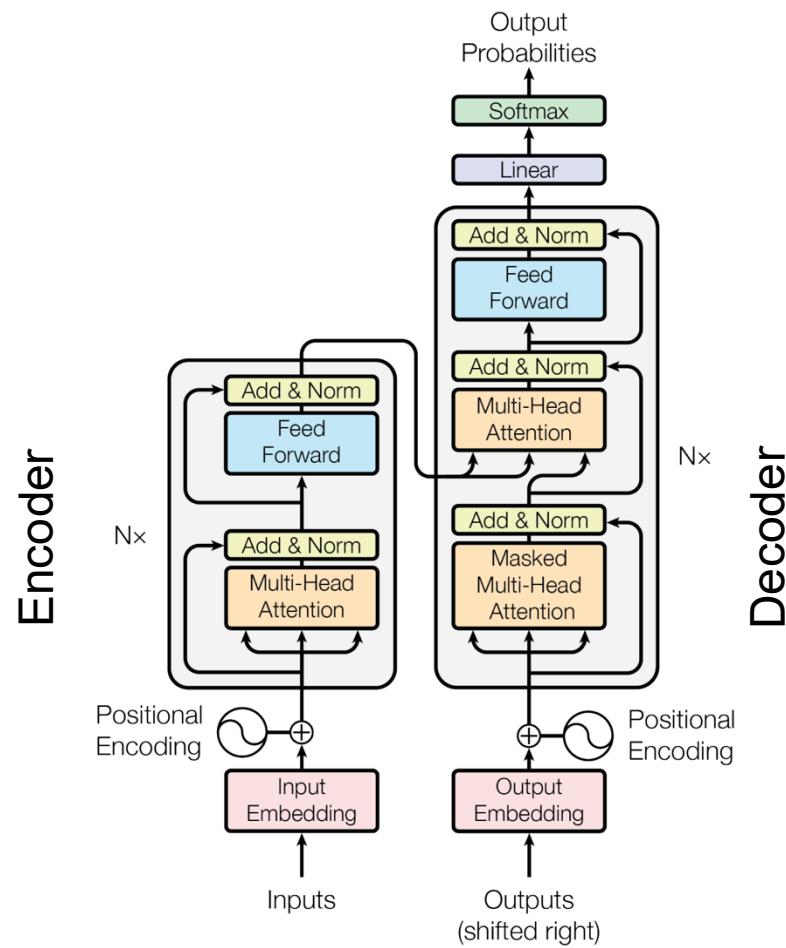
$g_t$  : Gradient at time  $t$  along  $\omega^j$

$\nu_t$  : Exponential Average of gradients along  $\omega_j$

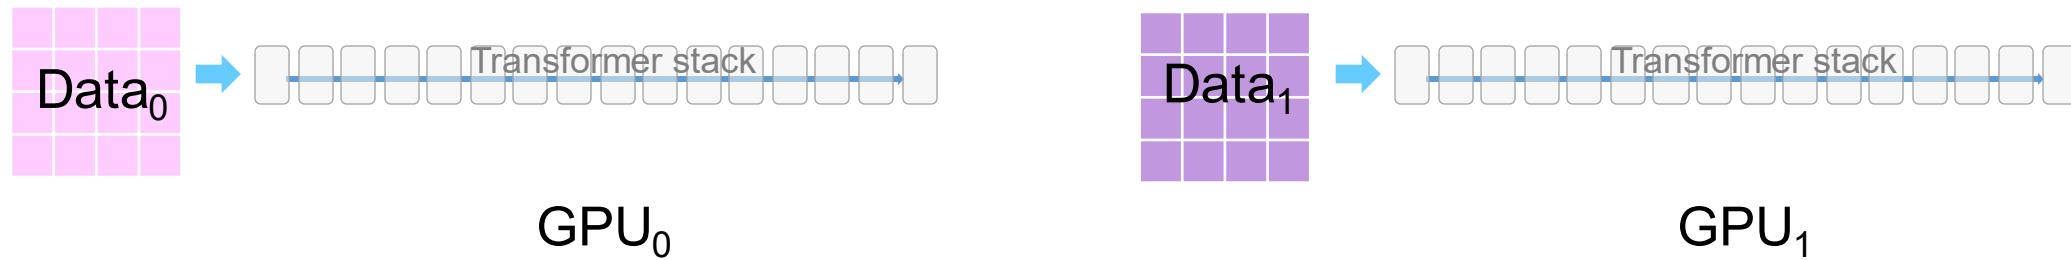
$s_t$  : Exponential Average of squares of gradients along  $\omega_j$

$\beta_1, \beta_2$  : Hyperparameters

# Recall: Transformer for Language Models

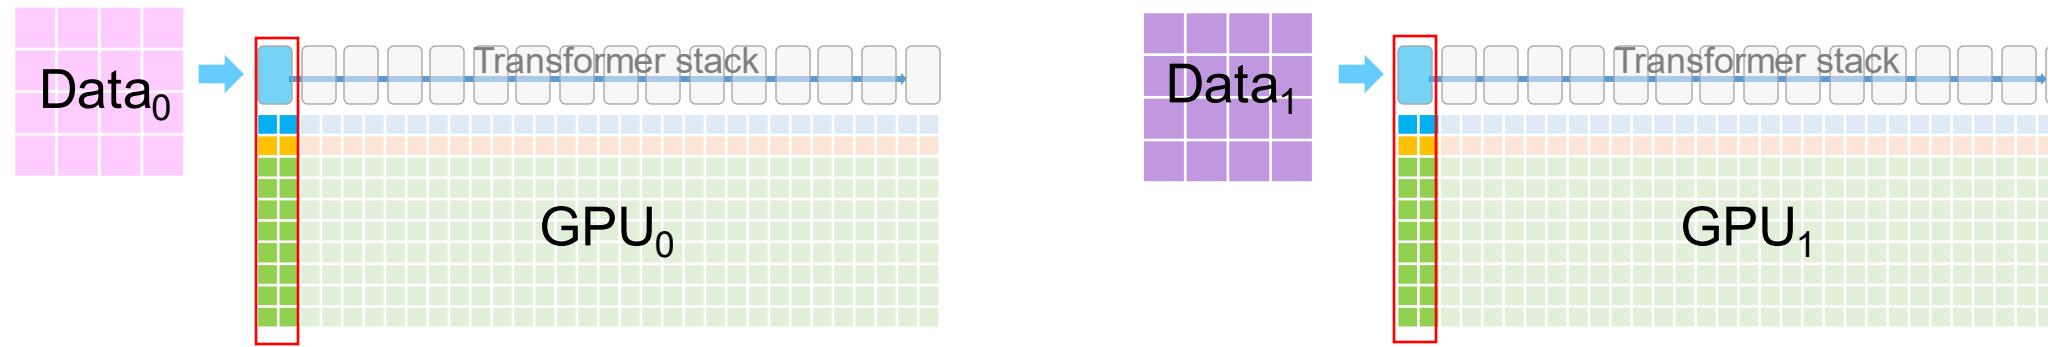


# Understanding Memory Consumption



A 16-layer transformer model = 1 layer

# Understanding Memory Consumption



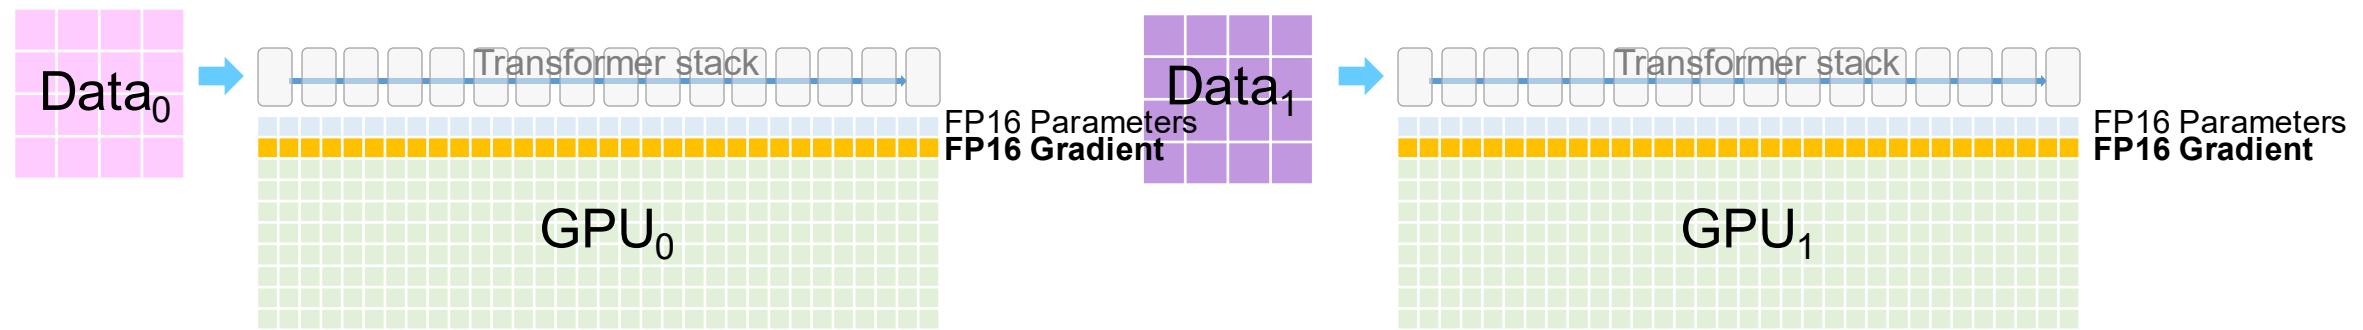
Each cell represents GPU memory used by its corresponding transformer layer

# Understanding Memory Consumption



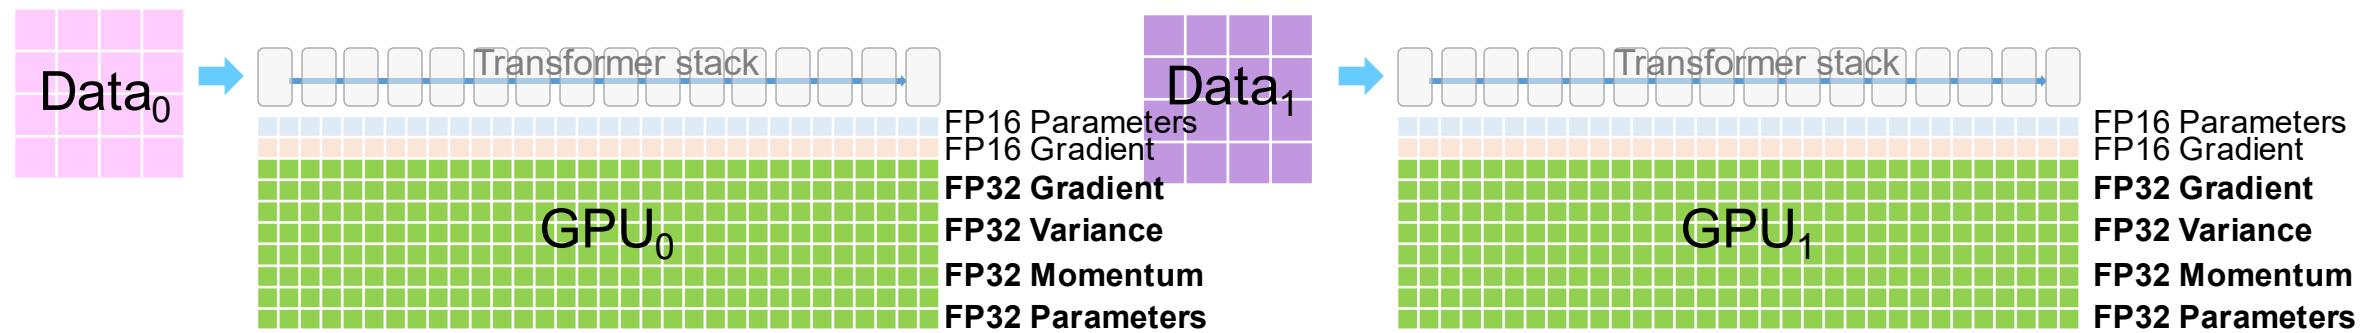
- FP16 parameter

# Understanding Memory Consumption



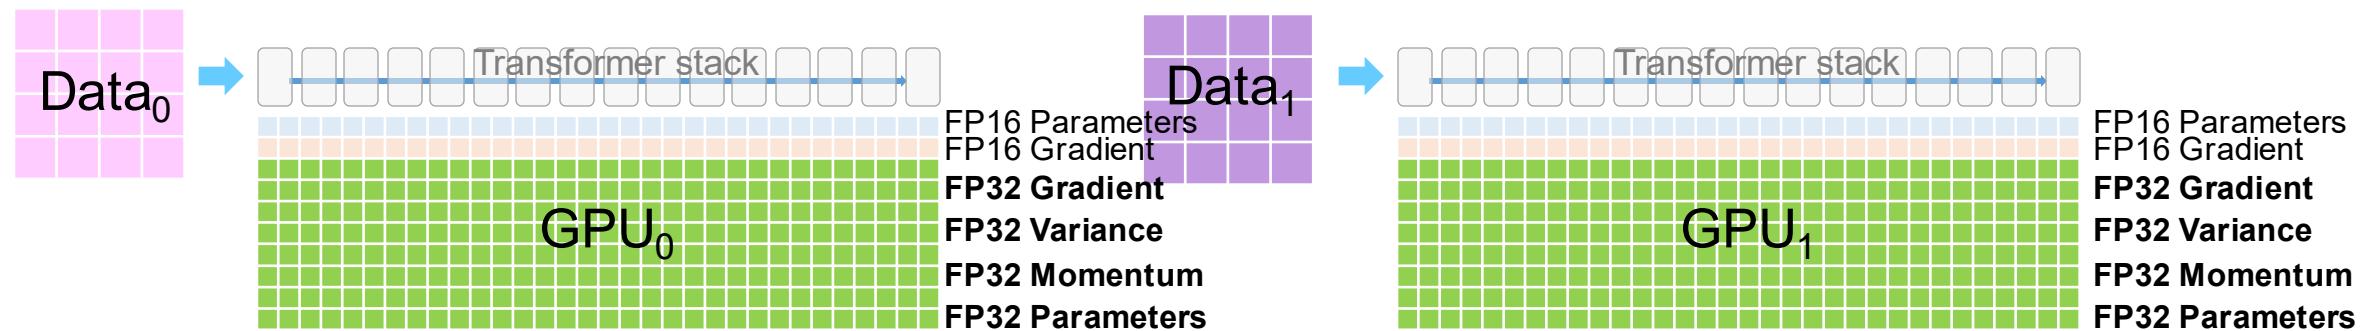
- FP16 parameter
- FP16 Gradients

# Understanding Memory Consumption



- FP16 parameter
- FP16 Gradients
- FP32 Optimizer States
  - Gradients, Variance, Momentum, Parameters

# Understanding Memory Consumption



- FP16 parameter : **2M bytes**
- FP16 Gradients : **2M bytes**
- FP32 Optimizer States : **16M bytes**
  - Gradients, Variance, Momentum, Parameters

Example 1B parameter model ->  
20GB/GPU

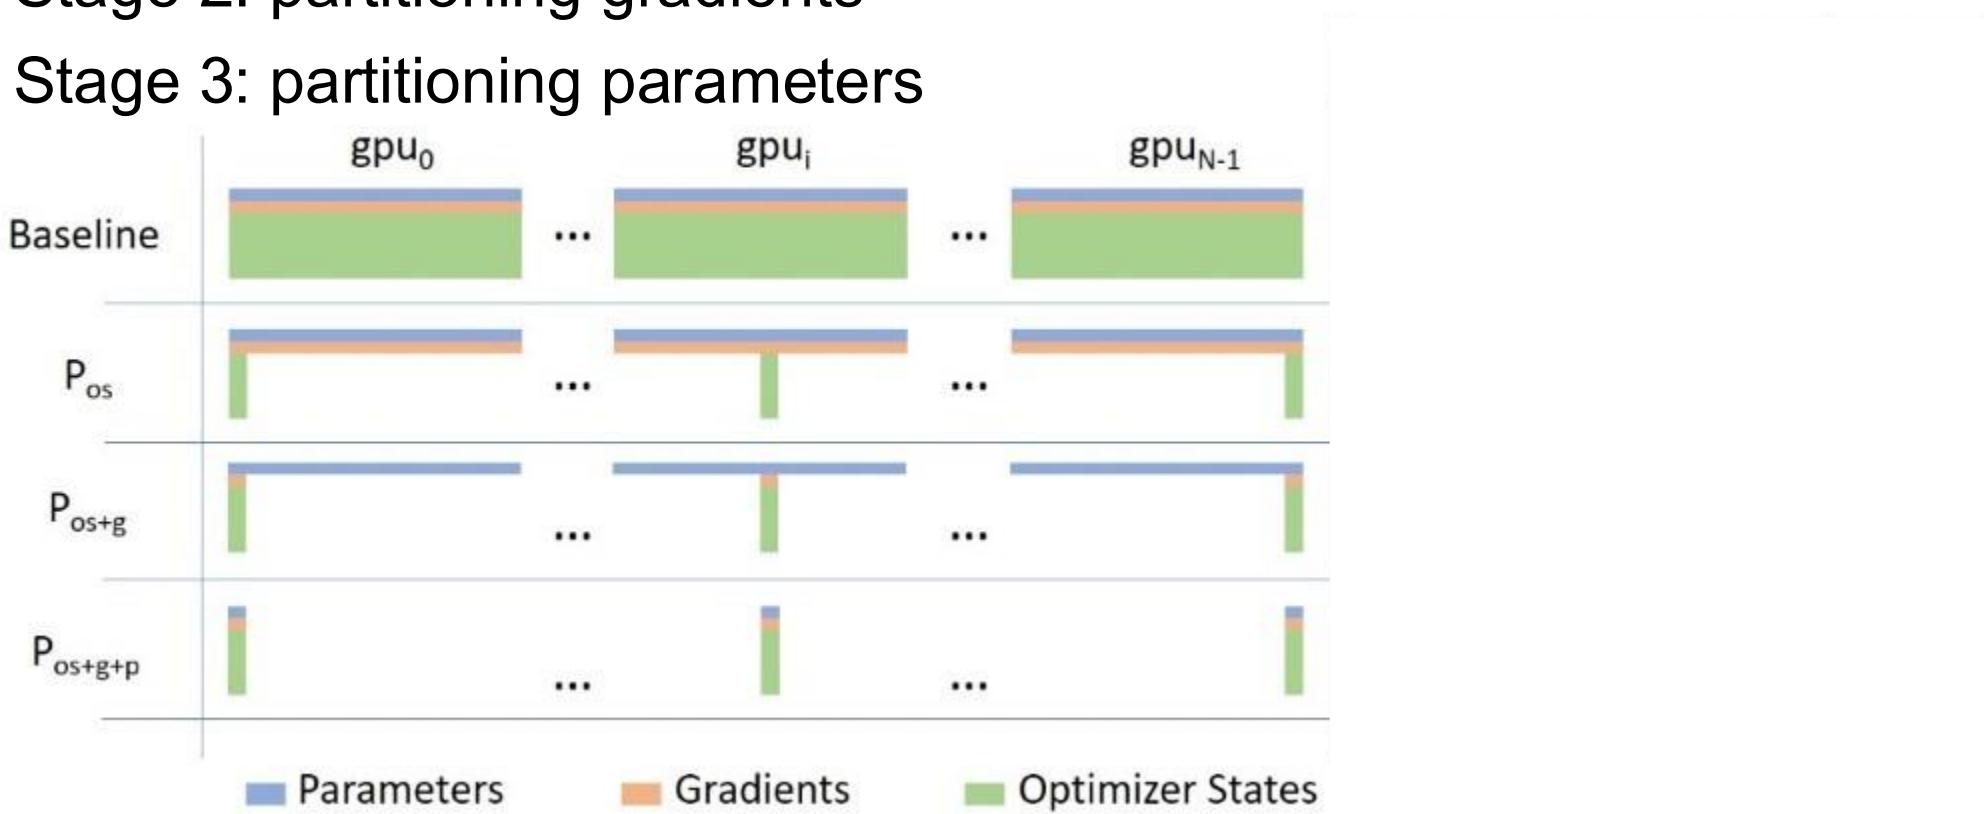
Memory consumption doesn't include:

- Input batch + activations

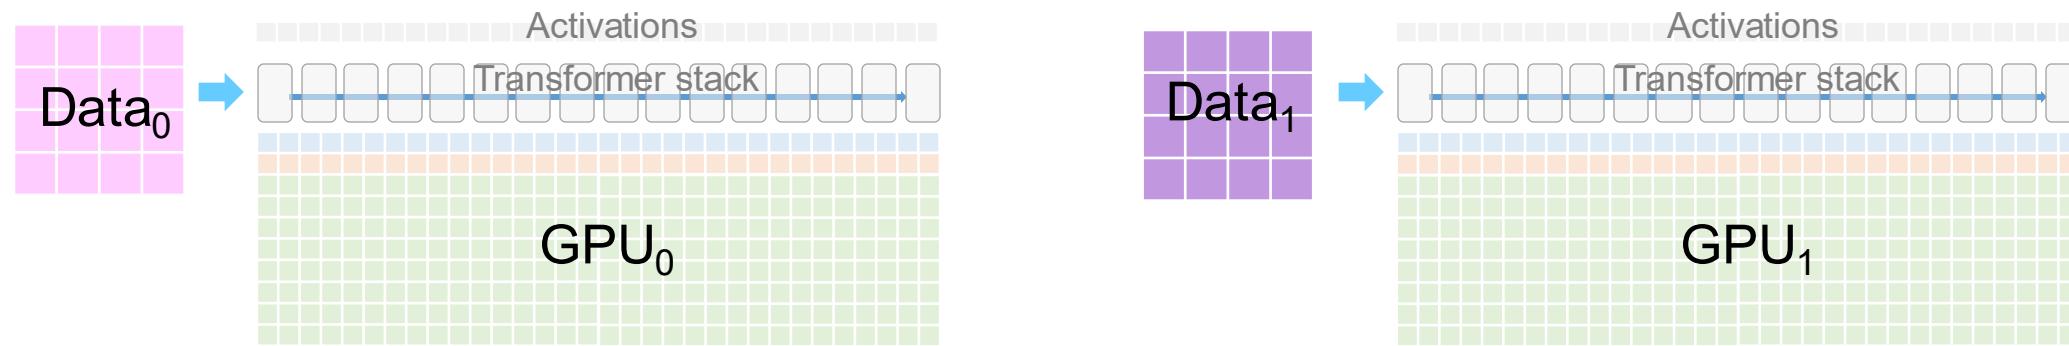
M = number of parameters in the model

# ZeRO-DP: ZeRO powered Data Parallelism

- ZeRO removes the redundancy across data parallel process
- Stage 1: partitioning optimizer states
- Stage 2: partitioning gradients
- Stage 3: partitioning parameters

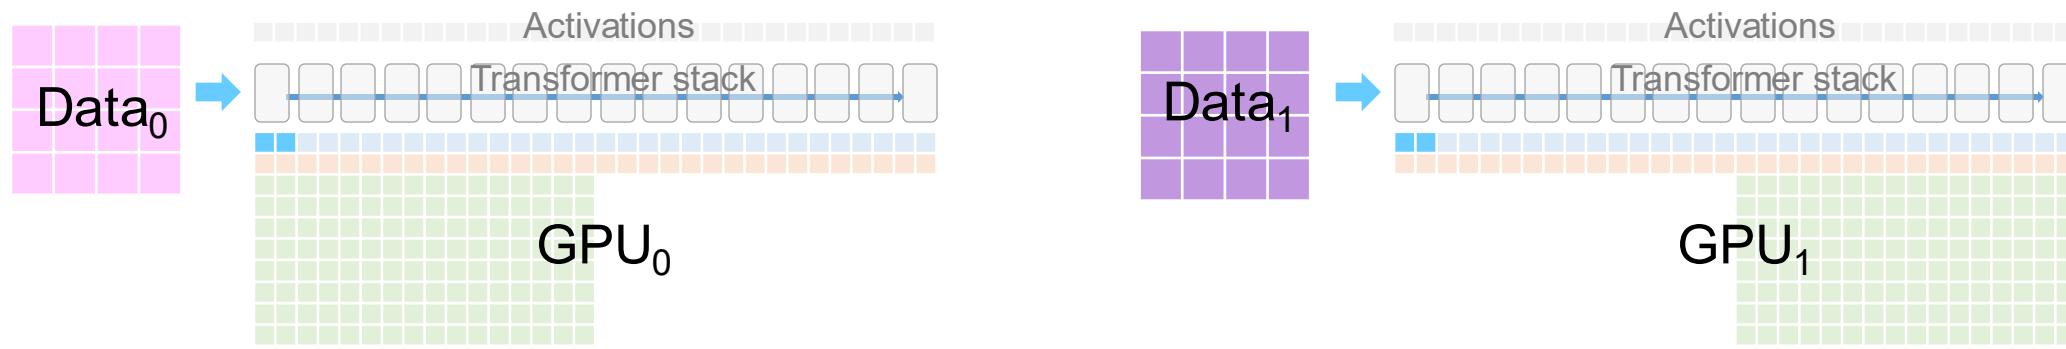


# ZeRO Stage 1: Partitioning Optimizer States



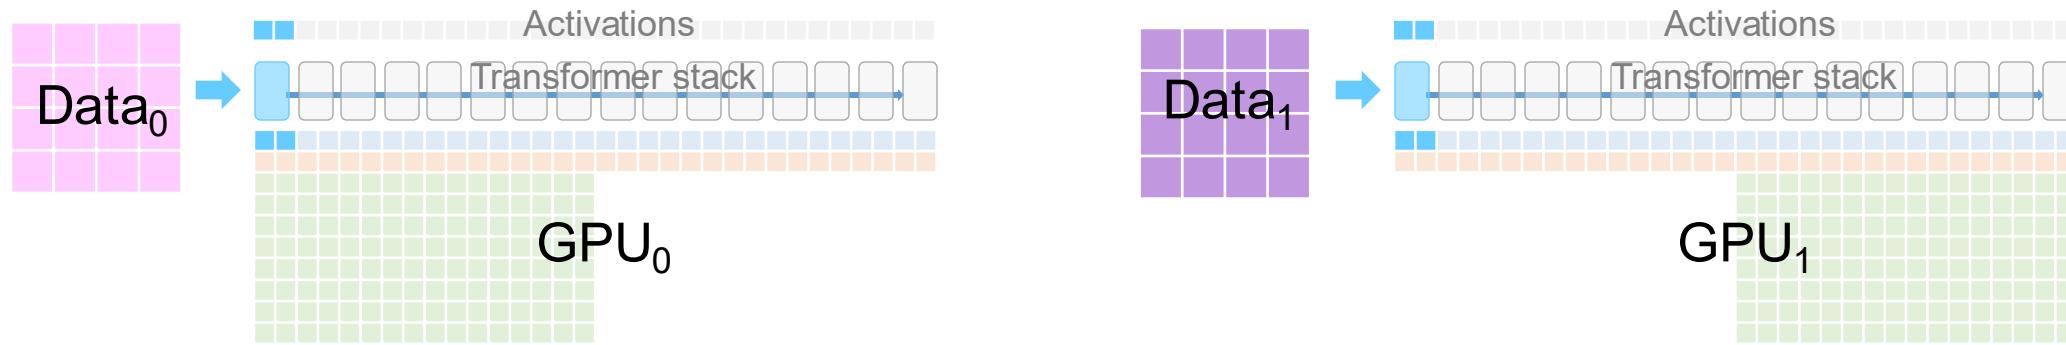
- ZeRO Stage 1

# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs

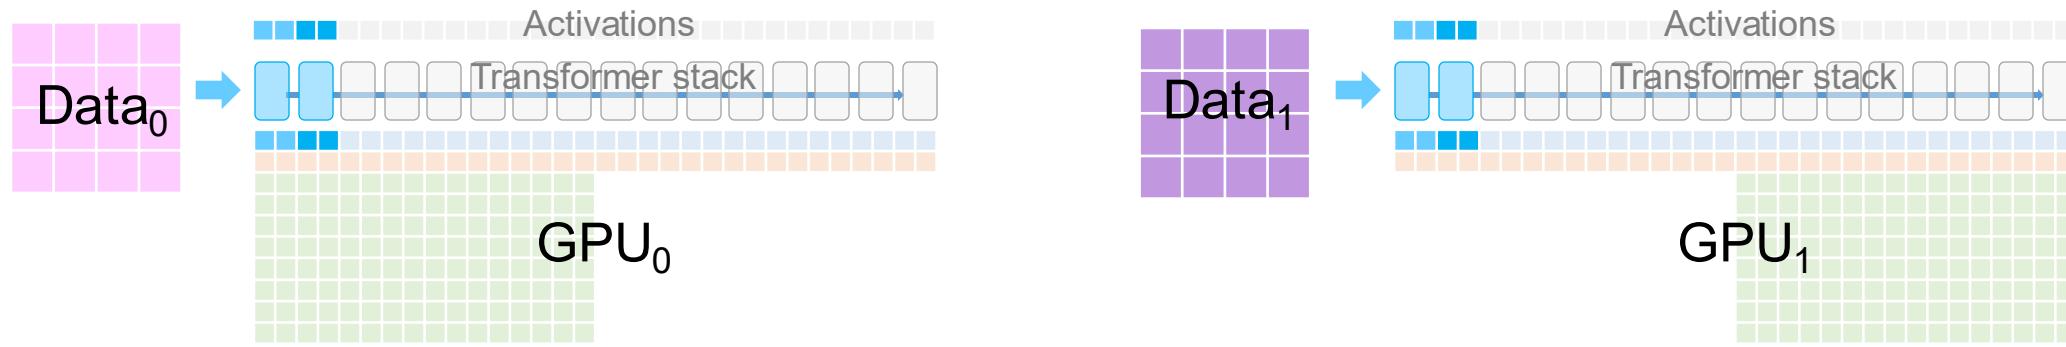
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

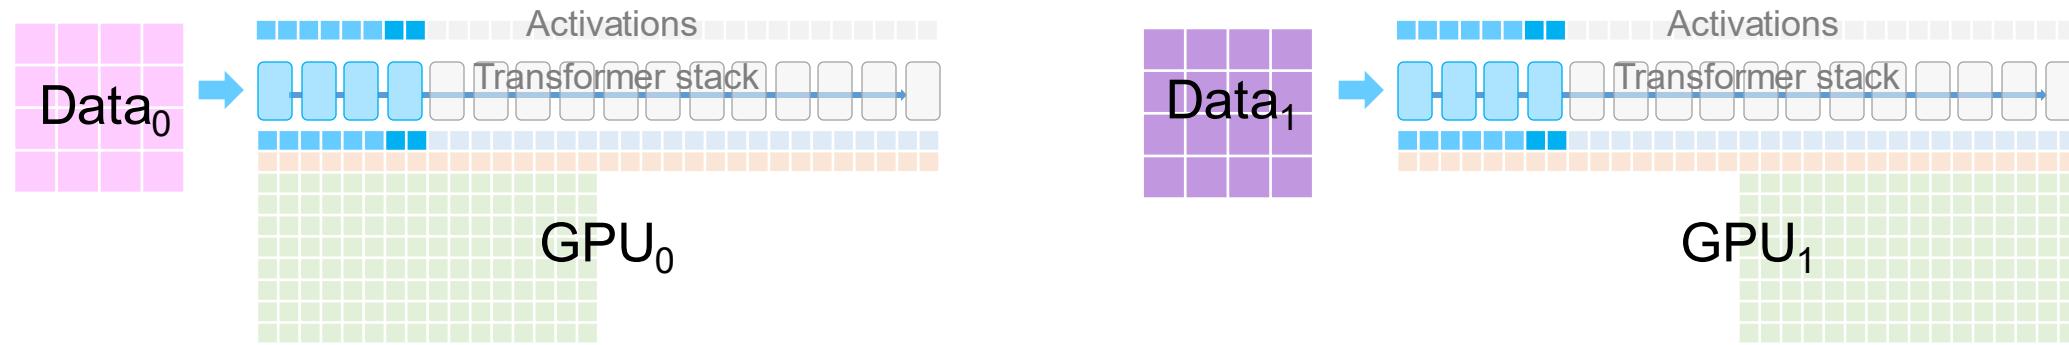
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

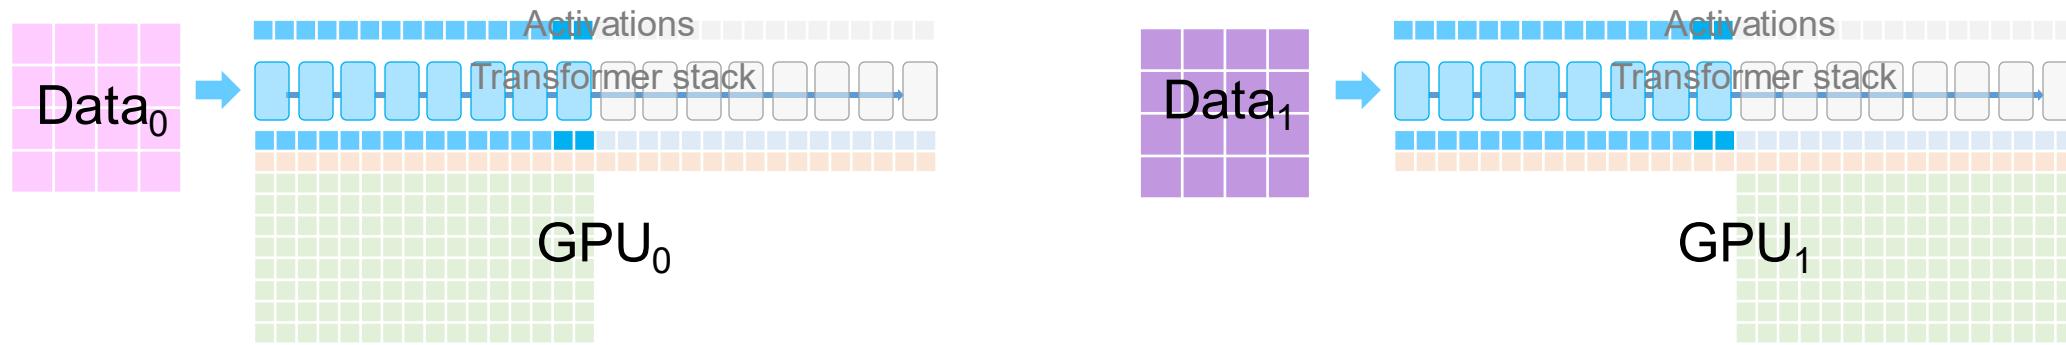
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

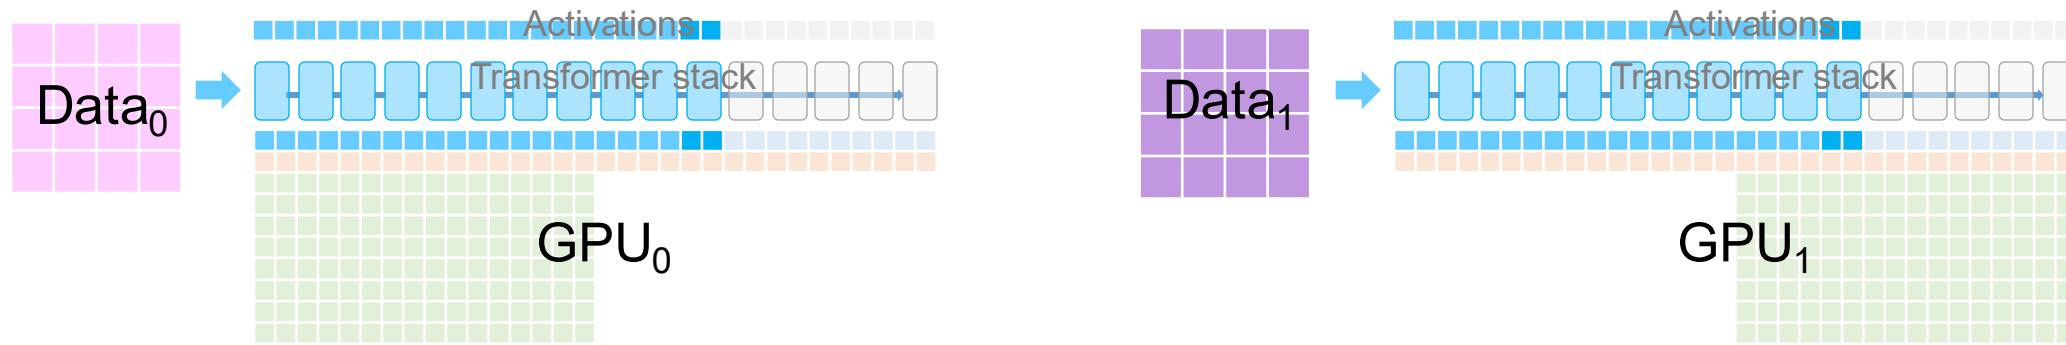
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

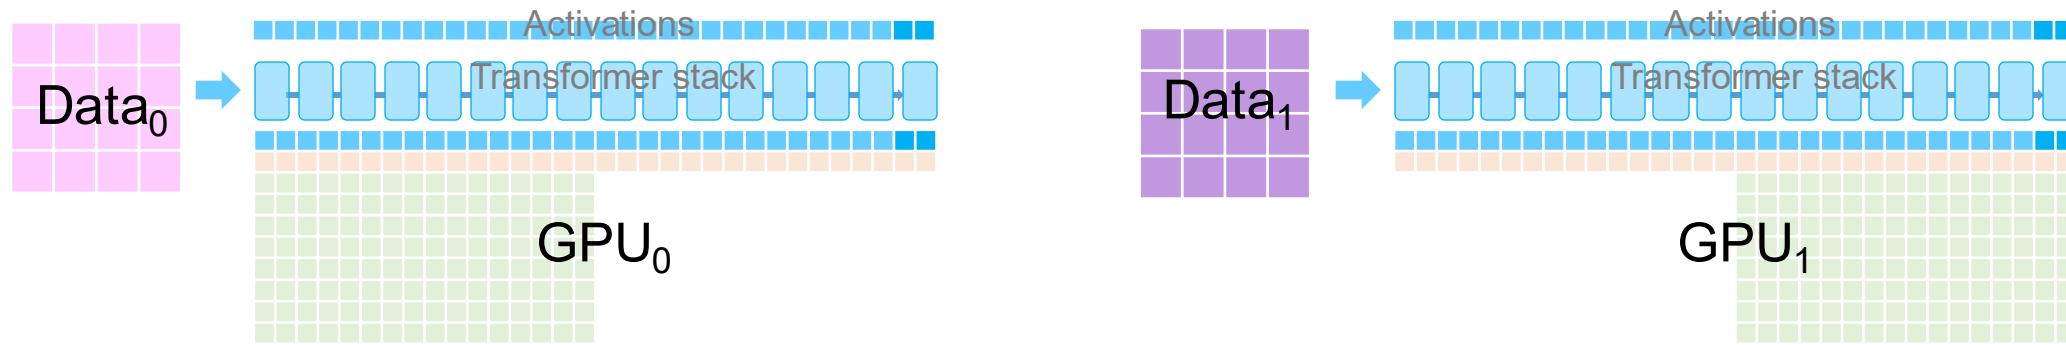
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

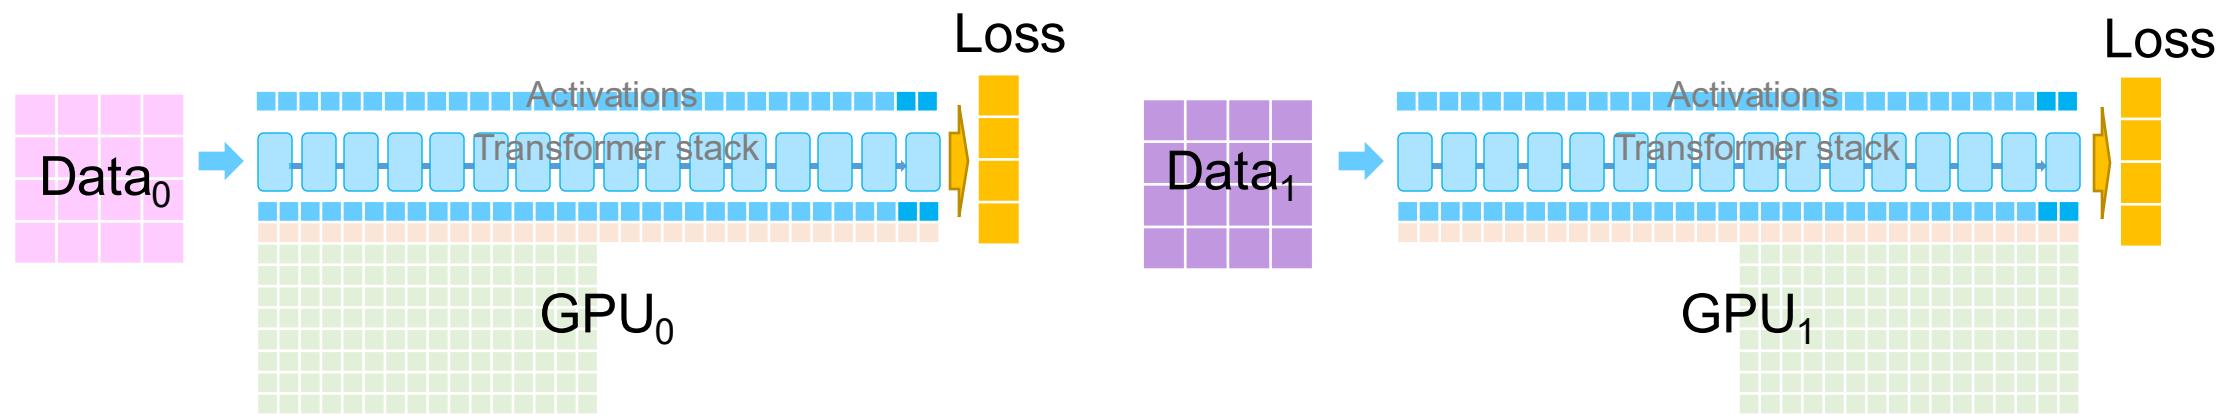
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

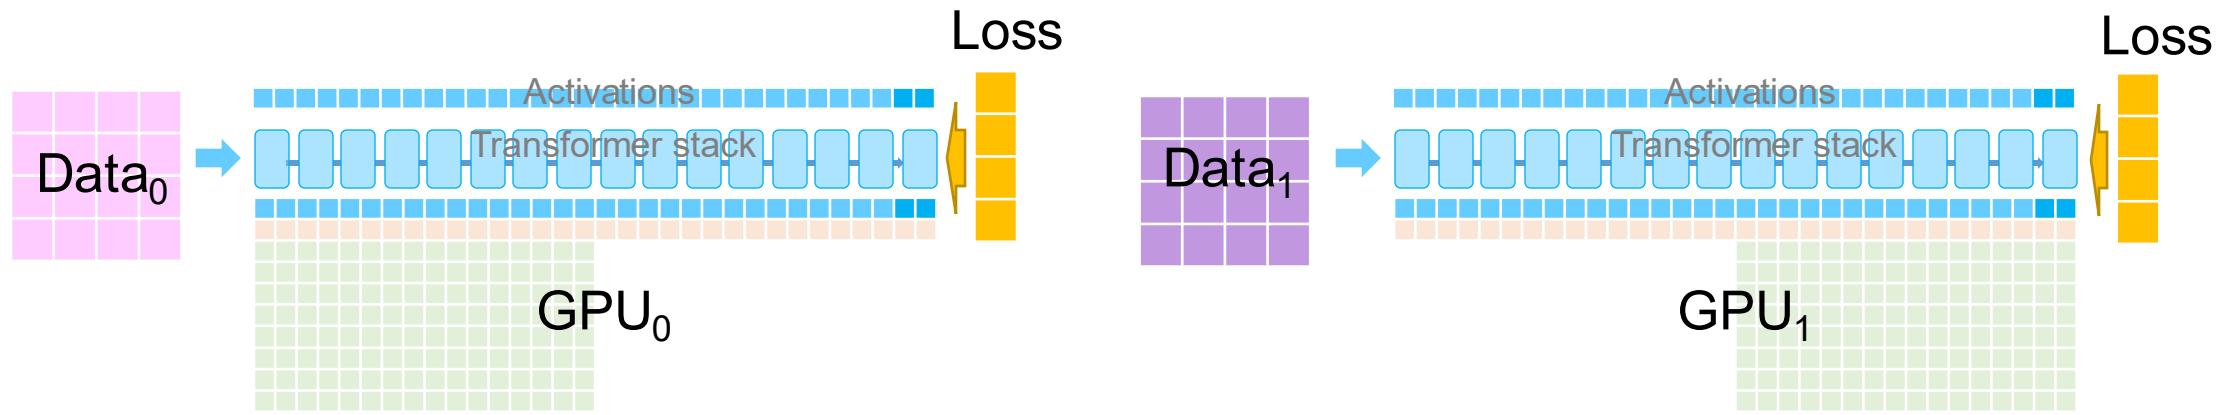
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks

■ Parameters ■ Gradients ■ Optimizer States

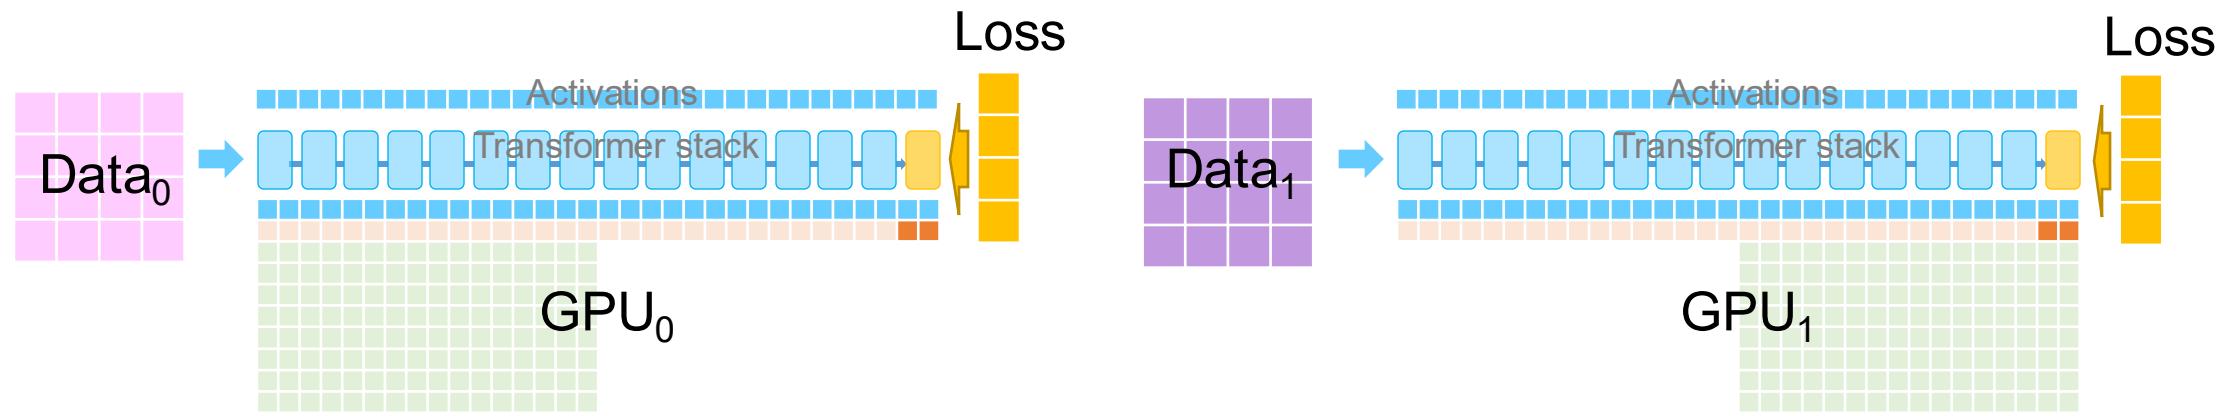
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

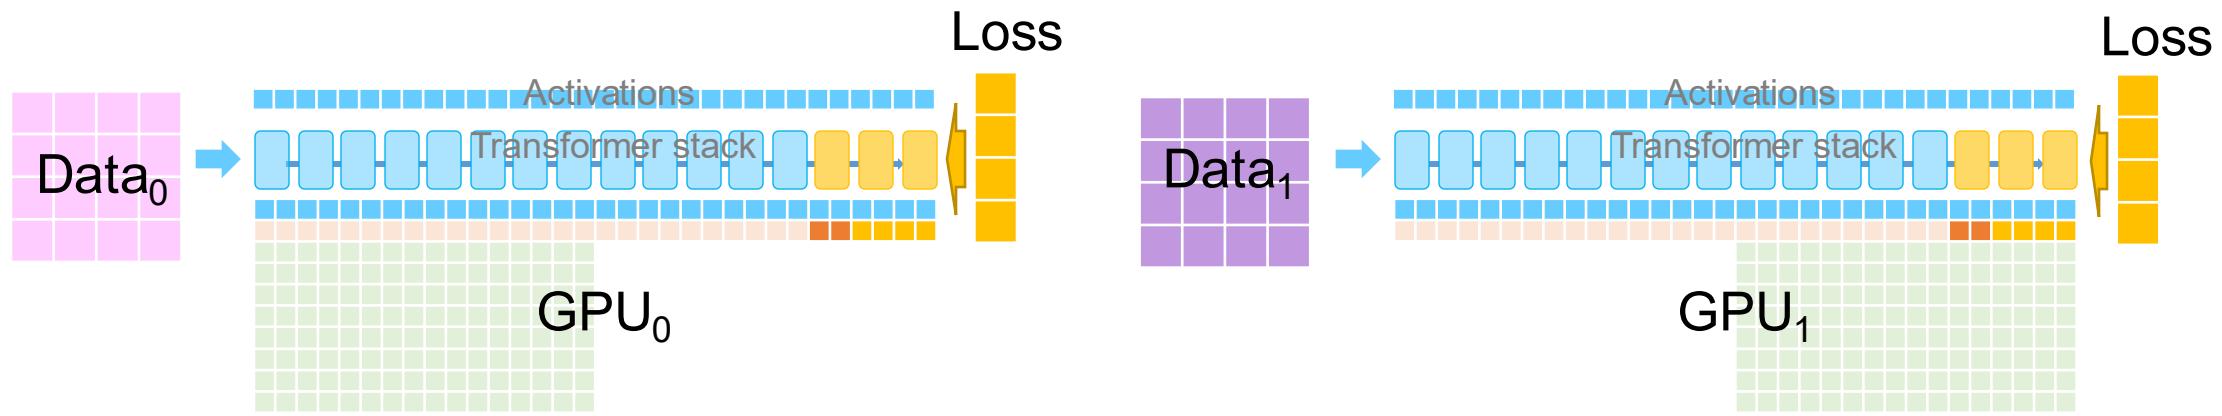
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

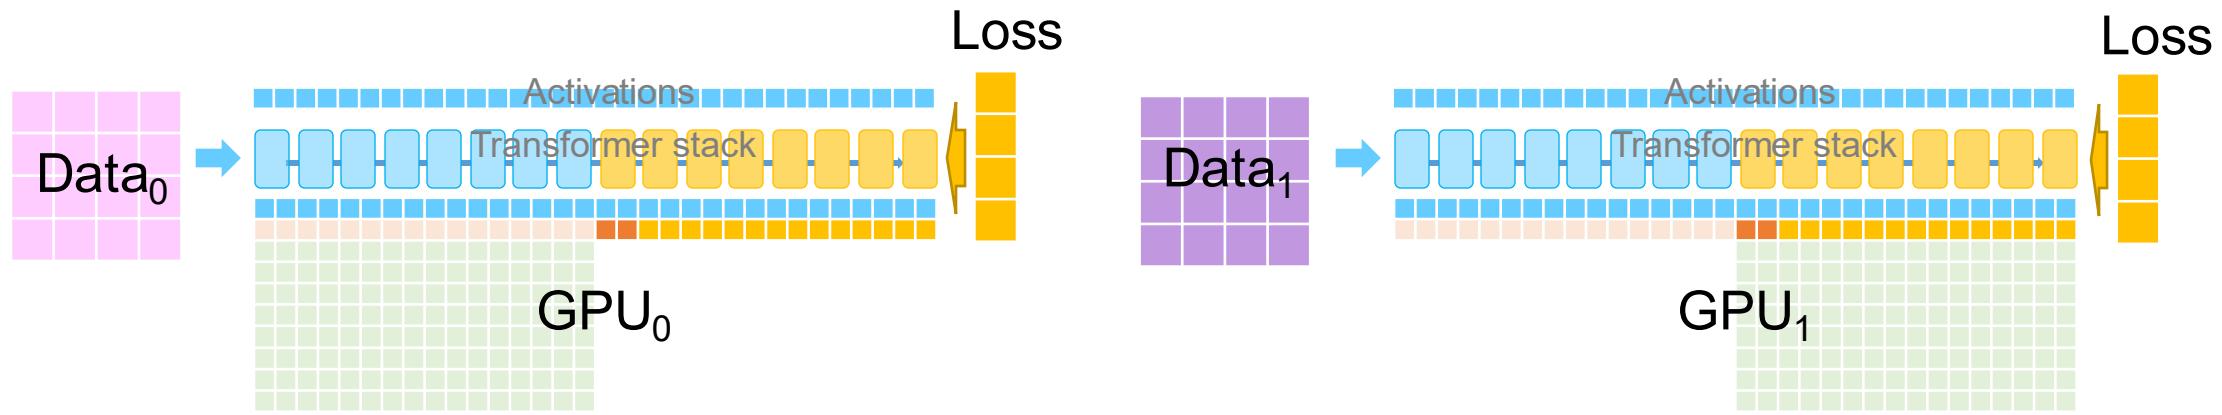
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

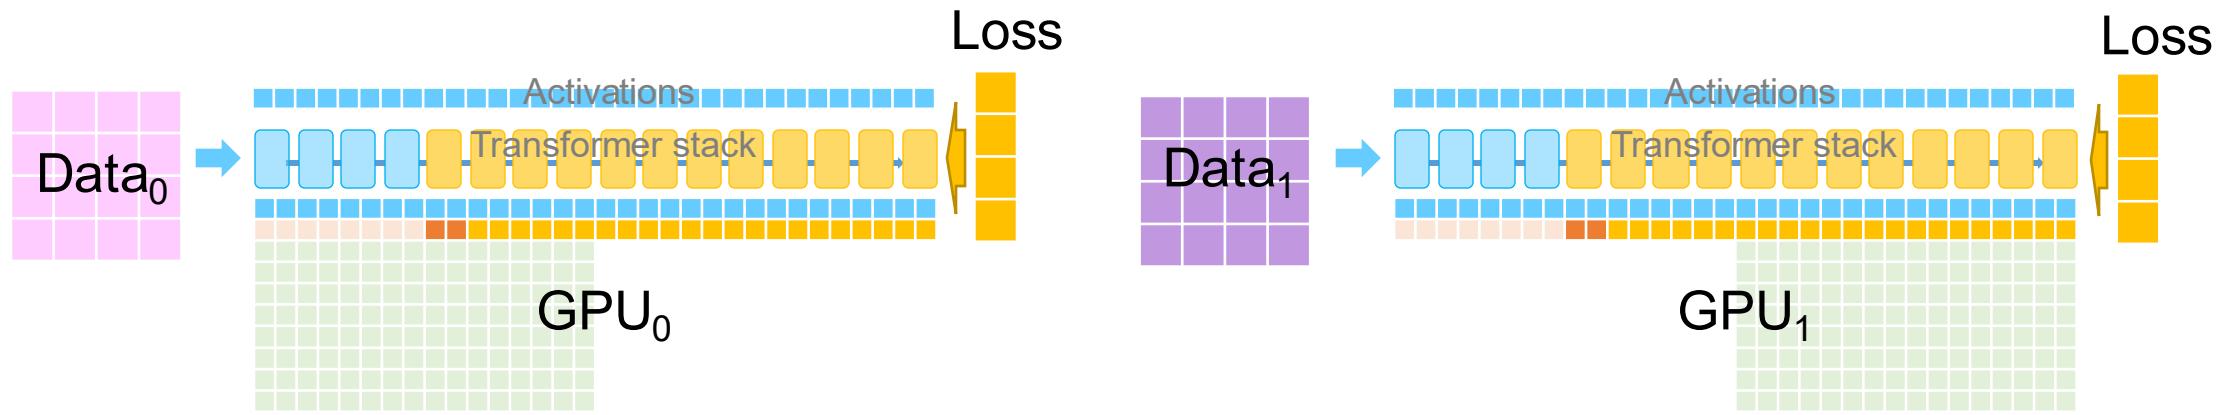
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

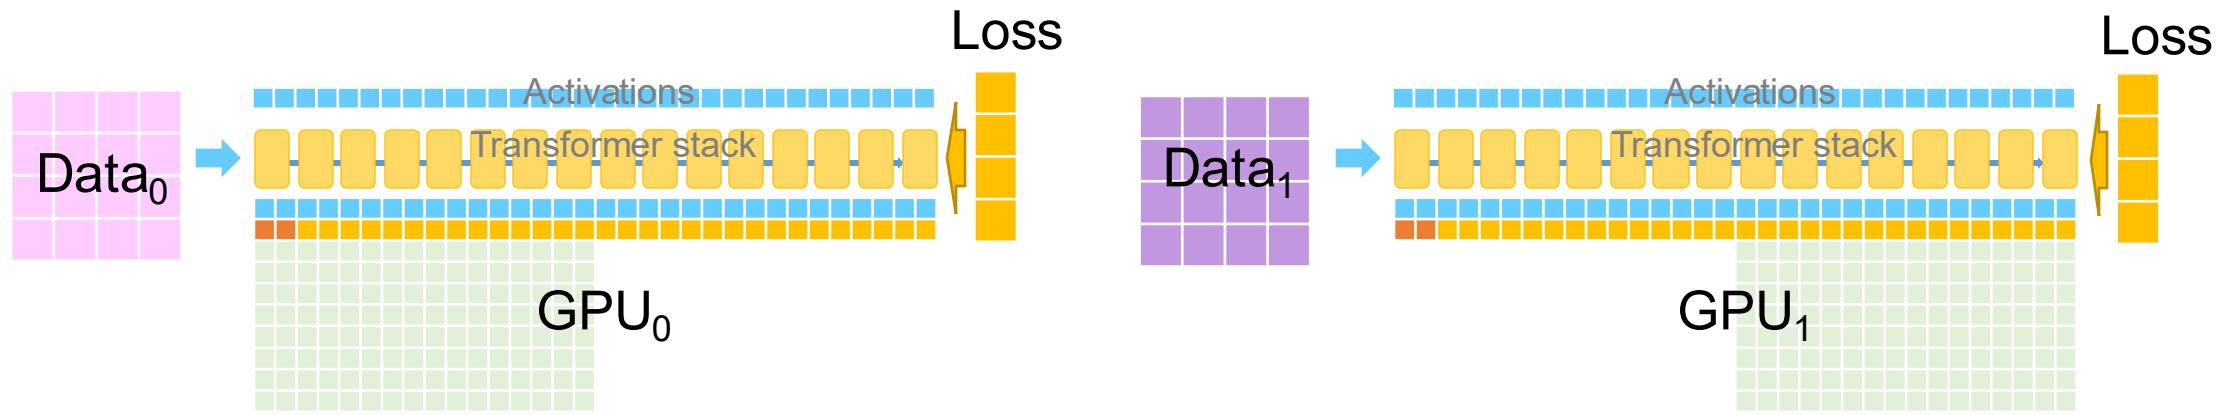
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

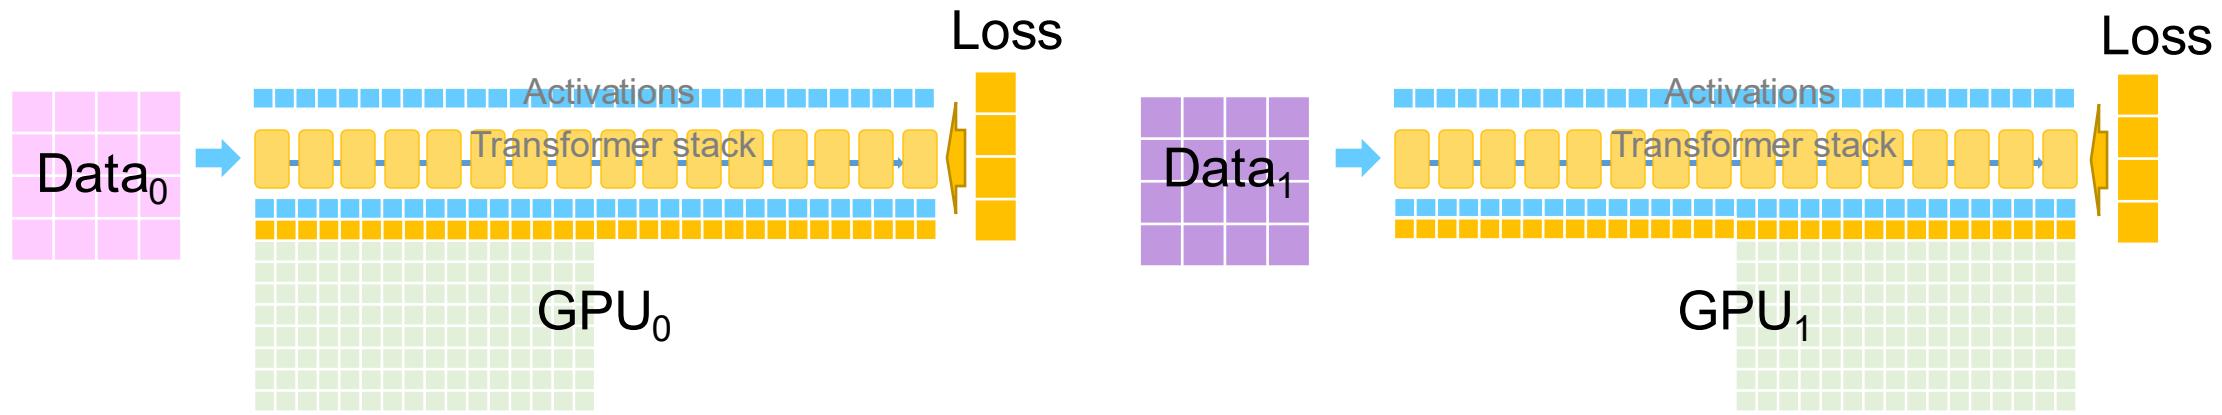
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients

■ Parameters ■ Gradients ■ Optimizer States

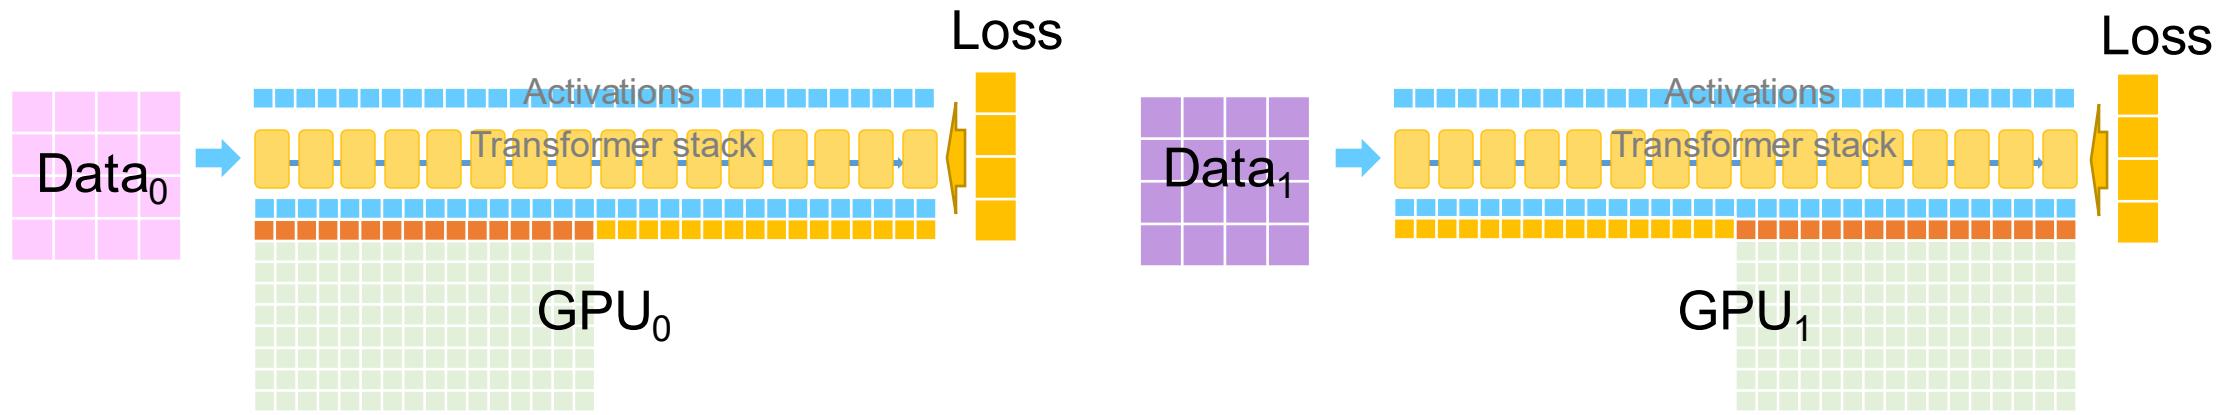
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average

■ Parameters ■ Gradients ■ Optimizer States

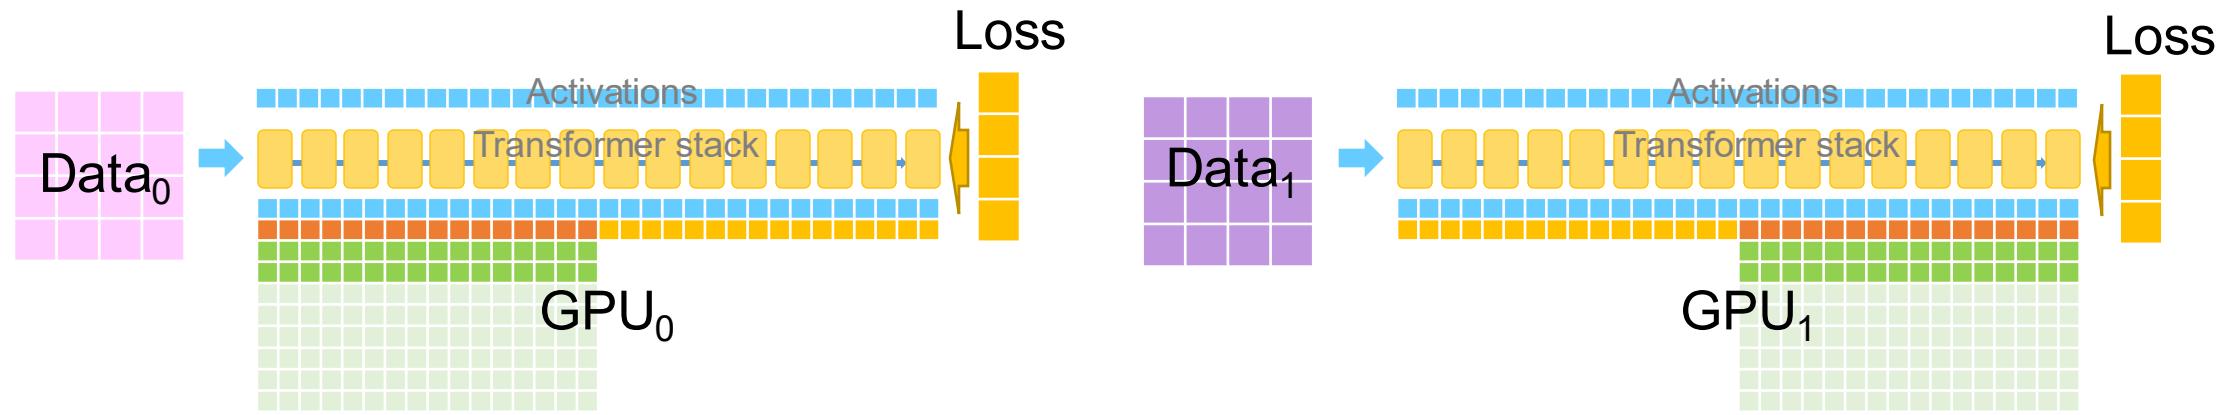
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average

■ Parameters ■ Gradients ■ Optimizer States

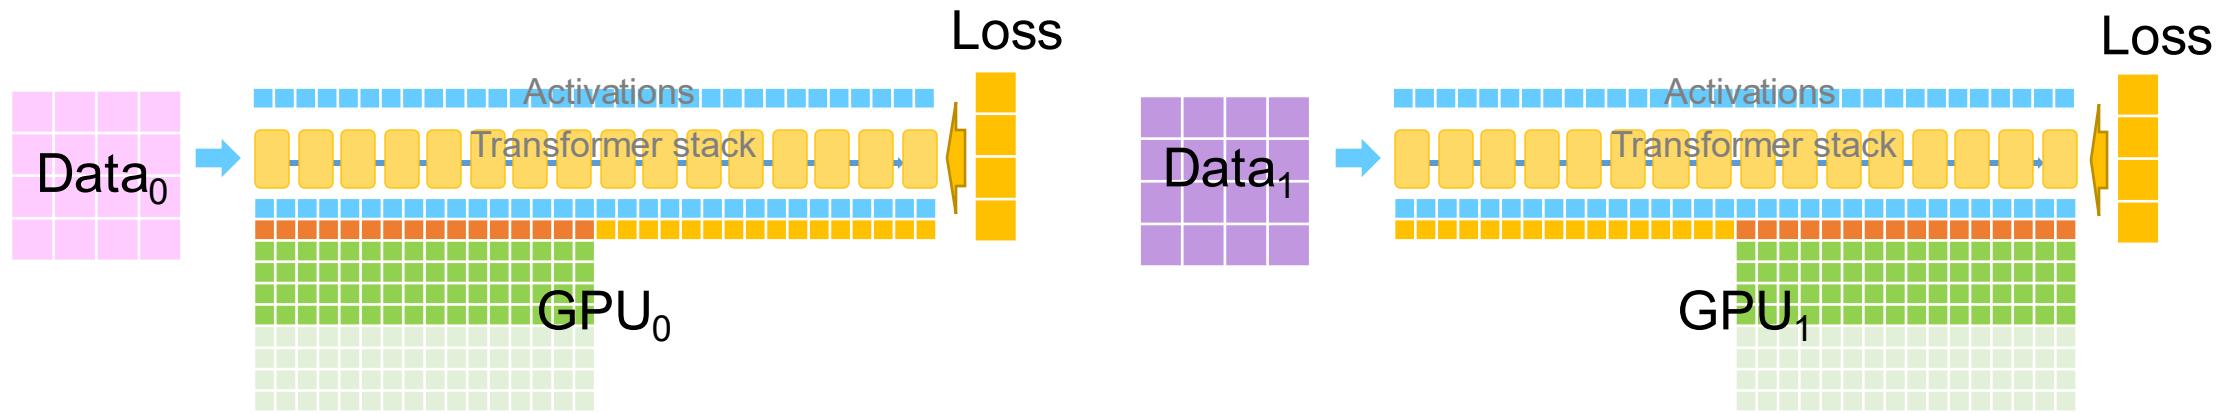
# ZeRO Stage 1: Partitioning Optimizer States



- ZeRO Stage 1
- Partitions optimizer states across GPUs
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer

■ Parameters ■ Gradients ■ Optimizer States

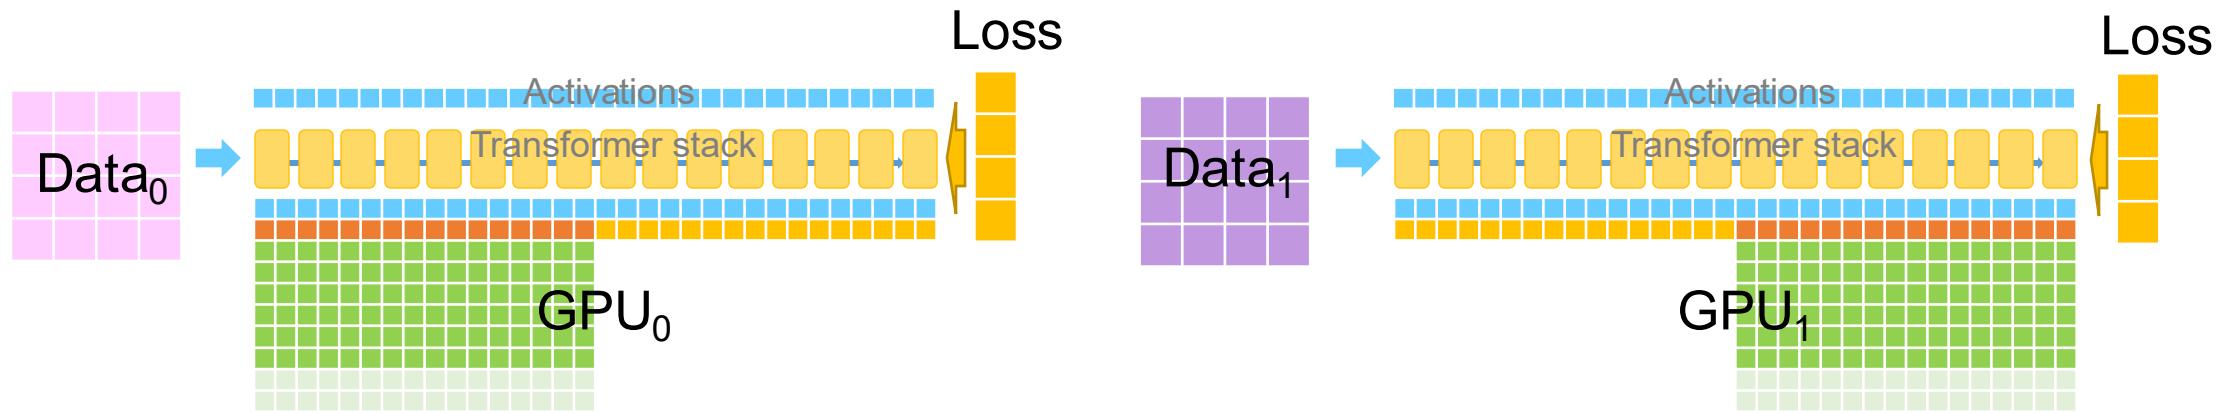
# ZeRO Stage 1: Partitioning Optimizer States



- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer

■ Parameters ■ Gradients ■ Optimizer States

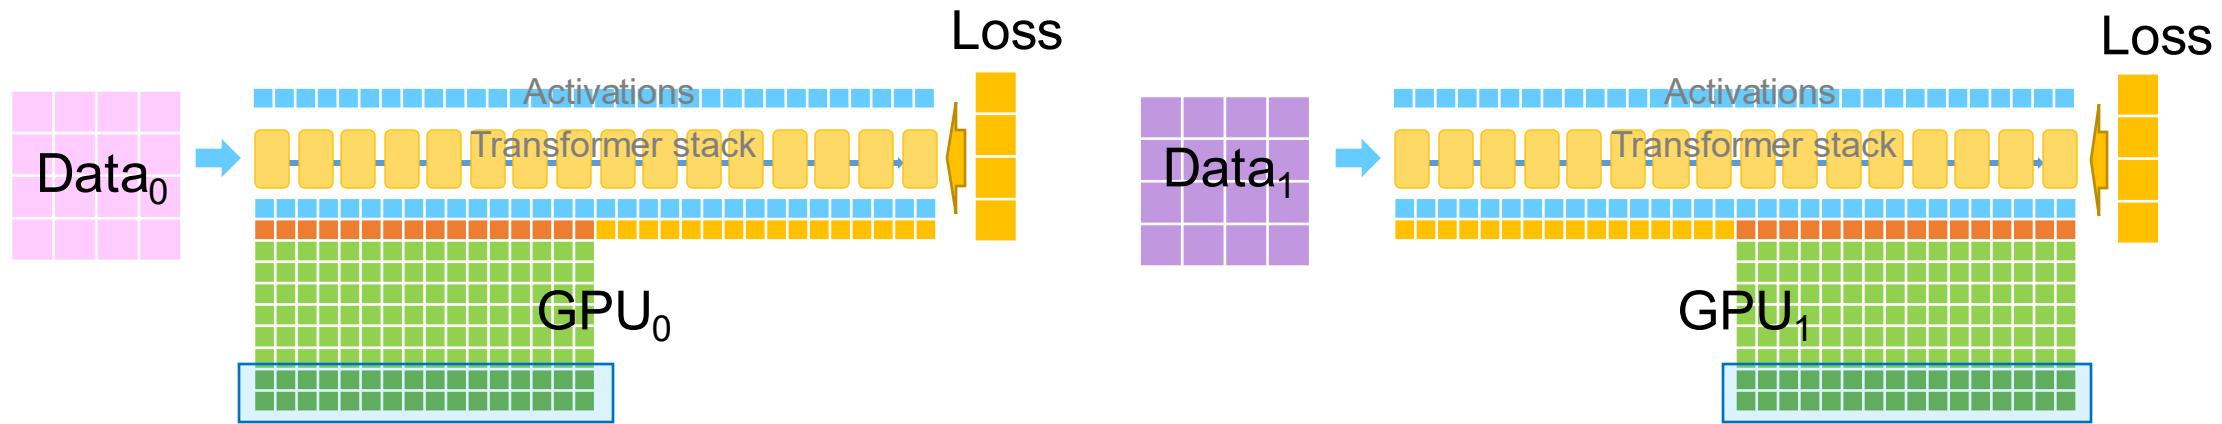
# ZeRO Stage 1: Partitioning Optimizer States



- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer

■ Parameters ■ Gradients ■ Optimizer States

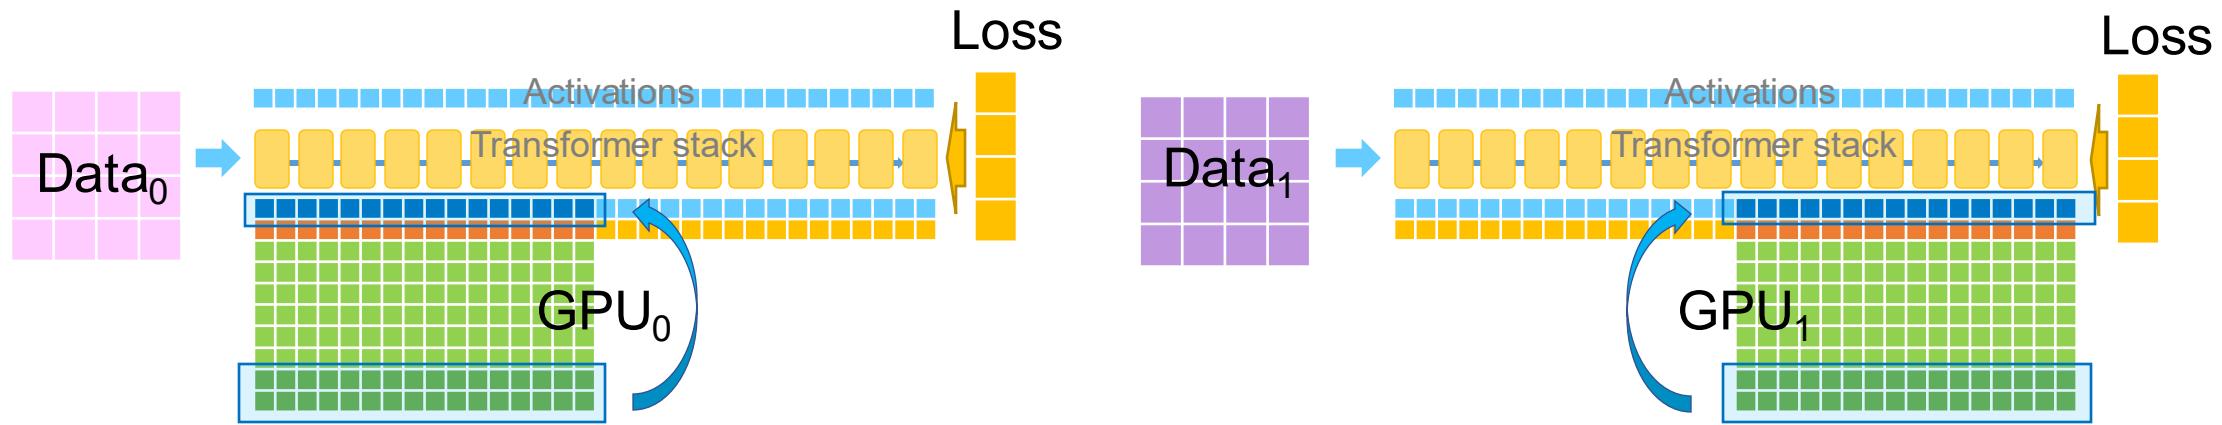
# ZeRO Stage 1: Partitioning Optimizer States



- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer

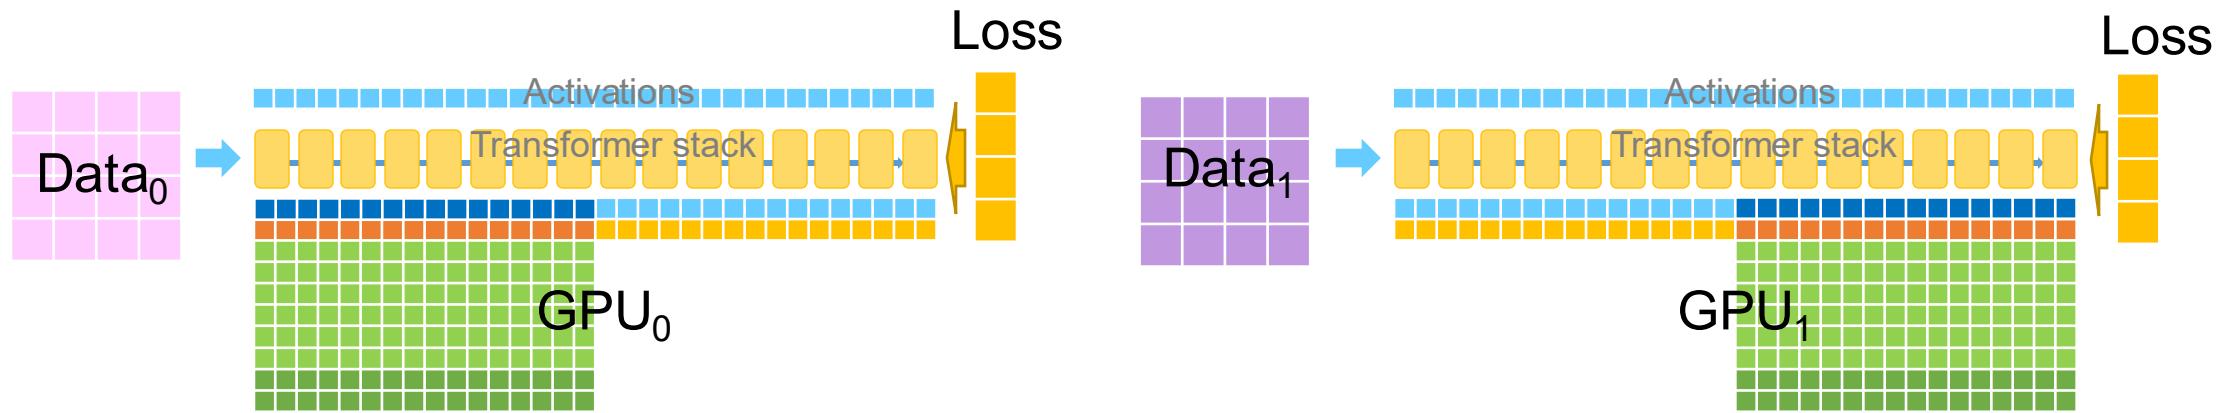
■ Parameters ■ Gradients ■ Optimizer States

# ZeRO Stage 1: Partitioning Optimizer States



- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer
- Update the FP16 weights

# ZeRO Stage 1: Partitioning Optimizer States



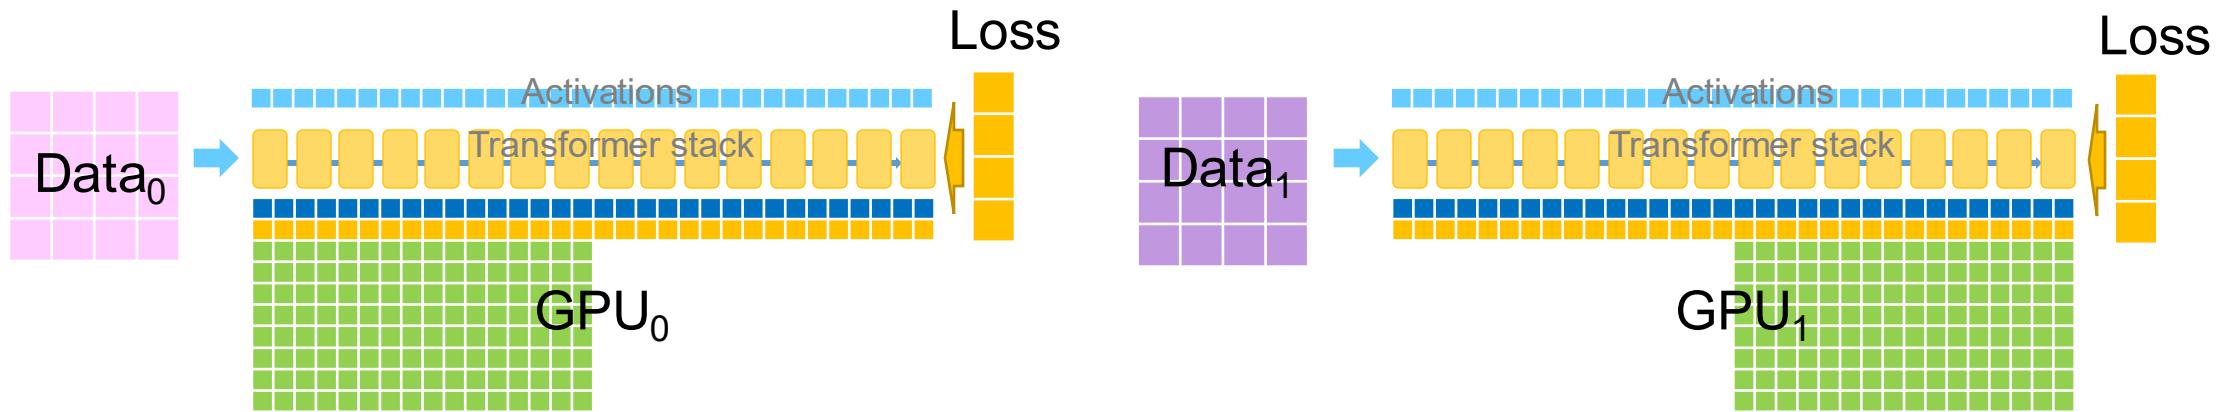
- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer
- Update the FP16 weights
- All Gather the FP16 weights to complete the iteration

Parameters

Gradients

Optimizer States

# ZeRO Stage 1: Partitioning Optimizer States



- Run Forward across the transformer blocks
- Backward propagation to generate FP16 gradients and AllReduce to average
- Update the FP32 weights with ADAM optimizer
- Update the FP16 weights
- All Gather the FP16 weights to complete the iteration

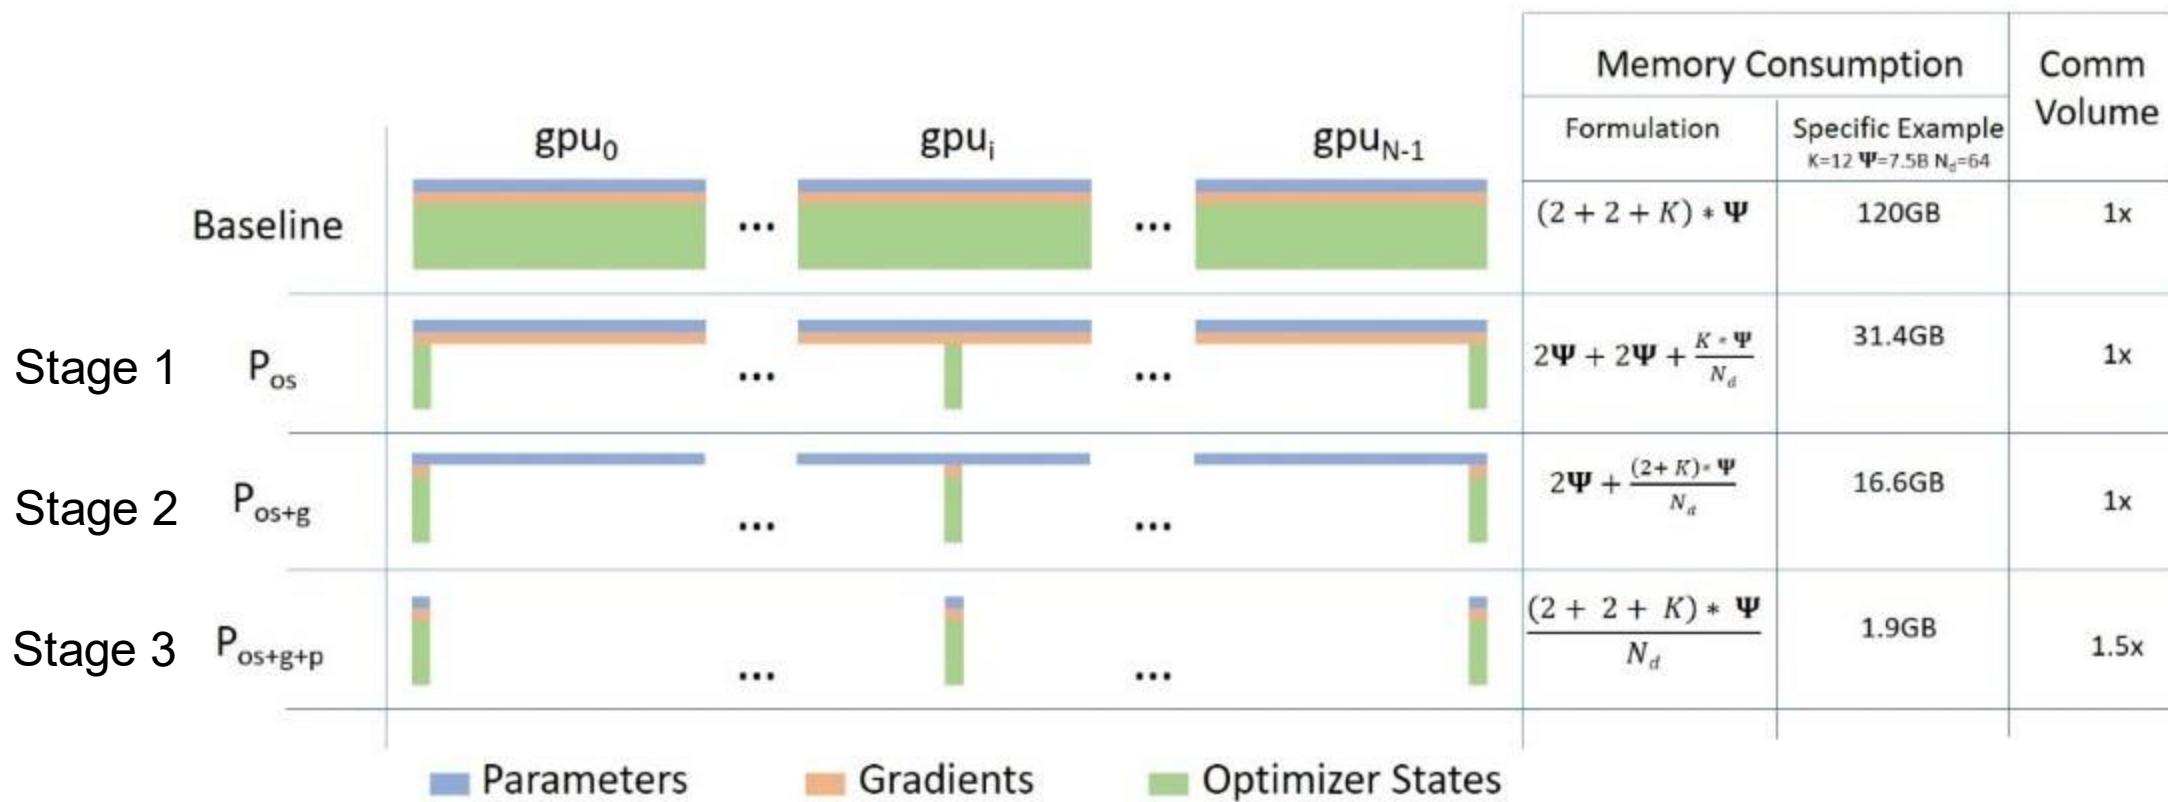
Parameters

Gradients

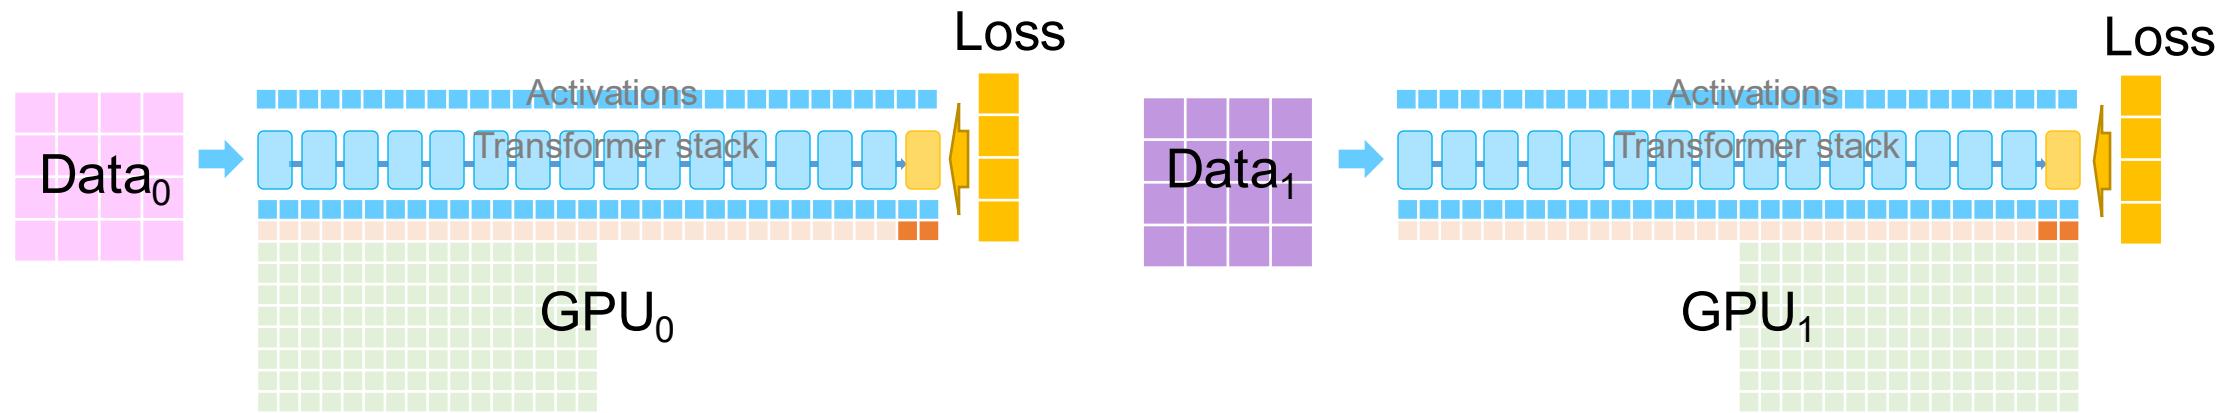
Optimizer States

# ZeRO: Zero Redundancy Optimizer

- Progressive memory savings and communication volume
- Turning NLP 17.2B is powered by Stage 1 and Megatron



# ZeRO Stage 2: Partitioning Gradients



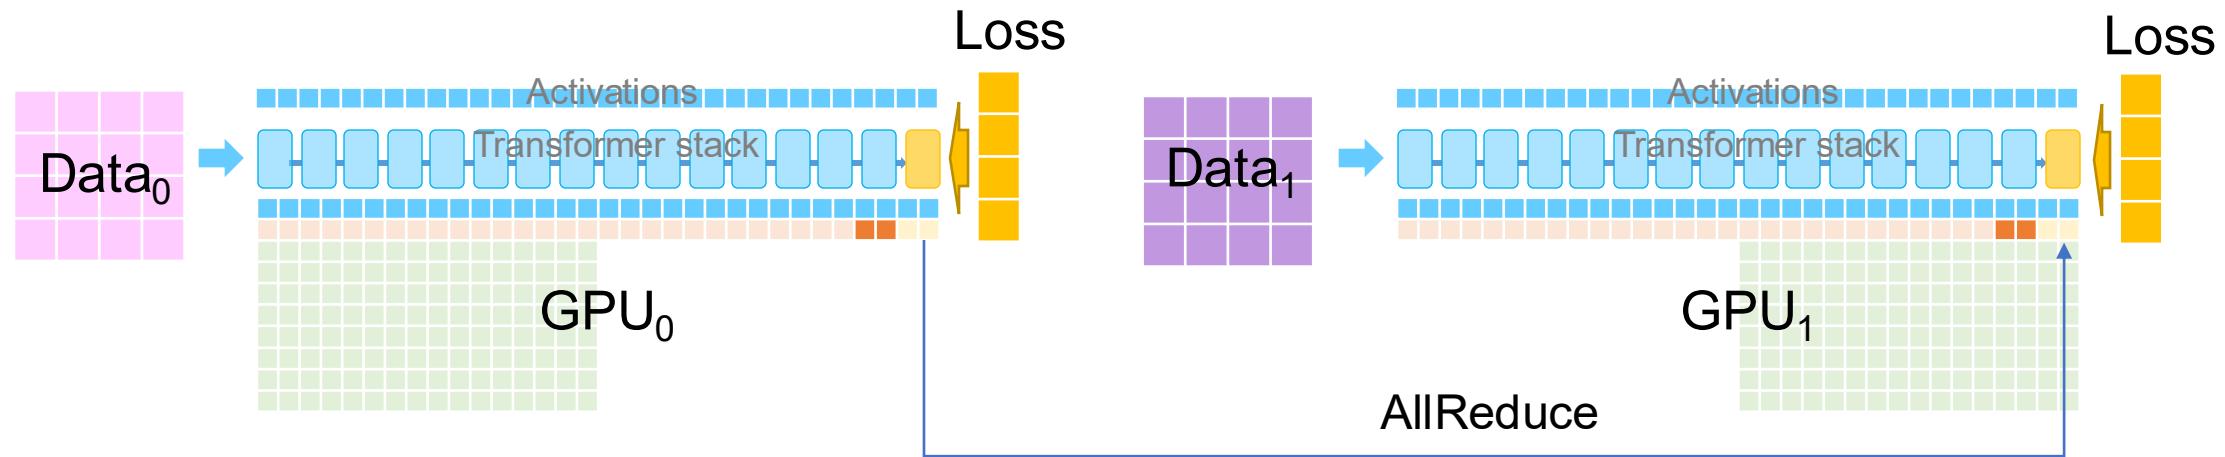
- Partitioning gradients across GPUs
- The forward process remains the same as stage 1

Parameters

Gradients

Optimizer States

# ZeRO Stage 2: Partitioning Gradients



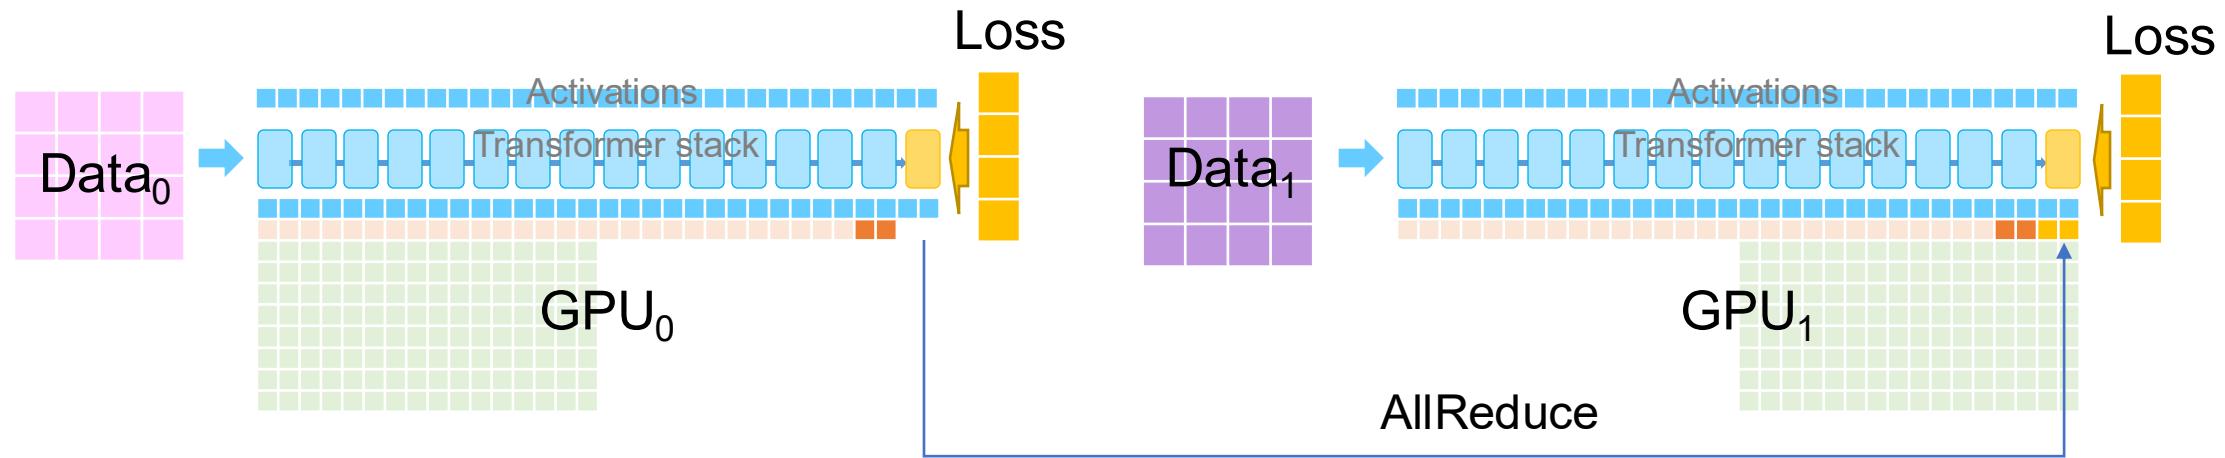
- Partitioning gradients across GPUs
- Perform AllReduce right after back propagation of each layer

Parameters

Gradients

Optimizer States

# ZeRO Stage 2: Partitioning Gradients



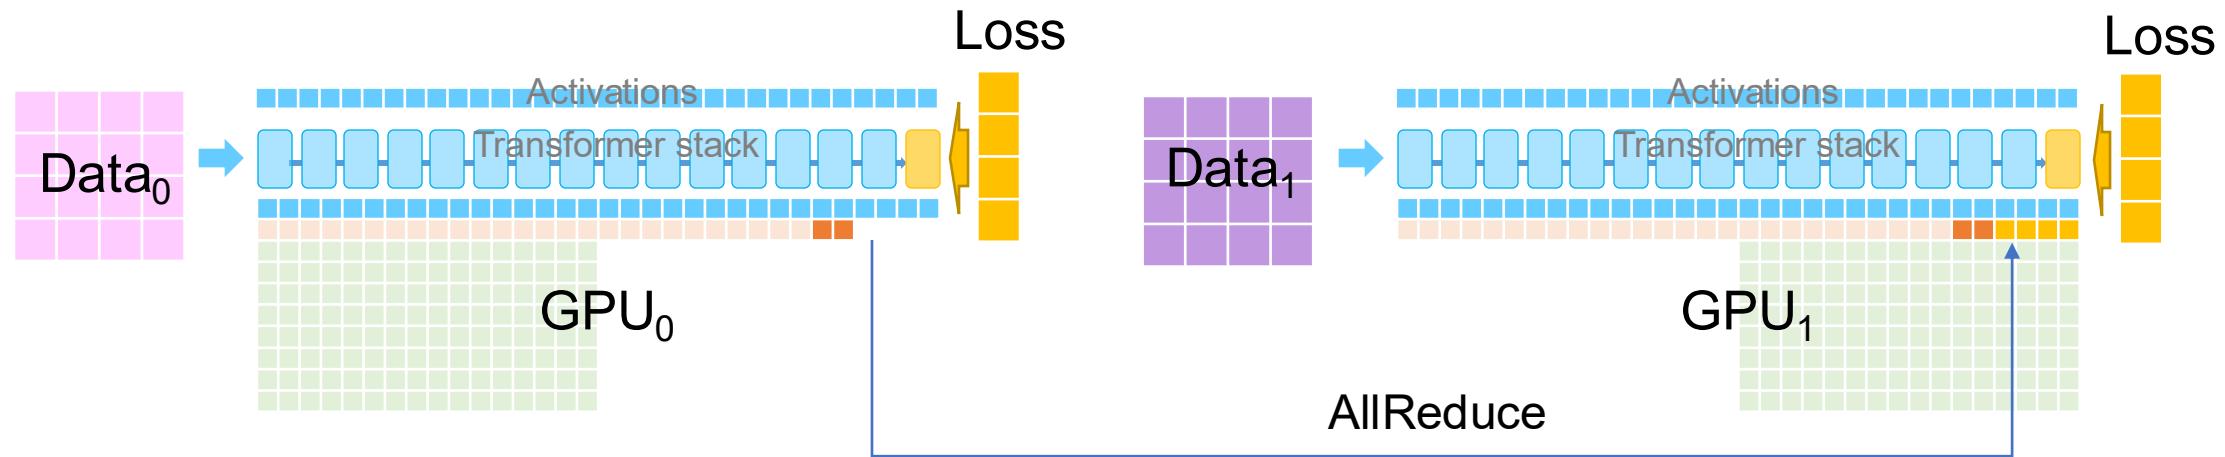
- Partitioning gradients across GPUs
- Only one GPU keeps the gradients after AllReduce

Parameters

Gradients

Optimizer States

# ZeRO Stage 2: Partitioning Gradients



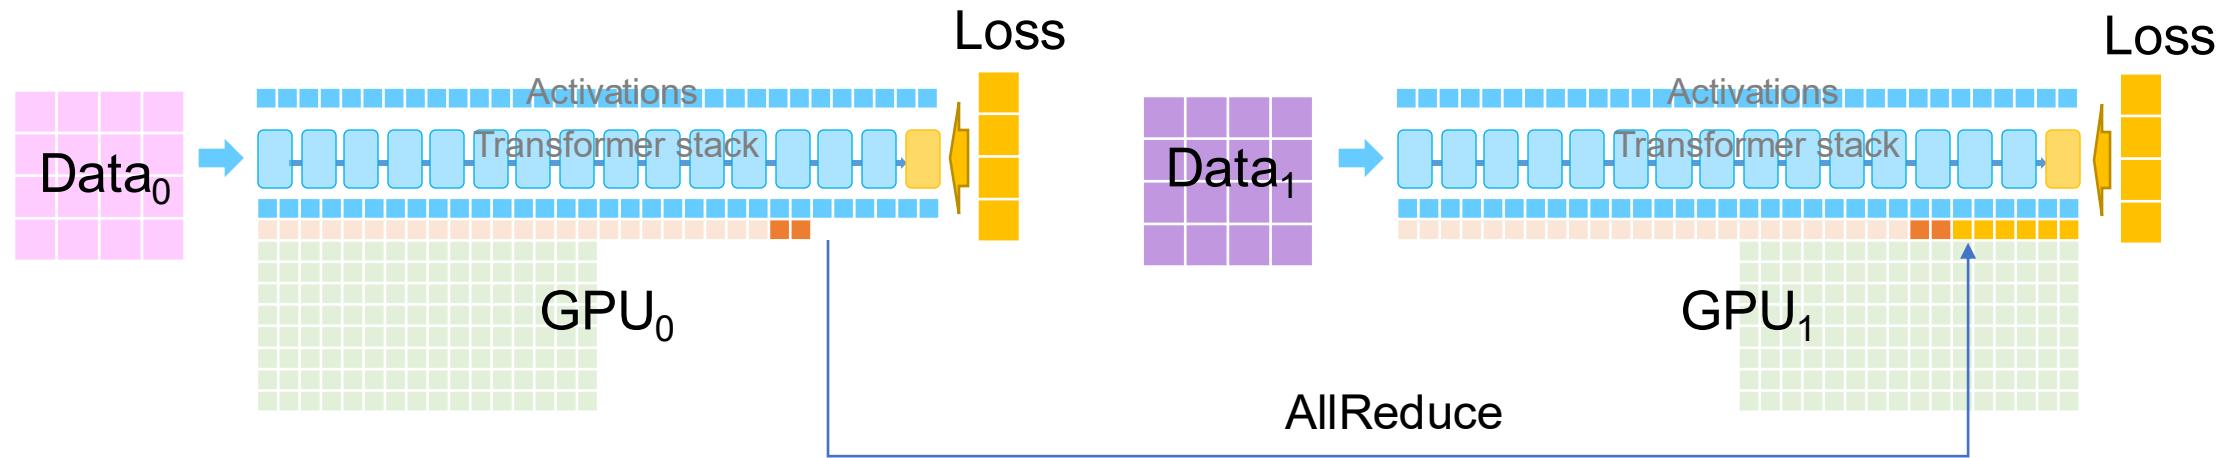
- Partitioning gradients across GPUs
- Reduce gradients on GPUs responsible for updating parameters

Parameters

Gradients

Optimizer States

# ZeRO Stage 2: Partitioning Gradients



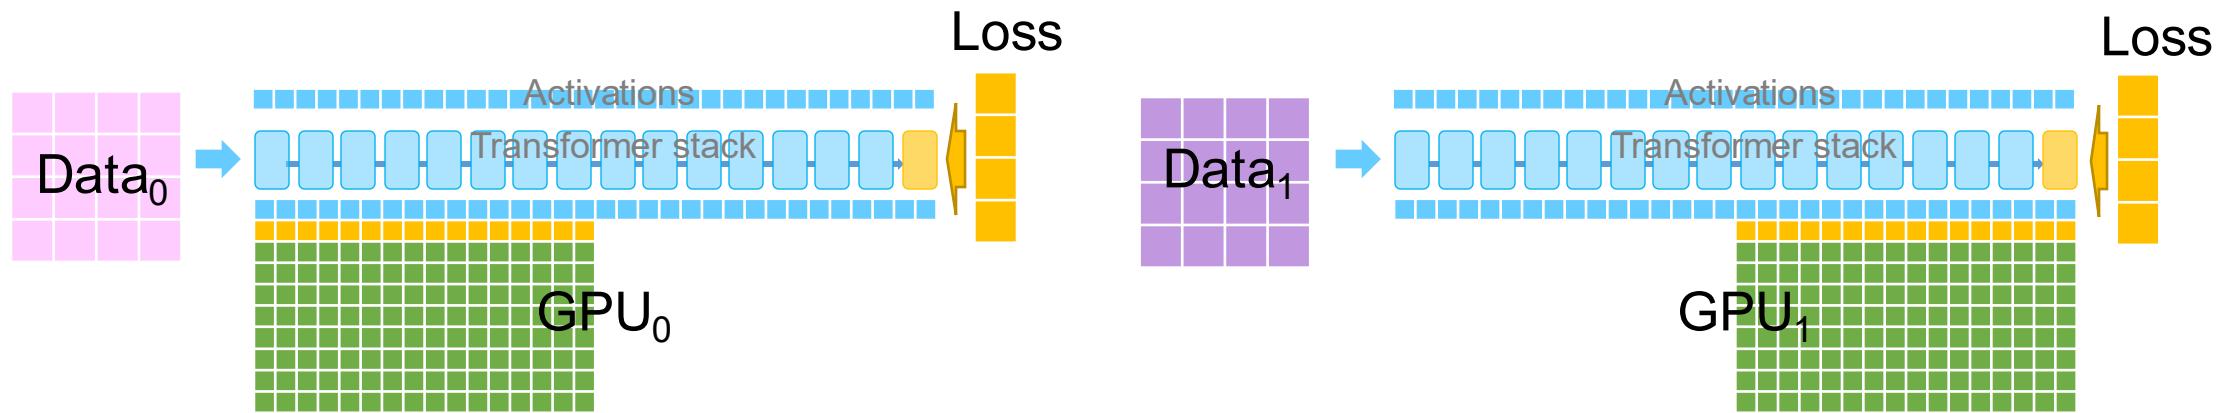
- Partitioning gradients across GPUs
- Reduce gradients on GPUs responsible for updating parameters

■ Parameters

■ Gradients

■ Optimizer States

# ZeRO Stage 2: Partitioning Gradients



- Partitioning gradients across GPUs
- Reduce gradients on GPUs responsible for updating parameters

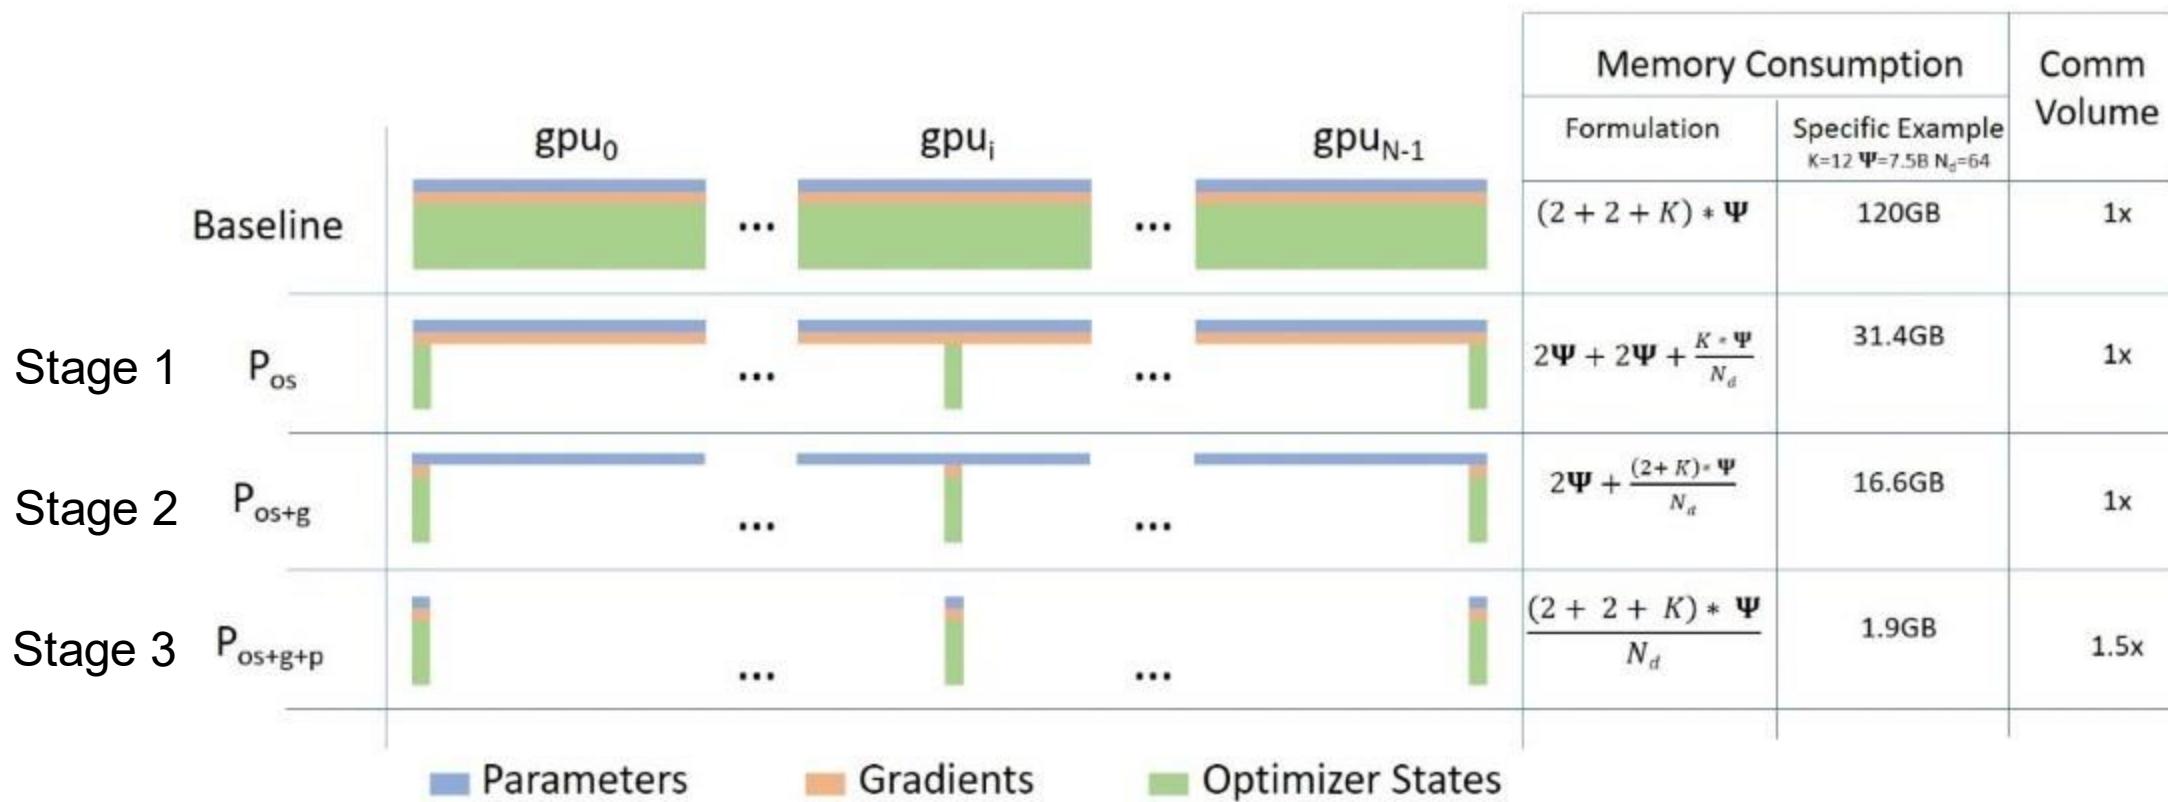
Parameters

Gradients

Optimizer States

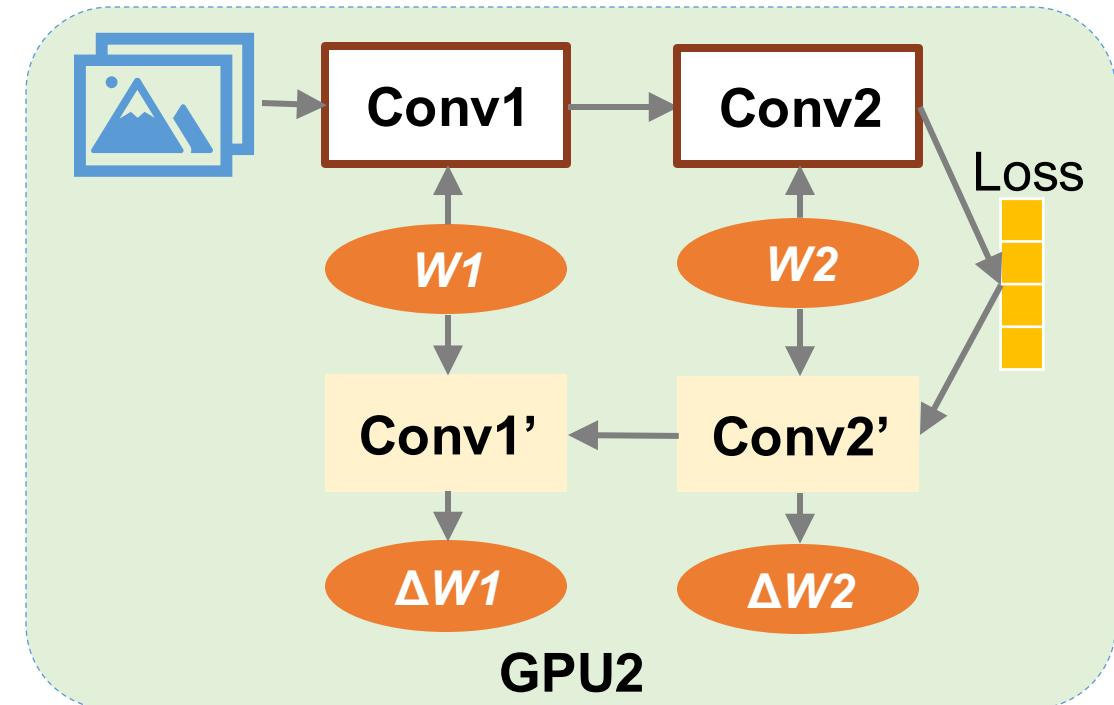
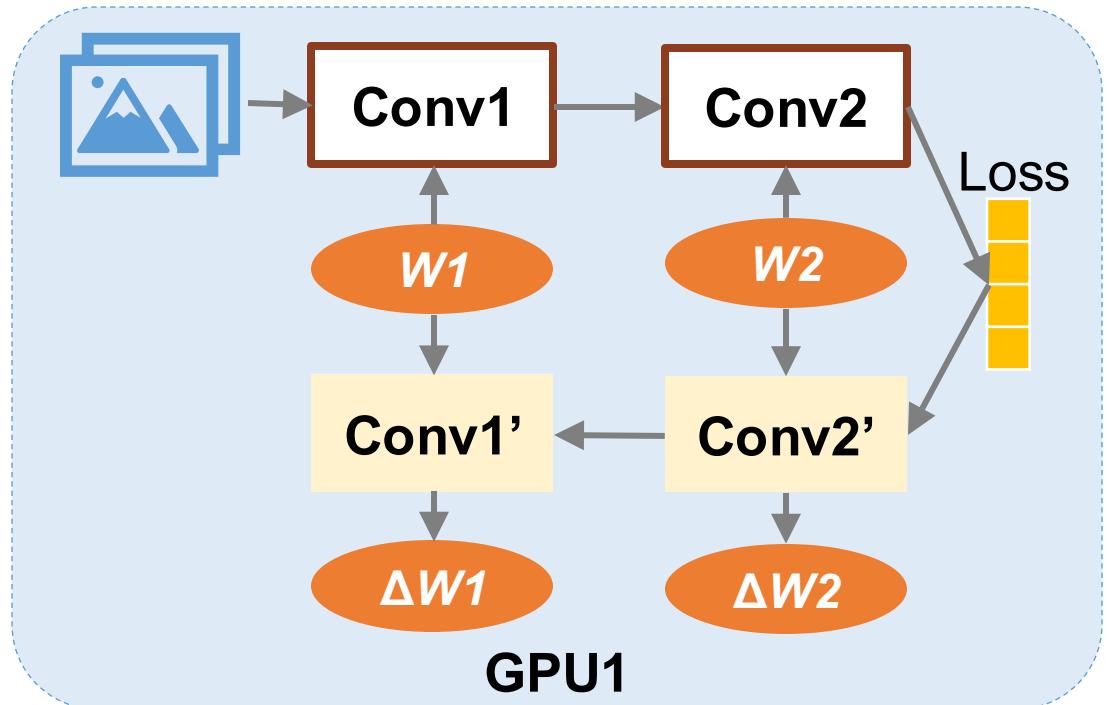
# ZeRO: Zero Redundancy Optimizer

- Progressive memory savings and communication volume
- Turning NLP 17.2B is powered by Stage 1 and Megatron



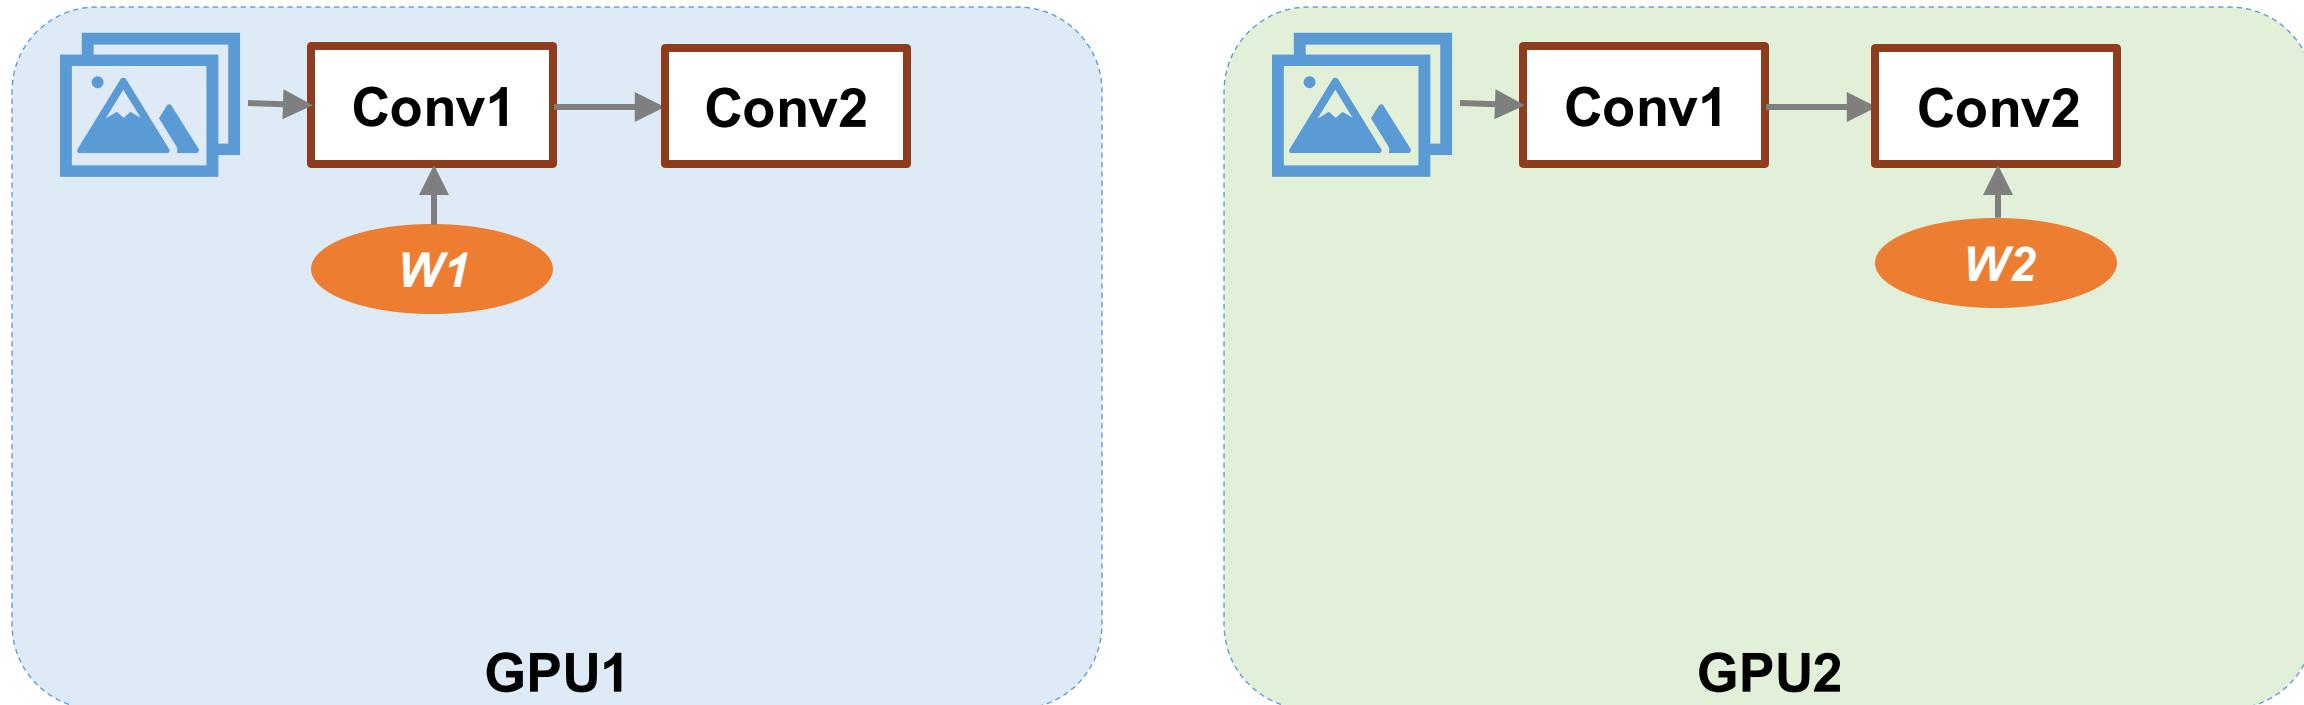
## ZeRO Stage 3: Partitioning Parameters

- In data parallel training, all GPUs keep all parameters during training



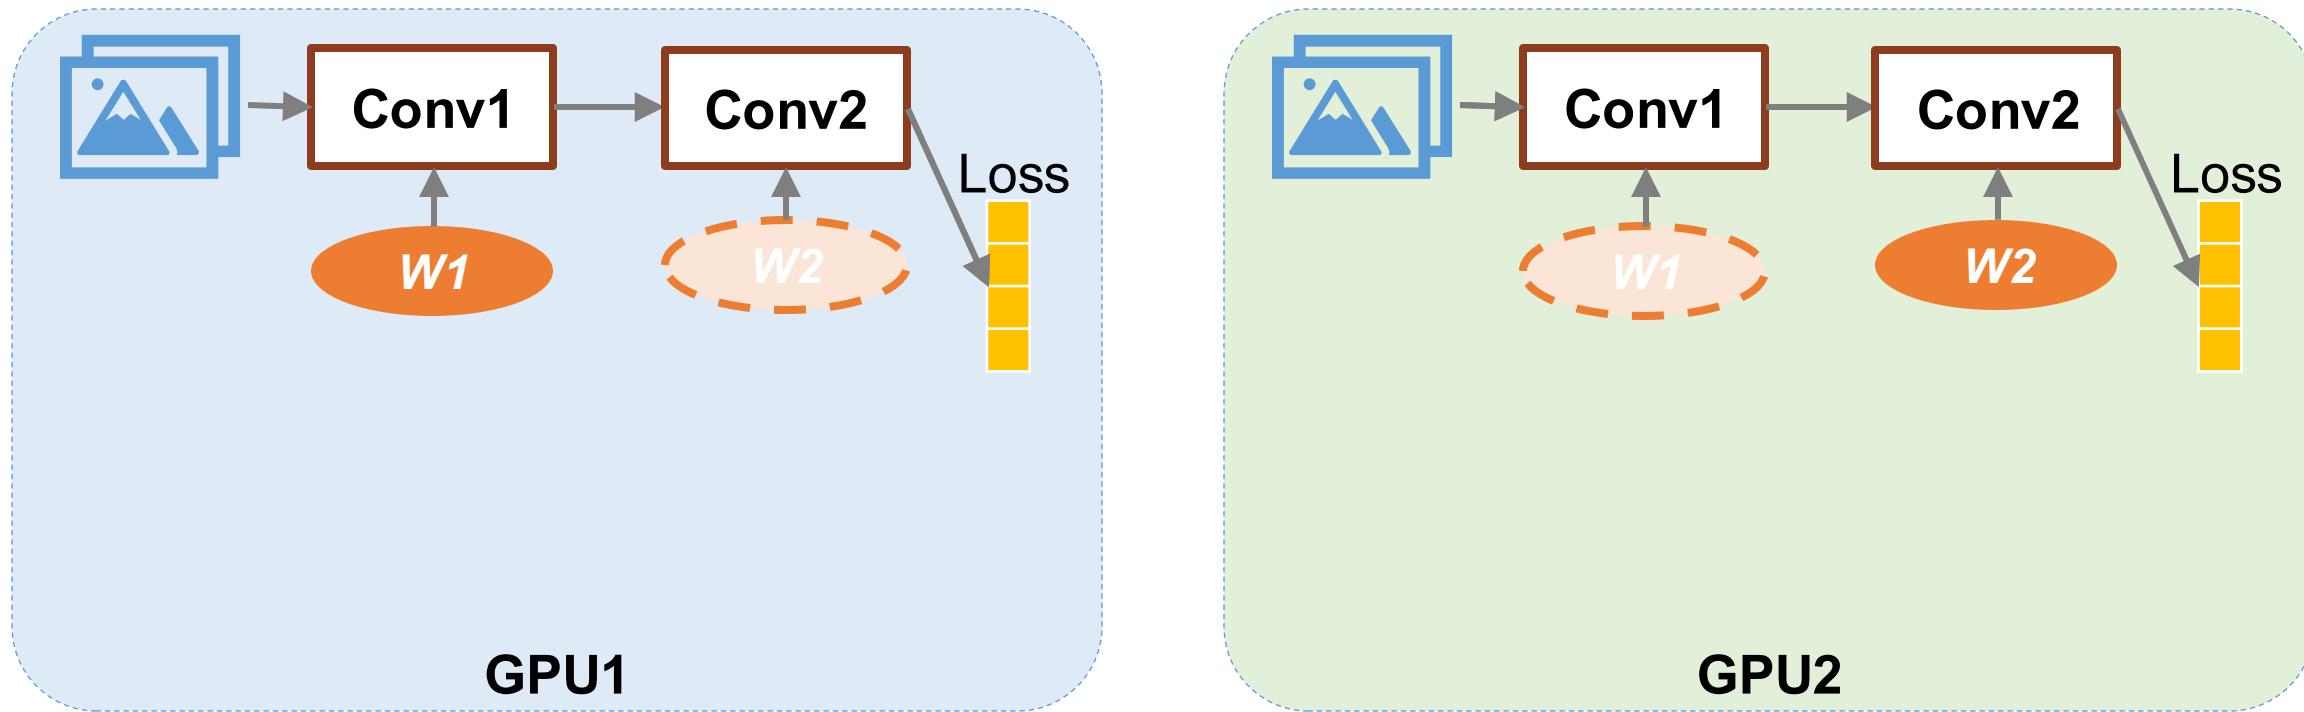
## ZeRO Stage 3: Partitioning Parameters

- In ZeRO, model parameters are partitioned across GPUs



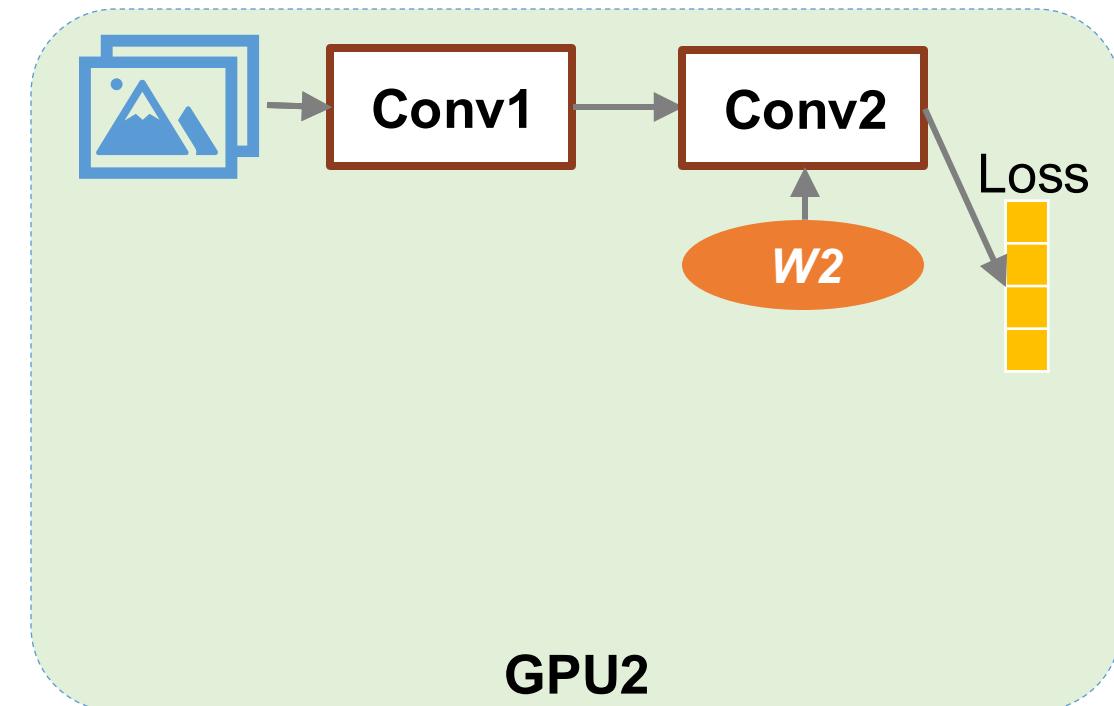
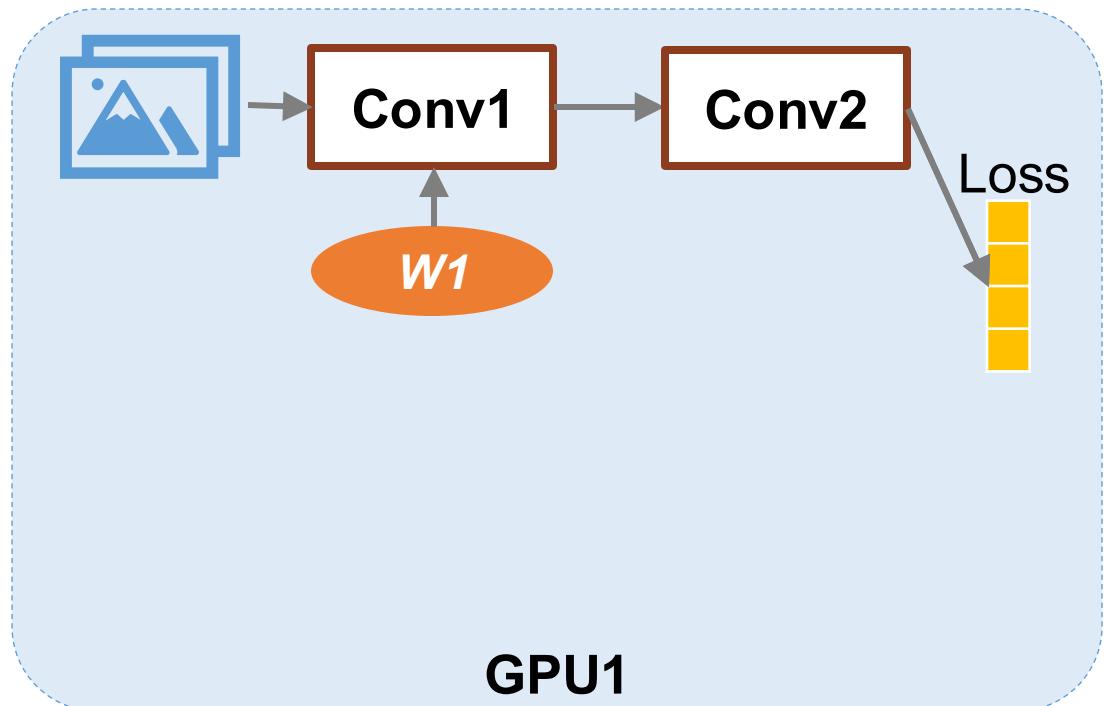
## ZeRO Stage 3: Partitioning Parameters

- In ZeRO, model parameters are partitioned across GPUs
- GPUs broadcast their parameters during forward



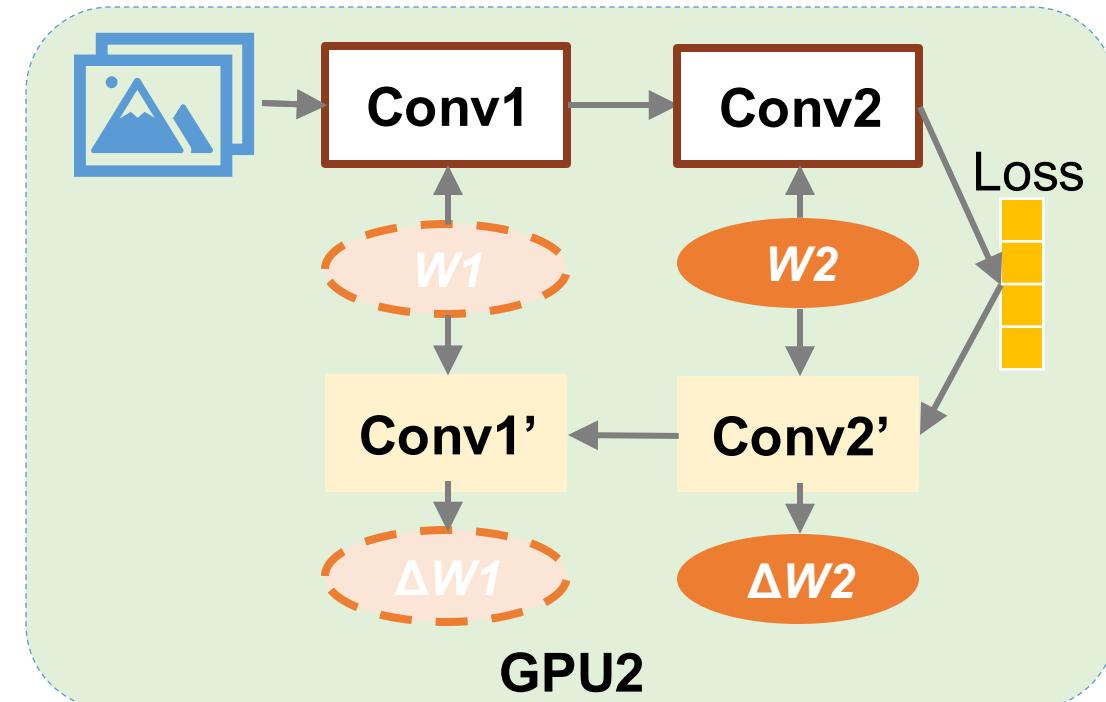
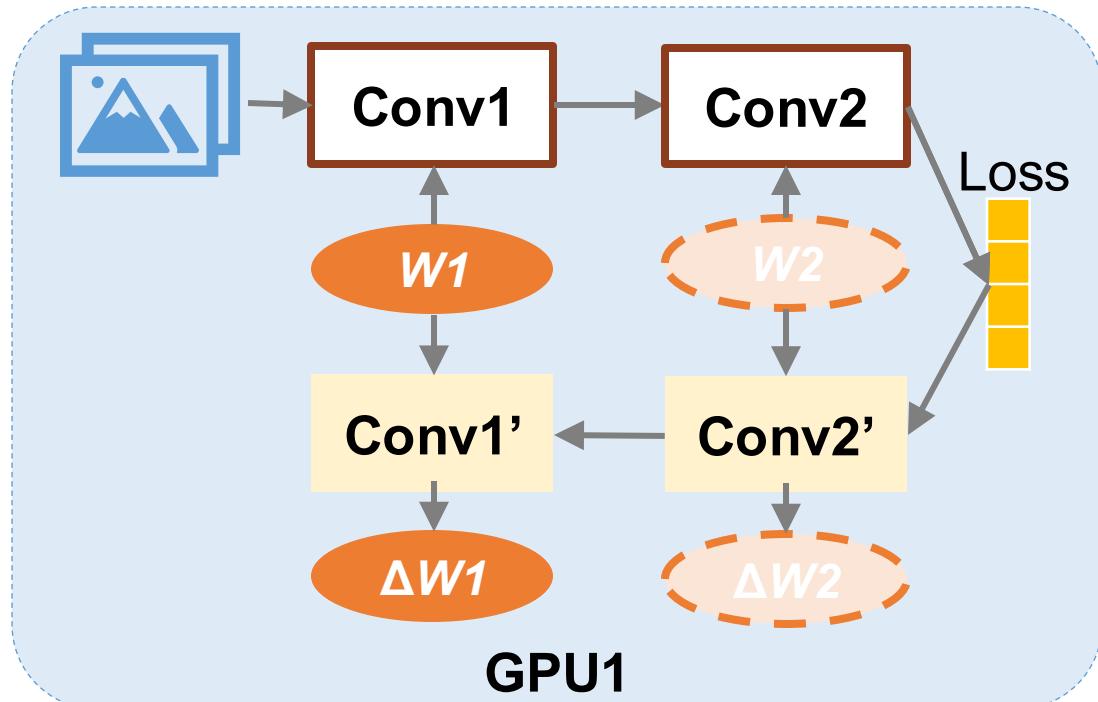
## ZeRO Stage 3: Partitioning Parameters

- In ZeRO, model parameters are partitioned across GPUs
- Parameters are discarded right after use



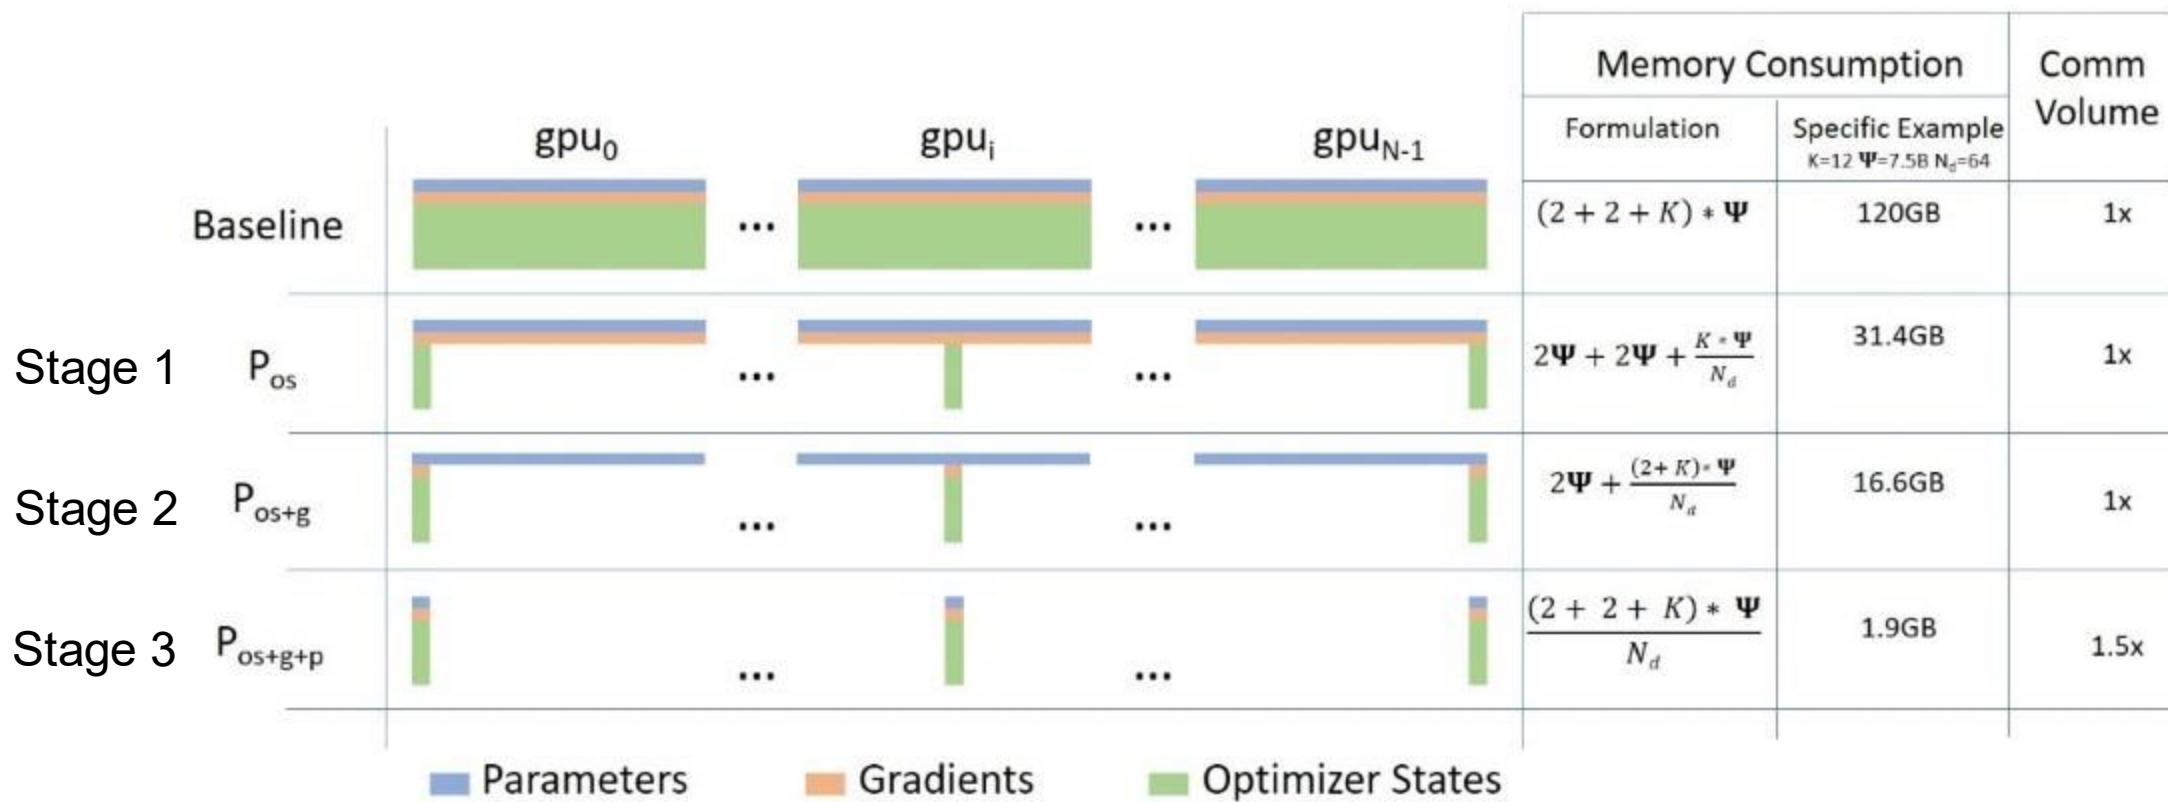
## ZeRO Stage 3: Partitioning Parameters

- In ZeRO, model parameters are partitioned across GPUs
- GPUs broadcast their parameters again during backward



# ZeRO: Zero Redundancy Optimizer

- ZeRO has three different stages
- Progressive memory savings and communication volume



# Summary

- Data-parallel training
  - Parameter server
  - Ring AllReduce
  - Tree AllReduce
  - Butterfly AllReduce
- ZeRO: zero redundancy optimizer