

15-780 – Graduate Artificial Intelligence: Markov Decision Processes

Aditi Raghunathan
Carnegie Mellon University

Search (recap)

- Want to reach goal from start state
- Each action takes us to a **deterministic successor** state
- Various algorithms
 - Uninformed search: DFS, BFS, iterative deepening, uniform cost search
 - Informed search: A*, heuristics, relaxations

Decision making under uncertainty

Real-world is not deterministic, there is **randomness** or **uncertainty**

Markov Decision Processes (MDPs) and their extensions provide a general way to think about how we can act optimally under uncertainty

MDPs introduced in 1950s – 60s

Applications: Robotics, self-driving cars, video games, robot soccer, scheduling and many more..

Example MDP



Example: dice game

For each round $r = 1, 2, \dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
 - If the dice results in 1 or 2, we end the game.
 - Otherwise, continue to the next round.

What policy should you follow?

MDP intuition

States encode all the information of a system needed to determine how it will evolve when taking actions

The system is governed via *transition probabilities*: they only depend on the current state and action; not on the history

Agent starts at a start state, and the goal is to take actions that maximize *expected reward*

MDPs formal definition

States: $s \in S$ assumed to be discrete

Actions: $a \in \mathcal{A}$ (assumed to be discrete)

Transition probabilities: distribution over next states given current state and current action

$T(s, a, s')$ is probability of next state being s' when taking action a at state s

Rewards: mapping states or transitions to reals $R(s, a, s')$

Start state, end state, discount factor γ (default 1)

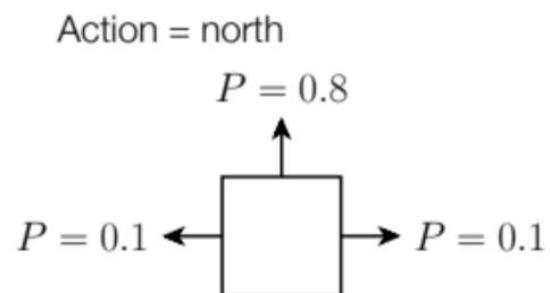
Gridworld domain

Simple grid world with a goal state with reward and a “bad state” with reward -100

Actions move in desired direction with probability 0.8 and one of two perpendicular directions with probability 0.1 each

Taking an action that bumps into a wall leaves agent where it is

0	0	0	1
0	-	0	-100
0	0	0	0



Policy and policy evaluation

A **policy** π is a mapping from each state $s \in S$ to an action $a \in \mathcal{A}$ (or distribution of actions)

Suppose you followed a path $s_0, a_1, r_1, s_1, a_2, r_2, s_2, a_2, r_2, s_3 \dots$; the expected sum of discounted rewards is $r_1 + \gamma r_2 + \gamma^2 r_3 + \dots$

Value function $V_\pi: S \rightarrow R$ such that $V_\pi(s)$ gives the expected sum of discounted rewards when following policy π from state s

Q-value of a policy $Q_\pi: S \times A \rightarrow R$ such that $Q_\pi(s, a)$ is the expected sum of discounted rewards when taking action a from state s and then following π

Recursive definitions

$$V_\pi(s) = Q_\pi(s, \pi(s)) \text{ otherwise } (= 0 \text{ if } s \text{ is the end-state})$$

$$Q_\pi(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_\pi(s')]$$

Algorithm for policy evaluation: start with arbitrary initialization $V_\pi(s) = 0 \forall s$ and then apply the recursion

Example MDP



Example: dice game

For each round $r = 1, 2, \dots$

- You choose **stay** or **quit**.
- If **quit**, you get \$10 and we end the game.
- If **stay**, you get \$4 and then I roll a 6-sided dice.
 - If the dice results in 1 or 2, we end the game.
 - Otherwise, continue to the next round.

Evaluate the “always stay” policy

What is $V_\pi(\text{stay})$?

Dice game

For *always stay* policy π

$$\begin{aligned} V_\pi(\text{end}) &= 0 \\ V_\pi(\text{stay}) &= \frac{1}{3}(4 + V_\pi(\text{end})) + \frac{2}{3}(4 + V_\pi(\text{stay})) \end{aligned}$$

Solve the recurrence

$$V_\pi(\text{stay}) = 12$$