

15-213 Recitation Caches & C Review

Your TAs
Friday, September 26th

Reminders

- **attacklab** was due *yesterday*.
- **cachelab** was released yesterday, and is due ***Thursday, October 9th.***
- Part 1 of the Midterm (take-home) is coming out October 1st, and is due ***October 8th***

Agenda

- **Intro to cachelab**
- **Review: Cache Concepts**
- **Review: Programming in C**
- **Parsing Command-Line Arguments with getopt ()**
- **Cache Practice Problems**

cachelab

cachelab: Overview

- First project-based assignment:
 - You'll write a *cache simulator* in C from scratch!
- Take in parameters defining the cache structure (**s**, **E**, **b**).
- Read a “trace file” of memory accesses and simulate them.
- After simulating those accesses, return the number of hits, misses, evictions, etc.

Using Git

- This is the first lab where we will use Git!
 - Note you will download the assignment through Github
- Try to keep good version control practice!
 - Commit after every “good fix”!
 - Good version control can help you recover from “bad mistakes”!

Review: Cache Concepts

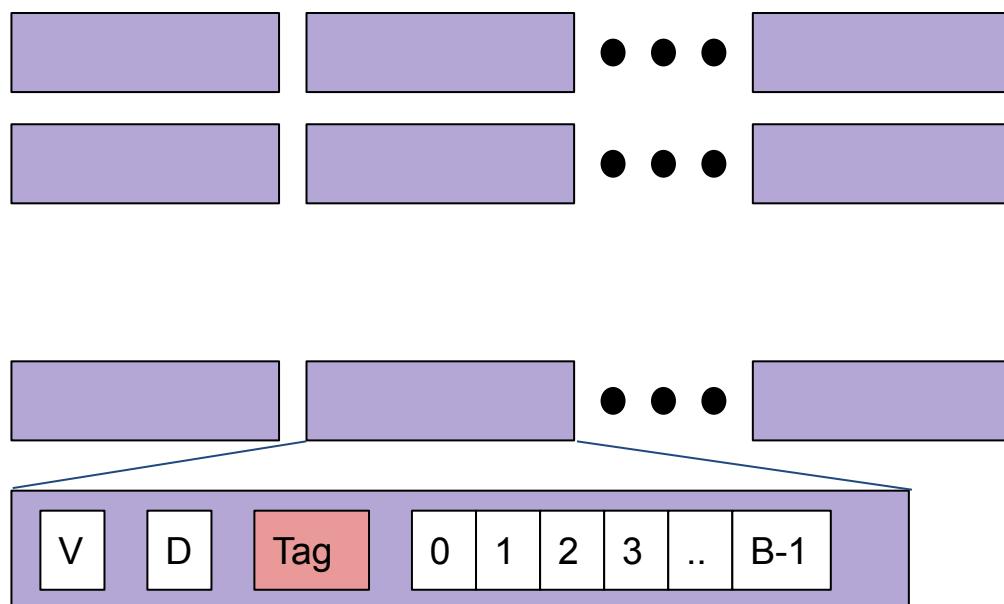
Cache Concepts: Configurations

- Your cache simulators need to support *parameters* (**s**, **E**, **b**) that allow the user to configure the layout of the cache.
- But what do these parameters mean?
- Let's review how a cache is organized!

Cache Organization

- A cache is composed of *sets*
- Each set is composed of some number of *lines*
- Each line stores the cached data itself, as well as information used by the cache.

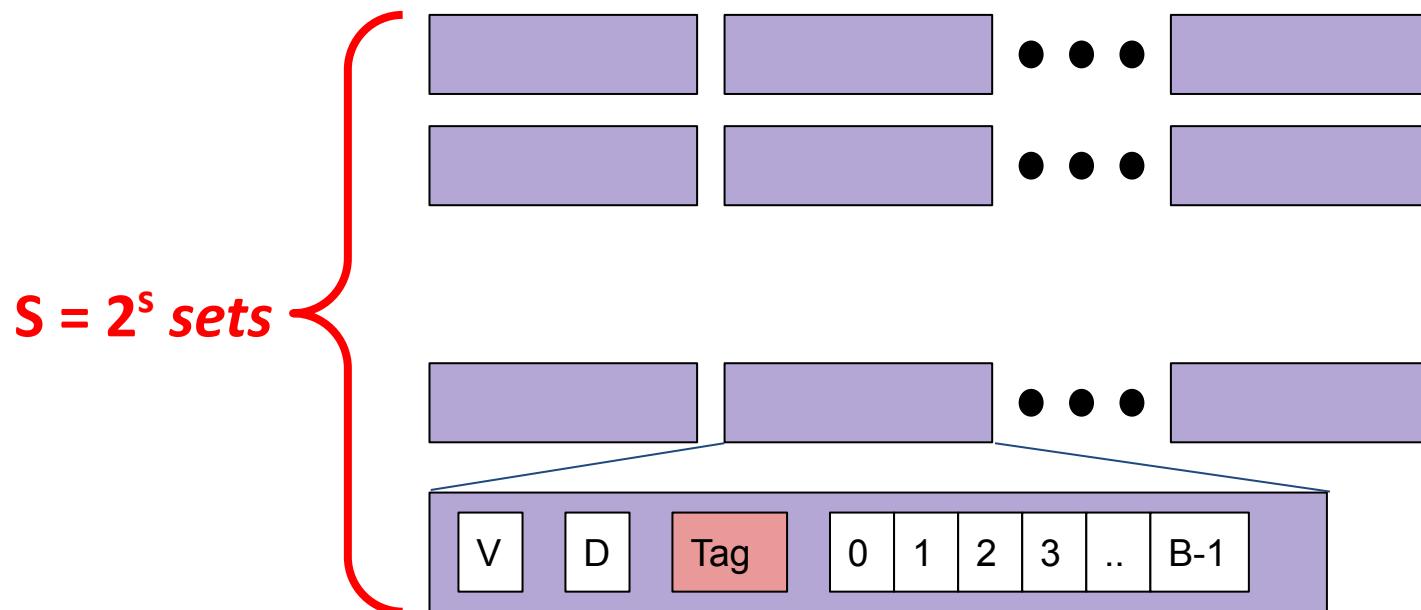
Note: don't need to store data for **cachelab**.



Cache Organization

- A cache is composed of sets

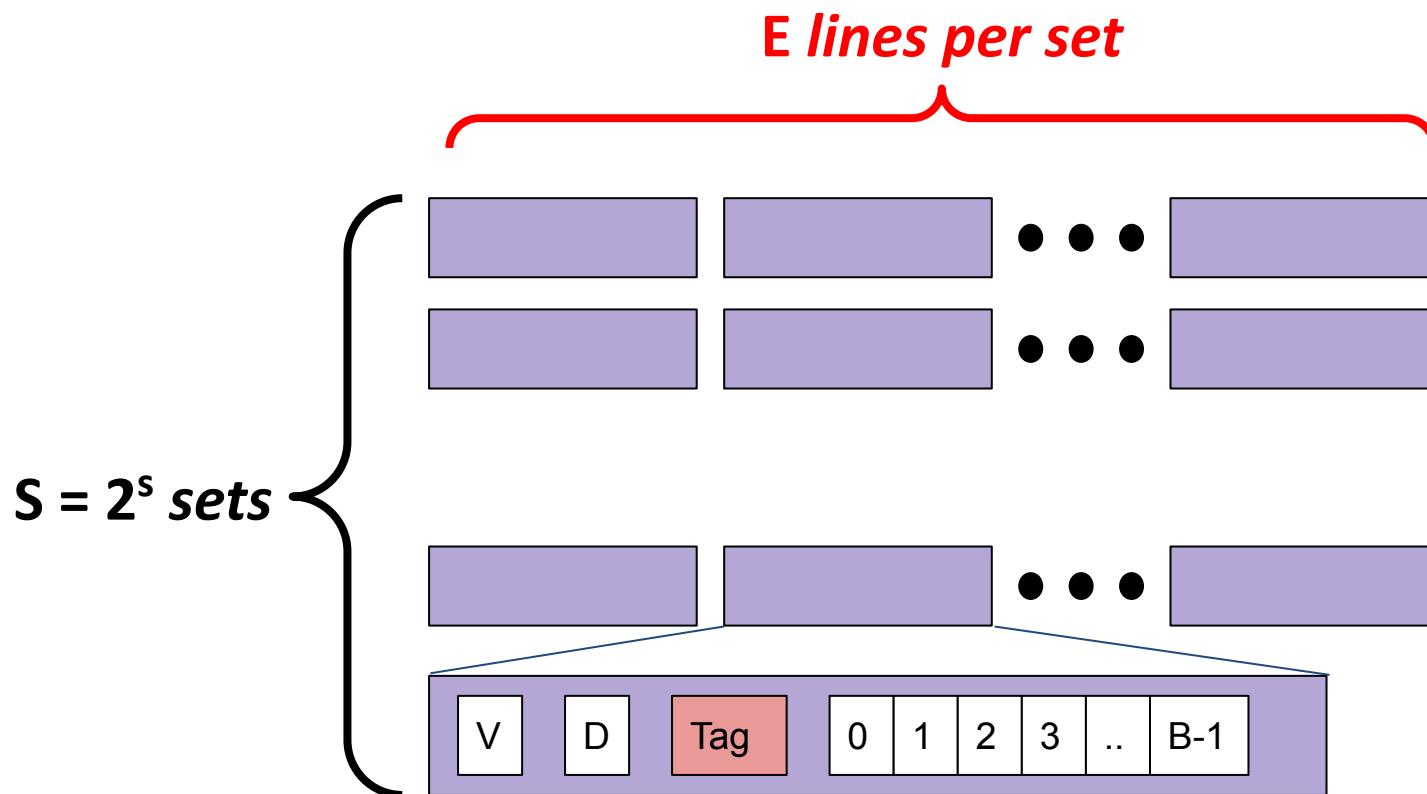
s – Number of set *bits*
 $S = 2^s$ – Number of sets



Cache Organization

- Each set is composed of *lines*

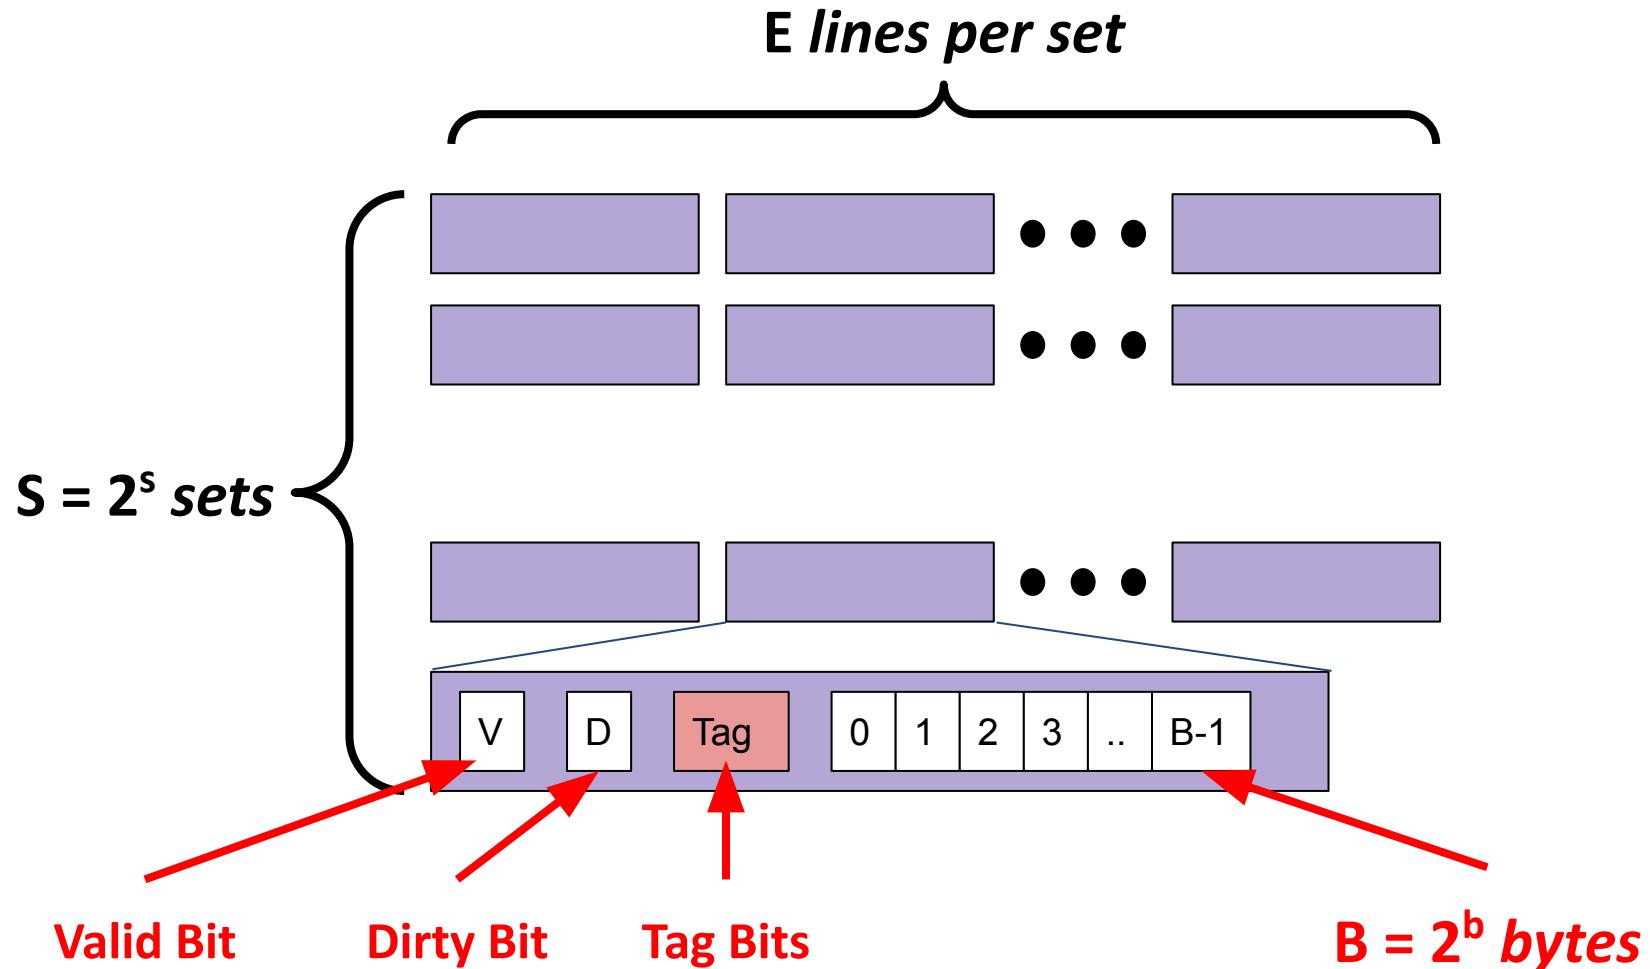
E – Number of *lines per set*



Cache Organization

- Each line stores data

b – Number of block offset bits
 $B = 2^b$ – Block Size



Cache Concepts: Cache Read

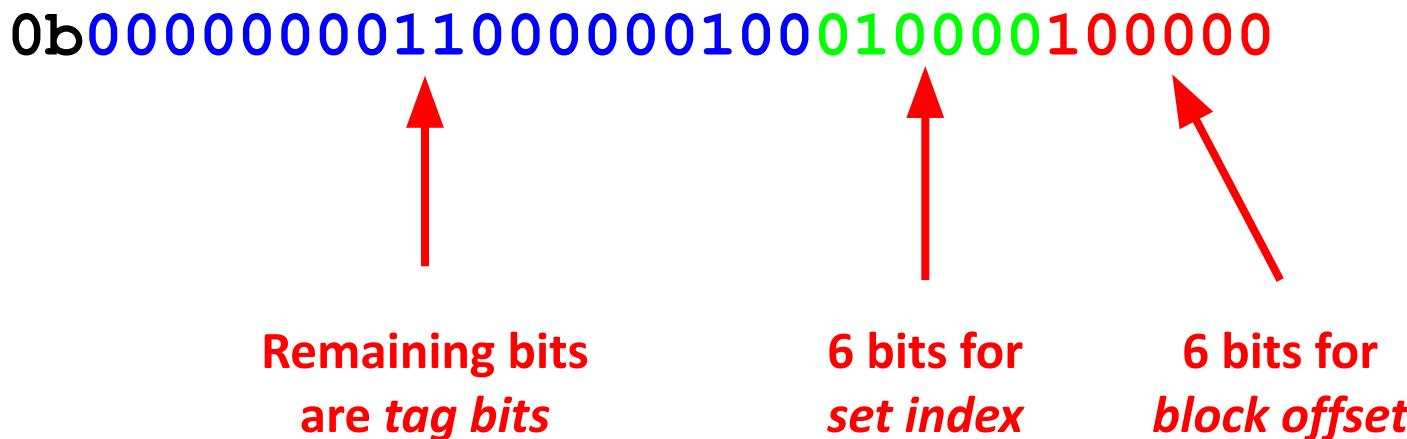
- We have an address that we want to look up in our cache.

0x00604420

- How do we search for it? Which set? Which line?
- Our *parameters* (**s** and **b**) determine how we partition the bits of our address.

Cache Concepts: Cache Read

- Our *parameters* (**s** and **b**) determine how we partition the bits of our address.
- Suppose **s** = 6 and **b** = 6



Cache Concepts: Cache Read

Tag: 0000000011000000100

Set: 010000

Block Offset: 100000

- These bits now tell us how to do the lookup in our cache!
- Use set index (**0b010000** = 16) to select the set
- Loop through lines in that set to find a matching tag
(**0b0000000011000000100**)
- If found and valid bit is set: ***Hit!***
 - Locate data starting at byte offset (**0b100000**)

Cache Concepts: Cache Miss

- But what happens if the cache doesn't have our data?
- We have a *cache miss*
- If we have a free line in the set, just load data into there
- Otherwise, the set is full!
 - We have to *evict* a line according to a *replacement policy*.
 - **cachelab**: LRU (Least Recently Used)
 - Other policies exist!
- Finally, load our new line into the free slot.

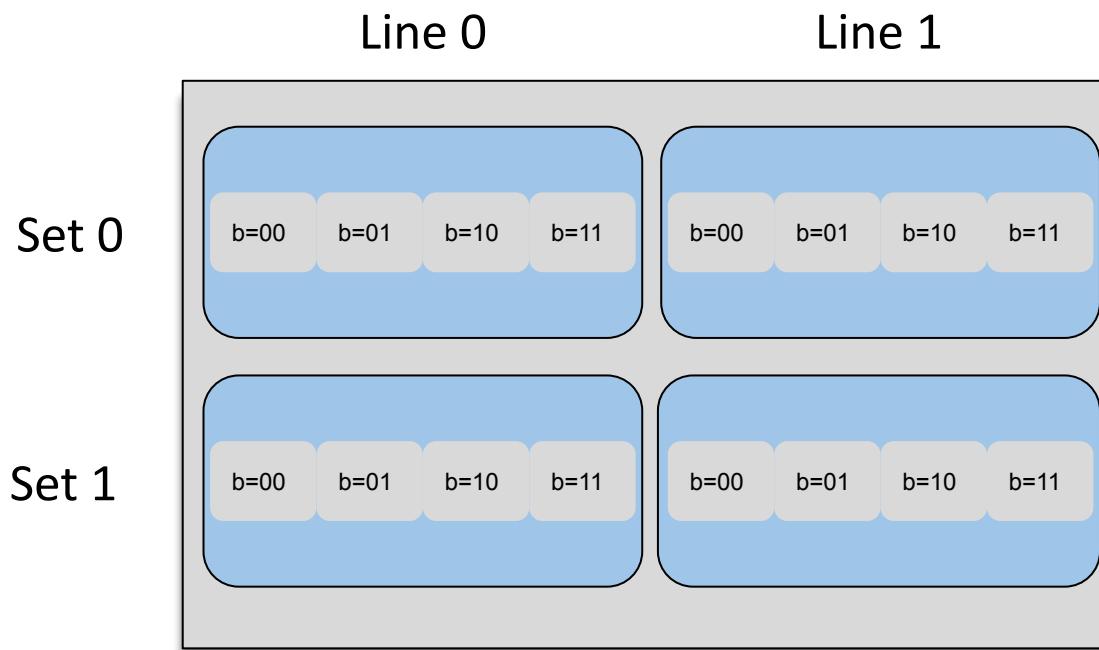
Cache Concepts: Dirty Bit

- You will implement a ***write-back, write-allocate*** policy for **cachelab**.
- ***Write-Allocate***: Write misses load the line into cache, update it in place.
- ***Write-Back***: Defer writing updates to memory until line is evicted.
 - Expensive to flush every evicted line to memory.
 - ***Dirty bit*** indicates whether cache line has been written to, and needs to be flushed to memory.

Example Trace

Example Trace

- We will use the following configuration:
 - $s = 1$, $E = 2$, $b = 2$



Example Trace: Reading a Trace

```
bpr.trace  
L 0,1  
L 0,1  
L 1,1  
S 2,1  
L 5,1  
L 4,1  
L 8,1  
L 0,1  
L 16,1  
L 9,1  
L 24,1  
L 32,1  
L 0,1
```

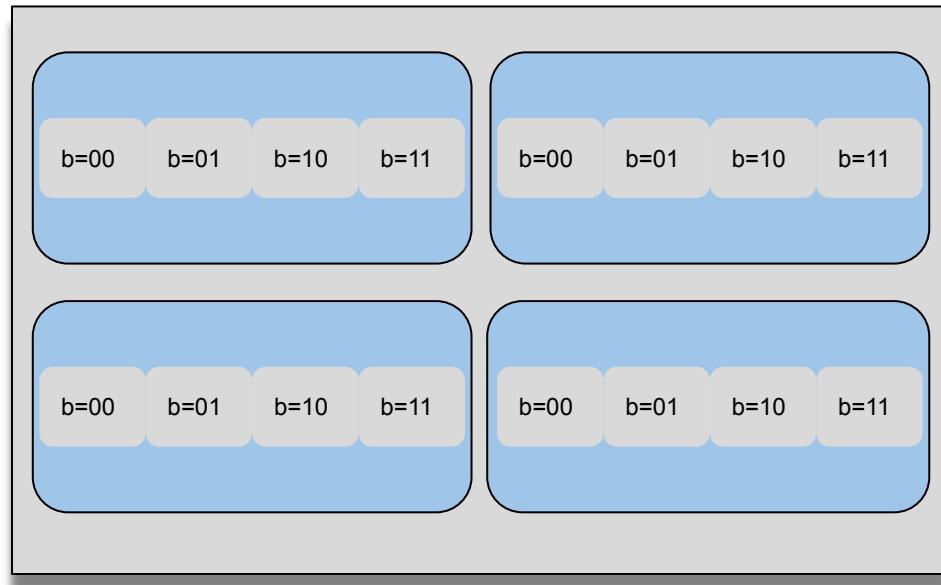
- **op <Addr>, <Size>**
- **op:**
 - **L – Load**
 - **S – Store**
- Note generally, we will **not** concern ourselves with the value being read/written
- For this trace however, we demonstrate write behavior through data values

Example Trace

...
L 0,1
...

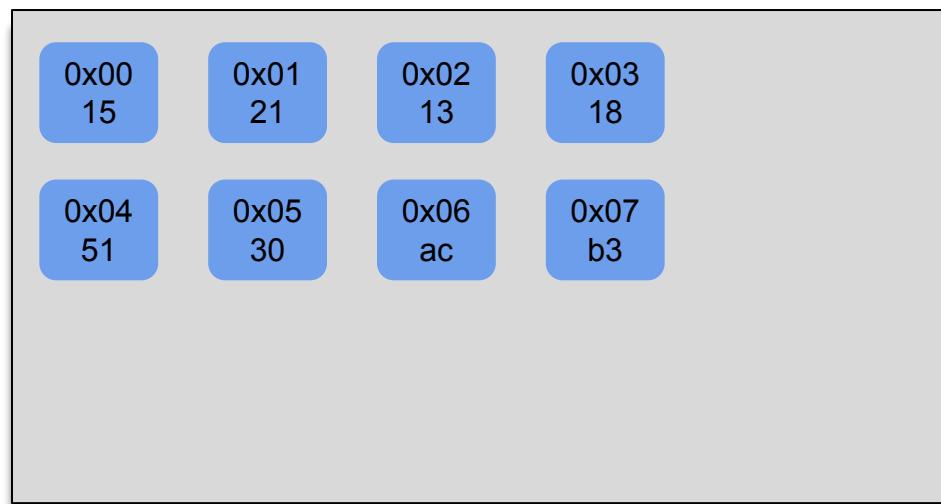


Cache



Will this instruction result in a hit or a miss?

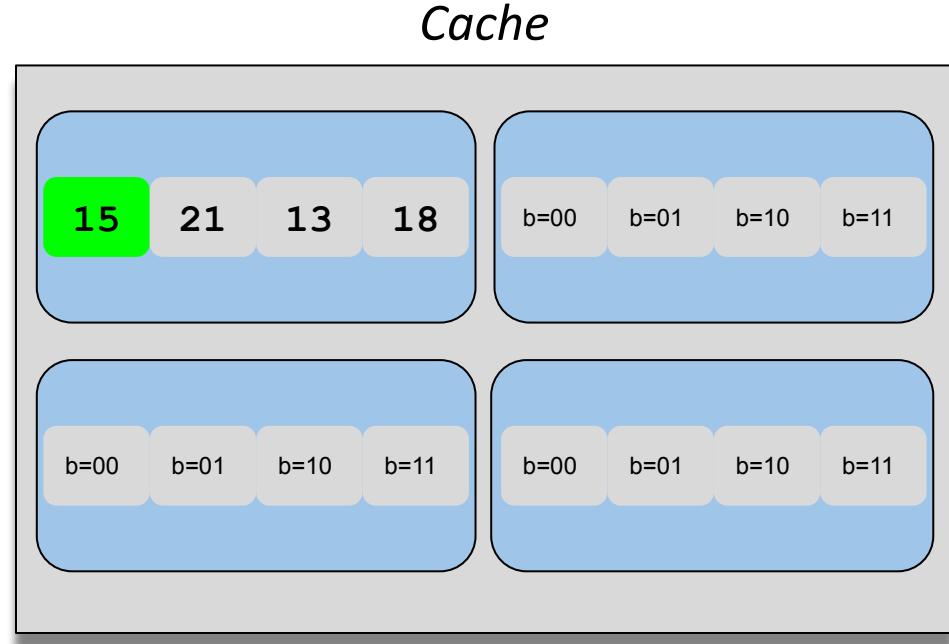
Memory



Example Trace

...
L 0,1
...

Miss

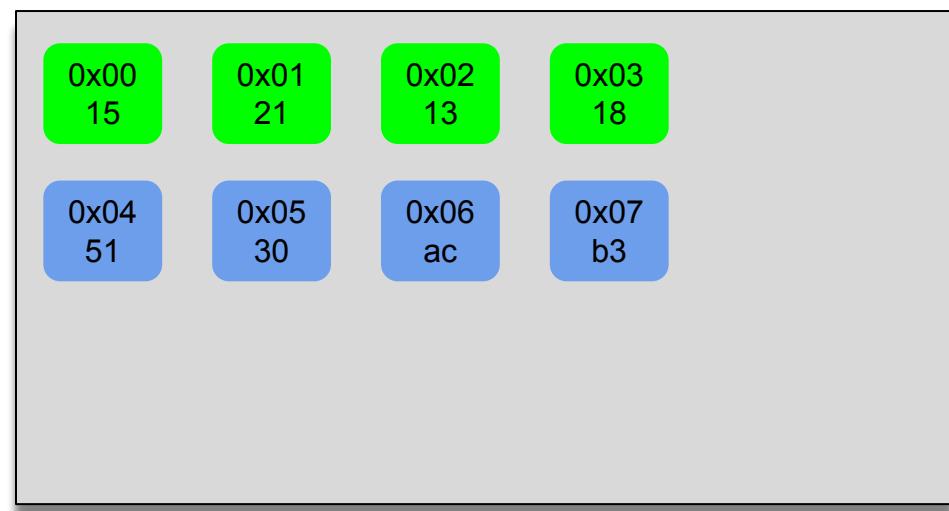


Why that line?

Where are those values from?

What kind of miss is this?

Memory



Example Trace

...

L 0,1

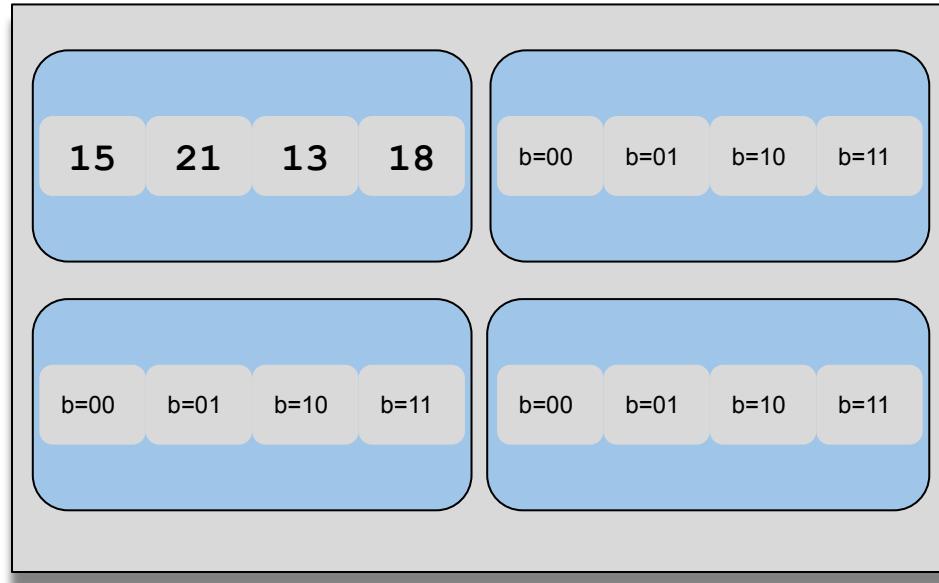
L 0,1

...

Miss
???

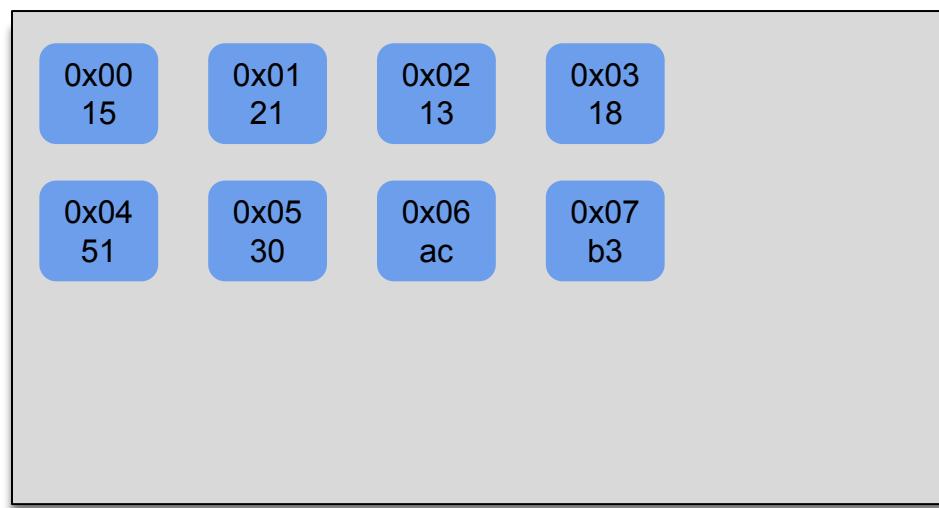


Cache



Memory

Will this instruction result in a hit or a miss?



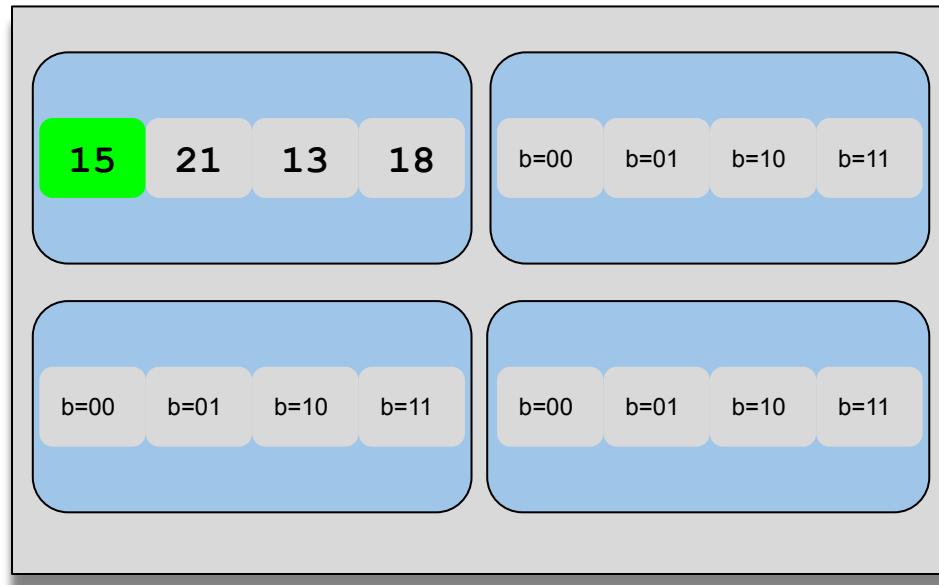
Example Trace

```
...  
L 0,1  
L 0,1  
...
```

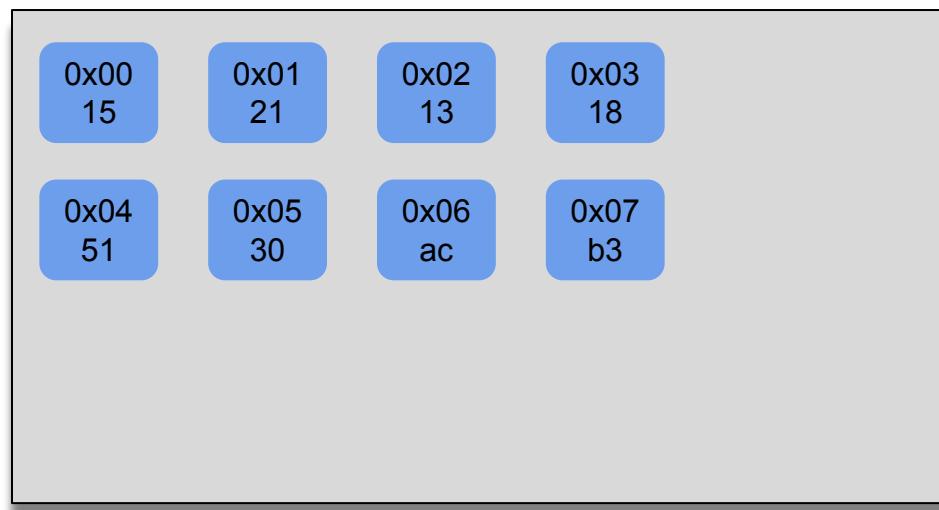
Miss
Hit!



Cache



Memory



Example Trace

...

L 0,1

L 0,1

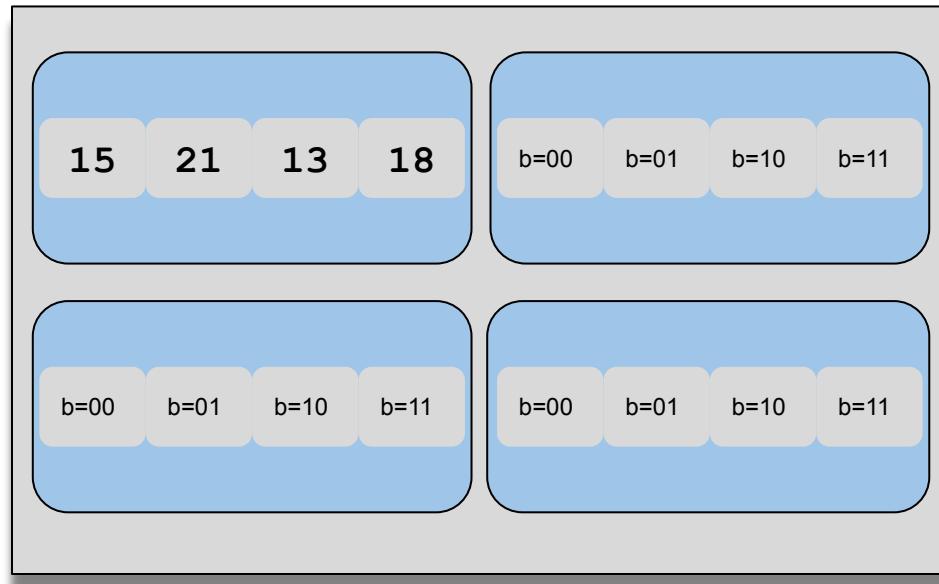
L 1,1

...

Miss
Hit!
???

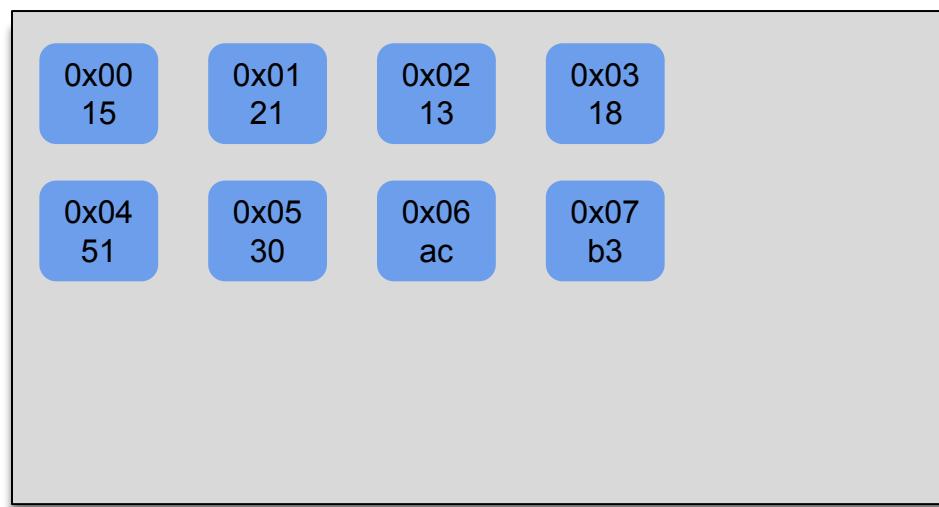


Cache



Memory

Will this instruction result in a hit or a miss?



Example Trace

...

L 0,1

L 0,1

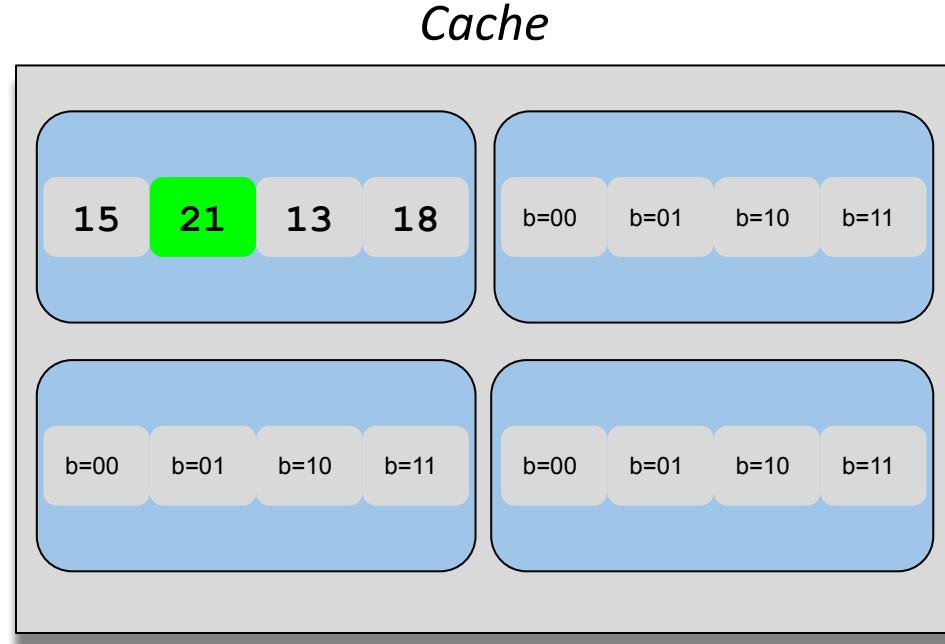
L 1,1

...

Miss

Hit!

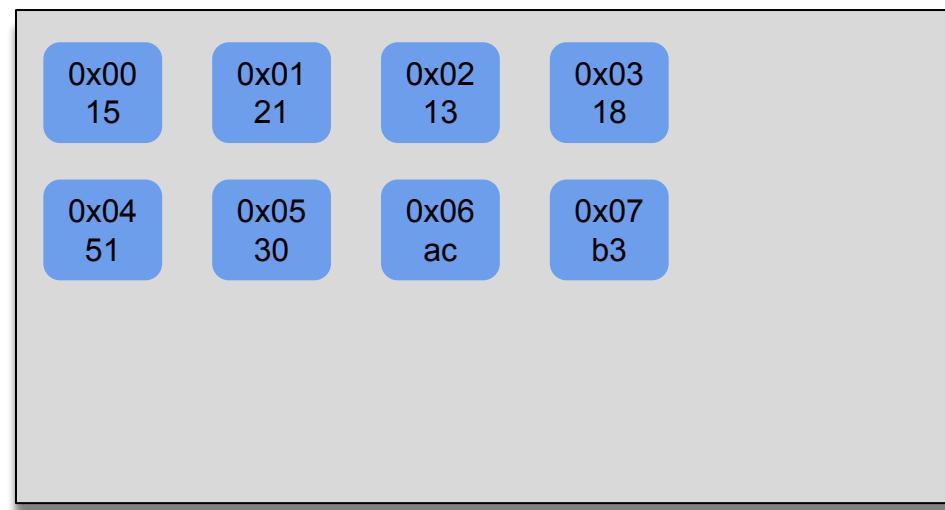
Hit!



Not a miss!

We had already loaded all four bytes of the line into cache. Why?

Memory



Example Trace

...

L 0,1

L 1,1

S 2,1

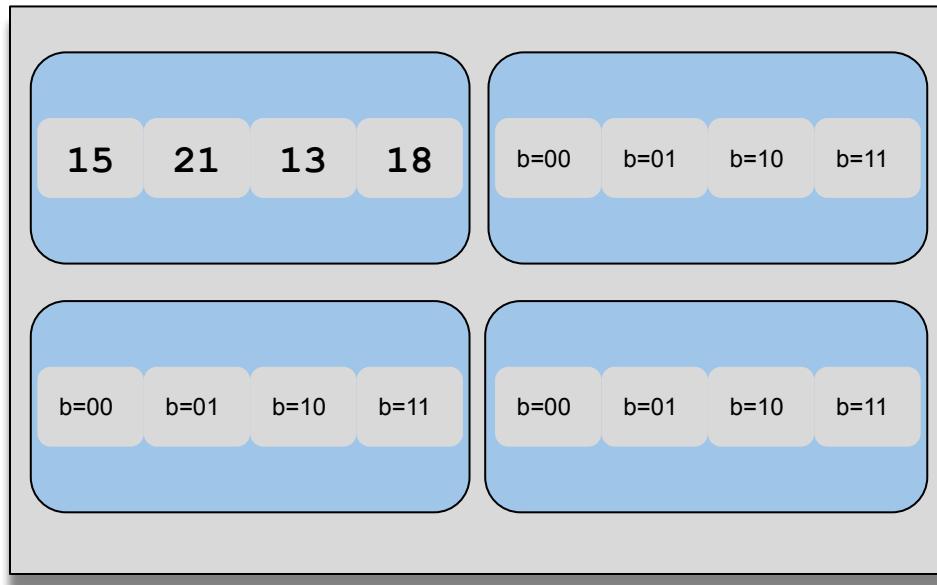
...

Hit!

Hit!

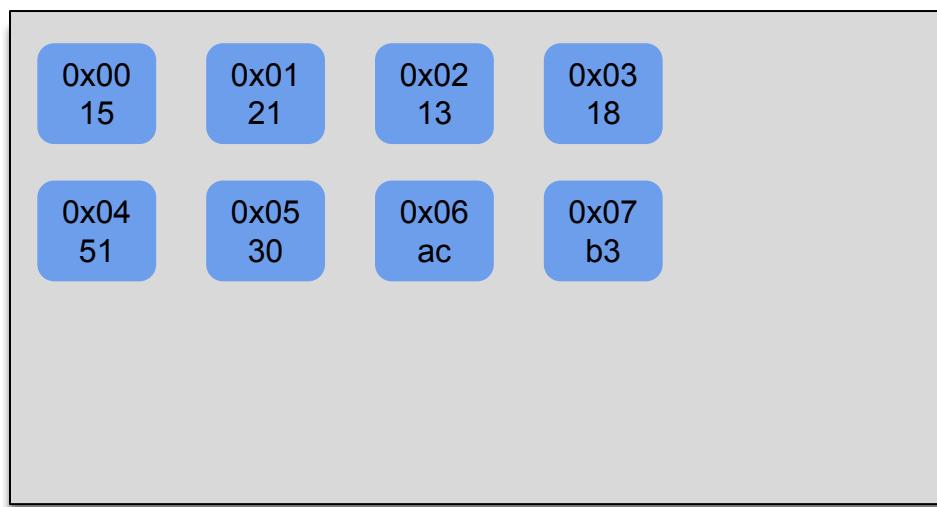
???

Cache



Will this instruction result in a hit or a miss?

Memory



Example Trace

...

L 0,1

L 1,1

S 2,1

...

Hit!

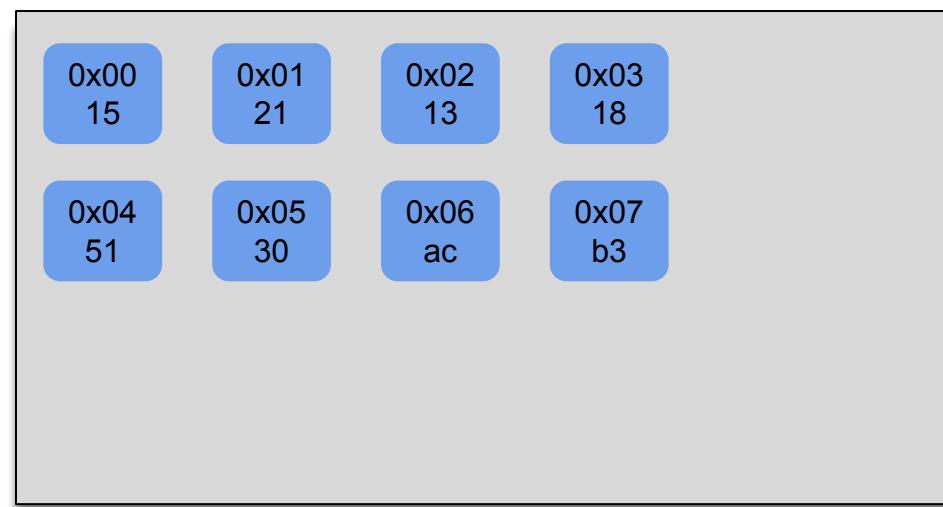
Hit!

Hit!



Write hit!

Set dirty bit.



Example Trace

...

L 1,1

S 2,1

L 5,1

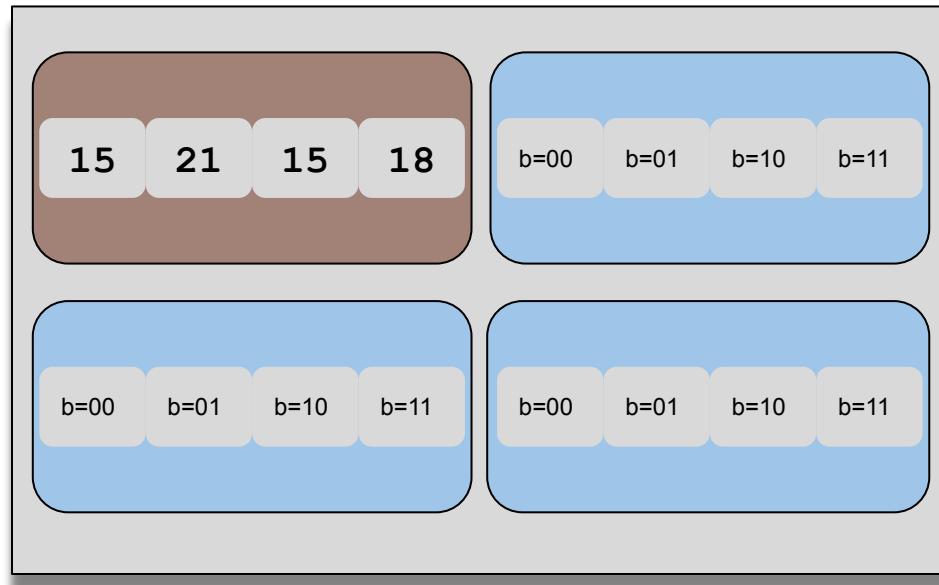
...

Hit!

Hit!

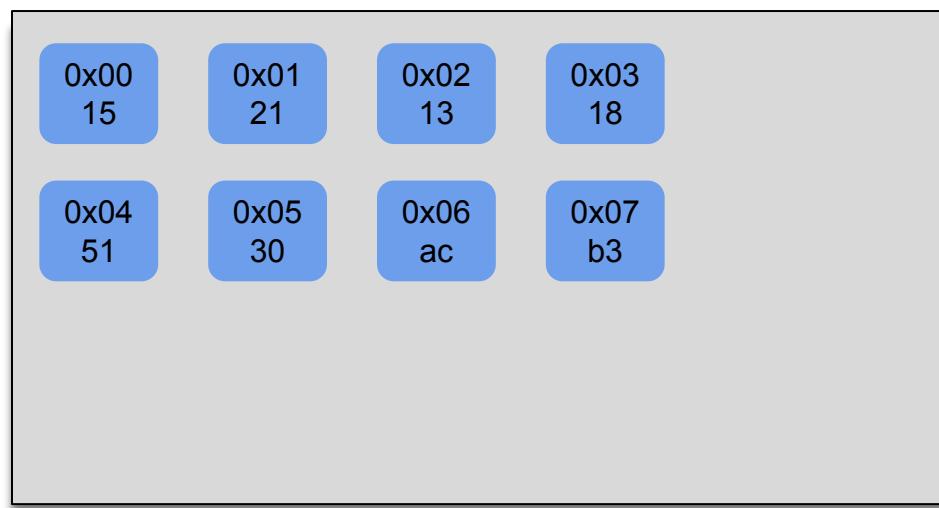
???

Cache



Will this instruction result in
a hit or a miss?

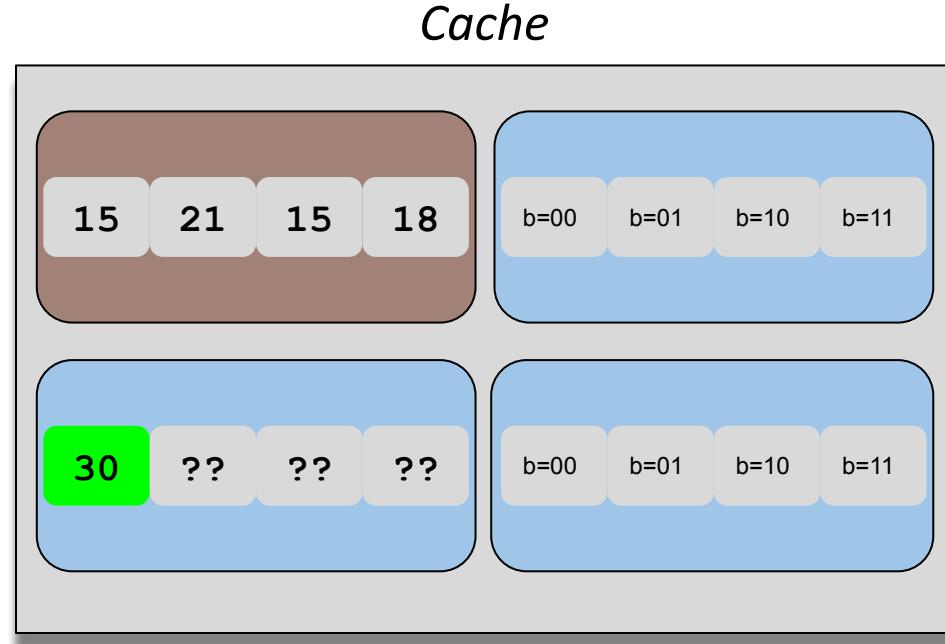
Memory



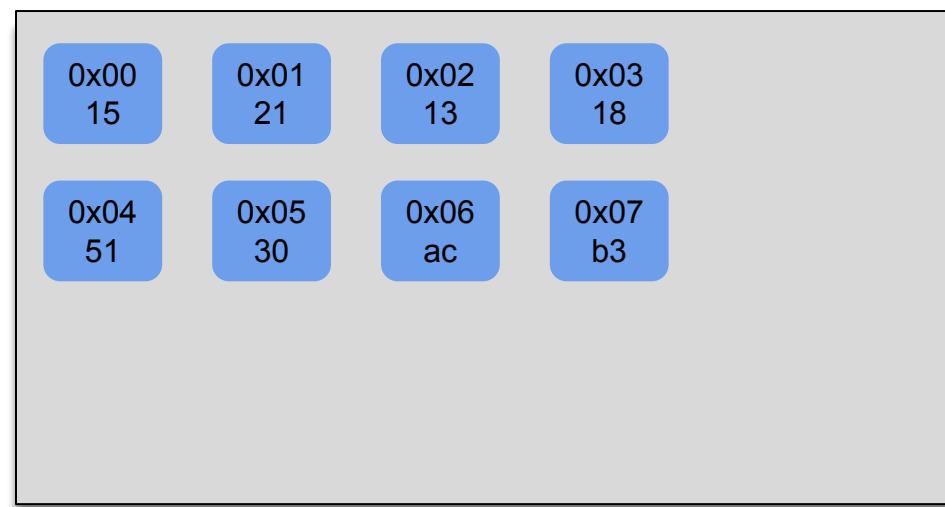
Example Trace

```
...  
L 1,1  
S 2,1  
L 5,1  
...
```

Hit!
Hit!
Miss



Do we load just one byte like this?



Example Trace

```
...  
L 1,1  
S 2,1  
L 5,1  
...
```

Hit!
Hit!
Miss

Do we load just one byte like this?

No!

Cache



Example Trace

...

L 1,1

S 2,1

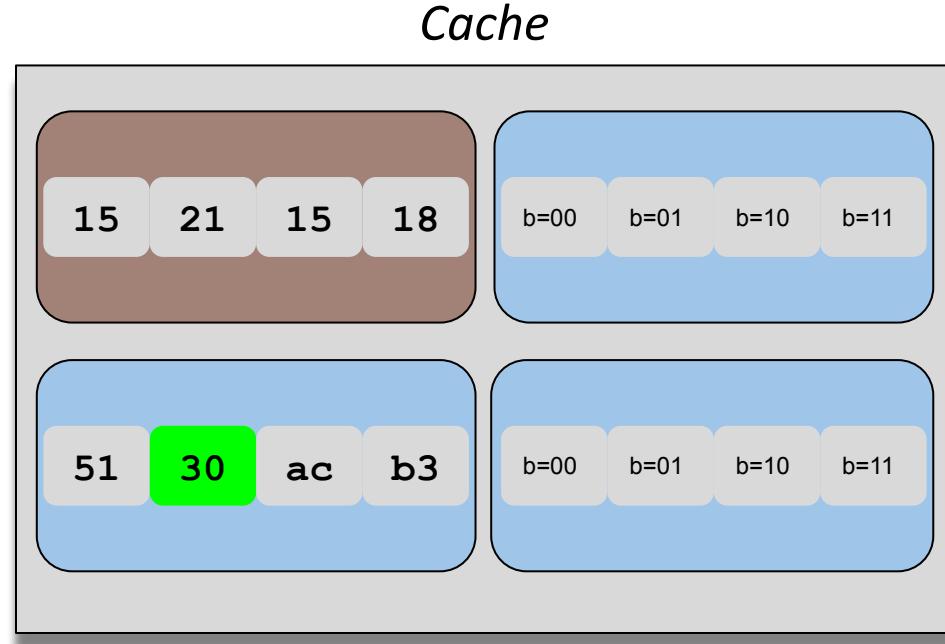
L 5,1

...

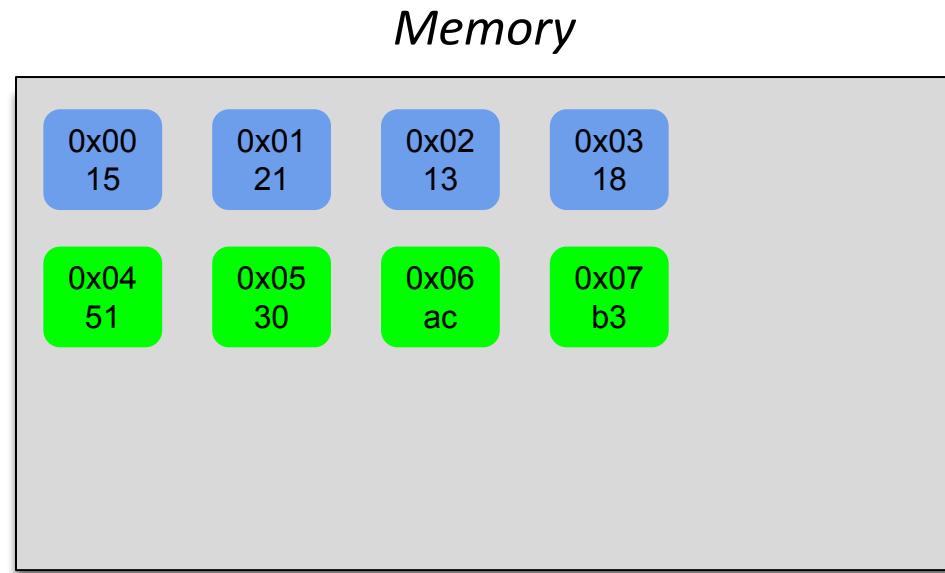
Hit!

Hit!

Miss



Why do we start with a byte from below address 5?



Example Trace

...

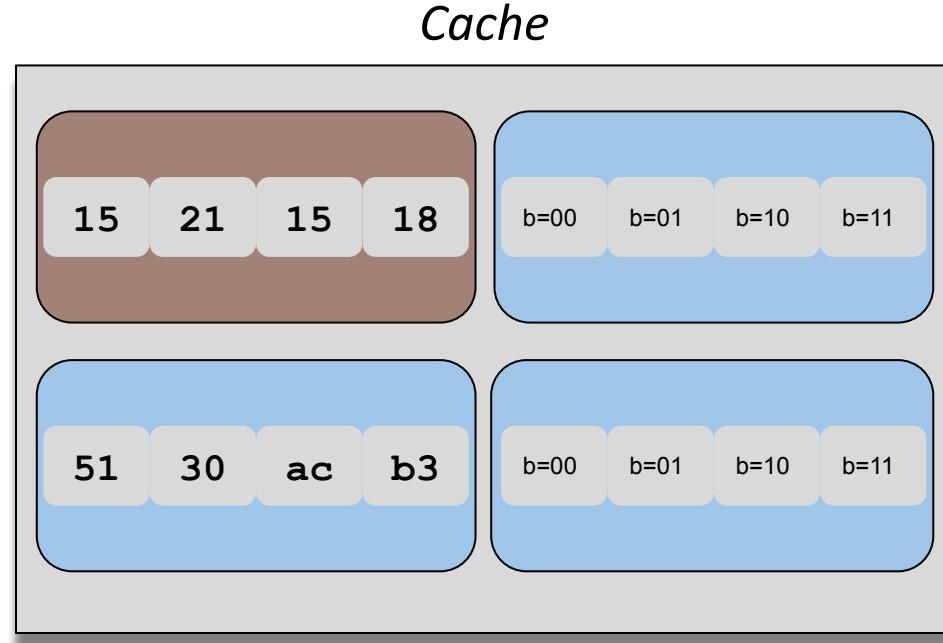
S 2,1

L 5,1

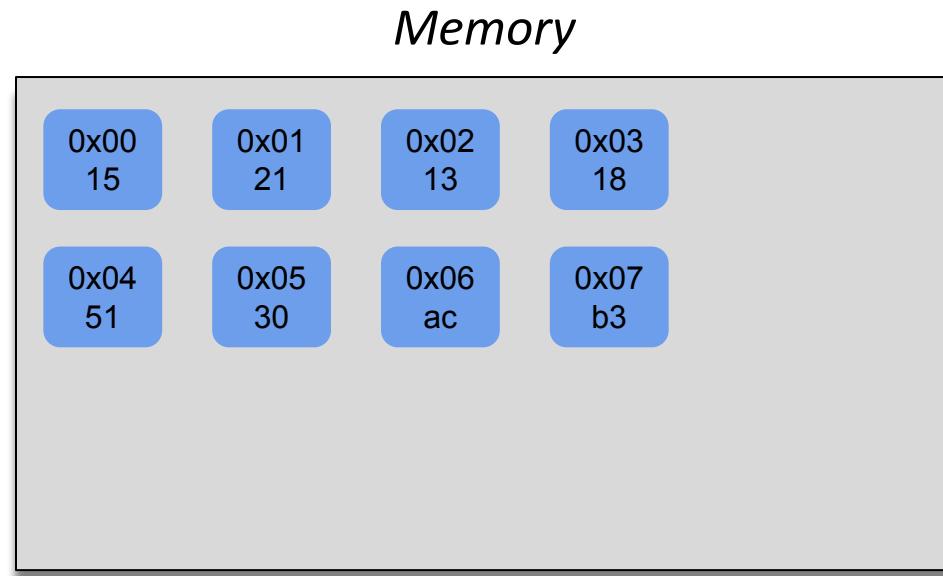
L 4,1

...

Hit!
Miss
???



Will this instruction result in a hit or a miss?

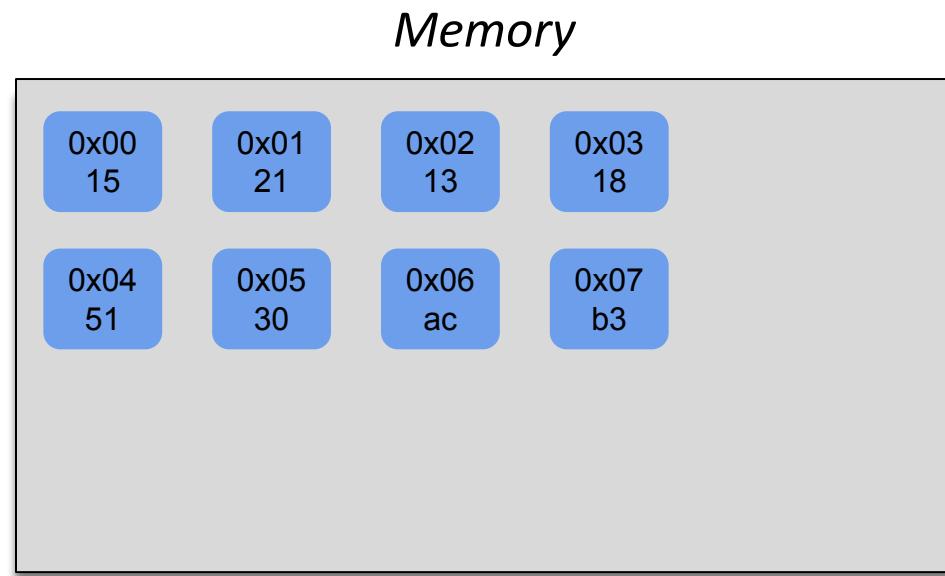
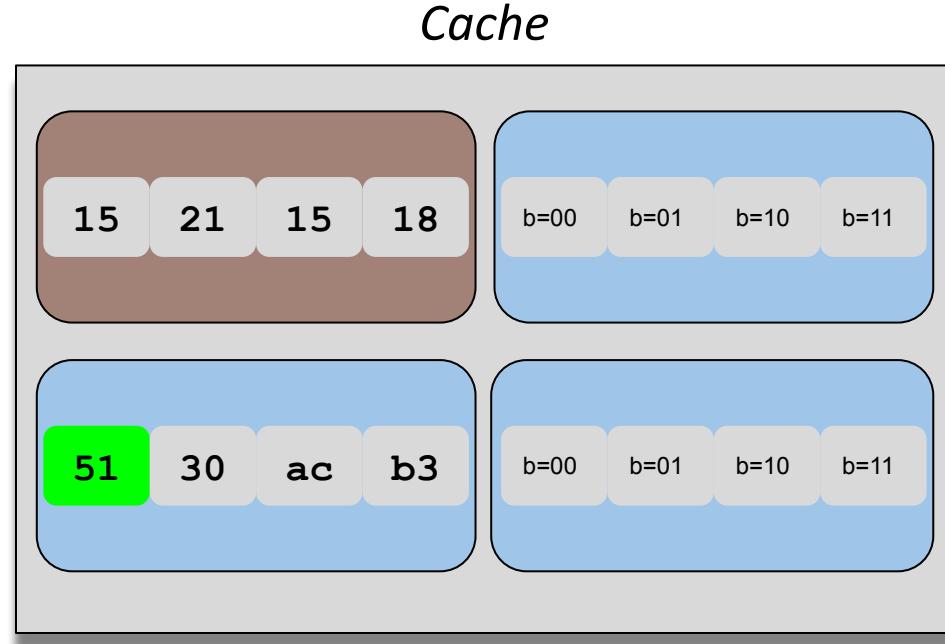


Example Trace

```
...  
S 2,1  
L 5,1  
L 4,1  
...
```



Hit!
Miss
Hit!



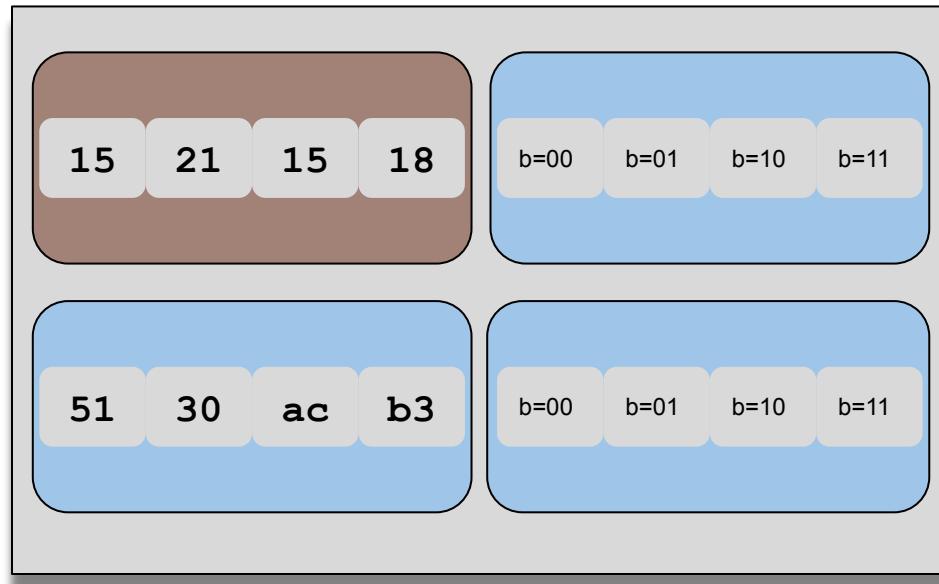
Example Trace

```
...  
L 5,1  
L 4,1  
L 8,1  
...
```

Miss
Hit!
???

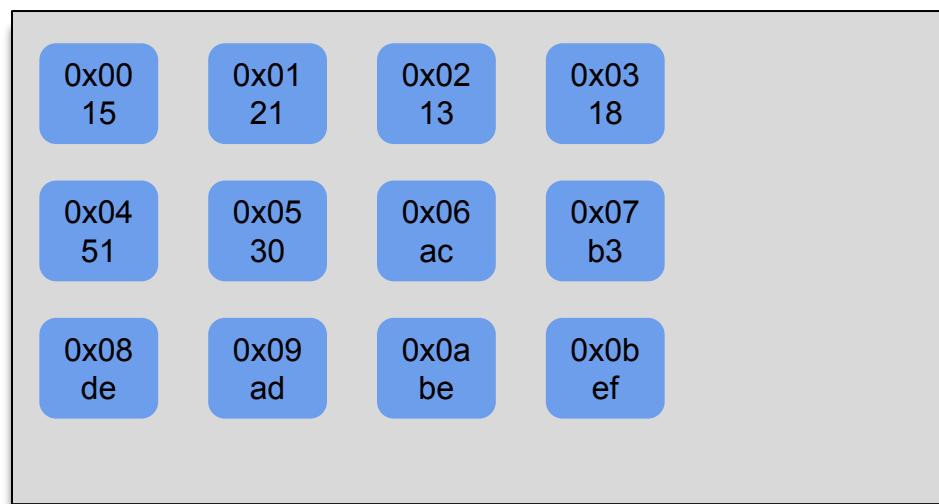


Cache



Memory

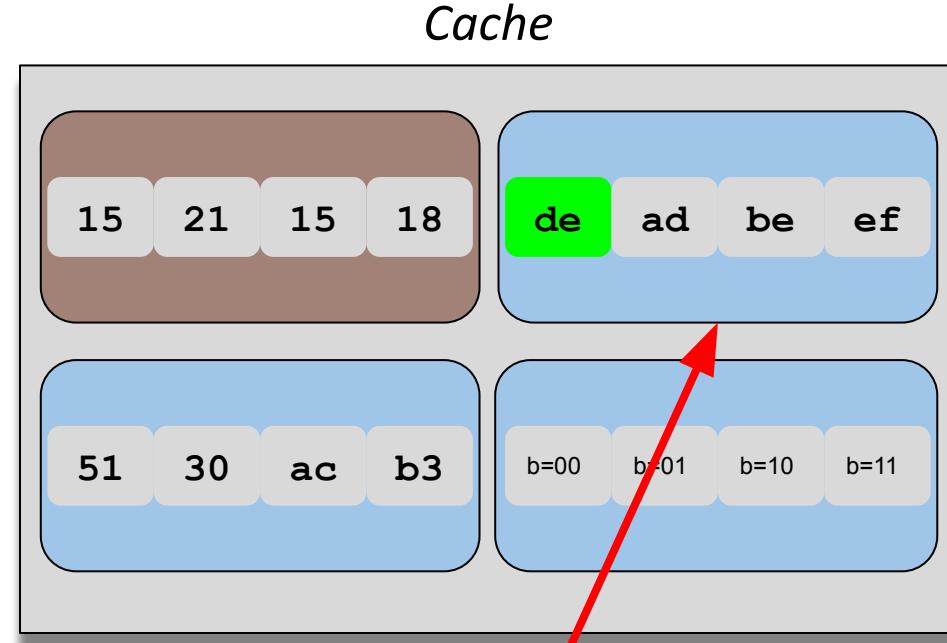
Will this instruction result in a hit or a miss?



Example Trace

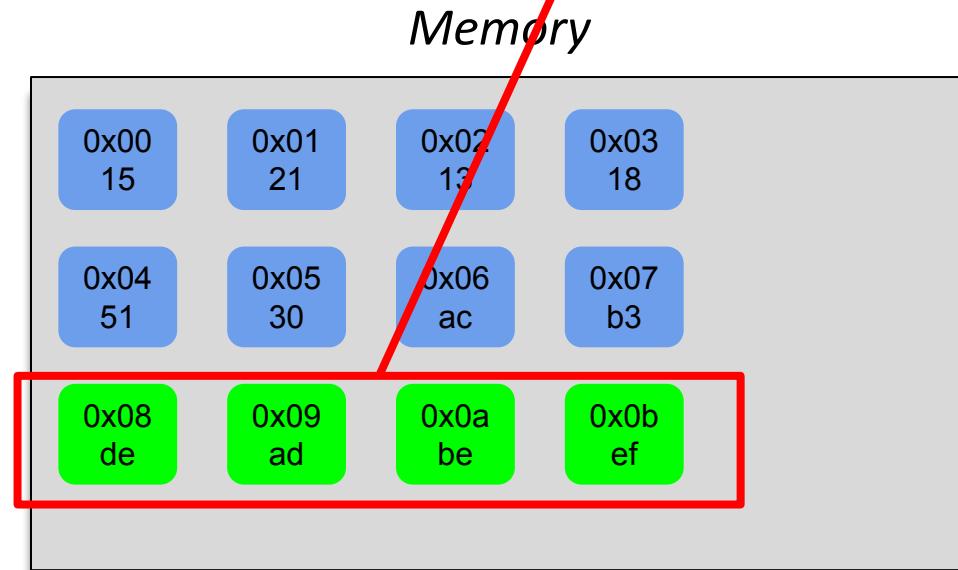
```
...  
L 5,1  
L 4,1  
L 8,1  
...
```

Miss
Hit!
Miss



Miss!

We had a free line, so just load the data into there.



Example Trace

...

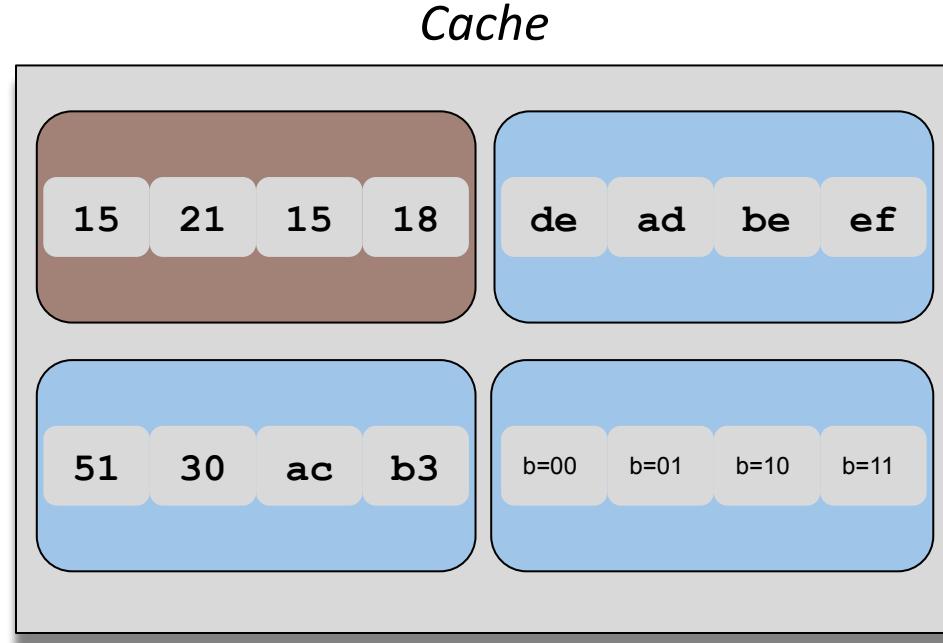
L 4,1

L 8,1

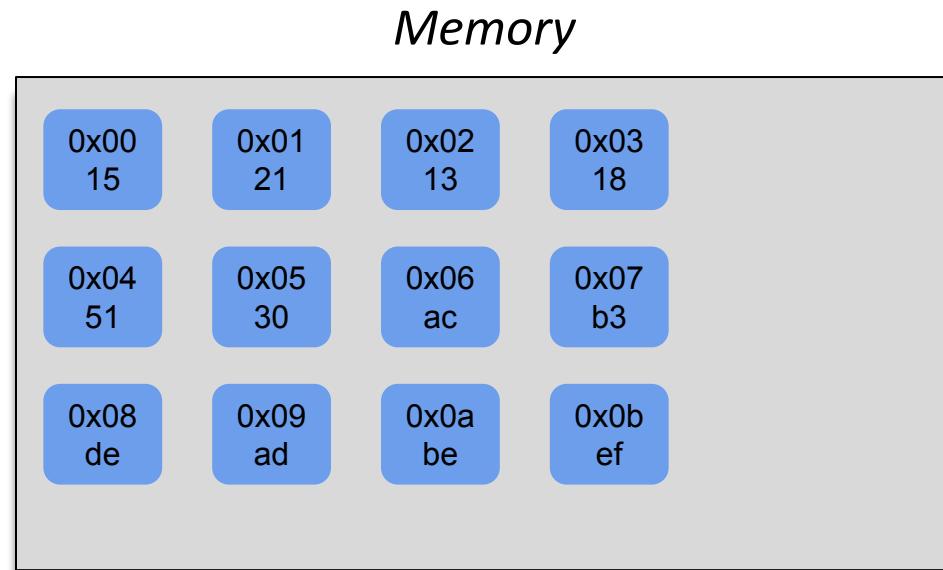
L 0,1

...

Hit!
Miss
???



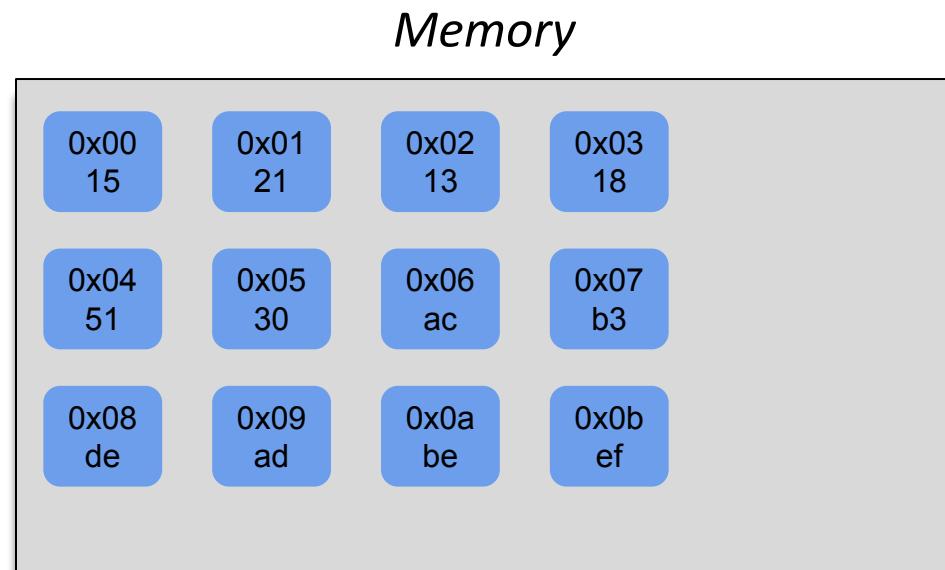
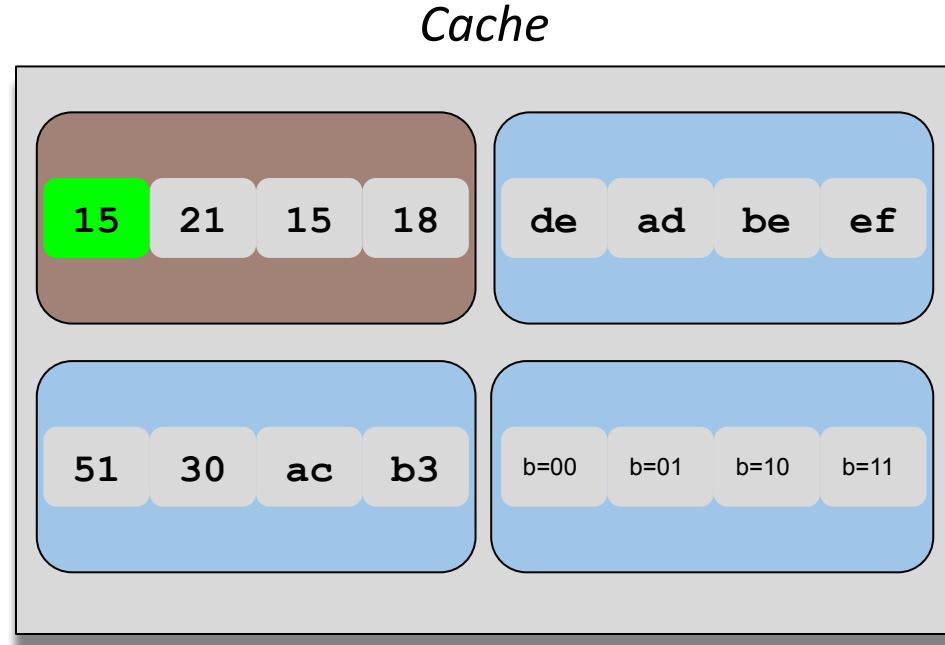
Will this instruction result in a hit or a miss?



Example Trace

```
...  
L 4,1  
L 8,1  
L 0,1  
...
```

Hit!
Miss
Hit!



Example Trace

...

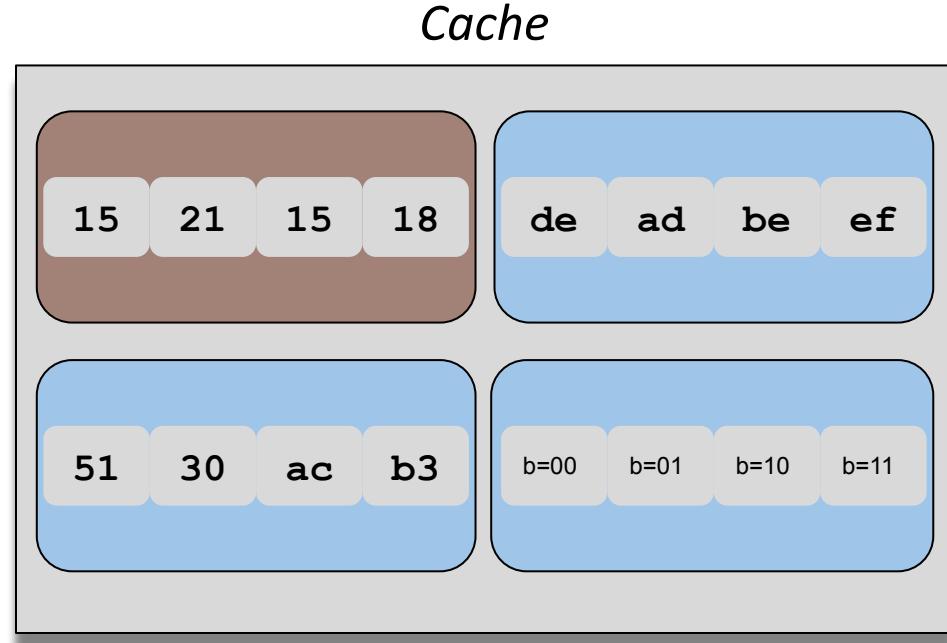
L 8,1

L 0,1

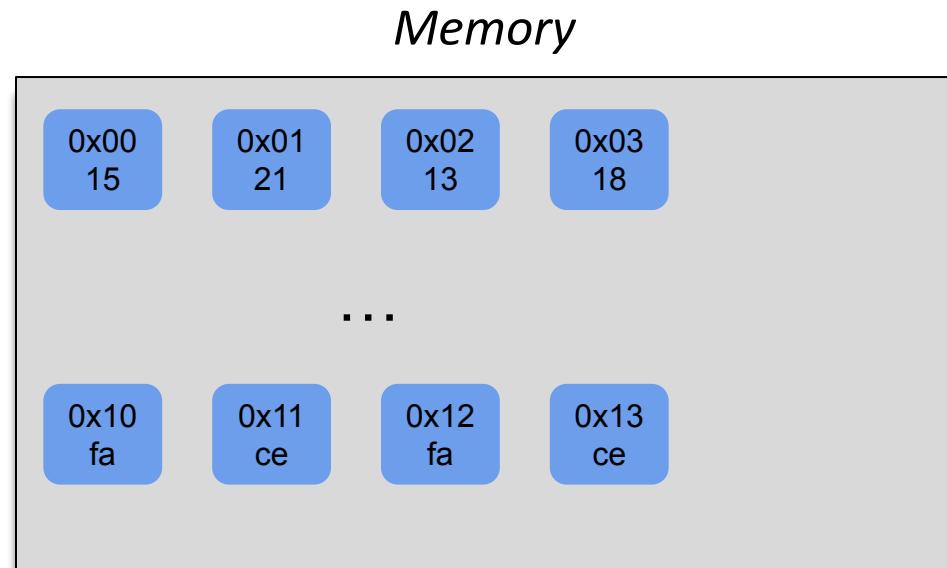
L 16,1

...

Miss
Hit!
???



Will this instruction result in a hit or a miss?

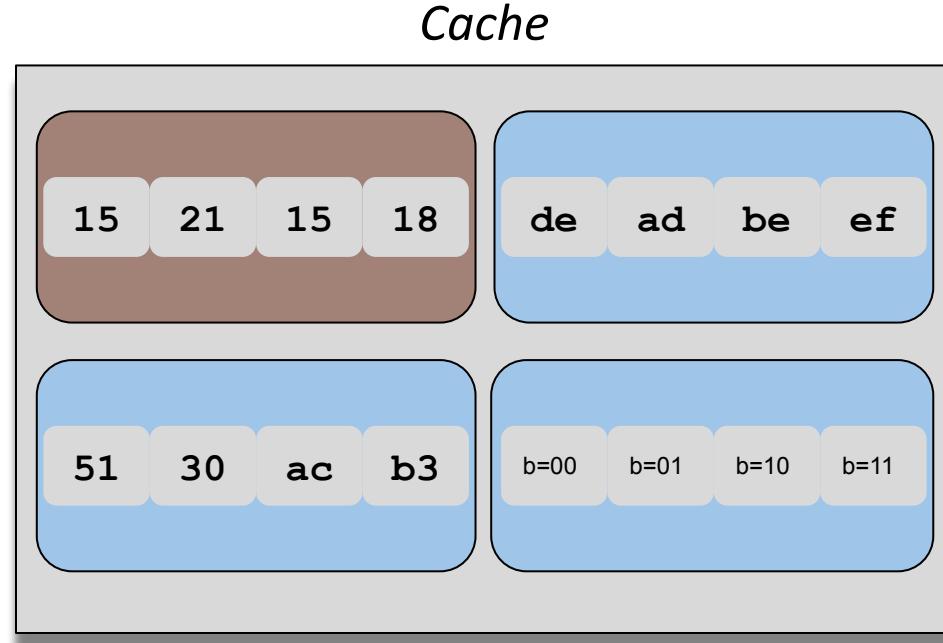


Example Trace

```

    ...
    L 8,1
    L 0,1
    L 16,1
    ...
  
```

Miss
Hit!
Miss



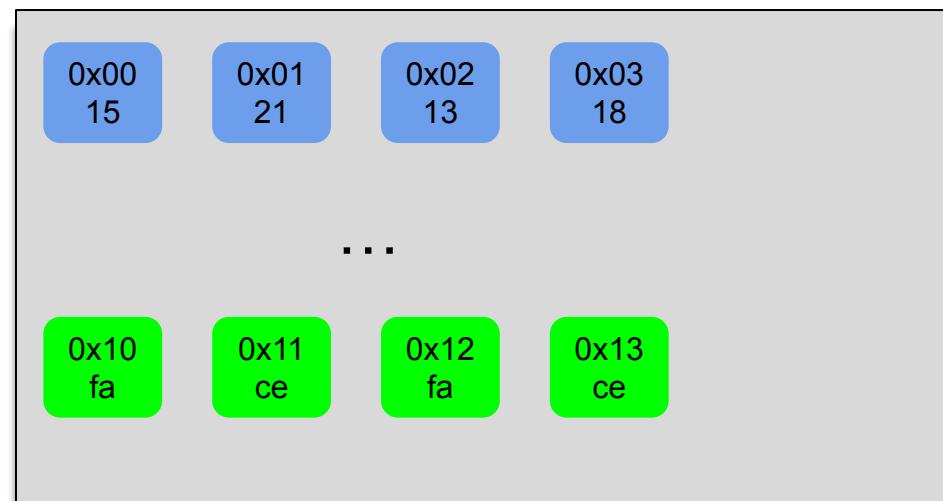
What kind of miss is this?

$16 = 0b10000$

=> Set Index 0

=> Have to evict!

Memory

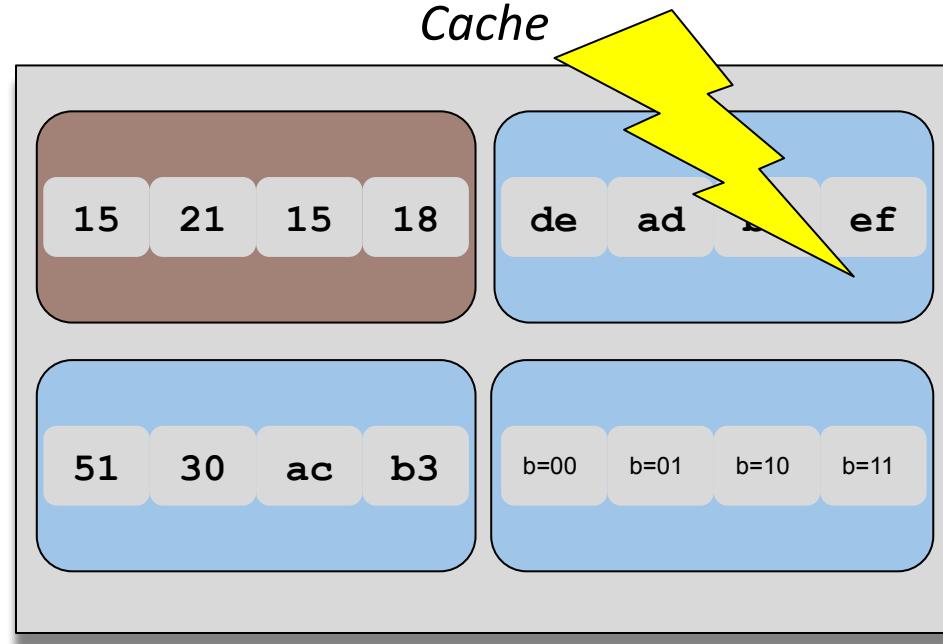


Example Trace

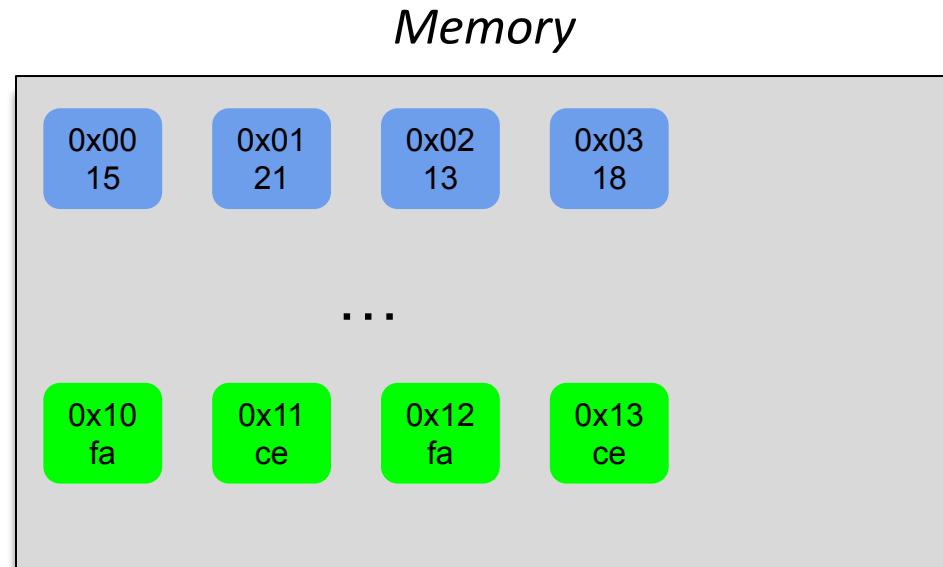
```

    ...
    L 8,1
    L 0,1
    L 16,1
    ...
  
```

Miss
Hit!
Miss



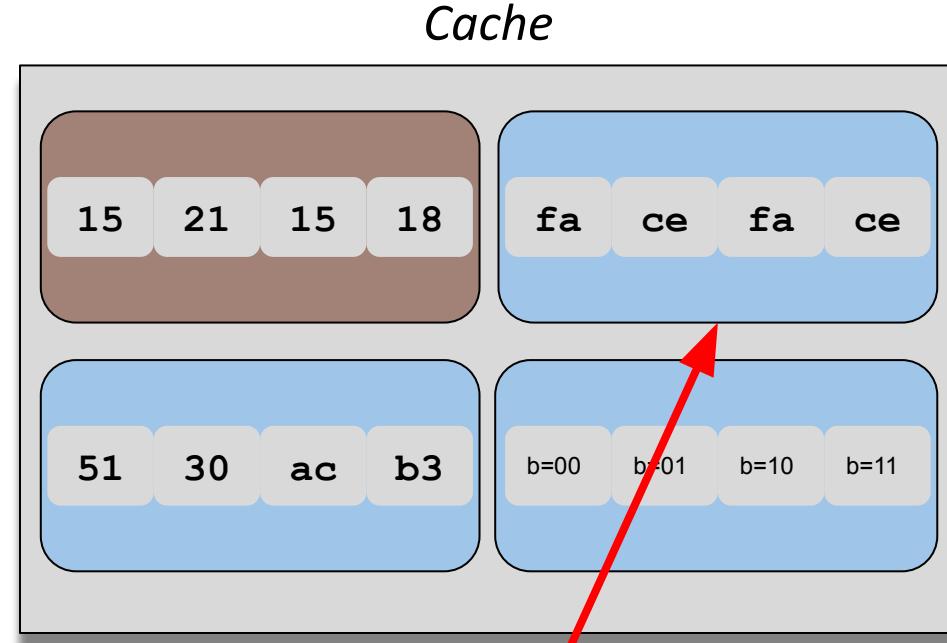
1. **Cold Miss** (first time seeing this block)
1. **Evict LRU** (Least Recently Used) line from set 0



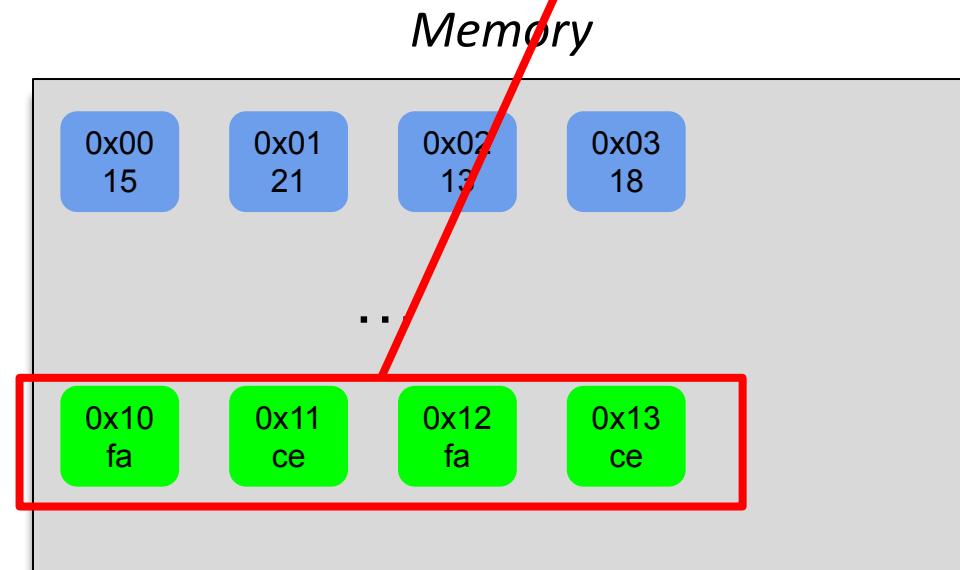
Example Trace

```
...  
L 8,1  
L 0,1  
L 16,1  
...
```

Miss
Hit!
Miss



Load new data into line



Example Trace

...

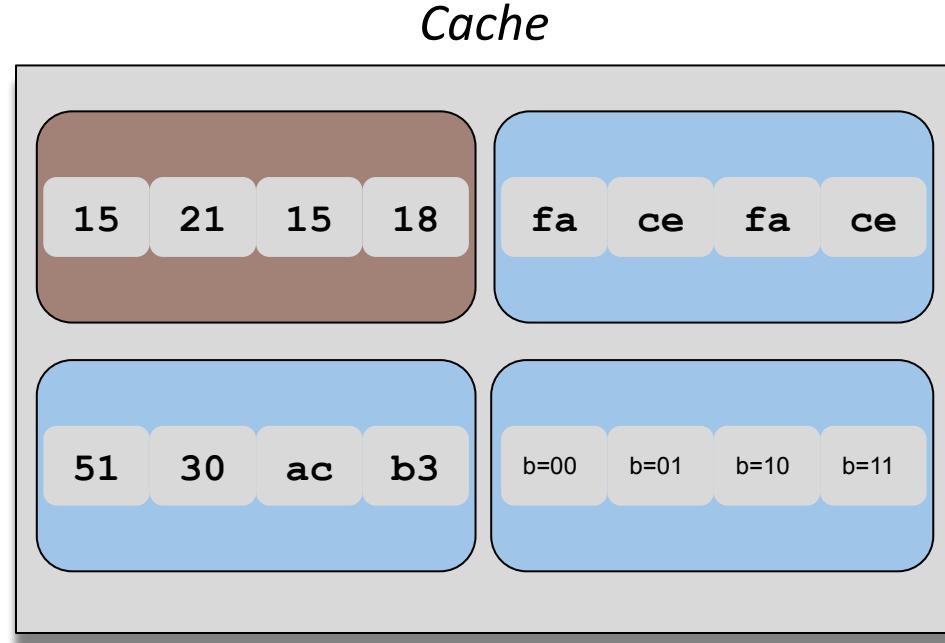
L 0,1

L 16,1

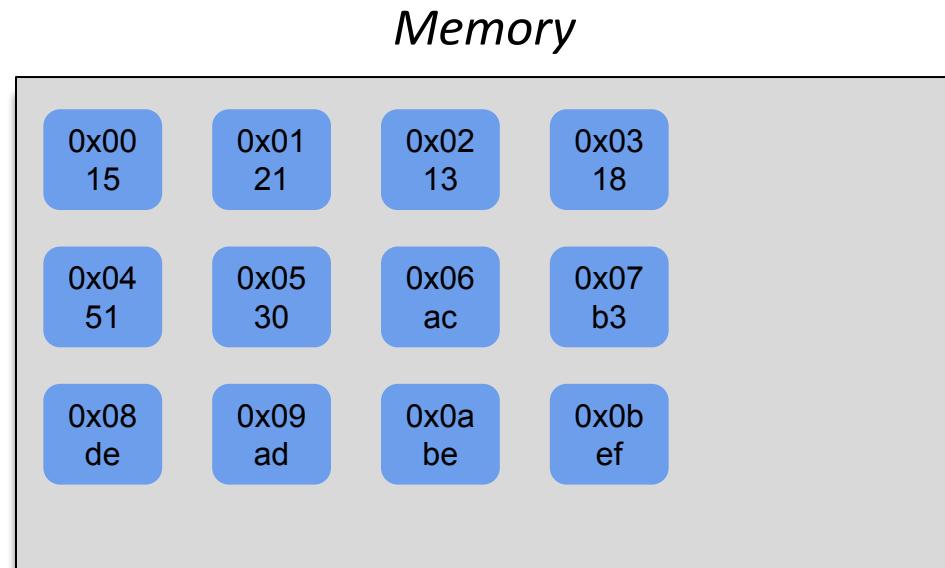
L 9,1

...

Hit!
Miss
???



Will this instruction result in a hit or a miss?



Example Trace

...

L 0,1

L 16,1

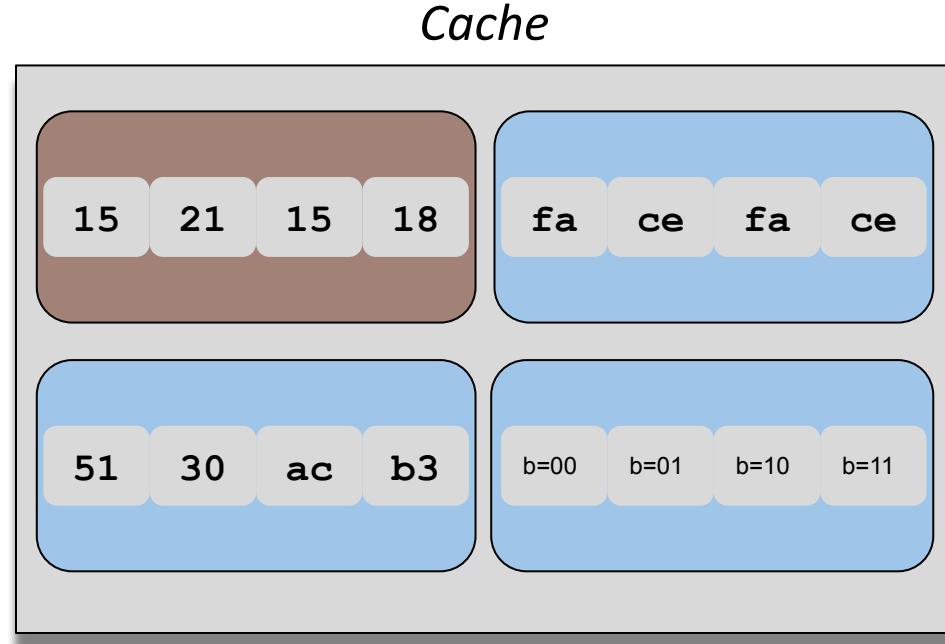
L 9,1

...

Hit!

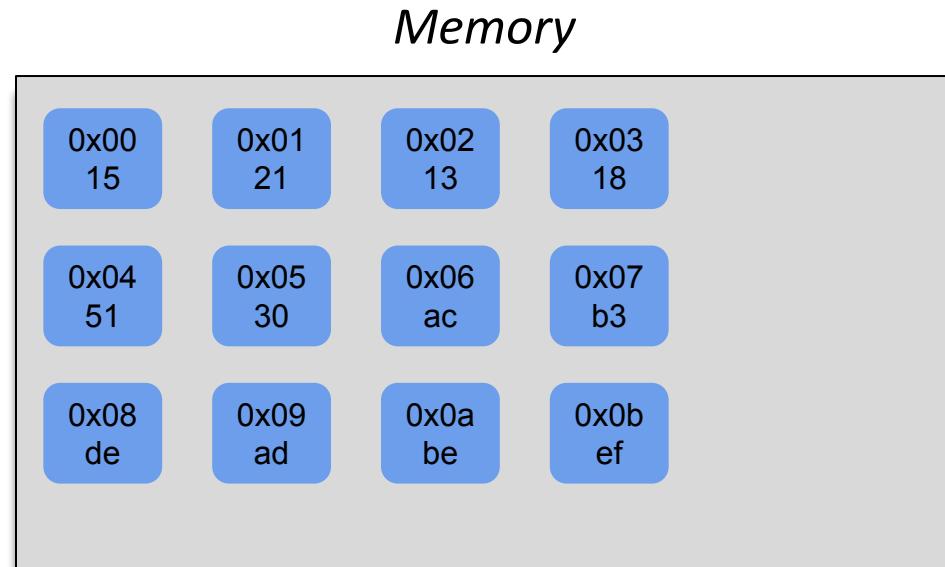
Miss

Miss

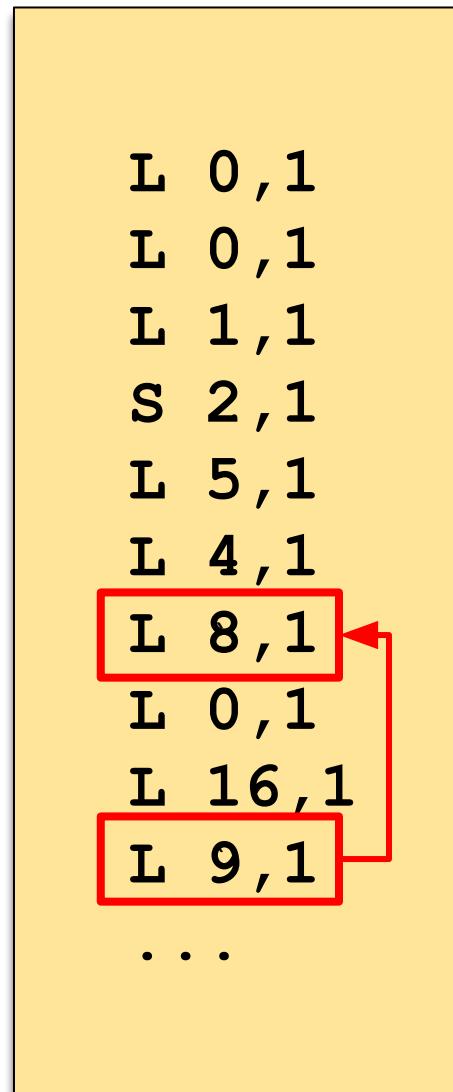


What kind of miss is this?

Has the block been in the cache before?

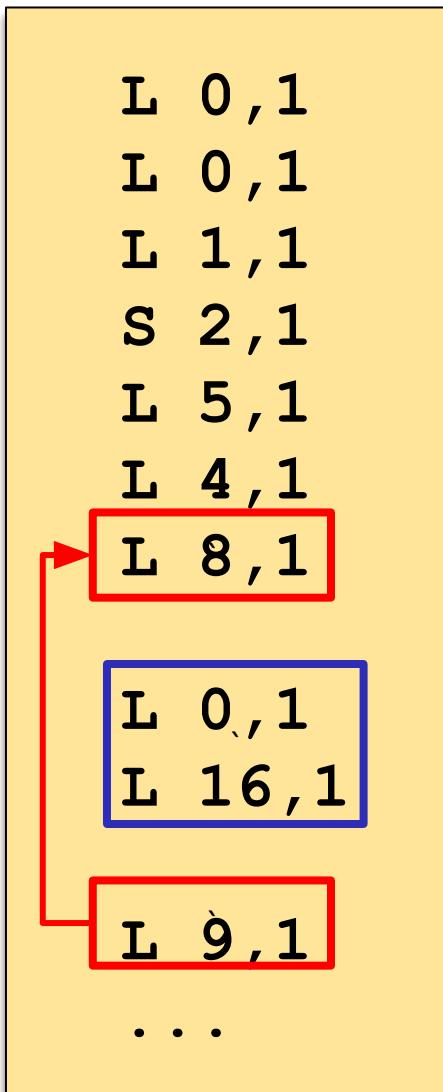


Cache Concepts: Conflict/Capacity Misses



- Has this block been in the cache before?
- Yes!
- If we've seen the block before:
 - Not a cold miss
 - Either a *conflict miss* or a *capacity miss*.

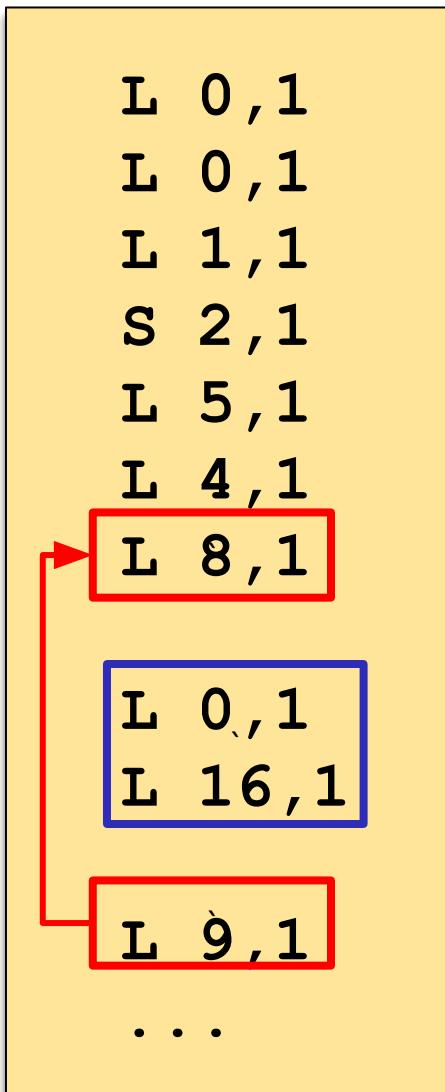
Cache Concepts: Conflict/Capacity Misses



How to distinguish between the two:

1. Find the last reference to that block in the trace.
2. Count the number of *unique* blocks referenced *in-between*:
 - a. If the number is greater than or equal to the total number of lines in the cache: ***Capacity Miss***
 - b. Otherwise: ***Conflict Miss***

Cache Concepts: Conflict/Capacity Misses



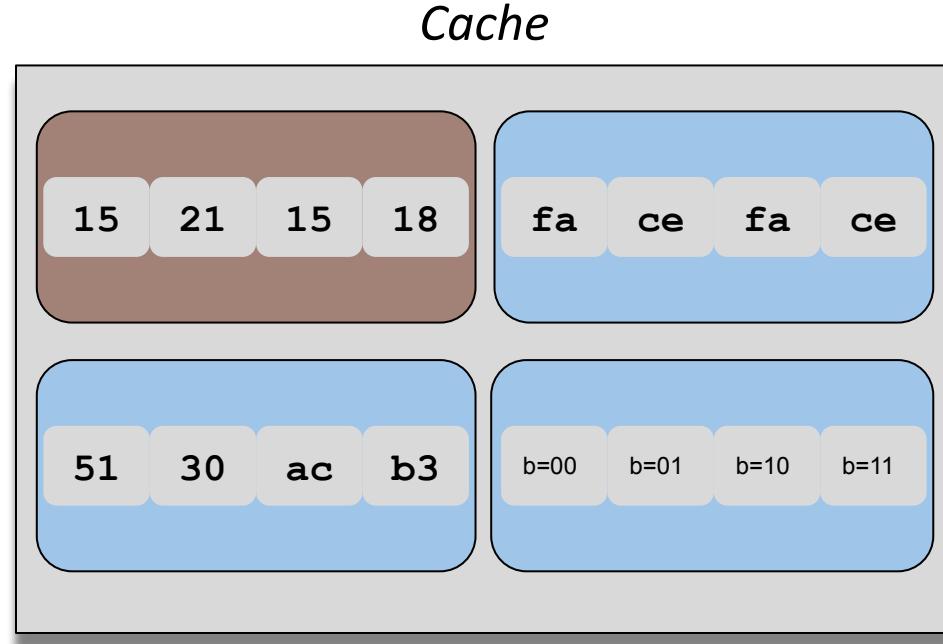
- In this case:
 - Two unique blocks in between current reference and last reference.
 - But we have *four* total lines in the cache
 - So we have a ***Conflict Miss***.

Example Trace

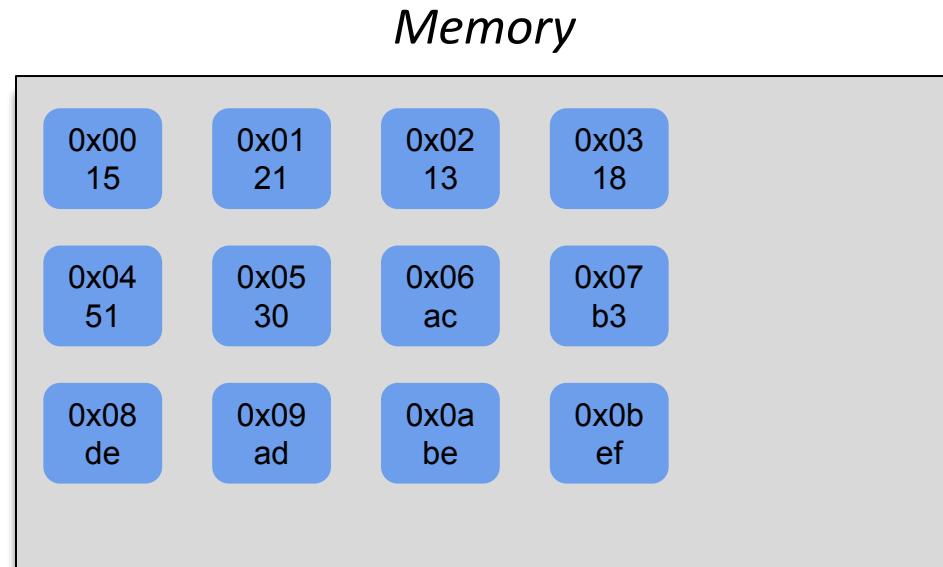
```

    ...
    L 0,1
    L 16,1
    L 9,1
    ...
  
```

Hit!
Miss
Miss



9 = 0b1001
=> Set Index 0
=> Have to evict!

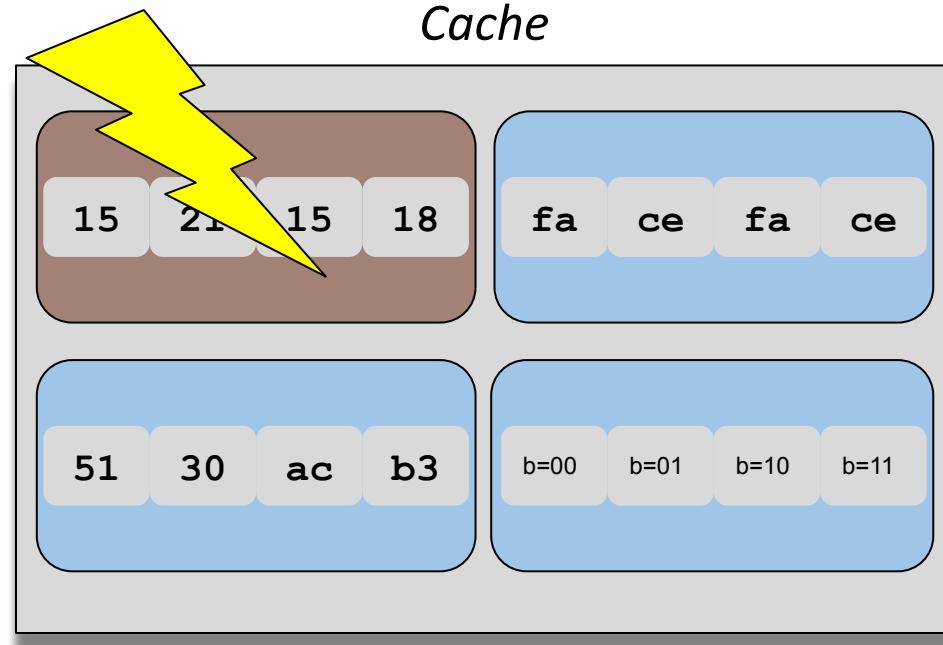


Example Trace

```

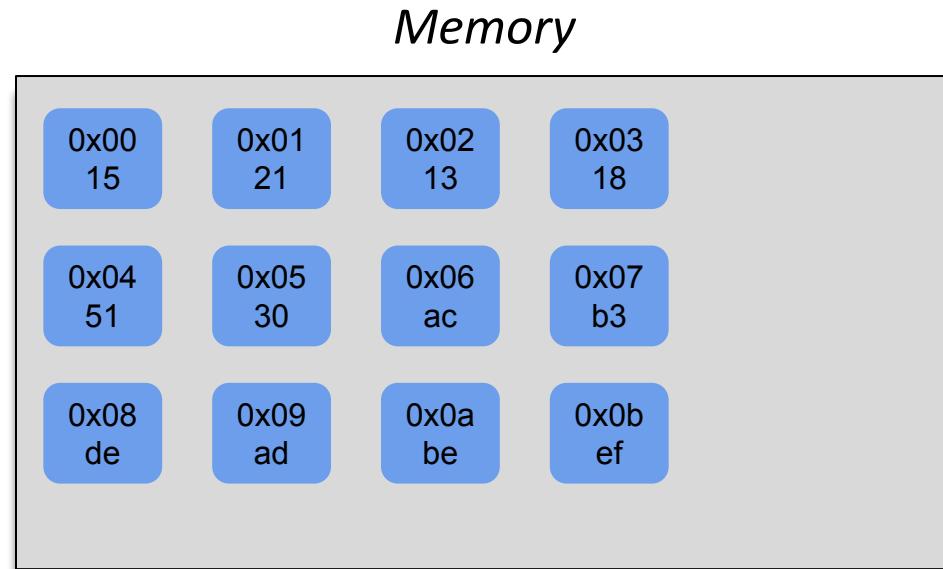
    ...
    L 0,1
    L 16,1
    L 9,1
    ...
  
```

Hit!
Miss
Miss



Evict least recently used line

Dirty bit set => Dirty Eviction



Example Trace

...

L 0,1

L 16,1

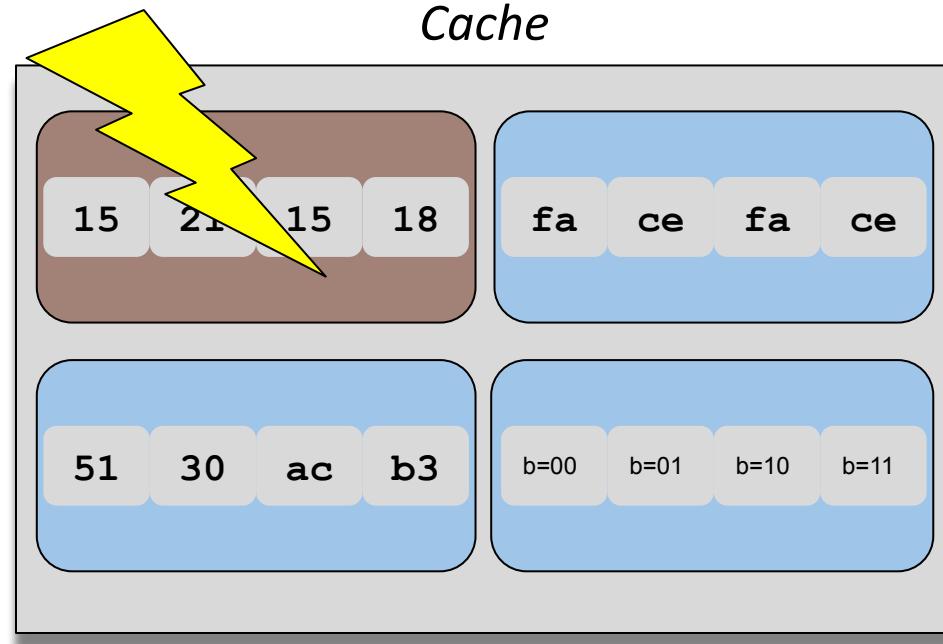
L 9,1

...

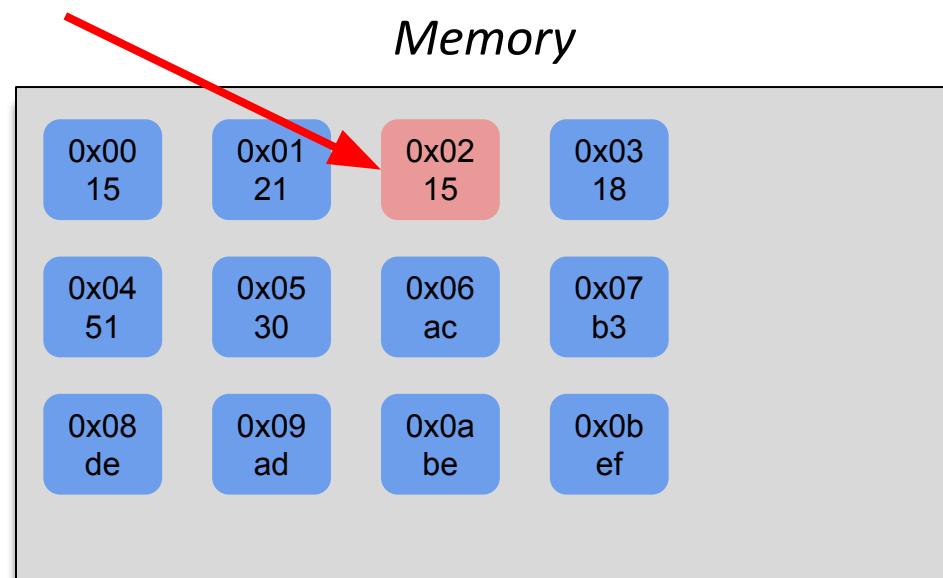
Hit!

Miss

Miss



Write-back policy!



Example Trace

...

L 0,1

L 16,1

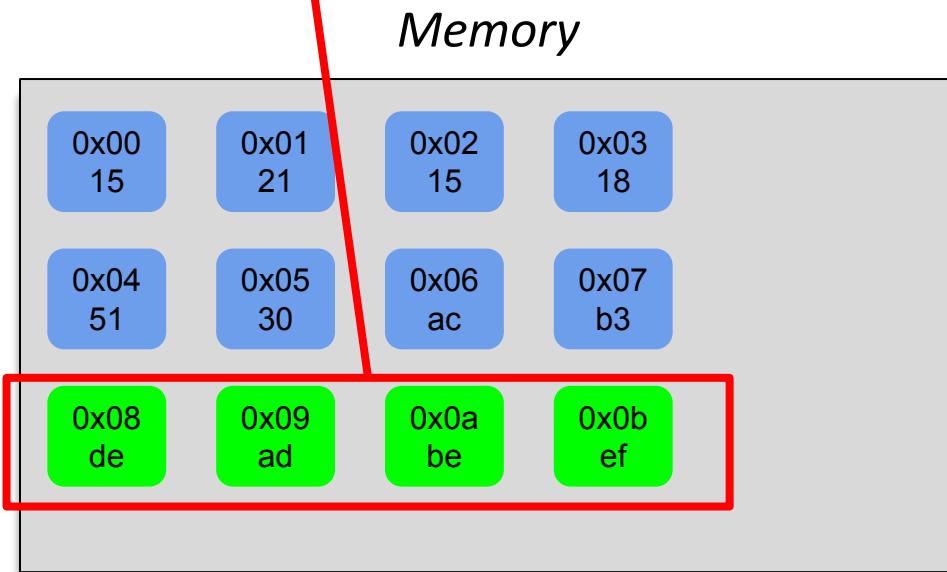
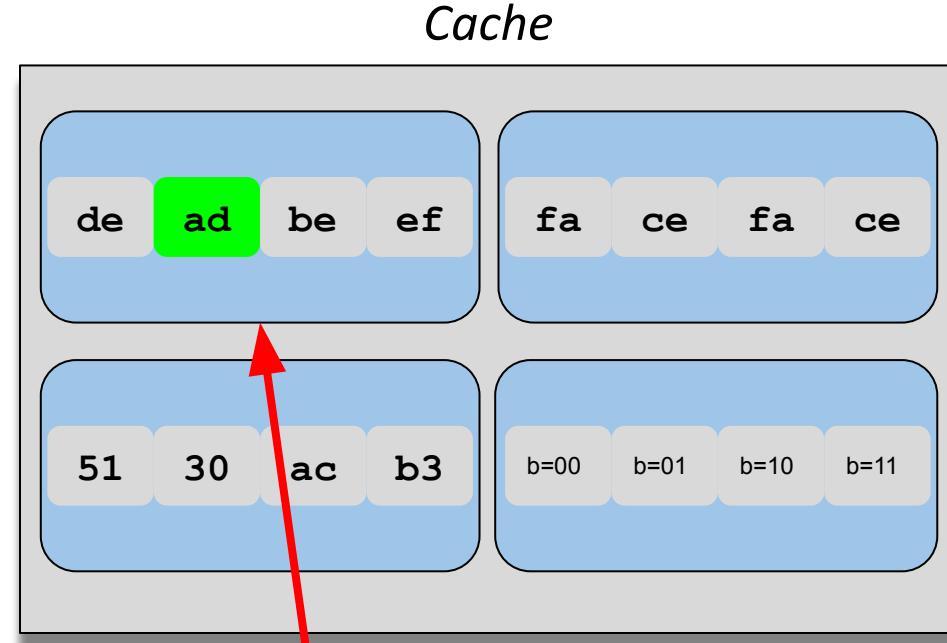
L 9,1

...

Hit!

Miss

Miss



Load new value into line

Example Trace

...

L 16,1

L 9,1

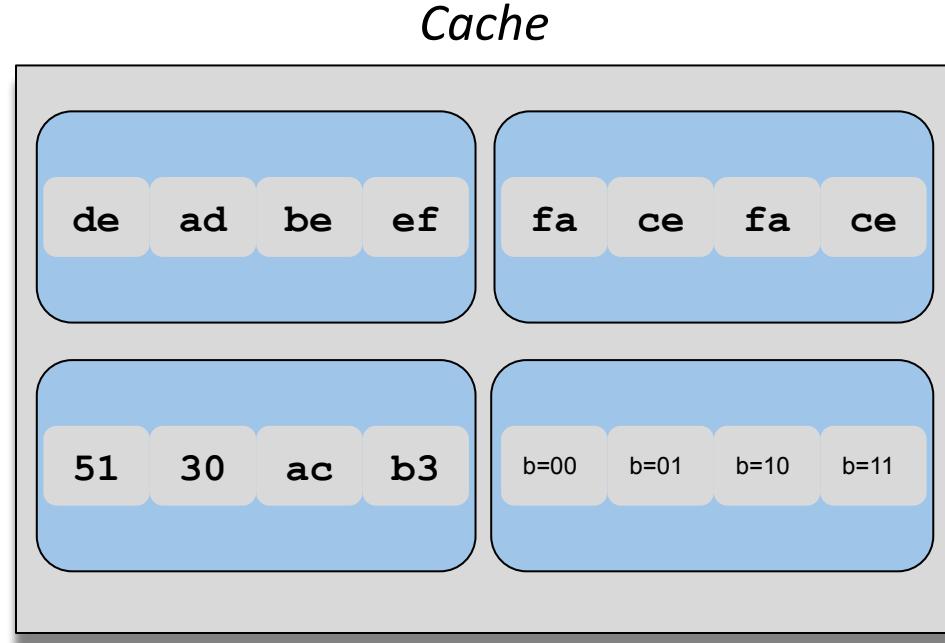
L 24,1

...

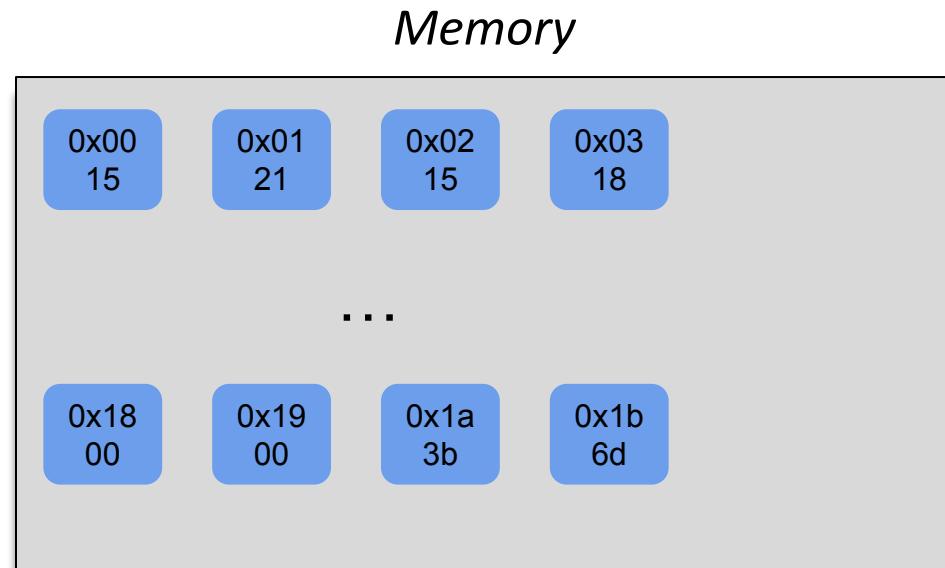
Miss

Miss

???



Will this instruction result in a hit or a miss?



Example Trace

...

L 16,1

L 9,1

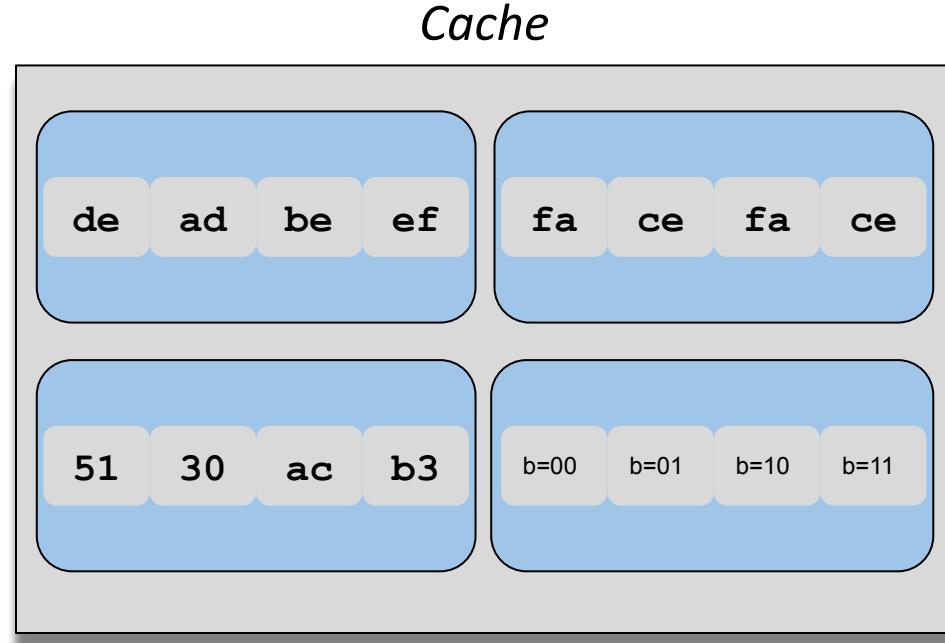
L 24,1

...

Miss

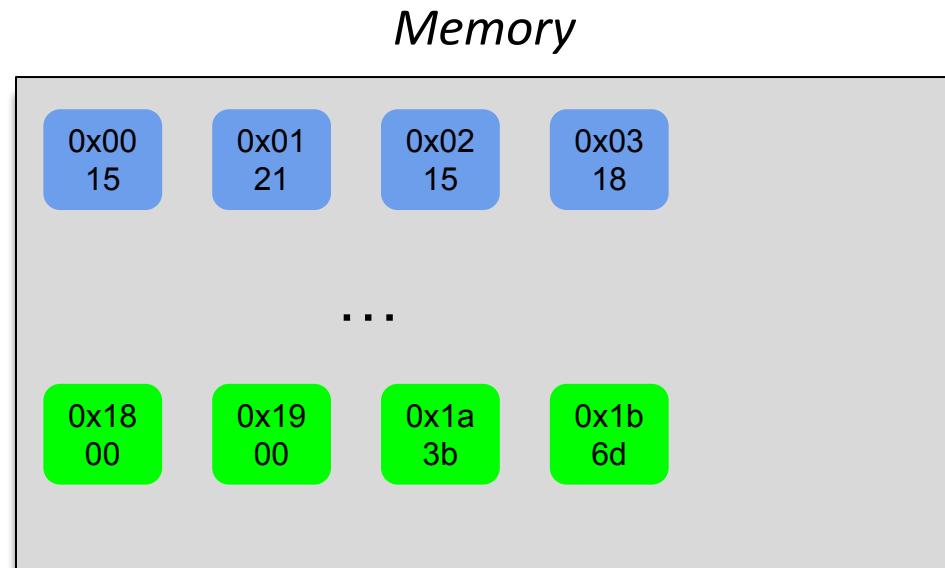
Miss

Miss



What type of miss is this?

Which line will get evicted?



Example Trace

...

L 16,1

L 9,1

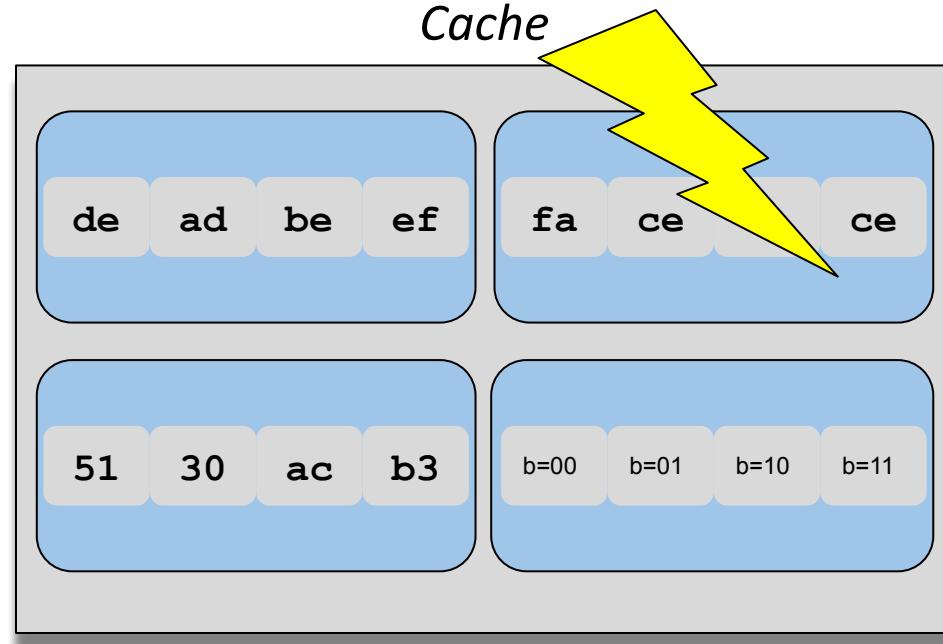
L 24,1

...

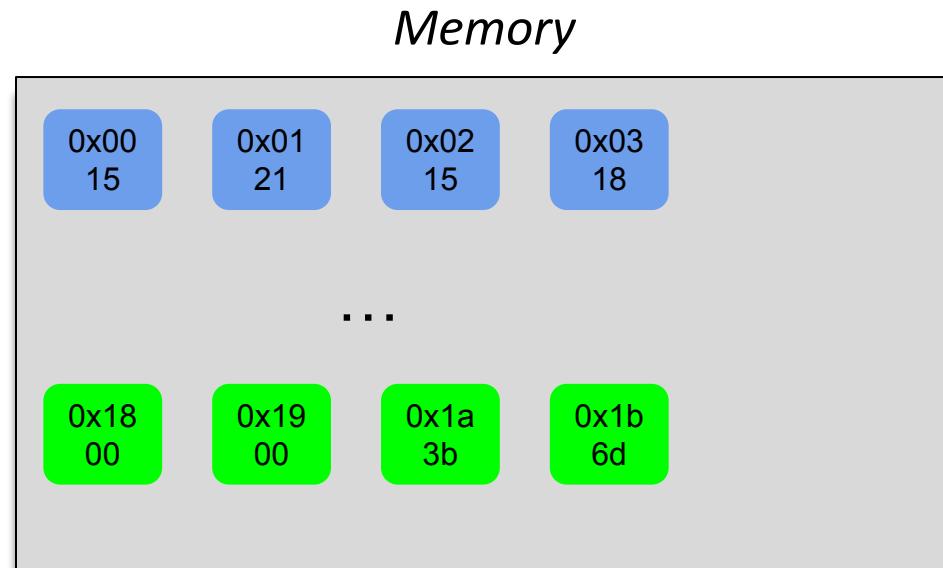
Miss

Miss

Miss



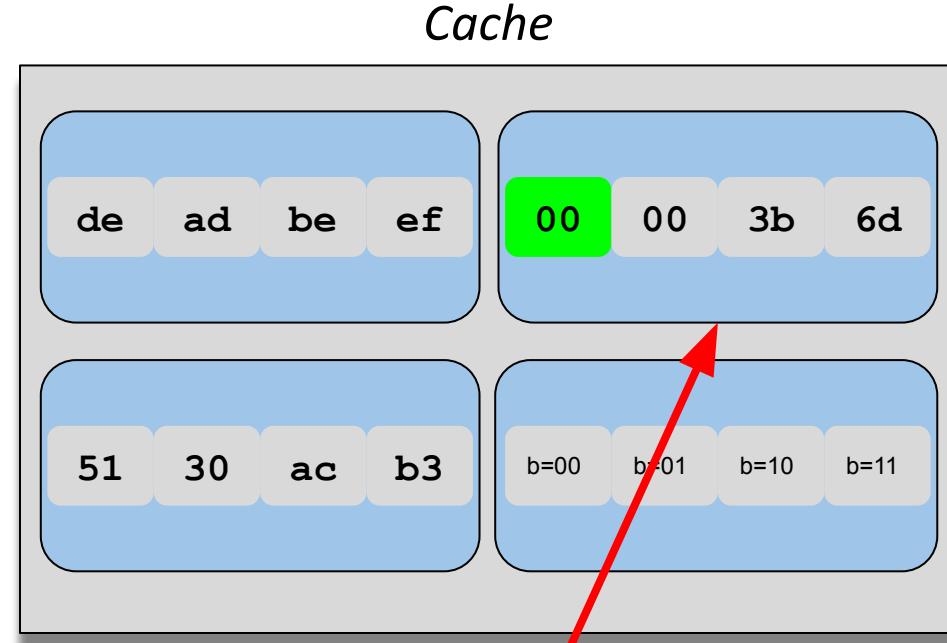
Evict least recently used line



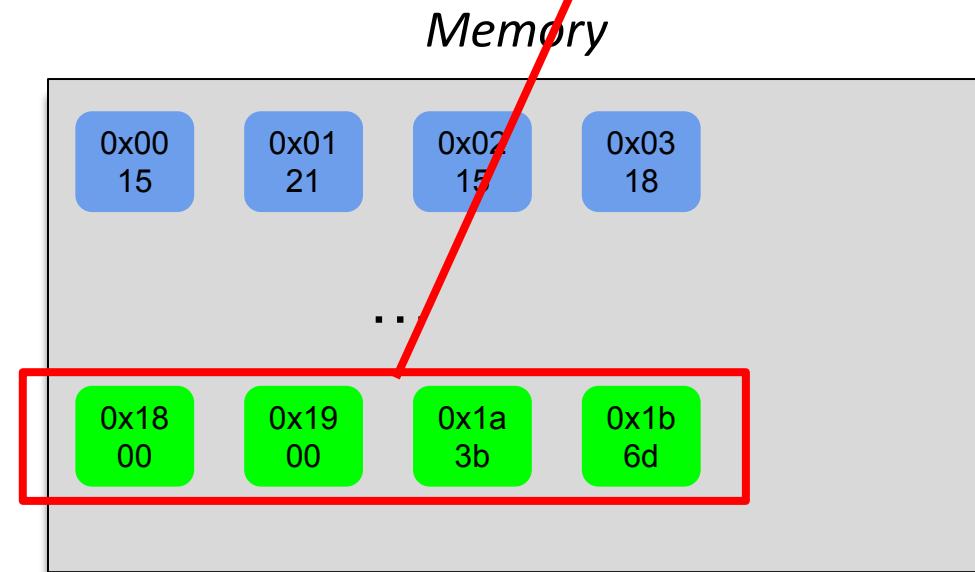
Example Trace

```
...  
L 16,1  
L 9,1  
L 24,1  
...
```

Miss
Miss
Miss



Load new value into line



Example Trace

...

L 9,1

L 24,1

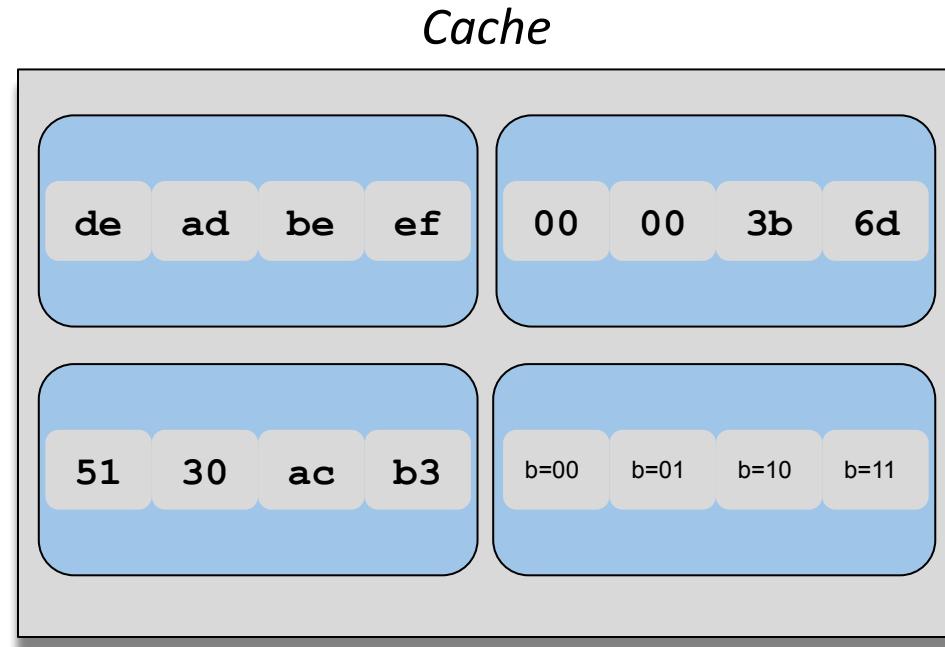
L 32,1

...

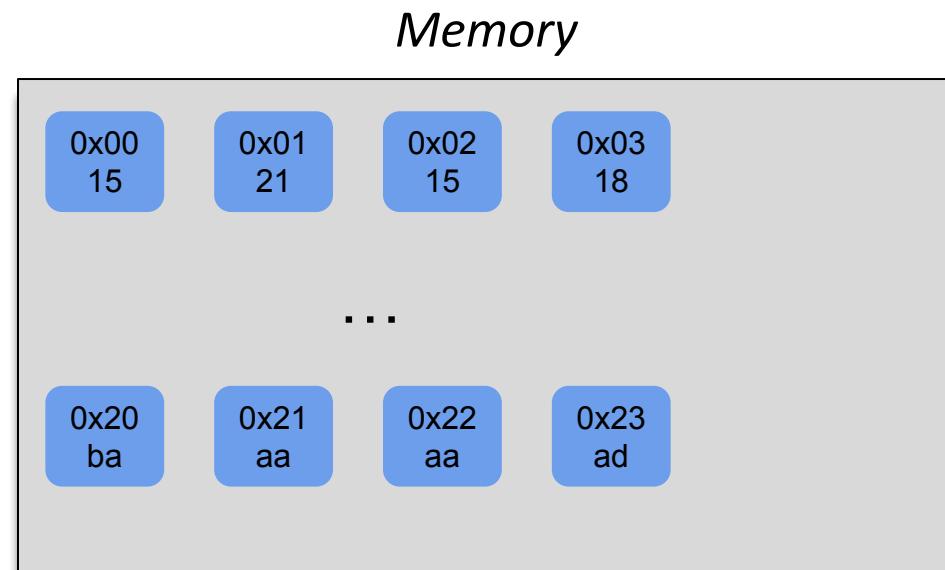
Miss

Miss

???



Will this instruction result in a hit or a miss?



Example Trace

...

L 9,1

L 24,1

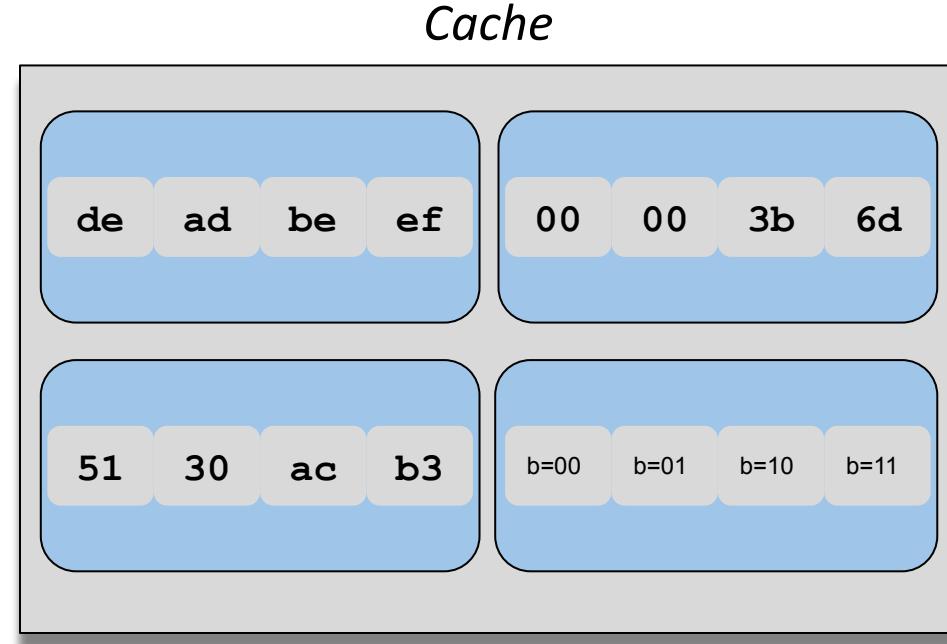
L 32,1

...

Miss

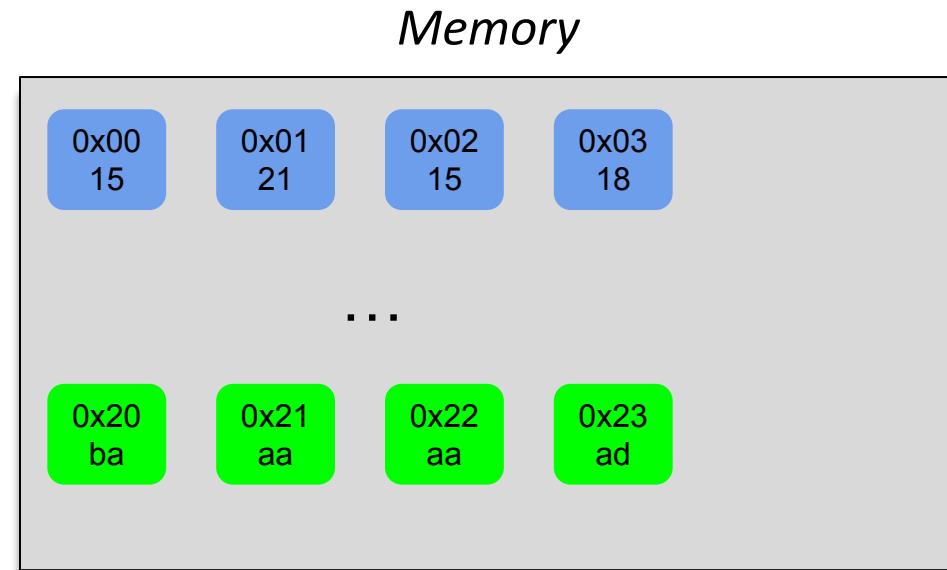
Miss

Miss



What type of miss is this?

Which line gets evicted?



Example Trace

...

L 9,1

L 24,1

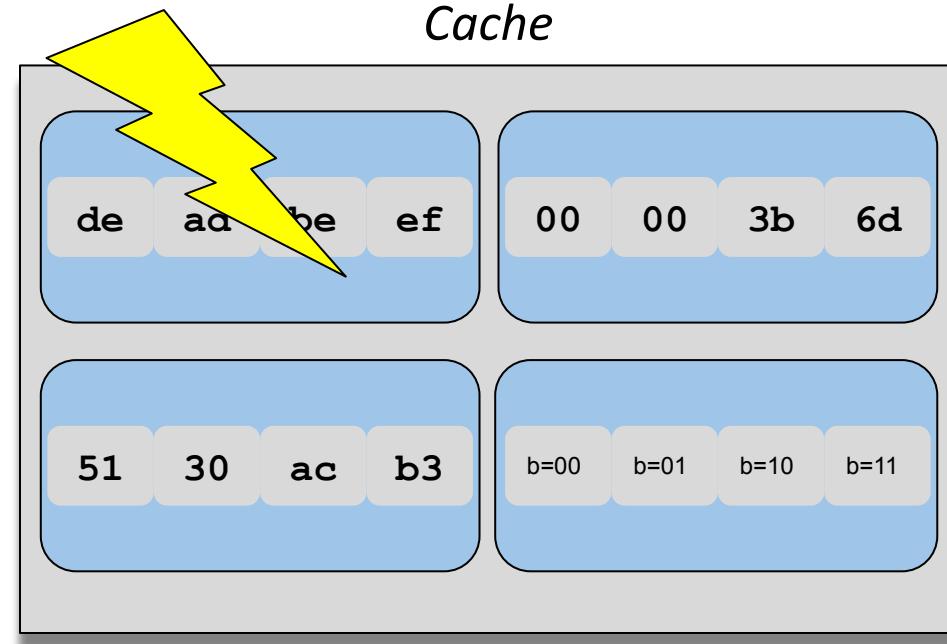
L 32,1

...

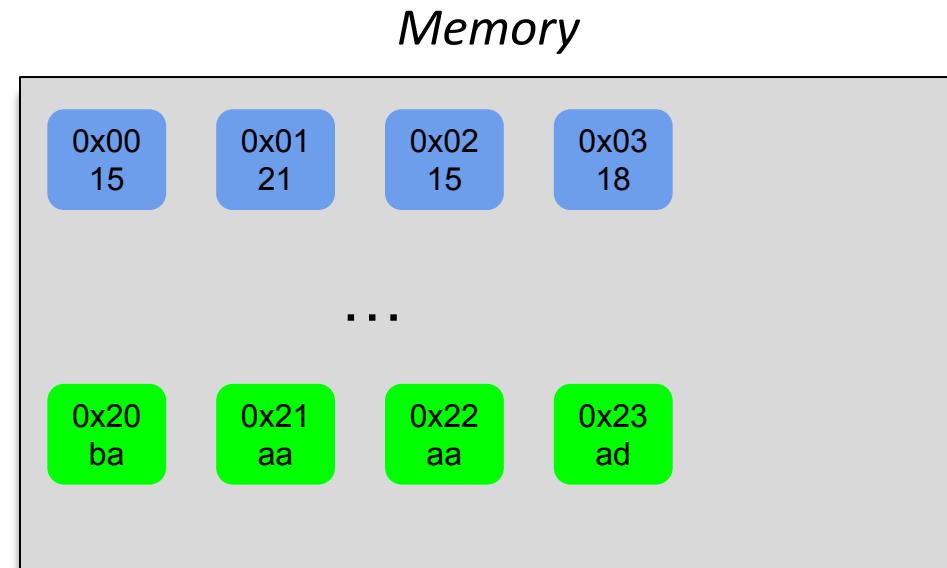
Miss

Miss

Miss



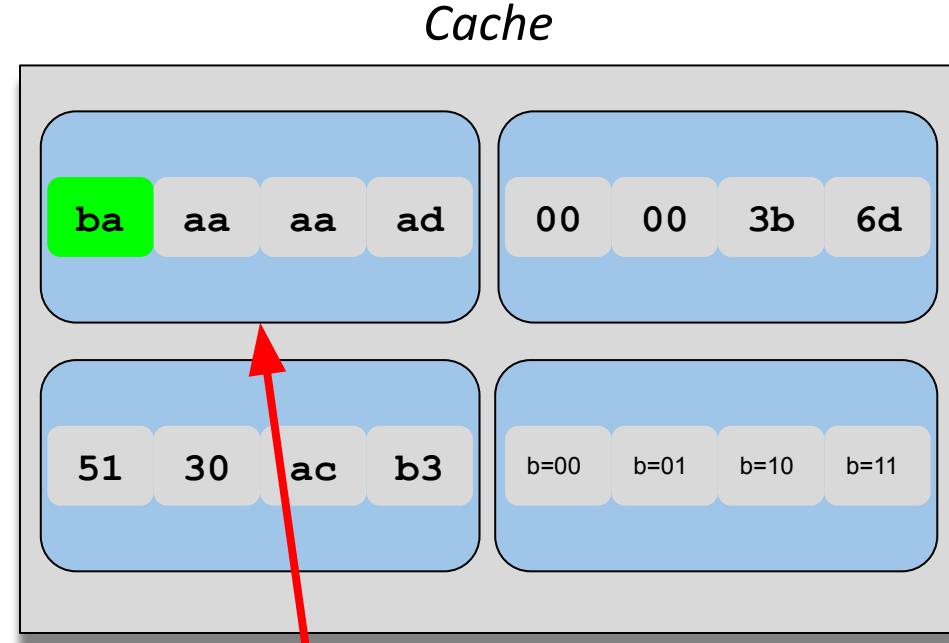
Evict least recently used line



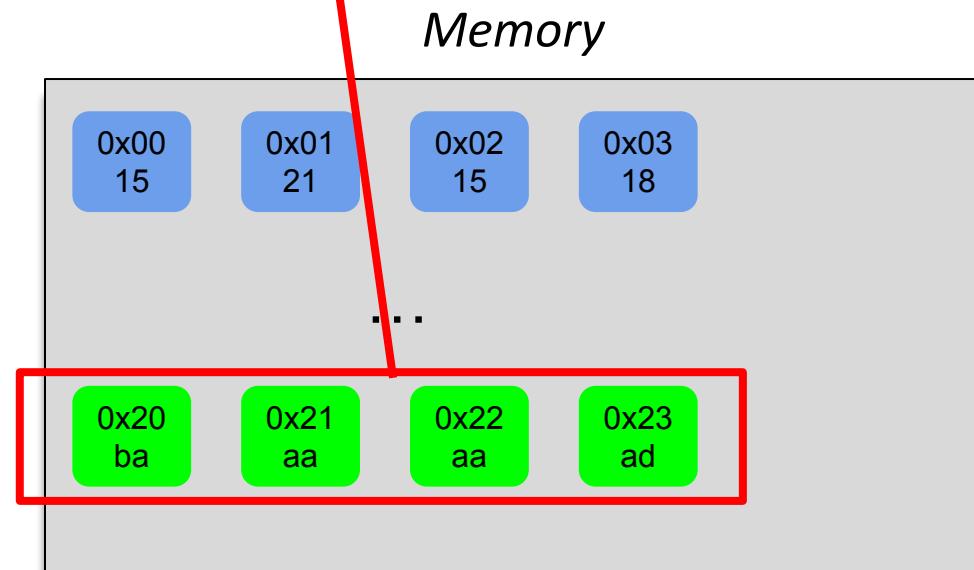
Example Trace

```
...  
L 9,1  
L 24,1  
L 32,1  
...
```

Miss
Miss
Miss



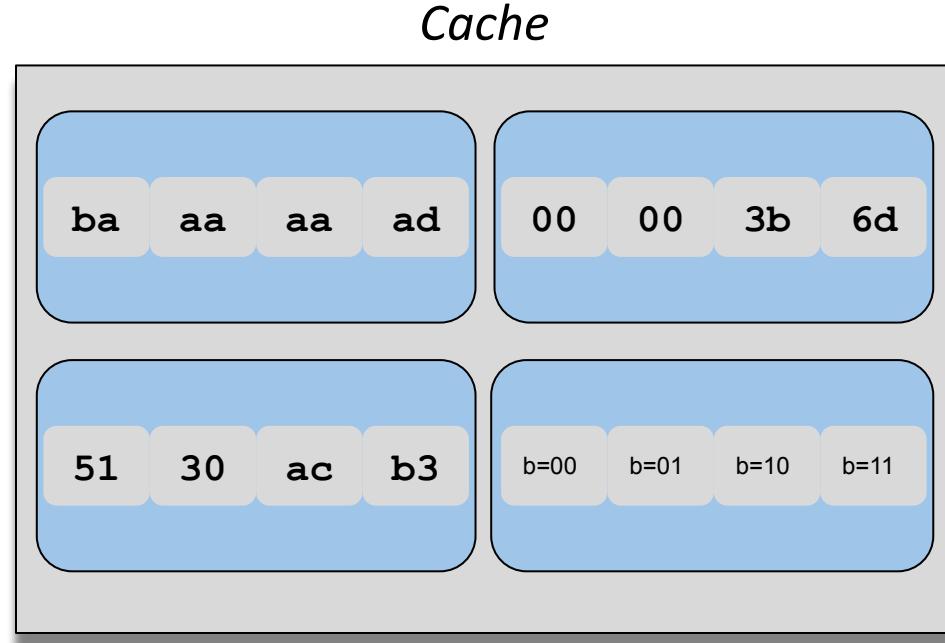
Load new value into line



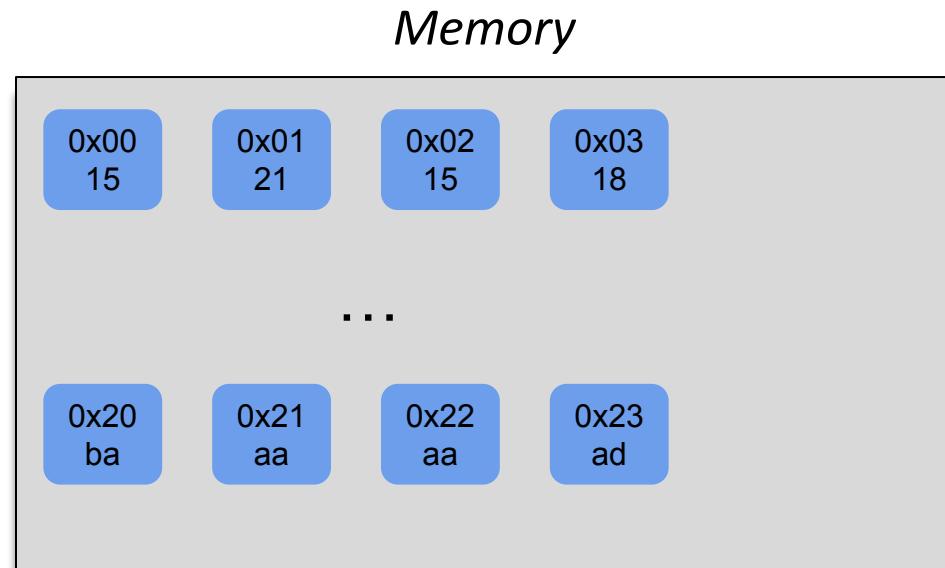
Example Trace

```
...  
L 24,1  
L 32,1  
L 0,1  
...
```

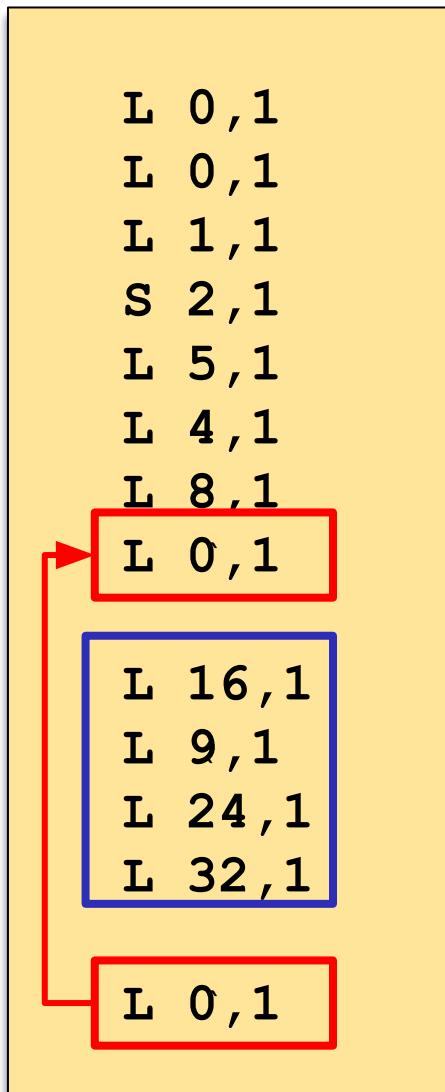
Miss
Miss
???



Will this instruction result in a hit or a miss?



Cache Concepts: Conflict/Capacity Misses

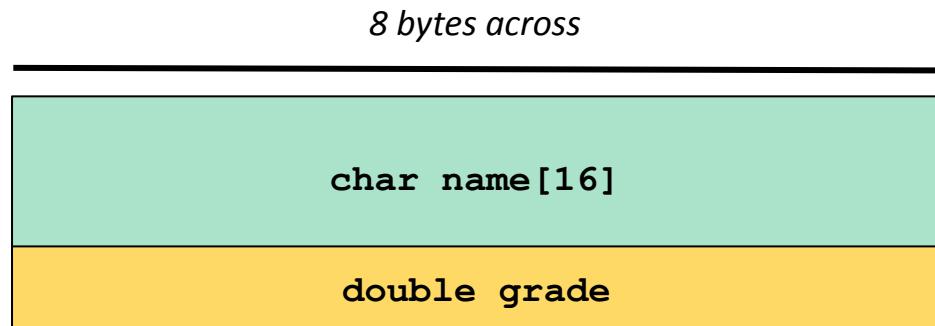


- In this case:
 - Number of unique blocks *in-between* current reference and most recent reference: 4
 - Our cache has 4 total lines
 - So: ***Capacity Miss***
- Note: the cache is not full!

Review: Programming in C

Programming in C: Structs

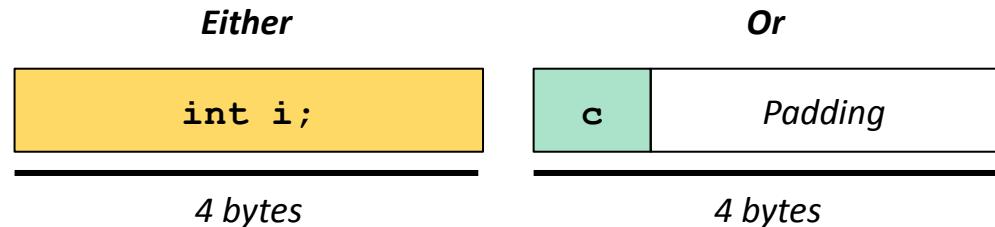
```
struct student {  
    char name[16];  
    double grade;  
};
```



- Group multiple related fields under one block of memory, at one address.
- Will probably be useful for **cachelab!**

Programming in C: Unions

```
union temp {  
    int i;  
    char c;  
};
```



- Store potentially different data types in the same region of memory.
- Specifies multiple ways to interpret data at the same memory location.

Programming in C: Style

- ***Code Reviews:*** `cachelab` will be the first lab graded for style by your TAs.
 - Comments
 - File Header
 - Modularity
 - Correctness:
 - `malloc()` can fail! Library functions can fail!
 - Memory leaks, File Descriptor leaks
 - Note: We will only review your final submission on autolab, not your github code

Using getopt()

- Let us first learn what `getopt()` does.
- Read the man pages!
 - `man 3 getopt`
 - <https://linux.die.net/man/3/getopt>

Activity: getopt_example.c

Returns **-1** when
done parsing!

```
while ((opt = getopt(argc, argv, "vn:")) != -1) {  
    <-- Omitted -->  
}
```

- Arguments are **-v** and **-n**
- Colon indicates option **-n** has required argument, which will get parsed into **optarg**.

Activity: getopt_example.c

```
while ((opt = getopt(argc, argv, "vn:")) != -1) {  
    switch (opt) {  
        case 'v':  
            verbose = 1;  
            break;  
        case 'n':  
            n = atoi(optarg);  
            break;  
        default:  
            fprintf(stderr, "usage: ...");  
            exit(1);  
    }  
}  
  
for (int i = 0; i < n; i++) {  
    if (verbose) printf("%d\n", i);  
}  
  
printf("Done counting to %d\n", n);
```

Define the optstring

Count up to -n argument

If -v argument is set,
print all numbers
before n

Cache Practice Problems

Cache Practice Problems

- We'll work through a series of questions together.
- Write down your answer to each question.
- Discuss with classmates!

Cache Practice Problem: Locality

- The following function exhibits which type of locality?

Consider ***only array accesses***.

```
void who(int *arr, int size) {  
    for (int i = 0; i < size-1; ++i)  
        arr[i] = arr[i+1];  
}
```

- A. Spatial
- B. Temporal
- C. Both spatial and temporal
- D. Neither

Cache Practice Problem: Locality

- The following function exhibits which type of locality?

Consider ***only array accesses.***

```
void who(int *arr, int size) {  
    for (int i = 0; i < size-1; ++i)  
        arr[i] = arr[i+1];  
}
```

- A. Spatial
- B. Temporal
- C. ***Both spatial and temporal***
- D. Neither

- *Spatial:* Items with nearby addresses tend to be referenced close together in time.
- *Temporal:* Recently accessed addresses tend to be accessed again in the near future.

Cache Practice Problem: Locality

- The following function exhibits which type of locality?

Consider ***only array accesses***.

```
void coo(int *arr, int size) {  
    for (int i = size-2; i >= 0; --i)  
        arr[i] = arr[i+1];  
}
```

- A. Spatial
- B. Temporal
- C. Both spatial and temporal
- D. Neither

Cache Practice Problem: Locality

- The following function exhibits which type of locality?

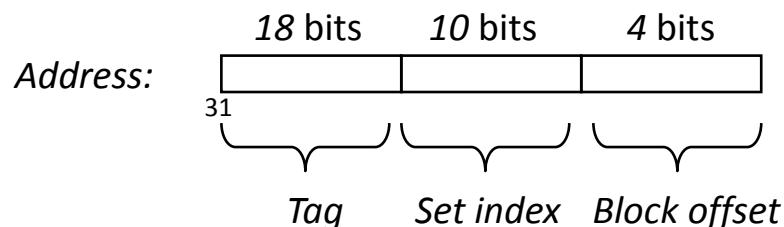
Consider ***only array accesses.***

```
void coo(int *arr, int size) {  
    for (int i = size-2; i >= 0; --i)  
        arr[i] = arr[i+1];  
}
```

- A. Spatial
- B. Temporal
- C. ***Both spatial and temporal***
- D. Neither

Cache Practice Problem: Cache Parameters

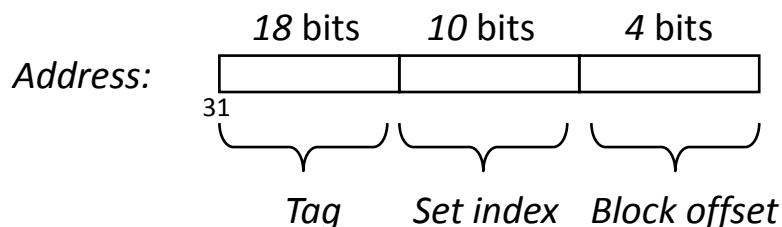
- Given the following address partition, how many int values fit in each block?



- A. 0
- B. 1
- C. 2
- D. 4
- E. Not enough information to determine

Cache Practice Problem: Cache Parameters

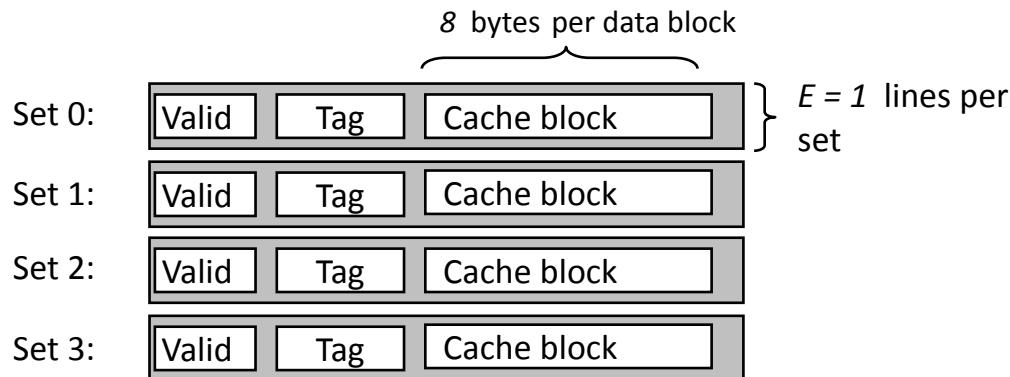
- Given the following address partition, how many int values fit in each block?



- A. 0
- B. 1
- C. 2
- D. 4
- E. Not enough information to determine

- ($b = 4$) Four Block Offset Bits
- So block size is $2^4 = 16$ bytes
- Integers are 4 bytes
- So we can fit four integers in each block.

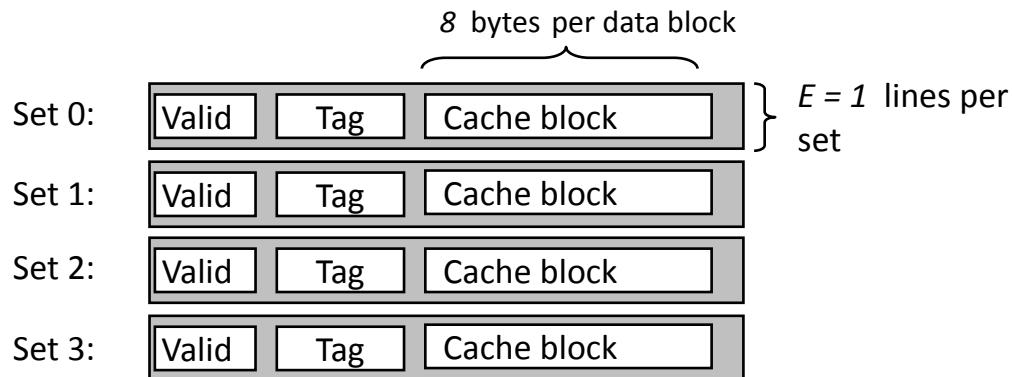
Cache Practice Problem: Cache Parameters



- What are the parameters corresponding to this cache organization?
(assume a 32 bit address space)

Option	t (# Tag Bits)	s	b
A	1	2	3
B	27	2	3
C	25	4	3
D	1	4	8
E	20	4	8

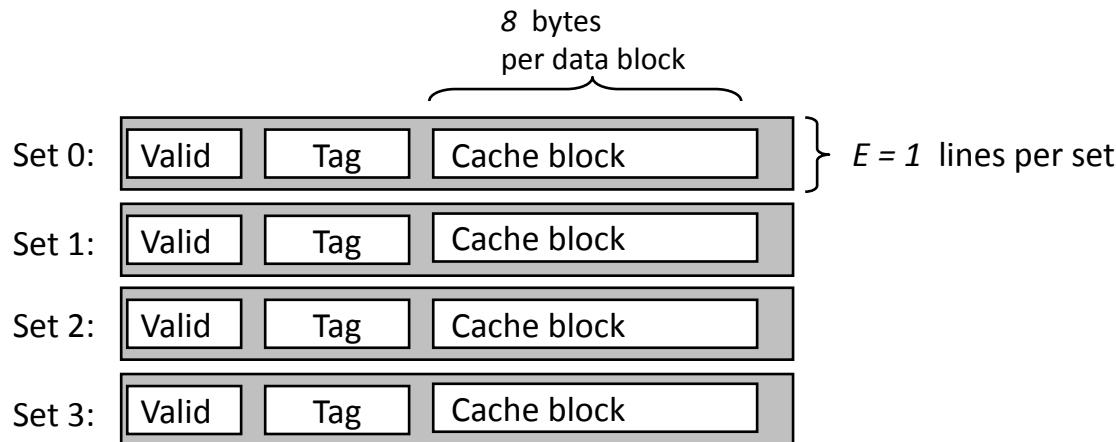
Cache Practice Problem: Cache Parameters



- What are the parameters corresponding to this cache organization?
(assume a 32 bit address space)

Option	t (# Tag Bits)	s	b
A	1	2	3
B	27	2	3
C	25	4	3
D	1	4	8
E	20	4	8

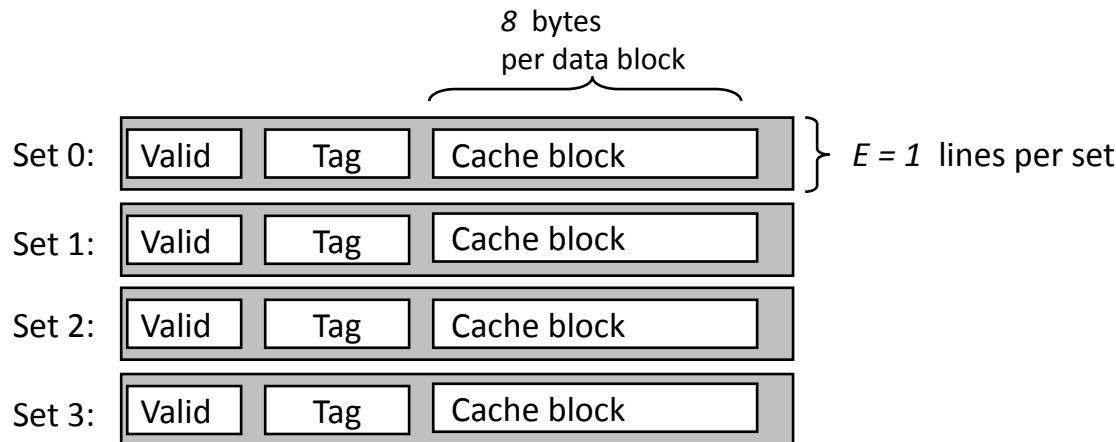
Cache Practice Problem: Which Set?



Which set does the address **0xfa1c** map to?

- A. 0
- B. 1
- C. 2
- D. 3
- E. None of the above

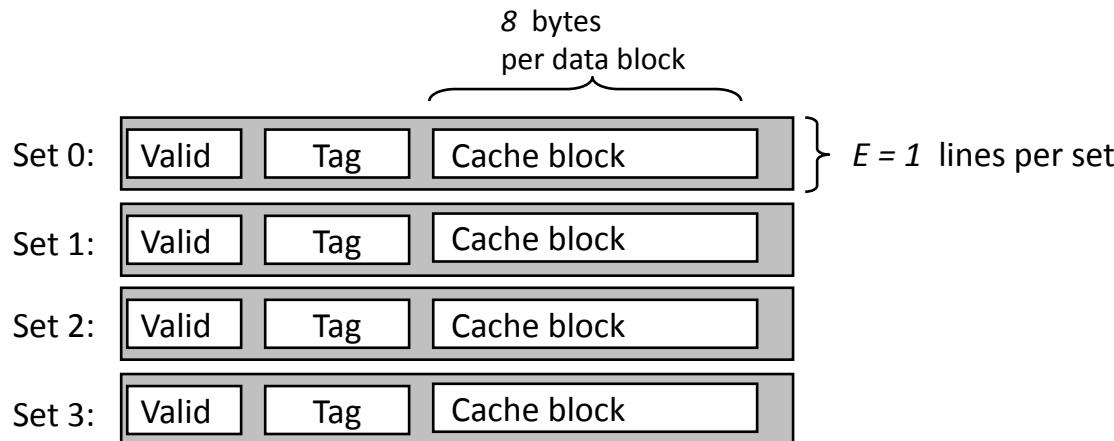
Cache Practice Problem: Which Set?



Which set does the address **0xfa1c** map to?

- A. 0
- B. 1
- C. 2
- D. 3
- E. None of the above

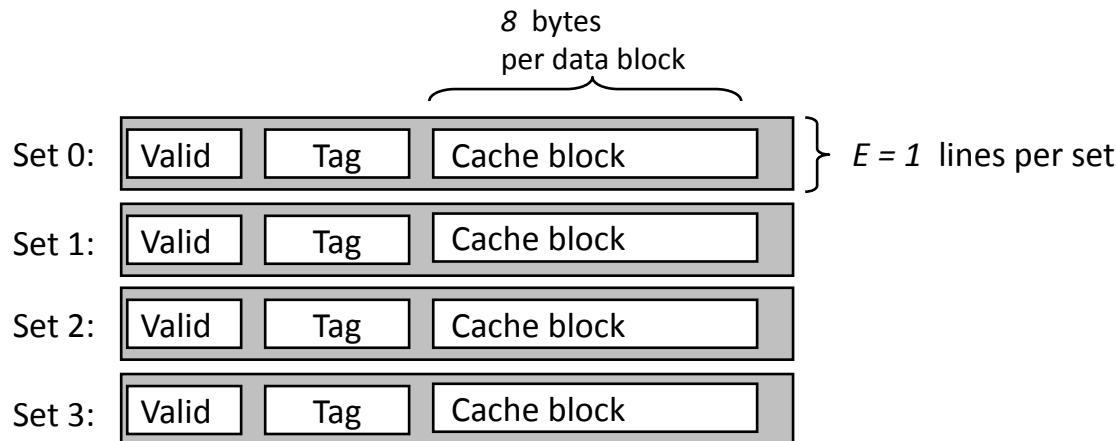
Cache Practice Problem: Range



Which range of addresses will be in the same block as **0xfa1c**?

- A. **0xfa1c**
- B. **0xfa1c-0xfa23**
- C. **0xfa1c-0xfa1f**
- D. **0xfa18-0xfa1f**
- E. It depends on the access size

Cache Practice Problem: Range



Which range of addresses will be in the same block as **0xfa1c**?

- A. **0xfa1c**
- B. **0xfa1c-0xfa23**
- C. **0xfa1c-0xfa1f**
- D. ***0xfa18-0xfa1f***
- E. It depends on the access size

Cache Practice Problem

- If $N = 16$, how many bytes of a does the loop access?

```
int foo(int* a, int N)
{
    int i;
    int sum = 0;
    for(i = 0; i < N; i++)
    {
        sum += a[i];
    }
    return sum;
}
```

- A. 4
- B. 16
- C. 64
- D. 256

Cache Practice Problem

- If $N = 16$, how many bytes does the loop access of a ?

```
int foo(int* a, int N)
{
    int i;
    int sum = 0;
    for(i = 0; i < N; i++)
    {
        sum += a[i];
    }
    return sum;
}
```

- A. 4
- B. 16
- C. 64
- D. 256

Cache Practice Problem: Cache Misses

```
void muchAccessSoCacheWow(int *bigArr) {
    // 48 KB array of ints
    int length = (48*1024)/sizeof(int);
    int access = 0;

    // traverse array with stride 8
    // pass 1
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }

    // pass 2
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }
}
```

- 32 KB cache
- 32 bit address space.
- 8-way associative
- 64 bytes per block.
- LRU

What is the miss rate on “Pass 1”?

- A. 0%
- B. 25%
- C. 33%
- D. 50%
- E. 66%

Cache Practice Problem: Cache Misses

```
void muchAccessSoCacheWow(int *bigArr) {
    // 48 KB array of ints
    int length = (48*1024)/sizeof(int);
    int access = 0;

    // traverse array with stride 8
    // pass 1
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }

    // pass 2
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }
}
```

- 32 KB cache
- 32 bit address space.
- 8-way associative
- 64 bytes per block.
- LRU

What is the miss rate on “Pass 1”?

- A. 0%
- B. 25%
- C. 33%
- D. **50%**
- E. 66%

Cache Practice Problem: Cache Misses

```
void muchAccessSoCacheWow(int *bigArr) {
    // 48 KB array of ints
    int length = (48*1024)/sizeof(int);
    int access = 0;

    // traverse array with stride 8
    // pass 1
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }

    // pass 2
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }
}
```

- 32 KB cache
- 32 bit address space.
- 8-way associative
- 64 bytes per block.
- LRU

What is the miss rate on “Pass 2”?

- A. 0%
- B. 25%
- C. 33%
- D. 50%
- E. 66%

Cache Practice Problem: Cache Misses

```
void muchAccessSoCacheWow(int *bigArr) {
    // 48 KB array of ints
    int length = (48*1024)/sizeof(int);
    int access = 0;

    // traverse array with stride 8
    // pass 1
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }

    // pass 2
    for(int i = 0; i < length; i+=8){
        access = bigArr[i];
    }
}
```

- 32 KB cache
- 32 bit address space.
- 8-way associative
- 64 bytes per block.
- LRU

What is the miss rate on “Pass 2”?

- A. 0%
- B. 25%
- C. 33%
- D. **50%**
- E. 66%

Wrapping Up

- **cachelab** tips:
 - Review Lectures
 - Start early! This lab can be challenging!
 - Don't get discouraged!
- C Programming Review materials:
 - Ed: #98
 - Keep an eye Ed for *Bootcamp 4: C Programming*.

The End!

